

TW128: A Tree-Structured Duplex Authenticated Encryption Mode

Coda Hale

Hale Institute for Bicycle Appreciation

Abstract. AEAD tags are often expected to be compact commitments to the full encryption context, an intuition reinforced by common experience with hash-based MACs such as HMAC. The AEAD literature has only recently begun to formalize this expectation through notions such as committing security and encryptment, but standard deployed modes such as AES-GCM and ChaCha20-Poly1305 do not provide it. We present TW128, a nonce-based authenticated encryption mode with associated data that adapts the KangarooTwelve tree-hashing pattern to keyed duplex encryption over KECCAK- p [1600, 12]. A serial root duplex binds the key, nonce, and associated data, encrypts the first message chunk, and aggregates tags from independently keyed leaf duplexes for the remaining chunks. The leaf decomposition is canonical, but the number of leaves evaluated at once is an implementation choice: scalar, SIMD, and multicore implementations produce identical ciphertexts and tags without making the degree of parallelism a mode parameter.

The final root tag is the authentication tag and is also analyzed as a digest-like commitment to the full (K, U, A, M) context. In the public permutation model, via an intermediate random oracle and the flat sponge lift, we prove concrete multi-user indistinguishability and derive all-in-one authenticated encryption as the single-user corollary. We also prove integrity under release of unverified plaintext in the sense of [ABL⁺14]. Separately, we prove that the wrapper returning the tag is a secure hash of the full context—collision-, second-preimage-, and preimage-resistant in the standard hash-function sense. The same tag analysis gives the AEAD-facing committing results, including full-input CMTDs-4 committing security. At the implemented parameters, all stated notions meet a 128-bit security target. Vectorized `amd64` and `arm64` implementations recover the scalable performance profile of the tree design: in bulk they reach 48 Gbit/s, with the AVX-512 and `arm64` paths delivering 51–76% of hardware-accelerated AES-128-GCM throughput.

Keywords: authenticated encryption · committing security · sponge constructions · duplex constructions · Keccak

1 Introduction

Authenticated encryption with associated data (AEAD) is often treated in practice as if its tag were a compact commitment to the full encryption context: a short string that binds the key, nonce, associated data, and message. This expectation is natural for users who first encounter message authentication through hash-based MACs such as HMAC, but it is not a property of standard deployed AEADs. AES-GCM and ChaCha20-Poly1305, for example, are not key-committing: an adversary can craft a single ciphertext that decrypts to valid plaintexts under two or more keys [DGRW18, ADG⁺22]. The cause is structural. Their authenticity rests on algebraic MACs whose outputs can often be made to validate under a second context, whereas a digest is expected to resist both collisions and preimages.

The gap is visible in deployed protocols and concrete attacks, not only in definitions. The “invisible salamander” attack defeats Facebook’s attachment franking scheme [DGRW18]; related constructions produce ciphertexts that open to different well-formed files under different keys [ADG⁺22]; partitioning oracles recover passwords from systems built on non-committing AEAD [LGR21]; and context discovery attacks compute valid decryption contexts for ciphertexts under standard modes such as CCM, EAX, and SIV, defeating the preimage half of the digest expectation [MLGR23]. These attacks differ in surface form, but all expose the same mismatch: the tag is used as if it were a full-context commitment, while the mode only proves ordinary authenticity.

Recent work has begun to model this expectation directly. Compactly committing AEAD was introduced for message franking [GLR17], and the encryption approach of [DGRW18] builds commitment into a dedicated primitive. The committing hierarchy of [BH22] makes the collision-style requirement explicit: no two distinct contexts should open the same ciphertext and tag. These notions explain when a ciphertext can have another valid opening, but they are still phrased as AEAD games over ciphertexts. The assumption made by protocols is often more tag-centric: the tag itself is treated as a compact commitment string. This paper therefore analyzes the tag with hash notions directly, and then recovers the AEAD commitment notions as game-specific consequences of the same tag analysis.

TW128 (*Tree Wrap*, with 128 denoting the 128-bit security target) takes this tag-centric view as a design constraint. If the authentication tag is to behave like a compact digest of the encryption context, then the mode should have a digest-like root rather than a serial stream cipher followed by an algebraic MAC. TW128 adapts the KangarooTwelve [BDP⁺16] tree-hashing pattern to keyed duplex encryption over KECCAK- p [1600, 12]: a serial root duplex binds the key, nonce, associated data, first message chunk, and leaf tags, while independently keyed leaf duplexes encrypt the remaining chunks. As in KangarooTwelve’s *kangaroo hopping*, a single-chunk message uses only the root duplex and pays no tree overhead (Section 8.4).

The construction fixes the root/leaf transcript, but not the amount of parallelism used by an implementation. A scalar implementation, an AVX2 implementation, an AVX-512 implementation, and a NEON implementation all produce identical ciphertexts and tags, as would a multicore implementation. The degree of parallelism is therefore not a mode parameter and is not visible on the wire. This separates TW128 from prior duplex-based AEADs with parallel variants, where the parallelism degree is fixed by the mode instance or otherwise shared by the communicating parties.

We prove that the wrapper H_{TW128} returning the root tag is a secure hash of the full (K, U, A, M) input in the sense of [RS04], and use the same root tag analysis to obtain the AEAD-facing committing results. We also prove ordinary AEAD security: in the public permutation model, TW128 has a concrete multi-user indistinguishability bound [BT16], with all-in-one authenticated encryption [RS06] as the single-user corollary. For integrity, the same root-centered structure gives INT-RUP security [ABL⁺14]: even if an adversary can obtain unverified plaintexts, it cannot make verification accept a fresh ciphertext except within the same tag-guessing and collision terms. At the implemented parameters ($k = c = \nu = \tau_{\text{root}} = \tau_{\text{leaf}} = 256$), every stated advantage is at most 2^{-128} for adversaries running up to $N_{\text{tot}} \leq 2^{63}$ KECCAK- p [1600, 12] calls—on the order of 2^{71} bytes of processed data. The cap is a data-scale choice rather than a security ceiling: the bounds are birthday-type on the 256-bit capacity and become vacuous only near 2^{128} calls (Section 7.2). Optimized amd64 and arm64 implementations preserve the canonical transcript while evaluating leaf duplexes in batches; in bulk, TW128 reaches 48 Gbit/s with AVX-512 and 39 Gbit/s on an Apple M4 Pro—51–76% of hardware-accelerated AES-128-GCM throughput (Section 8).

1.1 Our Contributions

- **A scalable tree AEAD mode.** We specify TW128, a nonce-based AEAD with associated data built as a two-level tree of strict [BDPVA11] duplexes over KECCAK- p [1600, 12] (Section 3). The mode fixes the root/leaf transcript but delegates the number of leaves evaluated simultaneously to the implementation, so scalar and vectorized implementations interoperate without changing ciphertexts or tags.
- **Digest-like commitment.** We prove that the wrapper H_{TW128} returning the root tag is a secure hash of the full (K, U, A, M) input in the sense of [RS04]: collision-, second-preimage-, and preimage-resistant (Theorems 8 to 10 and Corollary 3). Thus the authentication tag behaves like the compact digest practitioners often expect an AEAD tag to be.
- **AEAD commitment consequences.** We instantiate the same root tag analysis in the existing AEAD vocabulary. The all-bodies-equal structure of CMTDs-4 gives the full-input s -way committing bound without the leaf-collision slack of the general hash collision theorem [BH22] (Corollary 5), and the remaining CMT/CMTD notions for $\ell \in \{1, 3, 4\}$ follow from the committing hierarchy and direct source-game arguments (Corollary 6).
- **Multi-user, AE, and RUP integrity.** We give a concrete multi-user indistinguishability bound [BT16] for D users (Theorem 2). The $D = 1$ case yields a concrete all-in-one AE bound [RS06] in the public permutation model (Corollary 1). We also prove INT-RUP integrity [ABL⁺14], with plaintext-computing oracle work charged through the same transcript resources (Theorem 4). The proof proceeds through a dedicated TW128 random oracle model and is transferred to the public permutation setting by the [BDPVA08] flat sponge lift.
- **Optimized implementations.** We describe and evaluate vectorized amd64 (AVX-512 and AVX2) and arm64 (NEON with the SHA-3 extension) implementations whose batched leaf cores realize the mode’s scalable parallelism in practice (Section 8).

1.2 Related Work

Committing AEAD. The study of committing authenticated encryption was initiated by [GLR17], who introduced compactly committing AEAD for message franking, and by [DGRW18], who built it from a dedicated primitive (encryption). [BH22] organized the collision-style notions into the CMT-1/3/4 and CMTD hierarchy and gave the UtC, RtC, and HtE transforms; [CR22] gave the CTX transform, which commits a nonce-based AE scheme to its full context by adding a hash. These transforms are efficient and broadly applicable, but they treat commitment as an external layer. TW128 instead makes the authentication tag itself the hash-like commitment: the root tag is the compact commitment, and the root tag analysis gives both the [BH22] full-input committing consequence. Closest in spirit is the recent work of [DHM⁺24], which also builds nonce-based AE natively on SHA-3 functions and obtains CMT-4 from collision resistance; it targets session-based communication through a sequential duplex chain, whereas TW128 is a single scalable tree for standalone AEAD (Section 9.3).

Sponge and duplex AEADs. Permutation-based authenticated encryption built on the sponge and duplex constructions of [BDPVA08, BDPVA11] is well established. Keyak [BDP⁺15] uses the Motorist mode over round-reduced KECCAK- p and already admits parallel duplex instances; Ascon [DEMS21] and Xoodyak [DHP⁺20] are serial duplex AEADs over smaller permutations. TW128 shares the single-permutation, constant-state

duplex foundation but differs in where parallelism lives. Keyak fixes the degree of parallelism as a mode parameter, while Ascon and Xoodyak expose no mode-level parallelism. In TW128, the decomposition into indexed leaves is fixed by the transcript, but the number of leaves evaluated at once is chosen by the implementation and is invisible to the wire format (Section 9.2). Farfalle [BDH⁺17] reaches implementation-chosen, wire-invisible parallelism by a different route: it builds a pseudorandom function whose compression and expansion layers apply the permutation in parallel under rolling masks, instantiates it over KECCAK- p as Kravatte with concrete security claims, and obtains authenticated encryption through the session mode Farfalle-SAE and the synthetic-IV mode Farfalle-SIV on top of the PRF. TW128 instead retains the strict [BDPVA11] duplex object and the permutation-model analysis that comes with it, drawing its parallelism from the tree topology rather than from the primitive construction; the committing behavior of the Farfalle modes has not been analyzed.

Release of unverified plaintext. Andreeva et al. [ABL⁺14] formalized authenticated encryption in the release-of-unverified-plaintext setting by separating plaintext computation from verification. They define INT-RUP for integrity and plaintext awareness notions for privacy, and show that ordinary ciphertext integrity does not imply INT-RUP. TW128 targets the integrity side of this model: the released stream plaintext is not claimed to be private, but it does not reveal the hidden tag values needed to make a fresh ciphertext verify.

Tree hashing. TW128’s two-level shape is inspired by KangarooTwelve [BDP⁺16], a tree hash over the same permutation. TW128 adopts its choice of permutation and final-node/leaf split and adapts them to authenticated encryption: each leaf chunk is encrypted and tagged independently, leaf tags replace unkeyed chaining values, and the root is keyed by (K, U) . The construction also inherits KangarooTwelve’s short-message behavior: inputs fitting in one chunk use only the root duplex. Because TW128 has a fixed two-level topology with independently keyed, indexed leaves and phase-separated duplex calls, it does not need KangarooTwelve’s Sakura tree encoding [BDPVA14] to make its transcripts unambiguous (Section 3).

1.3 Paper Organization

Section 2 recalls the notation, sponge and duplex constructions, nonce-based AEAD and RUP integrity notions, committing security [BH22], and the hash notions of [RS04]. Section 3 specifies the TW128 construction and parameters. Section 4 fixes the scope of the security claims, the public permutation games, intermediate random oracle model, resource accounting, and loss terms. Section 5 proves multi-user indistinguishability, all-in-one AE security, and INT-RUP integrity. Section 6 proves root tag hash security and committing security, and Section 7 summarizes the concrete bounds. Section 8 reports implementation performance, and Section 9 compares TW128 to prior AEAD and committing constructions and discusses limitations outside the formal scope. The appendices collect structural lemmas, the flat sponge lift, vectorized kernel details, and test vectors.

2 Preliminaries

2.1 Notation

Strings. Write $\{0, 1\}$ for the binary alphabet. For an integer $n \geq 0$, $\{0, 1\}^n$ denotes the set of n -bit strings, $\{0, 1\}^* = \bigcup_{n \geq 0} \{0, 1\}^n$ the set of all finite-length bit strings, and ε the empty string. For $x \in \{0, 1\}^*$, $|x|$ is the bit length and $x \parallel y$ is the concatenation;

the bitwise XOR of two equal-length strings $x, y \in \{0, 1\}^n$ is $x \oplus y$. A *byte string* is an element of $(\{0, 1\}^8)^*$; its byte length is denoted $\|x\| = |x|/8$. For an integer $0 \leq n < 2^{64}$, $\langle n \rangle_8 \in \{0, 1\}^{64}$ is its eight-byte little-endian encoding. For a finite or infinite bit string x and $0 \leq \tau \leq |x|$ when x is finite, $\lfloor x \rfloor_\tau \in \{0, 1\}^\tau$ denotes the leading τ bits of x .

Sampling and probability. We write $x \leftarrow \$ S$ for sampling x uniformly at random from a finite set S , with implicit coins independent across samples. For an event E in a probability experiment, $\Pr[E]$ denotes its probability.

Sets of maps. For a set S , $\text{Perm}(S)$ denotes the set of permutations on S ; for a positive integer b , $\text{Perm}(\{0, 1\}^b)$ is the set of b -bit permutations.

Games and advantages. We write $\mathcal{A}^{O_1, \dots, O_k}$ for a (possibly randomized) adversary with oracle access to $O_1, \dots, O_k$, and $\Pr[\mathbf{G}^{\mathcal{A}} \Rightarrow 1]$ for the probability that a game $\mathbf{G}$ —an experiment that runs $\mathcal{A}$ against specified oracles and returns a Boolean outcome—returns 1. For a distinguishing experiment between two games $\mathbf{G}_0, \mathbf{G}_1$, the advantage of $\mathcal{A}$ against a scheme Π in notion $\mathbf{X}$ is the absolute difference

$$\text{Adv}_{\Pi}^{\mathbf{X}}(\mathcal{A}) = |\Pr[\mathbf{G}_0^{\mathcal{A}} \Rightarrow 1] - \Pr[\mathbf{G}_1^{\mathcal{A}} \Rightarrow 1]|.$$

For a win-event experiment without a hidden challenge bit it is instead $\Pr[\mathbf{G}^{\mathcal{A}} \text{ wins}]$, the probability the game returns 1; a bit-guessing experiment with a hidden challenge bit uses the rescaled form $|\Pr[\mathbf{G}^{\mathcal{A}} \text{ wins}] - 1/2|$, as in the multi-user notion of Section 2.4.

Concrete security. All bounds in this paper are concrete: each advantage is an explicit function of the resource caps introduced in Section 4.4. We make no use of asymptotic or negligibility conventions.

Reference summary. Table 1 summarizes the resource caps and loss terms used throughout. The caps are defined formally in Section 4.4, the loss terms in Section 4.5.

Table 1: Resource caps and loss terms.

Symbol	Meaning	Defined
N	Construction interface cost (KECCAK- p [1600, 12] calls)	§ 4.4
N_{prim}	Direct primitive-interface queries	§ 4.4
N_{tot}	$N + N_{\text{prim}}$ (total query-sequence cost)	§ 4.4
q_v, Q_v	Valid verification queries / cap	§ 4.4
$q_{\text{RO}}, Q_{\text{RO}}$	Direct random oracle queries / cap (RO-model)	§ 4.4
$n_{\text{leaf}}, N_{\text{leaf}}$	Distinct construction-generated final leaf transcripts / cap	§ 4.4
$\text{Lift}_c(N_{\text{tot}})$	[BDPVA08] flat sponge loss	§ 4.5
$\text{KeyGuess}_k(Q)$	Ideal system key guess	§ 4.5
$\text{RootColl}_\tau(Q, s)$	s -way root-collision	§ 4.5
$\text{RootPre}_\tau(Q)$	Root-preimage	§ 4.5
$\text{LeafColl}_\tau(Q)$	Known-key leaf-collision	§ 4.5

2.2 Sponges and Duplexes over Keccak- p

The Keccak- p permutation. The KECCAK- p family of permutations, defined in FIPS 202 [Nat15] as the generalization of the Keccak design’s Keccak- f permutations to a free round count, is parameterized by a state width b and a round count n_r ; we write $\text{KECCAK} - p[b, n_r] \in \text{Perm}(\{0, 1\}^b)$ for the corresponding permutation. TW128 uses

214 KECCAK- p [1600, 12], the same primitive as KangarooTwelve [BDP⁺16], fixed throughout.
 215 In the public permutation security model of Section 4.1 and the flat sponge lift of Section B,
 216 this permutation is modeled as a uniformly random element of $\text{Perm}(\{0, 1\}^{1600})$.

217 **The sponge construction.** Following [BDPVA08], fix a permutation $\pi \in \text{Perm}(\{0, 1\}^b)$,
 218 a *sponge-compliant* padding rule pad , and a rate $1 \leq r < b$ with capacity $c = b - r$. A
 219 padding rule appends $\text{pad}[r](|M|)$, a bit string determined by the input length and the
 220 rate, to an input M , extending it to a sequence of r -bit blocks; it is sponge-compliant
 221 [BDPVA11, Definition 1] if $M \parallel \text{pad}[r](|M|)$ is never empty and, for all $M \neq M'$ and all
 222 $n \geq 0$, $M \parallel \text{pad}[r](|M|) \neq M' \parallel \text{pad}[r](|M'|) \parallel 0^{nr}$. The sponge function $\text{Sponge}[\pi, \text{pad}, r] :$
 223 $\{0, 1\}^* \times \mathbb{N} \rightarrow \{0, 1\}^*$ maps an input M and an output length ℓ to an ℓ -bit string by
 224 absorbing $M \parallel \text{pad}[r](|M|)$ into the all-zero b -bit state in r -bit blocks (each separated by an
 225 application of π) and squeezing ℓ bits in r -bit blocks from the resulting state. [BDPVA08,
 226 Theorem 2] shows that $\text{Sponge}[\pi, \text{pad10}^*, r]$, where the simple rule pad10^* appends a single
 227 1 bit followed by the minimum number of 0 bits reaching a multiple of r , is indistinguishable
 228 from a random oracle when π is uniformly random. We bound its distinguishing advantage
 229 by $\text{Lift}_c(N_{\text{tot}}) = N_{\text{tot}}(N_{\text{tot}} + 1)/2^{c+1}$, derived in Lemma 20 from the product upper bounds
 230 of [BDPVA08, Lemmas 4 and 5].

231 **The duplex construction.** The *duplex* construction of [BDPVA11] extends the sponge
 232 to a stateful interface. A duplex object $\text{Duplex}[\pi, \text{pad}, r]$ is initialized to the all-zero state.
 233 Each call $\text{duplexing}(\sigma, \ell)$, with $\ell \leq r$, absorbs an input string σ of at most $\rho_{\max}(\text{pad}, r)$
 234 bits via π and returns the first ℓ bits of the resulting state. Here $\rho_{\max}(\text{pad}, r) = \min\{x :$
 235 $x + |\text{pad}[r](x)| > r\} - 1$ is the largest input length for which the padded input fits in a single r -
 236 bit block; for the pad10^*1 rule, $\rho_{\max}(\text{pad10}^*1, r) = r - 2$ [BDPVA11]. [BDPVA11, Lemma 3]
 237 establishes the duplex-to-sponge equality: for any call sequence $((\sigma_0, \ell_0), \dots, (\sigma_i, \ell_i))$, the
 238 i -th output equals

$$\text{Sponge}[\pi, \text{pad}, r](\sigma_0 \parallel \text{pad}_0 \parallel \dots \parallel \sigma_i, \ell_i),$$

240 where $\text{pad}_j = \text{pad}[r](|\sigma_j|)$ is the padding the duplex appends in its j -th call; the flattened
 241 input carries these interior paddings explicitly, and the sponge appends only the final pad_i .
 242 The flattening map is injective by [BDPVA11, Lemma 4]: the duplex input-string sequence
 243 is recoverable from the flattened sponge input.

244 **TW128 instantiation.** TW128 fixes the padding rule pad10^*1 and the rate $r = 1344$,
 245 so the capacity is $c = 256$ bits. The largest input length per duplex call is therefore
 246 $\rho_{\max}(\text{pad10}^*1, 1344) = 1342$ bits.

247 **Bridge to the H -input domain.** The indistinguishability theorem of [BDPVA08] is stated
 248 for the simpler pad10^* padding over its unpadded H -input domain, while TW128 uses
 249 pad10^*1 internally. [BDPVA11, Theorem 3] supplies an injective preprocessing bridge
 250 $I[r, r_{\max}]$ between the two domains. For TW128, $r = r_{\max} = 1344$, so the bridge reduces
 251 to appending $1 \parallel 0^q$ with q determined by the input length modulo $r_{\max}$; on the bridge
 252 image, q is recovered as the trailing-zero count. Section 4.3 instantiates this bridge on
 253 typed TW128 duplex transcript prefixes as the map $\text{flat}(\cdot)$ used by the random oracle
 254 proofs throughout the paper.

255 2.3 Authenticated Encryption with Associated Data

256 **Syntax.** A *nonce-based authenticated encryption scheme with associated data* (AEAD
 257 scheme) is a pair $\text{SE} = (\text{Enc}, \text{Dec})$ of deterministic algorithms with associated key space
 258 $\mathcal{K} \subseteq (\{0, 1\}^8)^*$, nonce space $\mathcal{U} \subseteq (\{0, 1\}^8)^*$, associated data space $\mathcal{A} \subseteq (\{0, 1\}^8)^*$, message
 259 space $\mathcal{M} \subseteq (\{0, 1\}^8)^*$, and ciphertext space $\mathcal{C} \subseteq (\{0, 1\}^8)^*$. The encryption algorithm

260 $\text{Enc} : \mathcal{K} \times \mathcal{U} \times \mathcal{A} \times \mathcal{M} \rightarrow \mathcal{C}$ produces a ciphertext $C = \text{Enc}_K(U, A, M)$. The decryption
 261 algorithm $\text{Dec} : \mathcal{K} \times \mathcal{U} \times \mathcal{A} \times \mathcal{C} \rightarrow \mathcal{M} \cup \{\perp\}$ returns either a plaintext or the rejection
 262 symbol $\perp$. The *tag length in bytes* is the constant $\|\text{Enc}_K(U, A, M)\| - \|M\|$, fixed by the
 263 scheme. We write nonce values as U to keep N available for the construction-cost resource
 264 variable used in the concrete bounds. In this generic syntax, C denotes the full AEAD
 265 ciphertext, including any tag material. The TW128 algorithm section (Section 3) writes
 266 this same value as $C \parallel T$, with C the ciphertext body and T the root tag.

267 **Correctness.** For every $(K, U, A, M) \in \mathcal{K} \times \mathcal{U} \times \mathcal{A} \times \mathcal{M}$, $\text{Dec}_K(U, A, \text{Enc}_K(U, A, M)) = M$.
 268 An AEAD scheme is *tidy* if every accepting decryption is consistent with encryption:
 269 whenever $\text{Dec}_K(U, A, C) = M \neq \perp$, one has $\text{Enc}_K(U, A, M) = C$.

270 **All-in-one AE.** We use the all-in-one authenticated encryption notion of [RS06], instan-
 271 tiated for nonce-based AE following [BT16]. The experiment grants either the real pair
 272 $(\text{Enc}_K, \text{Ver}_K)$ or the ideal pair $(\text{Rand}, \text{Replay})$, where Rand returns a fresh uniform byte
 273 string of the ciphertext length on each encryption query and Replay accepts exactly the
 274 tuples Rand produced. The advantage $\text{Adv}_{\text{SE}}^{\text{ae}}(\mathcal{A})$ is the distinguishing advantage between
 275 these two worlds over nonce-respecting adversaries. This all-in-one notion is instantiated
 276 for TW128 in Section 4.1.

277 **Integrity under release of unverified plaintext.** For the RUP setting of [ABL⁺14], a
 278 conventional decryption algorithm is separated into a plaintext-computing algorithm and
 279 a verification algorithm. The integrity notion INT-RUP gives the adversary access to
 280 encryption, plaintext computation, and verification, and asks whether verification accepts
 281 a fresh ciphertext and tag not returned by encryption. As in the nonce-IV setting of
 282 [ABL⁺14], encryption queries are nonce-respecting, while plaintext computation and
 283 verification may be queried on arbitrary nonces. This notion protects only ciphertext
 284 integrity; privacy under release of unverified plaintext requires plaintext awareness in
 285 addition to ordinary encryption privacy.

286 2.4 Multi-User Indistinguishability

287 The multi-user indistinguishability notion of [BT16, Fig. 3] extends the same real-vs-ideal
 288 oracle pair to an adversary-selected set of users. The number of users u is adversary-
 289 determined: it is the number of **New** queries the adversary makes, i.e., the final value of
 290 the user counter v in the game below. A challenge bit $b \leftarrow \{0, 1\}$ selects between the real
 291 construction interface (real ciphertexts, real verification) and an ideal interface (uniform
 292 strings of matching length, verification that accepts only previously-issued encryption
 293 tuples). The adversary also accesses the public primitive interface (π, π^{-1}) , exposed
 294 identically in both worlds. This follows the same-oracle convention of [BT16], who work in
 295 the ideal-cipher model and give both worlds the same ideal-cipher oracles; here the shared
 296 primitive is the random permutation. The adversary outputs a guess b' and wins if $b' = b$.

297 **Game and advantage.** The game $\text{G}_{\text{SE}}^{\text{mu-ind}}(\mathcal{A})$ of [BT16, Fig. 3] samples $b \leftarrow \{0, 1\}$ and
 298 maintains a user counter v , a set V of (d, U) pairs already used by encryption, and a set W
 299 of recorded encryption tuples; its **New**, **Enc**, and **Vf** oracles are instantiated for TW128 in
 300 Section 4.2, together with the public primitive interface (π, π^{-1}) , which does not depend
 301 on the challenge bit. We write the user index as d and the associated data input as A ,
 302 where [BT16] writes i and H , and take the absolute value of their signed advantage so that

$$\begin{aligned} \text{Adv}_{\text{SE}}^{\text{mu-ind}}(\mathcal{A}) &= |2 \Pr[\text{G}_{\text{SE}}^{\text{mu-ind}}(\mathcal{A}) \text{ returns } 1] - 1| \\ &= |\Pr[\mathcal{A}^{\text{Enc}, \text{Vf}, b=1} \Rightarrow 1] - \Pr[\mathcal{A}^{\text{Enc}, \text{Vf}, b=0} \Rightarrow 1]| \end{aligned}$$

the second equality holding because b is uniform. Thus $\text{Adv}_{\text{SE}}^{\text{mu-ind}}$ is the two-game distinguishing advantage between the real (Enc, Ver) pair and the ideal (Rand, Replay) pair instantiated per user. The single-user case $u = 1$ recovers the all-in-one AE notion of Section 2.3, augmented with the public primitive interface (the real permutation in both worlds).

Per-user nonce uniqueness. Nonce uniqueness is enforced per user via the set V : the same nonce may appear under distinct users but not twice under the same user. This is strictly weaker than the global nonce-uniqueness requirement that would force each nonce value U to appear under at most one user.

2.5 Committing Notions and Root Tag Hash Security

[BH22] formalize a family of committing security notions for nonce-based AEAD schemes, parameterized by how much of the encryption input is committed and by whether commitment is established through decryption (the D-notions) or through encryption agreement (the E-notions). We recall their definitions and the s -way generalization, and then frame the root tag itself as a hash function whose [RS04] security—collision, second-preimage, and preimage resistance—captures the binding and preimage properties expected of a compact commitment. The AEAD committing notions are then proved from the same tag analysis in Section 6.

What-is-committed functions. For $\ell \in \{1, 3, 4\}$, the *what-is-committed* function WiC_ℓ projects an encryption input $(K, U, A, M) \in \mathcal{K} \times \mathcal{U} \times \mathcal{A} \times \mathcal{M}$ to the portion it should commit:

$$\begin{aligned}\text{WiC}_1(K, U, A, M) &= K, \\ \text{WiC}_3(K, U, A, M) &= (K, U, A), \\ \text{WiC}_4(K, U, A, M) &= (K, U, A, M).\end{aligned}$$

CMTDs- ℓ and CMTs- ℓ games. For $s \geq 2$, [BH22, Fig. 5] defines the s -way committing games on tuples $(K_i, U_i, A_i, M_i)_{i=1}^s$ with pairwise distinct WiC_ℓ -images, all entries in the scheme's syntactic domain and none equal to $\perp$. The s -way CMTDs- ℓ (decryption) game has the adversary output a ciphertext C alongside the tuples and returns 1 if $\text{Dec}_{K_i}(U_i, A_i, C) = M_i$ for every i ; the printed return line of [BH22, Fig. 5] writes C_i , but the game outputs a single ciphertext C . The s -way CMTs- ℓ (encryption) game has the adversary output the tuples alone and returns 1 if all encryptions $\text{Enc}_{K_i}(U_i, A_i, M_i)$ are equal. Advantages $\text{Adv}_{\text{SE}}^{\text{CMTDs-}\ell}(\mathcal{A})$ and $\text{Adv}_{\text{SE}}^{\text{CMTs-}\ell}(\mathcal{A})$ are the probabilities that the adversary wins the respective game; the parameter s is suppressed from the notation when clear from context. The two-tuple case $s = 2$ is the original CMTD- ℓ /CMT- ℓ pair of [BH22, Fig. 4]. These definitions are instantiated for TW128 in Section 4.1, with the public permutation lift deferred to that section.

Definition 1 (Root tag wrapper). For every TW128 encryption input $X = (K, U, A, M)$ in the scheme domain, define the root tag wrapper by

$$\text{H}_{\text{TW128}}(K, U, A, M) = T \quad \text{where} \quad \text{Enc}_K(U, A, M) = C \parallel T.$$

We also write $\text{H}_{\text{TW128}}(X)$ for $\text{H}_{\text{TW128}}(K, U, A, M)$. Equivalently, H_{TW128} runs TW128 encryption and discards the ciphertext body.

The root tag as a hash. Many uses of commitment store, transmit, or compare a short tag rather than the full ciphertext, and ask that the short string alone be binding—the compactly committing AEAD goal introduced by [GLR17] for message franking. TW128

ciphertexts are written $C \parallel T$ with T the τ_{root} -bit root tag, and we capture this goal directly by treating Definition 1 as a hash function. Keyed by the permutation, this is a hash family, and we ask that it satisfy the hash-function security notions of [RS04], over the TW128 input domain, together with the local s -way multicollision extension used for the committing theorems. For fixed byte lengths a and μ , write

$$\mathcal{X}_{a,\mu} = \{0,1\}^k \times \{0,1\}^\nu \times (\{0,1\}^8)^a \times (\{0,1\}^8)^\mu$$

for the fixed-length slice of valid inputs, with bit length $m_{a,\mu} = k + \nu + 8a + 8\mu$.

Collision resistance (Coll). No adversary finds distinct inputs $X_1, \dots, X_s$ (any $s \geq 2$) with $H_{\text{TW128}}(X_1) = \dots = H_{\text{TW128}}(X_s)$. The case $s = 2$ is the RS04 Coll notion; larger s is the corresponding multicollision extension.

Preimage resistance (Pre). For fixed a, μ , given $H_{\text{TW128}}(X^*)$ for a uniform $X^* \leftarrow \$ \mathcal{X}_{a,\mu}$, no adversary finds any valid input X with $H_{\text{TW128}}(X) = H_{\text{TW128}}(X^*)$.

Second-preimage resistance (Sec/eSec). The standard and everywhere forms of second-preimage resistance of [RS04], on each fixed input slice $\mathcal{X}_{a,\mu}$, are reduced explicitly to two-way collision resistance in Corollary 3.

Everywhere preimage resistance (ePre). For a target tag T^* fixed in advance, no adversary finds X with $H_{\text{TW128}}(X) = T^*$.

The advantages $\text{Adv}_{\text{TW128}}^{\text{Coll}}$, $\text{Adv}_{\text{TW128}}^{\text{Pre}[a,\mu]}$, and $\text{Adv}_{\text{TW128}}^{\text{ePre}}$ are the respective success probabilities, instantiated for TW128 in Section 4.1. The Pre guarantee is proved directly in Theorem 9, while the second-preimage guarantees are inherited from Coll in Corollary 3. Coll is the binding property of the tag taken as a compact commitment to (K, U, A, M) , Pre is the random-opening preimage property on fixed input slices, and ePre is the fixed-target preimage property for query-independent tags. The [BH22] committing notions are the collision-facing AEAD projection of the same tag analysis: a winning opening gives distinct inputs sharing a tag, while the dedicated CMTDs-4 proof uses the all-bodies-equal game structure to avoid the leaf-collision slack of the general hash collision theorem (Section 6).

Relations. [BH22, Figs. 4 and 5] record the implications among the notions. For the D-notions we use:

$$\text{CMTDs-4} \Rightarrow \text{CMTDs-}\ell \text{ for } \ell \in \{1, 3\}$$

(both immediate: pairwise distinct keys ($\ell = 1$) or pairwise distinct (K, U, A) triples ($\ell = 3$) are *a fortiori* pairwise distinct full tuples, so every CMTDs- ℓ winner is a CMTDs-4 winner). The E-notions are bounded directly in Corollary 6 by the same final-root candidate-collision argument, rather than through an encryption reduction.

3 The TW128 Construction

3.1 Design Overview

TW128 is a nonce-based authenticated encryption scheme with associated data, built as a two-level tree of duplex objects over the KECCAK- $p[1600, 12]$ permutation. A single *root* duplex, initialized with the key K and nonce U , absorbs the associated data A when it is nonempty, emits keystream for the first plaintext chunk M_0 , whose maximum size β is defined in Table 2, and absorbs the resulting ciphertext chunk C_0 ; when A is empty the associated data phase is elided and the initialization call itself emits the initial M_0 keystream; the remaining plaintext is partitioned into *leaf* chunks $M_1, \dots, M_n$, each handled by an independent leaf duplex initialized with (K, U, j) . The leaf duplex emits




Figure 1: Schematic of TW128 for $n \geq 1$ leaves. The root duplex (bottom row) is initialized with (K, U) and absorbs in turn the associated data A , the root ciphertext chunk C_0 , each leaf tag $T_1, \dots, T_n$, and the eight-byte little-endian leaf count $\langle n \rangle_8$; it emits the root tag T . Each leaf duplex (top row) is initialized with (K, U, j) , emits keystream for plaintext chunk M_j , absorbs the resulting ciphertext chunk C_j , and produces T_j . Leaves can be evaluated in parallel. The ciphertext chunks $C_0, C_1, \dots, C_n$ are obtained by XORing the keystream emitted by the corresponding duplex with the plaintext chunks $M_0, M_1, \dots, M_n$. When $n = 0$ the diagram reduces to the bottom row without any T_j boxes or aggregation block: the aggregation phase is elided and the closing message phase call on C_0 emits the root tag T directly. When the associated data is empty the A block is absent: the associated data phase is elided and the initialization call emits the initial M_0 keystream directly.

the keystream for M_j , absorbs the resulting ciphertext chunk C_j , and produces a τ_{leaf} -bit leaf tag T_j . When $n \geq 1$ the leaf tags, together with the eight-byte little-endian leaf count $\langle n \rangle_8$, are absorbed into the root duplex, which emits a single τ_{root} -bit root tag T ; when $n = 0$ the aggregation phase is elided and the closing message phase call of the root chunk emits T directly, mirroring the way a leaf duplex emits its leaf tag. The construction is depicted in Figure 1.

Three design properties motivate the tree shape. First, the leaf duplexes operate on disjoint state and consume disjoint inputs once (K, U) is fixed, so leaves $1, \dots, n$ can be evaluated in parallel. The transcript fixes the leaf decomposition and order, but not the number of leaves an implementation evaluates at once; scalar, SIMD, and multicore implementations therefore produce the same ciphertexts and tags. Section 8 quantifies the resulting throughput on AMD64 and ARM64 hardware. Second, every duplex operates on a bounded 1600-bit state regardless of message length, and the leaf tags are aggregated as a stream rather than buffered in full (Section 3.5), so the working memory is constant in $\|M\|$ at a fixed degree of parallelism, scaling only with the number of leaves evaluated simultaneously. Third, the explicit aggregation of leaf tags through a single deterministic root transcript is the AEAD analog of KangarooTwelve’s final-node hashing of per-chunk chaining values. In TW128, the chaining values are the leaf tags and the final digest is the root tag. This makes the root tag the natural hash point for the full (K, U, A, M) context; the formal hash and committing analysis appears in Section 6.

3.2 Parameters and Domain

Table 2 lists the user-visible and derived scheme parameters. The implementation uses the same byte length for the root tag and each leaf tag, but the proofs of Sections 5 and 6 keep τ_{root} and τ_{leaf} as separate symbols.

The bound $M_{\text{max}} = 8183 \cdot 2^{64}$ bytes is the largest plaintext whose leaf index fits in $\langle \cdot \rangle_8$: writing $\text{leaves}(M) = \max(0, \lceil \|M\|/\beta \rceil - 1)$, every supported plaintext satisfies

Table 2: TW128 parameters.

Parameter	Value	Description
k	256	Key length (bits)
ν	256	Nonce length (bits)
τ_{root}	256	Root tag length (bits)
τ_{leaf}	256	Leaf tag length (bits)
β	8183	Plaintext chunk size (bytes)
M_{max}	$8183 \cdot 2^{64}$ bytes	Maximum plaintext length
A_{max}	$8183 \cdot 2^{64}$ bytes	Maximum associated data length
b	1600	Permutation state width (bits)
n_r	12	Permutation rounds
r	1344	Sponge rate (bits)
c	256	Sponge capacity (bits)
ρ	167	Maximum byte-aligned absorption per duplex call (bytes)

415 $0 \leq \text{leaves}(M) \leq 2^{64} - 1$, so leaf indices lie in $\{1, \dots, 2^{64} - 1\}$. The associated data bound
 416 A_{max} is a fixed scheme parameter, taken equal to M_{max} . Inputs exceeding these bounds
 417 are outside the scheme domain.

418 The derived per-call absorption bound ρ follows the [BDPVA11] analysis: with $r = 1344$
 419 and `pad10*1` padding, $\rho_{\text{max}}(\text{pad10*1}, r) = r - 2 = 1342$ bits. The byte-aligned data field
 420 shares this allotment with the four-bit phase suffix of Section 3.4, so ρ is the largest integer
 421 satisfying $8\rho + 4 \leq 1342$, giving $\rho = 167$ data bytes (and $8\rho + 4 = 1340 \leq 1342$).

422 3.3 Parameter Choice Rationale

423 The four free parameters c , k , $\tau_{\text{root}} = \tau_{\text{leaf}}$, and β (the rate $r = b - c$ is then determined
 424 by the capacity) are chosen to give 128-bit single-key security against every term of the
 425 concrete bounds derived in Section 7.1 while keeping per-byte cost low; the nonce width ν
 426 is fixed separately, by the headroom the wide permutation affords. Specifically:

- 427 • $c = 256$. The dominant single-key term is the [BDPVA08] flat sponge loss

$$428 \quad \text{Lift}_c(N_{\text{tot}}) = N_{\text{tot}}(N_{\text{tot}} + 1)/2^{c+1},$$

429 which reaches the 2^{-128} threshold at $N_{\text{tot}} \approx 2^{(c-127)/2} \approx 2^{64.5}$ KECCAK- $p[1600, 12]$
 430 calls; the work cap $N_{\text{tot}} \leq 2^{63}$ used in Section 7.2 keeps this term at most $2^{-131} + 2^{-194}$,
 431 and hence below 2^{-130} . Setting $c = 256$ thereby delivers 128-bit single-key security on
 432 this term. Under the mu-ind bound of Section 5 the per-user security margin degrades
 433 by $\log D$ bits in D users; the concrete D -indexed bound appears in Section 7.2.

- 434 • $r = 1344$. With state width $b = 1600$ fixed by the KECCAK- $p[1600, 12]$ choice,
 435 $r = b - c = 1344$ is the largest rate compatible with the chosen capacity. The 12-
 436 round KECCAK- $p[1600, 12]$ permutation matches KangarooTwelve’s choice [BDP⁺16];
 437 TW128 adopts that established tree-hashing parameter rather than introducing a
 438 new reduced-round setting (Section 9.4 surveys the cryptanalytic record behind this
 439 round count).
- 440 • $\tau_{\text{root}} = \tau_{\text{leaf}} = 256$. Matching the capacity ensures that the per-tag guessing terms
 441 in the authenticity bounds and the tag collision terms in the hash and committing
 442 bounds (Section 7.1) are dominated by the capacity birthday rather than the tag
 443 length.

- $k = 256$. Matching the capacity ensures that the per-user key guess

$$\text{KeyGuess}_k(N_{\text{prim}}) = N_{\text{prim}}/2^k$$

remains below the capacity birthday at the target D of Section 7.2.

- $\beta = 8183$. The chunk size is set to $49\rho = 49 \cdot 167$ bytes, the largest multiple of the per-call absorption bound ρ not exceeding 2^{13} bytes. Taking β to be an exact multiple of ρ makes every full leaf split into exactly 49 fully utilized message duplex calls with no short final block, so a full leaf costs one initialization duplex call plus 49 message duplex calls. This amortizes the fixed per-leaf initialization cost over a large chunk while keeping the leaf state working set ($8183 < 2^{13}$ bytes) small enough for L1 cache; Section 8.3 reports the resulting cycles-per-byte figures.
- $\nu = 256$. The nonce width is not constrained by the concrete bounds, which assume a nonce-respecting adversary; it is set generously because the wide permutation makes a wide nonce nearly free. The root initialization absorbs $p_{\text{root}} \parallel K \parallel U$, 560 bits in all, in a single duplex call, comfortably within the 1336-bit ($\rho = 167$ -byte) data field (Section 3.2), so widening the nonce to 256 bits adds neither a permutation call nor any per-message cost. The width also removes any practical concern over random nonce collisions: a uniformly random 256-bit nonce repeats over q messages with probability at most $\binom{q}{2}/2^{256}$, negligible at the work cap of Section 7.2. Deployments that do not require a full-width random nonce lose nothing by the choice: a narrower nonce—a 96-bit packet counter, say—is carried in the fixed ν -bit field by zero-padding, so a single mode serves both random nonce and counter-based settings with no change of parameters.

3.4 Domain Separation

Distinct duplex instances and distinct phases within a duplex are kept separate at the transcript level by two mechanisms. First, the initialization strings carry a 6-byte ASCII prefix and, for leaves, an eight-byte little-endian chunk index:

$$\begin{aligned} p_{\text{root}} &= \text{"TW128R"} \in \{0, 1\}^{48}, \\ p_{\text{leaf}} &= \text{"TW128L"} \in \{0, 1\}^{48}. \end{aligned}$$

The root duplex's first absorbed string is $p_{\text{root}} \parallel K \parallel U$, and the j -th leaf duplex's first absorbed string is $p_{\text{leaf}} \parallel K \parallel U \parallel \langle j \rangle_8$. Together with the fixed k -bit and ν -bit widths of K and U , these initialization strings separate root from leaf transcripts and distinguish leaves with different chunk indices.

Second, every byte-aligned absorption carries a four-bit phase suffix. We identify each suffix with the four-bit string given by its Table 3 value written least significant bit first, so that a duplex call absorbing a byte-aligned block σ under a phase suffix σ_{ph} feeds the duplex interface of Section 2.2 the input $\sigma \parallel \sigma_{\text{ph}}$ before `pad10*1` padding. The seven phase suffixes are summarized in Table 3; they correspond to initialization, non-final and final associated data blocks, non-final and final message blocks, and non-final and final aggregation blocks. The seven values are pairwise distinct, so any two transcripts that differ in the phase of any block also differ in the flattened sponge input.

Table 4 summarizes the root and leaf transcript schedule. The algorithms of Sections 3.5 and 3.6 give the normative byte-level specification; the table records which phase is present and what output the final call of that phase supplies.

The combined effect of the initialization prefixes and the phase suffixes is that the flattened sponge input of any TW128 duplex call uniquely determines the duplex instance, the leaf index when applicable, the phase, and the duplex-call position within that phase.

Table 3: Phase suffixes (4 bits each, appended LSB-first).

Name	Hex	Use
σ_{init}	0xC	End of initialization (the sole init call)
σ_{ad}^-	0x9	Non-final associated data block
σ_{ad}^+	0xD	Final associated data block
σ_{msg}^-	0xA	Non-final message block
σ_{msg}^+	0xE	Final message block
σ_{agg}^-	0xB	Non-final aggregation block
σ_{agg}^+	0xF	Final aggregation block (emits root tag)

Table 4: TW128 transcript schedule. “First M_i keystream” means $8 \min(\|M_i\|, \rho)$ output bits for the first duplex block of that message chunk; subsequent message-phase outputs have the length of the next duplex block, as in Algorithm 1.

Instance	Phase	Suffixes	Absorbed string	Output supplied	When
Root	Init	σ_{init}	$p_{\text{root}} \parallel K \parallel U$	0 if $A \neq \varepsilon$; otherwise first M_0 keystream	Always
Root	AD	$\sigma_{\text{ad}}^-, \sigma_{\text{ad}}^+$	ρ -byte block partition of A	Final call emits first M_0 keystream	$A \neq \varepsilon$
Root	Message	$\sigma_{\text{msg}}^-, \sigma_{\text{msg}}^+$	ρ -byte block partition of C_0	Non-final calls emit next keystream; final call emits T if $n = 0$, and 0 otherwise	Always
Leaf j	Init	σ_{init}	$p_{\text{leaf}} \parallel K \parallel U \parallel \langle j \rangle_s$	First M_j keystream	$1 \leq j \leq n$
Leaf j	Message	$\sigma_{\text{msg}}^-, \sigma_{\text{msg}}^+$	ρ -byte block partition of C_j	Non-final calls emit next keystream; final call emits T_j	$1 \leq j \leq n$
Root	Aggregation	$\sigma_{\text{agg}}^-, \sigma_{\text{agg}}^+$	$T_1 \parallel \dots \parallel T_n \parallel \langle n \rangle_s$	Final call emits T	$n \geq 1$

490 Definition 2 (Well-formed TW128 transcript prefixes), Lemma 6 (Transcript Injectivity),
 491 and its structural corollaries formalize this parsing fact for the proofs of Sections 5 and 6
 492 and Corollary 6.

493 3.5 Encryption

494 Algorithm 2 specifies $\text{Enc}_K(U, A, M)$ in terms of two helper procedures, ABSORB and
 495 STREAMENC, defined in Algorithm 1. These helpers sit below the tree layer: only
 496 Algorithms 2 and 3 split messages or ciphertexts into β -byte root and leaf chunks, while
 497 the helpers split an already chosen byte string into ρ -byte duplex blocks. ABSORB drives
 498 a duplex through a sequence of such blocks with non-final σ^- and final σ^+ suffix tags,
 499 returning the output of the final call; the empty-input case yields a single empty duplex
 500 block carrying the σ^+ tag. The streaming helpers use the same empty-input convention, so
 501 a zero-length message string still performs one final message phase duplex call and returns
 502 a zero-length body. Algorithm 2 invokes ABSORB for the associated data phase only when
 503 $A \neq \varepsilon$; when $A = \varepsilon$ the phase is elided and the root σ_{init} call emits the initial M_0 keystream
 504 directly. STREAMENC encrypts a plaintext byte string one duplex block at a time, XORing
 505 each block with the current keystream, absorbing the resulting ciphertext duplex block
 506 back into the duplex with the appropriate phase tag, and using the duplex’s output as
 507 the next keystream. This is the duplex-level presentation of the overwrite mechanism of
 508 [BDPVA11, §6.2 and Algorithm 5]: when the current rate prefix emitted as keystream is
 509 Z , XOR-absorbing $m \oplus Z$ has the same update on the message bytes as overwriting that
 510 prefix with m , with the phase tag still carried as part of the duplex input. The decryption
 511 helper STREAMDEC uses the same schedule as STREAMENC with ciphertext duplex blocks
 512 c_i as input: it recovers each plaintext duplex block as $m_i = c_i \oplus Z$, absorbs the same c_i
 513 under the same phase suffix, and returns $(m_1 \parallel \dots \parallel m_t, Z')$; we omit its pseudocode. In
 514 both streaming helpers, each duplex call absorbs a duplex block followed by its phase

Algorithm 1 TW128 helper procedures.

```

1: procedure ABSORB( $\mathcal{D}, s, \sigma^-, \sigma^+, \ell$ )
2:   Partition  $s$  into duplex blocks  $s_1, \dots, s_t$  of exactly  $\rho$  bytes each, except that the
   last block may be shorter.
3:   Require  $t \geq 1$ ; the empty case yields the single empty duplex block  $s_1 = \varepsilon$ .
4:   for  $i \leftarrow 1$  to  $t - 1$  do
5:      $\mathcal{D}.\text{duplexing}(s_i \parallel \sigma^-, 0)$ 
6:   end for
7:   return  $\mathcal{D}.\text{duplexing}(s_t \parallel \sigma^+, \ell)$ 
8: end procedure

9: procedure STREAMENC( $\mathcal{D}, m, Z, \sigma^-, \sigma^+, \ell$ )
10:  Partition  $m$  into duplex blocks  $m_1, \dots, m_t$  of exactly  $\rho$  bytes each, except that the
   last block may be shorter.
11:  Require  $t \geq 1$ ; if  $m = \varepsilon$ , set  $t = 1$  and  $m_1 = \varepsilon$ .
12:  for  $i \leftarrow 1$  to  $t - 1$  do
13:     $c_i \leftarrow m_i \oplus Z$ 
14:     $Z \leftarrow \mathcal{D}.\text{duplexing}(c_i \parallel \sigma^-, 8\|m_{i+1}\|)$ 
15:  end for
16:   $c_t \leftarrow m_t \oplus Z$ 
17:   $Z' \leftarrow \mathcal{D}.\text{duplexing}(c_t \parallel \sigma^+, \ell)$ 
18:  return  $(c_1 \parallel \dots \parallel c_t, Z')$ 
19: end procedure

```

suffix— σ^- for a non-final block and σ^+ for the final block—and requests the output bits the caller consumes next: ℓ (possibly zero) on the final call, and the next duplex block’s keystream length on non-final calls. Each call is the duplex interface of Section 2.2 applied to the suffixed input of Section 3.4.

The encryption procedure (Algorithm 2) instantiates the root duplex and obtains the initial keystream for the root chunk M_0 : when $A \neq \varepsilon$ it absorbs the associated data A , and when $A = \varepsilon$ it elides the associated data phase and takes the keystream from the σ_{init} call itself. It then encrypts M_0 while absorbing the resulting ciphertext chunk C_0 . When $n = 0$ the closing σ_{msg}^+ call of this root chunk requests τ_{root} output bits, which are the root tag T , and the aggregation phase is elided. When $n \geq 1$ it independently encrypts each leaf chunk M_j while absorbing C_j , producing the leaf tag T_j ; the leaf tags together with $\langle n \rangle_8$ are then absorbed into the root duplex, and the final σ_{agg}^+ call emits the root tag T . The ciphertext is the concatenation $C_0 \parallel C_1 \parallel \dots \parallel C_n$. Thus this section writes C for the ciphertext body and $C \parallel T$ for the full AEAD ciphertext returned by $\text{Enc}_K(U, A, M)$.

The initial keystream request for a message string X is $8 \min(\|X\|, \rho)$, the bit length of its first duplex block. For the root chunk M_0 , that value is emitted by the closing σ_{ad}^+ call when $A \neq \varepsilon$, and by the root σ_{init} call itself when $A = \varepsilon$ and the associated data phase is elided. Each leaf’s σ_{init} call likewise emits $8 \min(\|M_j\|, \rho)$ keystream bits for the first duplex block of M_j , so the empty associated data root σ_{init} call plays exactly the role a leaf σ_{init} call already plays. When the corresponding message string is empty, the requested keystream length is zero and the next message phase duplex call absorbs the empty ciphertext duplex block carrying σ_{msg}^+ . The final σ_{msg}^+ call in the root chunk requests τ_{root} output bits when $n = 0$, namely the root tag T emitted by the elided-aggregation path, and zero output bits when $n \geq 1$ because the next duplex operation on the root is then the first aggregation block; the corresponding call on a leaf requests τ_{leaf} bits, namely the leaf tag. When $n \geq 1$ the final σ_{agg}^+ call requests τ_{root} bits, which is the root tag T .

The leaf loop is annotated **in parallel** to make explicit that the n leaf computations

Algorithm 2 $\text{Enc}_K(U, A, M)$.**Require:** $K \in \{0, 1\}^k$, $U \in \{0, 1\}^\nu$, A with $\|A\| \leq A_{\max}$, M with $\|M\| \leq M_{\max}$.**Ensure:** $C \parallel T$ with $\|C\| = \|M\|$ and $T \in \{0, 1\}^{\tau_{\text{root}}}$.

```

1:  $M_0 \leftarrow$  first  $\min(\|M\|, \beta)$  bytes of  $M$  (possibly  $\varepsilon$ ).
2: Partition the remaining suffix into chunks  $M_1, \dots, M_n$  of exactly  $\beta$  bytes each, except
   that the last chunk may be shorter but nonempty, so  $n = \text{leaves}(M)$ .
3:  $\mathcal{D}_{\text{root}} \leftarrow$  new Duplex
4: if  $A = \varepsilon$  then
5:    $Z \leftarrow \mathcal{D}_{\text{root}}.\text{duplexing}(p_{\text{root}} \parallel K \parallel U \parallel \sigma_{\text{init}}, 8 \min(\|M_0\|, \rho))$ 
6: else
7:    $\mathcal{D}_{\text{root}}.\text{duplexing}(p_{\text{root}} \parallel K \parallel U \parallel \sigma_{\text{init}}, 0)$ 
8:    $Z \leftarrow \text{ABSORB}(\mathcal{D}_{\text{root}}, A, \sigma_{\text{ad}}^-, \sigma_{\text{ad}}^+, 8 \min(\|M_0\|, \rho))$ 
9: end if
10: if  $n = 0$  then
11:    $(C_0, T) \leftarrow \text{STREAMENC}(\mathcal{D}_{\text{root}}, M_0, Z, \sigma_{\text{msg}}^-, \sigma_{\text{msg}}^+, \tau_{\text{root}})$ 
12:   return  $C_0 \parallel T$ 
13: end if
14:  $(C_0, \_) \leftarrow \text{STREAMENC}(\mathcal{D}_{\text{root}}, M_0, Z, \sigma_{\text{msg}}^-, \sigma_{\text{msg}}^+, 0)$ 
15: for  $j \leftarrow 1$  to  $n$  in parallel do
16:    $\mathcal{D}_j \leftarrow$  new Duplex
17:    $Z_j \leftarrow \mathcal{D}_j.\text{duplexing}(p_{\text{leaf}} \parallel K \parallel U \parallel \langle j \rangle_8 \parallel \sigma_{\text{init}}, 8 \min(\|M_j\|, \rho))$ 
18:    $(C_j, T_j) \leftarrow \text{STREAMENC}(\mathcal{D}_j, M_j, Z_j, \sigma_{\text{msg}}^-, \sigma_{\text{msg}}^+, \tau_{\text{leaf}})$ 
19: end for
20:  $G \leftarrow T_1 \parallel \dots \parallel T_n$ .
21:  $T \leftarrow \text{ABSORB}(\mathcal{D}_{\text{root}}, G \parallel \langle n \rangle_8, \sigma_{\text{agg}}^-, \sigma_{\text{agg}}^+, \tau_{\text{root}})$ 
22: return  $C_0 \parallel C_1 \parallel \dots \parallel C_n \parallel T$ 

```

are mutually independent; Section 8.1 discusses the SIMD realizations. The annotation is an implementation freedom, not a mode parameter: evaluating leaves serially, in SIMD batches, or across cores leaves the transcript and wire format unchanged.

Streaming aggregation. The string $G = T_1 \parallel \dots \parallel T_n$ in Algorithms 2 and 3 names the leaf tags only for specification. The closing ABSORB consumes $G \parallel \langle n \rangle_8$ as ρ -byte duplex blocks in leaf order and advances the root state after each, so an implementation need not hold G in full. It can absorb leaf tags in order as they become available, or buffer completed out-of-order leaves until the next in-order tag or SIMD batch can be absorbed, keeping only a bounded number of tags at a time. With the bounded 1600-bit state of every duplex, this keeps the working memory constant in $\|M\|$ for a fixed degree of parallelism, growing only with the number of leaves evaluated at once rather than with the message length; Section 8.1 describes the batched realization.

3.6 Decryption

Algorithm 3 specifies $\text{Dec}_K(U, A, C \parallel T)$. The procedure mirrors Enc at the duplex block level: each ciphertext duplex block c_i is absorbed into the appropriate duplex (root or leaf) under the corresponding σ_{msg}^- or σ_{msg}^+ suffix, and the plaintext duplex block is recovered as $m_i = c_i \oplus Z^{(i)}$, writing $Z^{(i)}$ for the value the running keystream variable Z of STREAMDEC holds when duplex block i is processed ($i < t$ inside the loop, $i = t$ after it), the same value the encryption procedure would have produced. The procedure absorbs ciphertext bytes before the final tag comparison, so the verification transcript is determined by (K, U, A, C) for every in-domain ciphertext, accepting or not. This property is used by the direct

Algorithm 3 $\text{Dec}_K(U, A, C \parallel T)$.

Require: K, U, A as in Algorithm 2, ciphertext C with $\|C\| \leq M_{\max}$, candidate tag $T \in \{0, 1\}^{\tau_{\text{root}}}$.

Ensure: M with $\|M\| = \|C\|$, or $\perp$.

```

1:  $C_0 \leftarrow$  first  $\min(\|C\|, \beta)$  bytes of  $C$  (possibly  $\varepsilon$ ).
2: Partition the remaining suffix into chunks  $C_1, \dots, C_n$  of exactly  $\beta$  bytes each, except
   that the last chunk may be shorter but nonempty, so  $n = \text{leaves}(C)$ .
3:  $\mathcal{D}_{\text{root}} \leftarrow$  new Duplex
4: if  $A = \varepsilon$  then
5:    $Z \leftarrow \mathcal{D}_{\text{root}}.\text{duplexing}(p_{\text{root}} \parallel K \parallel U \parallel \sigma_{\text{init}}, 8 \min(\|C_0\|, \rho))$ 
6: else
7:    $\mathcal{D}_{\text{root}}.\text{duplexing}(p_{\text{root}} \parallel K \parallel U \parallel \sigma_{\text{init}}, 0)$ 
8:    $Z \leftarrow \text{ABSORB}(\mathcal{D}_{\text{root}}, A, \sigma_{\text{ad}}^-, \sigma_{\text{ad}}^+, 8 \min(\|C_0\|, \rho))$ 
9: end if
10: if  $n = 0$  then
11:    $(M_0, T') \leftarrow \text{STREAMDEC}(\mathcal{D}_{\text{root}}, C_0, Z, \sigma_{\text{msg}}^-, \sigma_{\text{msg}}^+, \tau_{\text{root}})$ 
12:   return  $M_0$  if  $T' = T$  (compared in constant time), else  $\perp$ .
13: end if
14:  $(M_0, \_) \leftarrow \text{STREAMDEC}(\mathcal{D}_{\text{root}}, C_0, Z, \sigma_{\text{msg}}^-, \sigma_{\text{msg}}^+, 0)$ 
15: for  $j \leftarrow 1$  to  $n$  in parallel do
16:    $\mathcal{D}_j \leftarrow$  new Duplex
17:    $Z_j \leftarrow \mathcal{D}_j.\text{duplexing}(p_{\text{leaf}} \parallel K \parallel U \parallel \langle j \rangle_8 \parallel \sigma_{\text{init}}, 8 \min(\|C_j\|, \rho))$ 
18:    $(M_j, T_j) \leftarrow \text{STREAMDEC}(\mathcal{D}_j, C_j, Z_j, \sigma_{\text{msg}}^-, \sigma_{\text{msg}}^+, \tau_{\text{leaf}})$ 
19: end for
20:  $G \leftarrow T_1 \parallel \dots \parallel T_n$ .
21:  $T' \leftarrow \text{ABSORB}(\mathcal{D}_{\text{root}}, G \parallel \langle n \rangle_8, \sigma_{\text{agg}}^-, \sigma_{\text{agg}}^+, \tau_{\text{root}})$ 
22: if  $T' = T$  (compared in constant time) then
23:   return  $M_0 \parallel M_1 \parallel \dots \parallel M_n$ 
24: else
25:   return  $\perp$ 
26: end if

```

563 mu-ind proof to argue that encryption-first and verification-first transcript classes are well
 564 defined regardless of the tag-comparison outcome.

565 The $n = 0$ and empty-body cases mirror their encryption counterparts (Section 3.5):
 566 for $n = 0$ the aggregation phase is elided and the closing σ_{msg}^+ call of the root chunk emits
 567 the candidate tag T' directly, and an empty body still performs the single empty σ_{msg}^+ call
 568 of Algorithm 1, requesting τ_{root} output bits.

569 **Correctness.** For every $K \in \{0, 1\}^k$, $U \in \{0, 1\}^\nu$, A with $\|A\| \leq A_{\max}$, and M with
 570 $\|M\| \leq M_{\max}$, $\text{Dec}_K(U, A, \text{Enc}_K(U, A, M)) = M$. This follows by inspection: every duplex
 571 call in Dec_K receives the same input as the corresponding call in Enc_K (each ciphertext
 572 duplex block is absorbed under the same phase tag in both directions), so the duplex states
 573 and the recomputed root tag T' exactly match. The constant-time comparison $T' = T$
 574 succeeds, and the recovered block $c_i \oplus Z^{(i)} = (m_i \oplus Z^{(i)}) \oplus Z^{(i)} = m_i$ at every position.

575 **Lemma 1** (TW128 tidiness). *For every valid TW128 decryption input, if $\text{Dec}_K(U, A, C \parallel$
 576 $T) = M \neq \perp$, then $\text{Enc}_K(U, A, M) = C \parallel T$.*

577 *Proof.* An accepting decryption first passes the syntax and length checks of Algorithm 3,
 578 so the recovered message M has $\|M\| = \|C\|$ and the same canonical chunking as the
 579 ciphertext body. The decryption transcript absorbs K , U , A , and each ciphertext block

under the same duplex instances and phase suffixes used by encryption. Its stream outputs recover each plaintext duplex block as $m_i = c_i \oplus Z^{(i)}$, so re-encrypting the recovered message under the same transcript emits $c_i = m_i \oplus Z^{(i)}$ and absorbs the same ciphertext duplex block at every message position. The leaf tags and final root candidate tag recomputed by decryption are therefore exactly the values recomputed by this re-encryption. Since decryption accepts only when its final candidate tag equals the submitted tag T , re-encryption returns the submitted body and tag $C \parallel T$. $\square$

4 Security Model

This section sets up the two proof views used throughout the paper. For confidentiality and authenticity, TW128 is treated as a nonce-based AEAD in the lifted public permutation model: real encryption and verification interfaces are compared with ideal random and replay interfaces, first in the multi-user setting and then in the single-user all-in-one AE setting. For the digest-like tag claim, the same construction is viewed through the root tag wrapper H_{TW128} : adversaries supply full inputs or candidate openings, and the games ask whether the root tag collides or can be hit as a preimage. We instantiate these games for TW128 (Sections 4.1 and 4.2), define the intermediate random oracle world through which every bound is first proved (Section 4.3), and fix the resource accounting that all later bounds reference (Section 4.4).

Table 5 summarizes the proof obligations before the formal games. The loss terms named there are defined in Section 4.5, and the implemented-parameter instantiations appear in Section 7.

Scope of the security claims. All results below are public-permutation claims for TW128 with the $\text{KECCAK-}p[1600, 12]$ permutation modeled as a uniformly random permutation, obtained by first proving random oracle bounds for the flat duplex interface of Section 4.3 and then applying the flat sponge lift. In the AE-style distinguishing games, encryption queries are nonce-respecting and the verification oracle returns only acceptance or rejection. The separate INT-RUP game below gives integrity under a plaintext-computing oracle, but it does not give privacy under release of unverified plaintext or plaintext-awareness security. The claims also do not provide nonce-misuse resistance. The hash and committing games are adversarial-initialization games over the submitted inputs (K, U, A, M) or ciphertext openings and impose no nonce uniqueness condition. Across all games, the stated bounds are birthday-type, single-stage, and leakage-free: beyond-birthday and side-channel security are outside the model. Section 9.5 discusses the consequences of these exclusions and the corresponding open problems, and Section 9.4 reviews the third-party cryptanalysis supporting the random-permutation modeling of $\text{KECCAK-}p[1600, 12]$.

The proof dependency spine is as follows. Section A first establishes transcript decoding, bridge injectivity, output-point separation, and the local probabilistic accounting lemmas for key tests, leaf tag collisions, and root tag guesses. Section 5 uses those facts directly for the random oracle mu-ind and INT-RUP bounds. Section 6 then uses the same structural facts to build a chronological sequence of final root candidates and apply a single root tag hash bound; the hash and committing theorems differ only in which collision or preimage event is reduced to that candidate sequence. Finally, Section B transfers the random oracle bounds to the public permutation games defined below.

4.1 Public Permutation Games

Primitive convention. In every public permutation game below, the challenger samples $\pi \leftarrow \$ \text{Perm}(\{0, 1\}^{1600})$ and exposes the direct primitive interface (π, π^{-1}) to the adversary. The real-vs-ideal AE-style games compare the real construction interface with the ideal

Table 5: Proof obligations and the dominant loss shapes used for TW128. The table is only a roadmap; the formal games and resource accounting are defined in Sections 4.1, 4.2 and 4.4.

Notion	Adversary goal	Main result	Loss shape
AE / mu-ind	Distinguish real encryption and verification from random ciphertexts and replay-only verification, with nonce-respecting encryption queries.	Theorem 2 and Corollary 1	Capacity lift, key guess, leaf collision, tag guess
INT-RUP	Produce a fresh accepting ciphertext after encryption and plaintext-computing queries.	Theorem 4	Capacity lift, key guess, leaf collision, tag guess
H_{TW128} collision	Find distinct full inputs (K, U, A, M) with the same root tag.	Theorem 8	Capacity lift, root collision, leaf collision
H_{TW128} preimage	Find an input hitting a fixed target tag.	Theorem 10	Capacity lift, root preimage
H_{TW128} second preimage	Find a second preimage in the standard or everywhere [RS04] sense.	Corollary 3	Inherited from collision resistance
CMT / CMTD	Produce equal encryptions or one valid ciphertext opening under distinct committed inputs.	Corollaries 5 and 6	Capacity lift, root collision

interface specified by the game over this shared primitive, following the same-oracle convention of [BT16]; the hash and committing games below are one-world success probabilities over the same sampled permutation. The simulator $(S^{\text{RO}_{\text{flat}}}, S^{-1, \text{RO}_{\text{flat}}})$ of Section B is a proof device, entering through the lift from the intermediate random oracle model.

Validity convention. All construction interface inputs are checked against TW128’s public syntax and domain before any construction computation or construction-internal primitive lookup is performed. An encryption input is valid when $U \in \{0, 1\}^\nu$, A is a byte string with $\|A\| \leq A_{\max}$, and M is a byte string with $\|M\| \leq M_{\max}$. A verification or decryption input is valid when $U \in \{0, 1\}^\nu$, A is a byte string with $\|A\| \leq A_{\max}$, and the full ciphertext parses as $C \parallel T$ with $\|C\| \leq M_{\max}$ and $T \in \{0, 1\}^{\tau_{\text{root}}}$. In the hash and committing games, each submitted input (K, U, A, M) must be a valid encryption input including $K \in \{0, 1\}^k$, and a submitted CMTDs ciphertext must be a valid verification/decryption ciphertext. Invalid encryption queries are rejected and not recorded, invalid plaintext-computing queries return $\perp$, invalid verification queries return 0, and invalid committing-game outputs make the game return 0. Such rejected inputs contribute no construction cost and create no construction transcripts. Throughout, an *accepted* encryption query is one that passes this validity check and any game-specific nonce or user check.

All-in-one AE. The real all-in-one AE game additionally samples $K \leftarrow \$ \{0, 1\}^k$ and exposes the construction oracles $\text{Enc}_K[\pi]$ (the TW128 encryption algorithm of Algorithm 2 with the sampled π) and $\text{Ver}_K[\pi]$, defined by $\text{Ver}_K[\pi](U, A, C, T) = 1$ if $\text{Dec}_K[\pi](U, A, C \parallel T) \neq \perp$ and 0 otherwise. The all-in-one authenticated encryption game [RS06] replaces the construction interface by the ideal pair $(\text{Rand}, \text{Replay})$. On each accepted encryption query (U, A, M) , Rand samples a uniformly random byte string of length $\|M\| + \tau_{\text{root}}/8$, parses it as $C \parallel T$ with $T \in \{0, 1\}^{\tau_{\text{root}}}$, records (U, A, C, T) , and returns $C \parallel T$. The oracle Replay returns 1 exactly on recorded tuples (U, A, C, T) and 0 otherwise. The adversary is nonce-respecting on encryption queries.

$$\text{Adv}_{TW128}^{\text{ae}}(\mathcal{A}) = |\Pr[\mathcal{A}^{\text{Enc}_K[\pi], \text{Ver}_K[\pi], \pi, \pi^{-1}} \Rightarrow 1] - \Pr[\mathcal{A}^{\text{Rand}, \text{Replay}, \pi, \pi^{-1}} \Rightarrow 1]|.$$

This all-in-one AE advantage is the single-user case of the multi-user real-vs-ideal notion of Section 2.4; Corollary 1 therefore derives it from the direct mu-ind bound of Theorem 2.

RUP integrity. Following the separated AE syntax of [ABL⁺14], define $\text{PDec}_K[\pi]$ as the plaintext-computing oracle that maps (U, A, C, T) to the candidate plaintext computed by $\text{Dec}_K[\pi](U, A, C \parallel T)$ before the final root tag comparison is applied. The computation follows the same decryption transcript as Algorithm 3, but returns the recovered plaintext and never returns the recomputed tag or the verification bit.

The INT-RUP game samples $K \leftarrow \$ \{0, 1\}^k$ and $\pi \leftarrow \$ \text{Perm}(\{0, 1\}^{1600})$, exposes $\text{Enc}_K[\pi]$, $\text{PDec}_K[\pi]$, $\text{Ver}_K[\pi]$, and (π, π^{-1}) , and maintains the replay table W of encryption outputs. Accepted encryption queries are nonce-respecting: an encryption query is rejected if its nonce has appeared in a previous accepted encryption query. Plaintext-computing and verification queries may use arbitrary nonces. On an accepted encryption query (U, A, M) the game returns $\text{Enc}_K[\pi](U, A, M)$, parsed and recorded as (U, A, C, T) in W . On a valid plaintext-computing query (U, A, C, T) it returns $\text{PDec}_K[\pi](U, A, C, T)$, and on an invalid query it returns $\perp$. On a valid verification query (U, A, C, T) it returns $\text{Ver}_K[\pi](U, A, C, T)$; if this bit is 1 and $(U, A, C, T) \notin W$ when the query is asked, the game records a forgery. Invalid verification queries return 0. The game returns 1 iff a forgery is recorded, and

$$\text{Adv}_{\text{TW128}}^{\text{INT-RUP}}(\mathcal{A}) = \Pr[\text{G}_{\text{TW128}}^{\text{INT-RUP}}(\mathcal{A}) \Rightarrow 1].$$

Hash games. The hash games for the wrapper H_{TW128} (Definition 1) have no challenger-sampled key. In the s -way collision game, all inputs are adversary-supplied: the adversary outputs s inputs, and the challenger returns 1 iff they are pairwise distinct, in the TW128 domain, and share a root tag:

$\pi \leftarrow \$ \text{Perm}(\{0, 1\}^{1600})$, $(K_i, U_i, A_i, M_i)_{i=1}^s \leftarrow \$ \mathcal{A}^{\pi, \pi^{-1}}$,
 return 0 unless the inputs are pairwise distinct and in the TW128 domain,
 return $[\text{H}_{\text{TW128}}(K_1, U_1, A_1, M_1) = \dots = \text{H}_{\text{TW128}}(K_s, U_s, A_s, M_s)]$,

with advantage $\text{Adv}_{\text{TW128}}^{\text{Coll}}(\mathcal{A}) = \Pr[\text{the game returns 1}]$; no nonce uniqueness is imposed. The [RS04] standard and everywhere second-preimage notions are used on fixed input slices; in Corollary 3 they are reduced explicitly to this two-way collision game. For fixed byte lengths a, μ , the standard preimage game samples $X^* \leftarrow \$ \mathcal{X}_{a, \mu}$ independently of the primitive, computes $T^* = \text{H}_{\text{TW128}}(X^*)$, gives T^* to $\mathcal{A}^{\pi, \pi^{-1}}$, and returns 1 iff $\mathcal{A}$ outputs a valid input X in the TW128 domain with $\text{H}_{\text{TW128}}(X) = T^*$; the advantage is $\text{Adv}_{\text{TW128}}^{\text{Pre}[a, \mu]}(\mathcal{A})$. The everywhere preimage game fixes a target tag $T^* \in \{0, 1\}^{\tau_{\text{root}}}$ before the primitive sampling and returns 1 iff $\mathcal{A}(T^*)$ outputs an input X in the domain with $\text{H}_{\text{TW128}}(X) = T^*$, the advantage $\text{Adv}_{\text{TW128}}^{\text{ePre}}$ taken as the worst case over query-independent targets.

Committing games. The [BH22] committing games are instantiated as in Section 2.5, with no challenger-sampled key. The s -way CMTDs- ℓ (decryption) game has the adversary output a ciphertext $C \parallel T$ and s tuples (K_i, U_i, A_i, M_i) with pairwise distinct WiC_ℓ -images, and returns 1 iff $\text{Dec}_{K_i}[\pi](U_i, A_i, C \parallel T) = M_i$ for every i , with the eager length rejection justified by Algorithm 3 ($\text{Dec}_K(U, A, C \parallel T)$ returns $\perp$ or a message of length $\|C\|$); its advantage is $\text{Adv}_{\text{TW128}}^{\text{CMTDs-}\ell}(\mathcal{A})$. The CMTs- ℓ (encryption) game has the adversary output the s tuples alone and returns 1 iff the encryptions $\text{Enc}_{K_i}[\pi](U_i, A_i, M_i)$ are all equal; its advantage is $\text{Adv}_{\text{TW128}}^{\text{CMTs-}\ell}(\mathcal{A})$.

4.2 Multi-User Setting

This subsection instantiates the mu-ind game of [BT16] (Section 2.4) for TW128 in the public permutation model, under the primitive convention of Section 4.1. The challenger samples a challenge bit $b \leftarrow \$ \{0, 1\}$ and a permutation $\pi \leftarrow \$ \text{Perm}(\{0, 1\}^{1600})$, and exposes

(π, π^{-1}) as the primitive interface in both worlds. The adversary $\mathcal{A}$ creates users on demand via **New**, and the game maintains an active-user counter v . We write D for an execution-uniform cap on the number of **New** calls. Equivalently, the game may sample independent keys $K_1, \dots, K_D$ at the start of the experiment and let the d -th **New** call activate K_d ; inactive indices are never valid users.

The mu-ind construction oracles inherit the validity convention of Section 4.1. The encryption oracle on input (d, U, A, M) rejects, without recording, if (U, A, M) is outside the TW128 encryption domain, if $d \notin \{1, \dots, v\}$, or if $(d, U) \in V$; otherwise it sets $C = \text{Enc}_{K_d}[\pi](U, A, M)$ when $b = 1$, and samples a uniformly random byte string C of length $\|M\| + \tau_{\text{root}}/8$ when $b = 0$. The game records (d, U) in V and (d, U, A, C) in W , and returns C . The verification oracle on input (d, U, A, C) rejects if $d \notin \{1, \dots, v\}$ or if (U, A, C) is not a valid TW128 verification/decryption input; if $(d, U, A, C) \in W$ it returns true; if $b = 0$ it returns false; otherwise it returns $[\text{Dec}_{K_d}[\pi](U, A, C) \neq \perp]$. Nonce uniqueness is enforced per user by the set V .

After querying these oracles, $\mathcal{A}$ outputs $b' \in \{0, 1\}$. The mu-ind advantage of $\mathcal{A}$ against TW128 is

$$\text{Adv}_{\text{TW128}}^{\text{mu-ind}}(\mathcal{A}) = |2 \Pr[b' = b] - 1|.$$

A valid verification query is *non-replay* if its tuple is not in W when the query is asked; replay hits are answered by the recorded transcript and are not counted in the verification-query resource below.

Resource caps split per potential user $d \leq D$: $N^{(d)}$, $q_v^{(d)}$, $n_{\text{leaf}}^{(d)}$ are the per-user counts under Section 4.4, with inactive users contributing zero. The global totals $N = \sum_{d=1}^D N^{(d)}$, $q_v = \sum_{d=1}^D q_v^{(d)}$, $n_{\text{leaf}} = \sum_{d=1}^D n_{\text{leaf}}^{(d)}$ appear in the final bound: N through the lift term $\text{Lift}_c(N_{\text{tot}})$, and n_{leaf} and q_v through the leaf-collision and verification-tag terms. N_{prim} remains a single shared count over the primitive interface.

The mu-ind treatment applies only to indistinguishability. The committing security games already give the adversary control of all s keys, so the s -way CMTDs- ℓ statements of Section 4.1 are themselves the multi-key committing statement and need no separate mu-cmt notion.

4.3 The TW128 Random Oracle Model

The proofs of Sections 5 and 6 and Corollary 6 are carried out through an intermediate random oracle model whose construction interfaces use the [BDPVA11] bridge of Section 2.2 to address a single random oracle. The flat sponge lift of Section B then converts the resulting bounds to the public permutation form.

The model has a single random oracle $\text{RO}_{\text{flat}} : \{0, 1\}^* \rightarrow \{0, 1\}^\infty$ over the unpadded H -input domain of the [BDPVA08] simple-padded sponge interface, with lazy sampling: each distinct input receives an independent infinite bit string, and repeated queries return the same string. A direct adversary query is finite-output: on input (Y, ℓ) with finite ℓ , the oracle returns the first ℓ bits of the lazy-sampled value for Y .

Construction calls do not query RO_{flat} on arbitrary strings. Every root or leaf duplex output is read from RO_{flat} at the bridged flattened transcript prefix for that duplex call, and the caller takes only the requested output projection. Direct adversary queries to RO_{flat} are the only random oracle queries counted as direct oracle queries; construction lookups are accounted for through the construction cost N of Section 4.4.

For a well-formed TW128 duplex transcript prefix X in the language of Definition 2, at one of the construction transcript lookup points formalized by Lemma 8 and Table 12, let $\text{flat}(X) = I[1344, 1344](\text{raw_flat}(X))$ be the bridged image of the native flattened input under the [BDPVA11, Theorem 3] preprocessing map of Section 2.2, and define

$$\text{RF}(X) = \text{RO}_{\text{flat}}(\text{flat}(X)).$$

The construction interfaces Enc_K^{RO} , Dec_K^{RO} , $\text{PDec}_K^{\text{RO}}$, and Ver_K^{RO} are the TW128 algorithms of Section 3 with every root or leaf duplex output replaced by the corresponding requested projection of $\text{RF}(X)$; the projection length is the duplexing call's output parameter ℓ , which may be zero. The bridge and the typed-restriction convention are formalized in Lemma 8; the direct-query key-hit predicate ignores the adversary's requested output length and is formalized as **KeyedTWOut** in Lemma 8.

The random oracle model games are the construction interface variants of Sections 4.1 and 4.2: the primitive interface is replaced by direct adversary access to RO_{flat} , and each construction oracle uses the corresponding RO algorithm, the ideal oracles (**Rand**, **Replay**) being unchanged. Under this substitution $\text{Adv}_{\text{RO}}^{\text{ae}}$ is the all-in-one advantage of Section 4.1 with direct RO_{flat} access as the primitive interface in both worlds; the keyless hash advantages $\text{Adv}_{\text{RO}}^{\text{Coll}}(\mathcal{B})$, $\text{Adv}_{\text{RO}}^{\text{Pre}[a,\mu]}(\mathcal{B})$, and $\text{Adv}_{\text{RO}}^{\text{ePre}}(\mathcal{B})$ of H_{TW128} , the committing $\text{Adv}_{\text{RO}}^{\text{CMTDs-4}}(\mathcal{B})$, and, for $\ell \in \{1, 3, 4\}$, $\text{Adv}_{\text{RO}}^{\text{CMTs-}\ell}(\mathcal{B})$ are the probabilities that the corresponding random oracle hash and committing games return 1, the CMTs challenger computing its deterministic encryption checks with Enc_K^{RO} . For mu-ind, $\text{Adv}_{\text{RO}}^{\text{mu-ind}}(\mathcal{B})$ denotes the Section 4.2 game with the same **New** interface and per-user nonce discipline—user d using $(\text{Enc}_{K_d}^{\text{RO}}, \text{Ver}_{K_d}^{\text{RO}})$ when $b = 1$ and (**Rand**, **Replay**) when $b = 0$ —equivalently, in the two-game notation of Section 2.4,

$$\text{Adv}_{\text{RO}}^{\text{mu-ind}}(\mathcal{B}) = |\Pr[\mathcal{B}^{\text{Real}^{\text{mu}}, \text{RO}_{\text{flat}}} \Rightarrow 1] - \Pr[\mathcal{B}^{\text{Ideal}^{\text{mu}}, \text{RO}_{\text{flat}}} \Rightarrow 1]|,$$

where the superscripts name the entire multi-user construction interface in the corresponding world. The random oracle INT-RUP game is the same win-event game as above, with Enc_K^{RO} , $\text{PDec}_K^{\text{RO}}$, and Ver_K^{RO} replacing the public-permutation construction oracles and with direct adversary access to RO_{flat} replacing (π, π^{-1}) ; its advantage is $\text{Adv}_{\text{RO}}^{\text{INT-RUP}}(\mathcal{B})$. Adversary queries to RO_{flat} are finite-output and counted in $q_{\text{RO}}(\mathcal{B})$ of Section 4.4. Construction-internal $\text{RF}(\cdot)$ lookups — those made by Enc_K^{RO} , Dec_K^{RO} , $\text{PDec}_K^{\text{RO}}$, Ver_K^{RO} , or a committing challenger — are not.

4.4 Resources

Adversary resources are measured by fixed, execution-uniform caps on the number and size of oracle queries. An execution-uniform cap is one that holds on every execution path, including over adversary randomness, oracle randomness, and adaptive query choices. Proofs may also introduce local execution-uniform caps, such as q_e for the number of accepted encryption queries, solely to index a finite hybrid sequence. Such local caps are part of the adversary's fixed query bounds, but they are not listed as headline resource parameters unless they affect the final bound.

Table 6 gives the compact notation summary used in later theorems. Lowercase quantities are realized counts in a random oracle experiment, while uppercase quantities are execution-uniform caps used in the public-permutation bounds. The detailed definitions below fix exactly which computations contribute to each count.

N: the construction interface cost, measured as the number of $\text{KECCAK-}p[1600, 12]$ calls that the adversary's construction interface queries induce in the real TW128 computation after the validity convention above has been applied. In ideal games, the same accounting applies to the corresponding real TW128 computation on the valid query, even when an ideal oracle answers by table lookup or without running the construction. N counts at per-duplex-call granularity, summing each root initialization, associated data duplex block, root message duplex block, leaf initialization, leaf message duplex block, and aggregation duplex block across the entire query sequence. Valid plaintext-computing computations and valid AE-style non-replay verification computations are counted whether they accept or reject; replay-table

Table 6: Resource-accounting notation. Uppercase symbols are the caps that appear in public-permutation theorems; lowercase symbols are realized counts in random oracle executions or source games.

Symbol	Counted objects	Primary use
N	Construction-internal KECCAK- p [1600, 12] calls induced by valid construction oracle queries and deterministic challenger checks.	Public-permutation lift and concrete work cap.
N_{prim}	Direct adversary queries to the public primitive interface.	Key-guess, root-collision, root-preimage terms; direct-query part of the lift.
$N_{\text{tot}} = N + N_{\text{prim}}$	Converted primitive-query cost used by the flat sponge replacement.	Argument of $\text{Lift}_c(N_{\text{tot}})$.
$q_{\text{RO}}(\mathcal{B}), Q_{\text{RO}}$	Direct finite-output queries to RO_{flat} in the intermediate random oracle model; construction lookups are excluded.	Random oracle bounds; after the lift $q_{\text{RO}} \leq N_{\text{prim}}$.
$q_v(\mathcal{A}), Q_v$	Valid verification queries that are not replayable hits.	Verification-tag terms in AE, mu-ind, and INT-RUP bounds.
$n_{\text{leaf}}(\mathcal{B}), N_{\text{leaf}}$	Distinct construction-generated final leaf transcripts.	Leaf tag collision and tag-guessing terms; always at most N .
D	Execution-uniform cap on users activated in the mu-ind game.	Multi-user key-collision and key-guess terms.

798 verification hits are transcript-consistency checks and are not charged to N or to
 799 the verification-tag terms. In the hash and committing games there is no adversary
 800 construction oracle: for the H_{TW128} collision game, N counts the s deterministic
 801 encryptions on the submitted inputs; for the standard preimage game it counts both
 802 challenger encryptions: the source encryption computing $T^* = \text{H}_{\text{TW128}}(X^*)$ and the
 803 one deterministic encryption on the submitted input; for the everywhere-preimage
 804 game, the one deterministic encryption on the submitted input; for CMTDs- ℓ ,
 805 N uses conservative committing-game accounting: after the syntax, domain, and
 806 message-length checks, it charges the deterministic decryption checks on the common
 807 ciphertext whether or not the particular WiC_ℓ distinctness check would reject before
 808 those checks. Likewise, for the CMTs- ℓ source games of Corollary 6, after syntax
 809 and domain checks N charges the deterministic encryptions that would be used to
 810 compare the s ciphertexts whether or not the WiC_ℓ distinctness check would reject
 811 earlier. For the Sec and eSec games of Corollary 3, N and N_{prim} are those of the
 812 explicit induced two-way collision adversary in that corollary, whose source and
 813 candidate deterministic encryptions are counted exactly as in the $s = 2$ collision
 814 game.

815 N_{prim} : the number of direct adversary queries to the primitive interface, the pair (π, π^{-1})
 816 in public permutation games.

817 $N_{\text{tot}} = N + N_{\text{prim}}$: the total primitive-query cost used by the flat sponge lift. The lift of
 818 Section B is applied at N_{tot} because after the operational duplex-trace accounting of
 819 Lemma 19, the construction-internal duplex calls counted by N and the adversary's
 820 direct primitive-interface queries counted by N_{prim} form a primitive-query sequence
 821 of cost at most $N + N_{\text{prim}}$.

822 $q_v(\mathcal{A}), Q_v$: the number of verification queries with valid inputs that are not replay-table
 823 hits, and a cap on it. Used in the verification terms of the direct mu-ind, AE, and
 824 INT-RUP bounds of Section 5. Every counted verification query contributes at least
 825 one verification/decryption construction call in the N accounting, so $q_v(\mathcal{A}) \leq N$.

826 $q_{\text{RO}}(\mathcal{B}), Q_{\text{RO}}$: the number of direct finite-output queries to RO_{flat} by a random oracle
 827 model adversary $\mathcal{B}$, and a cap on it. Repeated direct queries on the same Y are

counted again. Construction-internal $\text{RF}(\cdot)$ lookups are not counted. After the derived-adversary construction, q_{RO} counts only simulator-induced RO_{flat} calls from the adversary’s direct primitive-interface queries, not the construction calls already counted in N for the flat sponge lift.

$n_{\text{leaf}}(\mathcal{B})$, N_{leaf} : the number of distinct construction-generated final leaf transcripts in the execution of $\mathcal{B}$, and a cap on it. Distinct final leaf transcripts are generated by valid encryption computations, by valid plaintext-computing computations, by valid AE-style verification computations (whether they accept or reject), and by hash and committing challenger construction computations that pass the early validity checks, including the preimage game’s source encryption.

For a byte string X ,

$$\text{leaves}(X) = \max(0, \lceil \|X\|/\beta \rceil - 1)$$

counts the leaf chunks induced by a body byte string of length $\|X\|$. Thus n_{leaf} is bounded by the sum of $\text{leaves}(X)$ over the valid construction computations that create final leaf transcripts, where X is the plaintext body for encryption computations and the submitted ciphertext body for plaintext-computing, verification, and committing decryption checks. Each distinct final leaf transcript contains a final leaf duplex call counted in N , so $n_{\text{leaf}}(\mathcal{B}) \leq N$ unconditionally. The ratio n_{leaf}/N is largest when most processed bytes lie in leaf chunks, where each full leaf contributes $1 + \beta/\rho = 50$ duplex calls to N per final leaf transcript.

All theorems are stated for adversaries whose execution-uniform caps hold at the stated bounds.

4.5 Loss Terms

Five closed-form expressions recur across the security bounds— Lift_c throughout, with KeyGuess_k , RootColl_τ , RootPre_τ , and LeafColl_τ in some. We define them here, after the resource caps that parameterize them.

Flat sponge loss. The flat sponge differentiating loss is

$$\text{Lift}_c(N_{\text{tot}}) = N_{\text{tot}}(N_{\text{tot}} + 1)/2^{c+1}.$$

For $N_{\text{tot}} < 2^c$ it upper-bounds the [BDPVA08] sponge-differentiating advantage; Lemma 20 carries out the derivation from the [BDPVA08, Lemmas 4 and 5] product bounds through the Weierstrass product inequality. For $N_{\text{tot}} \geq 2^c$, $\text{Lift}_c(N_{\text{tot}}) \geq 1$ and the advantage bound is trivial without any appeal to indistinguishability, and already at $N_{\text{tot}} \geq 2^{c/2}$ the bound exceeds $1/2$ and provides no meaningful security.

Key-guessing term. The ideal system key-guessing term of [BDPVA11] is

$$\text{KeyGuess}_k(Q) = Q/2^k.$$

It accounts for direct RO_{flat} queries whose inputs satisfy the decoder KeyedTWOOut of Table 12. Each such query is an adaptive test of the decoded candidate key against the hidden key K , independent of the query’s requested output length. Lemma 10 bounds the probability that any of Q direct queries tests the right key by $Q/2^k$.

Root-collision term. The s -way root tag multi-collision term is

$$\text{RootColl}_\tau(Q, s) = \binom{Q+s}{s} / 2^{\tau(s-1)},$$

with the abbreviation $\text{RootColl}_{\tau_{\text{root}}}(Q, s) = \text{RootColl}_\tau(Q, s)$ at $\tau = \tau_{\text{root}}$. The $+s$ is candidate-sequence slack for the s final root transcript inputs generated by the challenger after the adversary outputs its tuples. These lookups are construction-internal and are not added to $q_{\text{RO}}(\mathcal{B})$; after the flat sponge lift of Section B, the public-permutation bounds instantiate Q as N_{prim} .

Root-preimage term. The root tag preimage term is

$$\text{RootPre}_\tau(Q) = (Q+1)/2^\tau,$$

with the abbreviation $\text{RootPre}_{\tau_{\text{root}}}(Q) = \text{RootPre}_\tau(Q)$ at $\tau = \tau_{\text{root}}$. It is the preimage analog of RootColl_τ : a single query-independent target tag replaces the s -way internal collision, one challenger lookup replaces the s final-root lookups, and the exponent drops from $\tau(s-1)$ to τ . As in the root-collision term, the flat sponge lift gives $Q = N_{\text{prim}}$.

Leaf-collision term. The known-key leaf tag collision term is

$$\text{LeafColl}_\tau(Q) = \binom{Q}{2} / 2^\tau.$$

It is the pairwise birthday bound on τ -bit leaf tag projections used by the collision-resistance theorem (Theorem 8). In this game the keys are adversary-supplied, so direct primitive queries can pre-read candidate leaf tag points; after the flat sponge lift this gives $Q = N_{\text{prim}} + N_{\text{leaf}}$. The hidden-key AE-style bounds instead carry the linear same-slot term $n_{\text{leaf}}/2^{\tau_{\text{leaf}}}$ of Lemma 11, written inline.

5 Authenticated Encryption Security

This section proves the AEAD security of TW128. The first theorem is a random oracle multi-user indistinguishability bound. The public-permutation theorem then follows by the flat sponge lift of Section B, whose real-vs-ideal form adds the ideal-side proximity term $\text{Lift}_c(N_{\text{prim}})$, and the all-in-one AE bound is the single-user corollary. The section then proves INT-RUP integrity in the separated-interface model of [ABL⁺14]. The root tag hash and committing results are proved separately in Section 6; all concrete instantiations are summarized in Section 7.

All random oracle adversaries in this section are flat-RO adversaries (Section 4.3) with the displayed resources, and the flat sponge lift transfers each random oracle bound using the derived-adversary cap $q_{\text{RO}}(\mathcal{B}) \leq N_{\text{prim}}$.

Theorem 1 (Random Oracle Multi-User Indistinguishability). *For every flat-RO mu-ind adversary $\mathcal{B}$ with execution-uniform user cap D and resource counts of Section 4.2,*

$$\text{Adv}_{\text{RO}}^{\text{mu-ind}}(\mathcal{B}) \leq \frac{D(D-1)}{2^{k+1}} + D \cdot \text{KeyGuess}_k(q_{\text{RO}}(\mathcal{B})) + \frac{n_{\text{leaf}}(\mathcal{B})}{2^{\tau_{\text{leaf}}}} + \frac{q_v(\mathcal{B})}{2^{\tau_{\text{root}}}}.$$

Here q_v is the non-replay verification-query count of Section 4.4.

Proof. Use the pre-sampled-key formulation of Section 4.2: sample keys $K_1, \dots, K_D$ at the start of the experiment, and let the d -th New call activate K_d , unused keys being ignored. Let coll denote the event that two of the pre-sampled keys $K_1, \dots, K_D$ collide;

905 $\Pr[\text{coll}] \leq \binom{D}{2}/2^k = D(D-1)/2^{k+1}$ by union bound. In the random oracle model, use a
 906 standard hybrid over the D users: H_d answers users $1, \dots, d$ with (Rand, Replay) and users
 907 $d+1, \dots, D$ with the real construction interface, so H_0 is the real mu-ind game and H_D
 908 is the ideal one. Let H_d^* be H_d with a fixed output bit on coll. Since coll has the same
 909 probability in all hybrids,

$$910 \quad \text{Adv}_{\text{RO}}^{\text{mu-ind}}(\mathcal{B}) \leq \Pr[\text{coll}] + |\Pr[\mathcal{B}^{H_0^*} \Rightarrow 1] - \Pr[\mathcal{B}^{H_D^*} \Rightarrow 1]|.$$

911 By the triangle inequality, it remains to bound each adjacent stopped hybrid. For each d ,
 912 Lemma 17 gives

$$913 \quad |\Pr[\mathcal{B}^{H_{d-1}^*} \Rightarrow 1] - \Pr[\mathcal{B}^{H_d^*} \Rightarrow 1]| \leq \text{KeyGuess}_k(q_{\text{RO}}(\mathcal{B})) + \frac{n_{\text{leaf}}^{(d)}}{2^{\tau_{\text{leaf}}}} + \frac{q_v^{(d)}}{2^{\tau_{\text{root}}}}.$$

914 In outline, Lemma 17 proceeds through the explicit game sequence of Definition 5: an
 915 instrumented implementation of H_{d-1}^* and an all-reject game share all code except the
 916 return statement of a successful non-replay user- d verification comparison, which sets the
 917 flag `forge`, while a direct RO_{flat} query decoding to the hidden key K_d sets the flag `badkey`.
 918 The real pair agrees pointwise until `forge`. Claim 3 in the proof of Lemma 17 couples the
 919 all-reject game to H_d^* until `badkey`: Encryption Transcript Marginal Uniformity (Lemma 12)
 920 supplies freshness and distinctness for newly generated user- d encryption points, while the
 921 all-reject structure ensures that any reused user- d root-tag entry was previously read only
 922 through constant reject answers and remains uniform conditioned on the visible transcript.
 923 Environment lookups for the other users decode to their own keys and touch neither flag
 924 (Lemma 13). A first-occurrence decomposition of `forge` $\vee$ `badkey` in the all-reject game then
 925 yields the three terms: `badkey` is charged at most $\text{KeyGuess}_k(q_{\text{RO}}(\mathcal{B}))$ by the foreign-key
 926 key-test claim of Section A.9; a same-slot collision between two user- d final leaf tags—at
 927 least one reached by an encryption computation—is charged at most $n_{\text{leaf}}^{(d)}/2^{\tau_{\text{leaf}}}$ by the
 928 same-slot case of Lemma 11; and a successful comparison before both events must lie in a
 929 verification-first transcript class and is charged at most $q_v^{(d)}/2^{\tau_{\text{root}}}$ by Lemma 15.

930 Summing across $d = 1, \dots, D$ gives $D \text{KeyGuess}_k(q_{\text{RO}}(\mathcal{B}))$ for the key-test contribution.
 931 The leaf tag contribution sums exactly:

$$932 \quad \sum_d \frac{n_{\text{leaf}}^{(d)}}{2^{\tau_{\text{leaf}}}} = \frac{n_{\text{leaf}}(\mathcal{B})}{2^{\tau_{\text{leaf}}}},$$

933 since $\sum_d n_{\text{leaf}}^{(d)} = n_{\text{leaf}}(\mathcal{B})$. The verification-tag contribution satisfies

$$934 \quad \sum_d q_v^{(d)}/2^{\tau_{\text{root}}} = q_v(\mathcal{B})/2^{\tau_{\text{root}}}.$$

935 Combining these estimates with the coll penalty gives the stated RO bound. $\square$

936 **Theorem 2** (Multi-User Indistinguishability). *For every mu-ind adversary $\mathcal{A}$ against*
 937 *TW128 with execution-uniform user cap D and the global resources of Section 4.2,*

$$938 \quad \text{Adv}_{\text{TW128}}^{\text{mu-ind}}(\mathcal{A}) \leq \text{Lift}_c(N_{\text{tot}}) + \text{Lift}_c(N_{\text{prim}}) + \frac{D(D-1)}{2^{k+1}} + D \cdot \text{KeyGuess}_k(N_{\text{prim}}) + \frac{N_{\text{leaf}}}{2^{\tau_{\text{leaf}}}} + \frac{Q_v}{2^{\tau_{\text{root}}}}.$$

939 *Proof.* By the Lift Application corollary (Corollary 9) for the mu-ind game—a real-vs-ideal
 940 game, so the ideal-side proximity term applies—there is a random oracle mu-ind adversary
 941 $\mathcal{B} = \mathcal{B}_{\text{derived}}(\mathcal{A})$ with $q_{\text{RO}}(\mathcal{B}) \leq N_{\text{prim}}$, $n_{\text{leaf}}(\mathcal{B}) \leq N_{\text{leaf}}$, $q_v(\mathcal{B}) \leq Q_v$, and the per-user
 942 construction-side caps of Section 4.2 such that

$$943 \quad \text{Adv}_{\text{TW128}}^{\text{mu-ind}}(\mathcal{A}) \leq \text{Lift}_c(N_{\text{tot}}) + \text{Adv}_{\text{RO}}^{\text{mu-ind}}(\mathcal{B}) + \text{Lift}_c(N_{\text{prim}}).$$

944 Theorem 1 bounds the random oracle term. Substituting $q_{\text{RO}}(\mathcal{B}) \leq N_{\text{prim}}$, $n_{\text{leaf}}(\mathcal{B}) \leq N_{\text{leaf}}$,
 945 and $q_v(\mathcal{B}) \leq Q_v$ gives the stated bound. $\square$

946 **Corollary 1** (All-in-One AE Security). *For every nonce-respecting AE adversary $\mathcal{A}$*
 947 *against TW128 with the resources of Section 4.4,*

$$948 \quad \text{Adv}_{\text{TW128}}^{\text{ae}}(\mathcal{A}) \leq \text{Lift}_c(N_{\text{tot}}) + \text{Lift}_c(N_{\text{prim}}) + \text{KeyGuess}_k(N_{\text{prim}}) + \frac{N_{\text{leaf}}}{2^{\tau_{\text{leaf}}}} + \frac{Q_v}{2^{\tau_{\text{root}}}}.$$

949 *Proof.* Build a mu-ind adversary $\mathcal{A}^\mu$ that makes one New query, runs $\mathcal{A}$, and forwards
 950 each encryption and verification query of $\mathcal{A}$ to user 1. The primitive interface is forwarded
 951 unchanged. The real and ideal views of $\mathcal{A}$ in the all-in-one AE experiment are exactly
 952 the real and ideal views of $\mathcal{A}^\mu$ in the mu-ind experiment with $D = 1$, and the resource
 953 counts are unchanged. On a replayed encryption tuple the two real verification oracles
 954 agree: the mu-ind oracle short-circuits through its replay table, while the AE oracle
 955 re-runs decryption, which accepts by TW128 decryption correctness (Section 3.6); all other
 956 queries are answered by identical code. Applying Theorem 2 with $D = 1$ eliminates the
 957 key-collision term and gives the stated bound. $\square$

958 **Theorem 3** (Random Oracle INT-RUP Integrity). *For every flat-RO INT-RUP adversary*
 959 *$\mathcal{B}$ against TW128 with nonce-respecting encryption queries and the resources of Section 4.4,*

$$960 \quad \text{Adv}_{\text{RO}}^{\text{INT-RUP}}(\mathcal{B}) \leq \text{KeyGuess}_k(q_{\text{RO}}(\mathcal{B})) + \frac{n_{\text{leaf}}(\mathcal{B})}{2^{\tau_{\text{leaf}}}} + \frac{q_v(\mathcal{B})}{2^{\tau_{\text{root}}}}.$$

961 *Plaintext-computing queries contribute to N and n_{leaf} , but not to q_v .*

962 *Proof.* Work in the flat-RO INT-RUP game of Section 4.3. Let **forge** be the event that a
 963 valid non-replay verification query accepts. Let **bad_{key}** be the event that a direct RO_{flat}
 964 query decodes under **KeyedTWOut** to the hidden key K and a counted output-point class,
 965 and let **bad_{leaf}** be the event that some construction-generated final leaf transcript receives
 966 the same τ_{leaf} -bit leaf tag projection as a distinct final leaf transcript that shares its key,
 967 nonce, and leaf index and is reached by an encryption computation (the same-slot event
 968 of Lemma 11). Final leaf transcripts generated by encryption, plaintext-computing, and
 969 verification computations are all included in n_{leaf} .

970 The key-test term is unchanged by the plaintext-computing oracle: Lemma 14, stated
 971 for the INT-RUP extension, shows that before **bad_{key}** the hidden key remains uniform
 972 outside the candidate keys already tested by direct queries, and Lemma 10 gives

$$973 \quad \Pr[\text{bad}_{\text{key}}] \leq \text{KeyGuess}_k(q_{\text{RO}}(\mathcal{B})).$$

974 For the leaf term, before **bad_{key}** no direct query has pre-read a hidden-key final leaf
 975 tag point: such a query would decode under **KeyedTWOut** to K and the corresponding
 976 $\mathcal{L}_{\text{tag}}(j)$ class. Encryption queries are nonce-respecting. Couple the real INT-RUP execution
 977 to the all-reject continuation used in the proof of Lemma 16. Up to the first accepting
 978 non-replay submission, the construction computations and leaf-tag table accesses agree;
 979 the continuation differs only by deferring the relevant final root-tag value until comparison.
 980 Thus the internal event **bad_{leaf}** transfers to the all-reject continuation, where the same-slot
 981 case of Lemma 11, applied to the enlarged set of construction computations, gives

$$982 \quad \Pr[\text{bad}_{\text{leaf}} \text{ before } \text{bad}_{\text{key}} \cup \text{forge}] \leq \frac{n_{\text{leaf}}(\mathcal{B})}{2^{\tau_{\text{leaf}}}}.$$

983 It remains to bound accepting verification before **bad_{key}** $\cup$ **bad_{leaf}**. The INT-RUP root-
 984 guessing lemma (Lemma 16) gives

$$985 \quad \Pr[\text{forge} \text{ before } \text{bad}_{\text{key}} \cup \text{bad}_{\text{leaf}}] \leq q_v(\mathcal{B})/2^{\tau_{\text{root}}}.$$

986 Combining the three bounds through the first-occurrence decomposition of $\Pr[\text{forge}]$ proves
 987 the theorem. $\square$

Theorem 4 (INT-RUP Integrity). *For every INT-RUP adversary $\mathcal{A}$ against TW128 with nonce-respecting encryption queries and the resources of Section 4.4,*

$$\text{Adv}_{\text{TW128}}^{\text{INT-RUP}}(\mathcal{A}) \leq \text{Lift}_c(N_{\text{tot}}) + \text{KeyGuess}_k(N_{\text{prim}}) + \frac{N_{\text{leaf}}}{2^{\tau_{\text{leaf}}}} + \frac{Q_v}{2^{\tau_{\text{root}}}}.$$

Proof. By the Lift Application corollary (Corollary 9) for the INT-RUP game, there is a random oracle INT-RUP adversary $\mathcal{B} = \mathcal{B}_{\text{derived}}(\mathcal{A})$ with $q_{\text{RO}}(\mathcal{B}) \leq N_{\text{prim}}$, $n_{\text{leaf}}(\mathcal{B}) \leq N_{\text{leaf}}$, and $q_v(\mathcal{B}) \leq Q_v$ such that

$$\text{Adv}_{\text{TW128}}^{\text{INT-RUP}}(\mathcal{A}) \leq \text{Lift}_c(N_{\text{tot}}) + \text{Adv}_{\text{RO}}^{\text{INT-RUP}}(\mathcal{B}).$$

Theorem 3 bounds the random oracle term, and substituting the derived-adversary resource caps gives the stated bound. $\square$

6 Root Tag Hash Security and Committing Security

This section proves the digest claim for the root of the TW128 tree. The construction adapts the KangarooTwelve final-node pattern to authenticated encryption: leaf tags play the role of chaining values, and the root emits the single tag returned by H_{TW128} . The proof therefore tracks the final root inputs that can matter in the hash and committing games. The Root Tag Candidate Sequence lemma (Lemma 3) collects those inputs into a chronological candidate sequence, and the Root Tag Hash and Target-Hit Bound (Lemma 4) turns a collision or target hit among the candidates into the matching adaptive bound of Section A, controlling the chance that distinct inputs land on the same τ_{root} -bit root tag projection.

The hash random oracle theorems—collision (Theorem 5) and preimage (Theorems 6 and 7)—are then short instantiations, differing only in which case of the bound they invoke, whether the target is external or sampled from a fixed input slice, and whether a leaf tag collision term is needed when distinct inputs may use different ciphertext bodies. CMTDs-4 (Theorem 11) and the remaining committing notions reuse the collision case. Section B transfers each instance to the public permutation setting, Section 7 records the concrete bounds, and Corollary 6 forwards the remaining [BH22] committing notions to the same root-collision bound.

The instantiations rest on two structural separation facts, recorded first: distinct contexts (K, U, A) reach distinct final root transcripts regardless of the leaf tags they aggregate (Opening Triple Separation, Proposition 1), and—absent a collision among those leaf tags—so do distinct full inputs (Final-Root Input Separation, Proposition 2).

6.1 Structural Separation

Proposition 1 (Opening Triple Separation). *If two TW128 root computations have distinct (K, U, A) triples, then their bridged final root transcript inputs under $\text{flat}(\cdot)$ are distinct, regardless of whether each computation is an encryption or a decryption check, and whether or not they act on a common ciphertext.*

Proof. Consider two root computations i and j with contexts $(K_i, U_i, A_i) \neq (K_j, U_j, A_j)$, reaching final root transcripts R_i and R_j . Suppose for contradiction that $\text{flat}(R_i) = \text{flat}(R_j)$. Since the bridge $\text{flat}(\cdot)$ is injective on well-formed transcript prefixes (Lemma 8), the native flattened inputs of the two final-root prefixes are equal. Then Lemma 6 recovers the full duplex-call sequence and reads off the pair $(\sigma, \sigma_{\text{ph}})$ for each call. The seven 4-bit phase suffixes are pairwise distinct, so the σ_{init} , $\sigma_{\text{ad}}^{\pm}$, $\sigma_{\text{msg}}^{\pm}$, and (when present) $\sigma_{\text{agg}}^{\pm}$ calls in the recovered sequence are unambiguously identified. The aggregation phase is present exactly when $n \geq 1$, and the two recovered sequences are equal, so they either both carry an

aggregation phase or both omit it; either way the recovery of (K, U, A) below uses only the σ_{init} and AD calls, which precede any message or aggregation block.

Equality of the recovered sequences forces equality of the σ_{init} call, whose σ is $p_{\text{root}} \| K \| U$ with fixed-width K and U , hence $K_i = K_j$ and $U_i = U_j$. It also forces equality of the recovered AD subsequences. By the AD bullet of Lemma 6 the associated data phase contributes a σ_{ad}^+ -tagged block exactly when the associated data is nonempty, so the recovered sequences agree on the presence of the AD phase. If exactly one of A_i, A_j were empty the recovered sequences would differ in the presence of a σ_{ad}^+ call, contradicting their equality; hence either both are empty, giving $A_i = A_j = \varepsilon$, or both are nonempty, in which case the phase-recovery sub-claim of Lemma 6 inverts the ρ -byte-block partition map on the AD byte string and equal AD σ -block sequences imply $A_i = A_j$. In all cases $(K_i, U_i, A_i) = (K_j, U_j, A_j)$, contradicting the hypothesis of distinct triples. $\square$

Triple Separation handles distinct contexts; when inputs share a context but differ in their messages, distinctness of the final root transcripts needs the absence of a leaf tag collision, recorded next.

Proposition 2 (Final-Root Input Separation). *Let $\text{bad}_{\text{leaf}}^*$ be the event that two distinct construction-generated final leaf transcripts receive the same τ_{leaf} -bit RF projection. On $\neg \text{bad}_{\text{leaf}}^*$, if two TW128 construction computations reach final root transcripts with $\text{flat}(R) = \text{flat}(R')$, then they share the same context (K, U, A) and the same full ciphertext C . Equivalently, on $\neg \text{bad}_{\text{leaf}}^*$ computations with distinct (K, U, A, C) reach distinct bridged final root transcripts.*

Proof. By bridge injectivity (Lemma 8), $\text{flat}(R) = \text{flat}(R')$ gives equal native flattened final-root inputs. Lemma 6 then recovers the same root initialization $p_{\text{root}} \| K \| U$ —hence the same K and U —the same associated data blocks—hence the same A —the same root ciphertext chunk, and the same leaf count and ordered leaf tag sequence. Under $\neg \text{bad}_{\text{leaf}}^*$, Leaf-Transcript Recovery (Lemma 7) identifies the equal leaf tag sequences with the same final leaf transcripts, hence the same leaf ciphertext bodies; with the common root ciphertext chunk this gives the same full C . $\square$

6.2 Root Tag Candidate Bound

The candidate machinery rests on a single probabilistic engine, stated here and proved in full in Section A.5: adaptively chosen distinct random oracle inputs whose values are still unread when they are chosen behave, for collision and target-hitting purposes, as fresh uniform draws.

Lemma 2 (Adaptive Candidate Collision and Preimage). *Consider any interaction with a lazy random oracle $\text{RO}_{\text{flat}} : \{0, 1\}^* \rightarrow \{0, 1\}^\infty$ that builds a sequence of distinct candidate inputs $Y_1, \dots, Y_m$. The process may make arbitrary non-candidate oracle queries and may choose each candidate adaptively from the full prior transcript and private state sampled before the candidate's unread oracle value. Whenever a new candidate Y_j is appended, two premises hold: predictability—both the decision to append and the string Y_j are determined by the state preceding the append step—and unrevealed value—no oracle lookup before the append step has requested a positive number of bits of $\text{RO}_{\text{flat}}(Y_j)$; zero-length lookups at Y_j are permitted. Fix a projection length τ , and suppose the sequence has at most L candidates ($m \leq L$).*

1. Collision. *For every $s \geq 2$, the probability $\Pr[\text{coll}_{s, \tau}]$ that some s distinct candidates have equal τ -bit projections,*

$$[\text{RO}_{\text{flat}}(Y_{i_1})]_\tau = \dots = [\text{RO}_{\text{flat}}(Y_{i_s})]_\tau \quad \text{for some } 1 \leq i_1 < \dots < i_s \leq m,$$

$$\text{satisfies } \Pr[\text{coll}_{s, \tau}] \leq \binom{L}{s} \cdot 2^{-\tau(s-1)}.$$

1079 2. Preimage. For every target $t \in \{0, 1\}^\tau$ chosen independently of RO_{flat} , the probability
 1080 $\Pr[\text{pre}_{\tau,t}]$ that some candidate's τ -bit projection equals t , i.e. $\lfloor \text{RO}_{\text{flat}}(Y_j) \rfloor_\tau = t$ for
 1081 some $1 \leq j \leq m$, satisfies $\Pr[\text{pre}_{\tau,t}] \leq L \cdot 2^{-\tau}$.

1082 The preimage case has a second-preimage corollary: for a designated candidate position j_0 ,
 1083 the probability $\Pr[\text{sec}_{\tau,j_0}]$ that some candidate $j \neq j_0$ has $\lfloor \text{RO}_{\text{flat}}(Y_j) \rfloor_\tau = \lfloor \text{RO}_{\text{flat}}(Y_{j_0}) \rfloor_\tau$,
 1084 i.e. collides with the designated candidate on its τ -bit projection, satisfies $\Pr[\text{sec}_{\tau,j_0}] \leq$
 1085 $(L - 1) \cdot 2^{-\tau}$.

1086 *Proof sketch.* The premises license lazy sampling at the append step: predictability fixes
 1087 the candidate string before its value is drawn, and the unrevealed-value premise guarantees
 1088 that no bit of that value has been exposed, so $\text{RO}_{\text{flat}}(Y_j)$ is uniform and independent of
 1089 the entire prior transcript. Conditioning position by position, each fixed position set then
 1090 contributes at most $2^{-\tau(s-1)}$ (each fixed position at most $2^{-\tau}$), and a union bound over
 1091 the $\binom{L}{s}$ position sets (the L positions) gives the bounds. The full argument, including the
 1092 second-preimage corollary, appears in Section A.5. $\square$

1093 The remaining work is structural: exhibiting, in each game, a candidate sequence that
 1094 satisfies these premises and contains the final root transcripts the win event constrains.

1095 **Lemma 3** (Root Tag Candidate Sequence). *Consider one of the hash or committing*
 1096 *random oracle games of Section 4.3. The adversary $\mathcal{B}$ makes direct RO_{flat} queries and has*
 1097 *no construction oracle. The only construction-internal positive-length lookups at root tag*
 1098 *inputs are the root tag samplings of $h \geq 1$ deterministic challenger encryption or decryption*
 1099 *checks at game-fixed points. Apart from the standard preimage game's designated sampled*
 1100 *source check, all check inputs and keys are adversary supplied. In that preimage game*
 1101 *the sampled source input, including its key, is private challenger state sampled before the*
 1102 *source tag is revealed, and the source computation counts as one of the h checks. Let the*
 1103 *i -th check reach a final root transcript R_i whose root tag is read by a single RO_{flat} lookup at*
 1104 *$\text{flat}(R_i)$. Call an RO_{flat} input a root tag input if it decodes under KeyedTWOut to a root*
 1105 *tag output-point class— $\mathcal{R}_{\text{msg}}^+$ (the final root message block, which carries the root tag only*
 1106 *on the no-leaf path) or $\mathcal{R}_{\text{tag}}$ (the aggregated root tag); multi-chunk computations also reach*
 1107 *$\mathcal{R}_{\text{msg}}^+$ with zero requested output before aggregation, so this predicate over-approximates*
 1108 *the root tag candidate set, which only loosens the count. Every execution then admits a*
 1109 *chronological candidate sequence in which each root tag input is appended exactly once,*
 1110 *immediately before the first event that is either a direct RO_{flat} query at it or a positive-length*
 1111 *RO_{flat} lookup reading it, and inputs never reached by such an event are never appended.*
 1112 *This sequence is such that*

- 1113 1. the candidates are pairwise distinct, and if the challenger reaches its checks, each
 1114 $\text{flat}(R_i)$ is in the sequence no later than the sampling of the i -th root tag answer;
- 1115 2. it satisfies the predictability and unrevealed-value premises of Lemma 2; and
- 1116 3. its length is at most $q_{\text{RO}}(\mathcal{B}) + h$.

1117 *Proof.* A direct query to a keyed transcript is one of $\mathcal{B}$'s visible RO_{flat} queries counted in
 1118 $q_{\text{RO}}(\mathcal{B})$, not a separate construction query. The only construction computations in the
 1119 game are the h deterministic challenger checks named in the statement.

1120 *Distinctness and coverage.* Each root tag input is appended at its first triggering event
 1121 and never again, so the candidates are pairwise distinct by construction, and the rule
 1122 is well defined however direct queries and challenger checks interleave. In particular, in
 1123 the standard preimage game the source check's root tag sampling may be the trigger for
 1124 $\text{flat}(R_1)$; a later direct query at the same address is not a first trigger and appends nothing.
 1125 Each $\text{flat}(R_i)$ is appended no later than the sampling of the i -th root tag answer, since
 1126 that sampling is a positive-length lookup at $\text{flat}(R_i)$, giving (1).

Predictability. The decoded root-tag-input predicate depends on the input string alone through the exact decoder of Lemma 8, which ignores the requested output length; by Output-Point Separation (Corollary 8) each input decodes to at most one output-point class. The append rule also uses the already fixed next event: a direct query triggers at the address, while a construction lookup triggers only when it requests positive output length. The string itself is fixed in the state preceding the trigger: a direct query's input is chosen by $\mathcal{B}$ before its lookup, and a check's final root transcript is determined by the check's input and the prior oracle answers—including any private challenger state sampled before the candidate's unread value, which Lemma 2 expressly allows.

Unrevealed value. The trigger is the first direct query or positive-length lookup at the address, so no earlier direct query at it exists, and it suffices that no earlier construction lookup has read a positive number of its bits. An earlier positive-length construction lookup is either a stream output point, a leaf-tag point $\mathcal{L}_{\text{tag}}(j)$, or a root-tag sampling. The first two are different output-point classes and hence, by Corollary 8, have different addresses from root tag inputs; the last would itself have been the first trigger at this address. Zero-requested-output lookups, including the $\mathcal{R}_{\text{msg}}^+$ lookup a multi-chunk computation makes before aggregation, are permitted by the unrevealed-value premise of Lemma 2, including when such a zero-length lookup lands at another check's future no-leaf final-root address. Together with predictability this gives (2).

Length. Appends are triggered only by direct queries—at most $q_{\text{RO}}(\mathcal{B})$ of them—or by the h checks' root tag samplings; zero-length construction lookups never trigger an append. Hence the sequence has length at most $q_{\text{RO}}(\mathcal{B}) + h$, giving (3). $\square$

Lemma 4 (Root Tag Hash and Target-Hit Bound). *Consider one of the hash or committing TW128 random oracle games (Section 4.3) for an adversary $\mathcal{B}$, satisfying the hypotheses of Lemma 3, with win event W . By the Root Tag Candidate Sequence lemma (Lemma 3) the root tag inputs form a chronological candidate sequence meeting the premises of Lemma 2 with $\tau = \tau_{\text{root}}$. Writing $q_{\text{RO}} = q_{\text{RO}}(\mathcal{B})$:*

1. Root-collision case. *If W forces s checks to reach pairwise distinct bridged final root transcripts carrying a common root tag, then, with $h = s$ checks, $\Pr[W] \leq \text{RootColl}_{\tau_{\text{root}}}(q_{\text{RO}}, s)$;*
2. Fixed-target case. *If W forces a single check ($h = 1$) to carry a root tag equal to a target fixed independently of RO_{flat} , then $\Pr[W] \leq \text{RootPre}_{\tau_{\text{root}}}(q_{\text{RO}})$;*
3. Sampled-source case. *If the game first runs a designated source check, reveals its root tag, and W forces one later check to reach a distinct bridged final root transcript carrying that designated tag, then, with $h = 2$ checks, $\Pr[W] \leq \text{RootPre}_{\tau_{\text{root}}}(q_{\text{RO}})$.*

Proof. In each case the candidate sequence has the stated length and meets the premises of Lemma 2. (1) The s distinct candidates carrying a common τ_{root} -projection are an s -way collision; the collision case with $L = q_{\text{RO}} + s$ gives $\binom{q_{\text{RO}}+s}{s} \cdot 2^{-\tau_{\text{root}}(s-1)} = \text{RootColl}_{\tau_{\text{root}}}(q_{\text{RO}}, s)$. (2) The candidate hitting the fixed target is a preimage; the preimage case with $L = q_{\text{RO}} + 1$ gives $(q_{\text{RO}} + 1) \cdot 2^{-\tau_{\text{root}}} = \text{RootPre}_{\tau_{\text{root}}}(q_{\text{RO}})$. (3) Let j_{src} be the position of the designated source check. The later check is required to be a distinct bridged final root transcript, so it is a candidate position $j \neq j_{\text{src}}$ with the same τ_{root} -bit projection. This is an application of the second-preimage corollary, not the fixed-target case, because the designated source tag is sampled from RO_{flat} rather than fixed independently of it. The candidate sequence has length at most $L = q_{\text{RO}} + 2$, and the second-preimage corollary of Lemma 2 gives $(L - 1)2^{-\tau_{\text{root}}} = (q_{\text{RO}} + 1)2^{-\tau_{\text{root}}} = \text{RootPre}_{\tau_{\text{root}}}(q_{\text{RO}})$. $\square$

6.3 Hash Random Oracle Bounds

Theorem 5 (Random Oracle Collision Resistance of H_{TW128}). *For every $s \geq 2$ and every s -way collision adversary $\mathcal{B}$ in the TW128 random oracle model,*

$$\text{Adv}_{\text{RO}}^{\text{Coll}}(\mathcal{B}) \leq \text{LeafColl}_{\tau_{\text{leaf}}}(q_{\text{RO}}(\mathcal{B}) + n_{\text{leaf}}(\mathcal{B})) + \text{RootColl}_{\tau_{\text{root}}}(q_{\text{RO}}(\mathcal{B}), s).$$

Proof. Run the random oracle collision game. Let $\text{bad}_{\text{leaf}}^*$ be the event that two distinct construction-generated final leaf transcripts receive the same τ_{leaf} -bit RF projection; by the general count of Lemma 11, $\Pr[\text{bad}_{\text{leaf}}^*] \leq \text{LeafColl}_{\tau_{\text{leaf}}}(q_{\text{RO}}(\mathcal{B}) + n_{\text{leaf}}(\mathcal{B}))$, so it remains to bound the win under $\neg \text{bad}_{\text{leaf}}^*$. After rejecting unless the s output inputs are pairwise distinct and in the TW128 domain, the challenger performs s deterministic encryptions $\text{Enc}_{K_i}^{\text{RO}}(U_i, A_i, M_i)$, reaching $R_1, \dots, R_s$, all carrying a common root tag T on a win. Pairwise distinct tuples (K_i, U_i, A_i, M_i) have pairwise distinct (K_i, U_i, A_i, C_i) : two tuples differing in their context already differ in (K, U, A) , while two sharing a context but differing in the message yield distinct ciphertexts, since for a fixed context encryption maps the message to the ciphertext bijectively. So on $\neg \text{bad}_{\text{leaf}}^*$ Final-Root Input Separation (Proposition 2) makes $\text{flat}(R_1), \dots, \text{flat}(R_s)$ pairwise distinct. The Root Tag Hash and Target-Hit Bound (Lemma 4, part 1) bounds the win under $\neg \text{bad}_{\text{leaf}}^*$ by $\text{RootColl}_{\tau_{\text{root}}}(q_{\text{RO}}(\mathcal{B}), s)$; adding $\Pr[\text{bad}_{\text{leaf}}^*]$ gives the theorem. $\square$

Theorem 6 (Random Oracle Preimage Resistance of H_{TW128}). *Fix byte lengths $0 \leq a \leq A_{\text{max}}$ and $0 \leq \mu \leq M_{\text{max}}$, and let $m_{a,\mu} = k + \nu + 8a + 8\mu$. For every Pre adversary $\mathcal{B}$ against H_{TW128} on the fixed input slice $\mathcal{X}_{a,\mu}$ in the TW128 random oracle model (Section 4.3),*

$$\begin{aligned} \text{Adv}_{\text{RO}}^{\text{Pre}[a,\mu]}(\mathcal{B}) &\leq \text{LeafColl}_{\tau_{\text{leaf}}}(q_{\text{RO}}(\mathcal{B}) + n_{\text{leaf}}(\mathcal{B})) \\ &\quad + \text{RootPre}_{\tau_{\text{root}}}(q_{\text{RO}}(\mathcal{B})) + 2^{\tau_{\text{root}} - m_{a,\mu}}. \end{aligned}$$

Proof. Run the random oracle Pre game. The challenger samples $X^* \leftarrow_{\$} \mathcal{X}_{a,\mu}$, writes $X^* = (K^*, U^*, A^*, M^*)$, computes $T^* = H_{TW128}(X^*)$, gives T^* to $\mathcal{B}$, and later checks the output $X = (K, U, A, M)$. Let same be the event $X = X^*$. For a fixed random oracle table and fixed adversary coins, the adversary's output is a function of the τ_{root} -bit challenge tag T^* , so over the uniform source X^* it can equal X^* with probability at most $2^{\tau_{\text{root}}}/2^{m_{a,\mu}} = 2^{\tau_{\text{root}} - m_{a,\mu}}$. Hence $\Pr[\text{same}] \leq 2^{\tau_{\text{root}} - m_{a,\mu}}$.

It remains to bound winning executions with $X \neq X^*$. Let $\text{bad}_{\text{leaf}}^*$ be the event that two distinct construction-generated final leaf transcripts in the source and output checks receive the same τ_{leaf} -bit RF projection. By the general count of Lemma 11—which charges every direct query and every construction-generated final leaf transcript regardless of key provenance, covering the challenger-sampled source key—

$$\Pr[\text{bad}_{\text{leaf}}^*] \leq \text{LeafColl}_{\tau_{\text{leaf}}}(q_{\text{RO}}(\mathcal{B}) + n_{\text{leaf}}(\mathcal{B})).$$

On $\neg \text{bad}_{\text{leaf}}^*$, the distinct tuples X^* and X have distinct final root transcripts: distinct contexts are separated by Proposition 1, while equal contexts and distinct messages produce distinct ciphertexts, and then Proposition 2 applies. A win with $X \neq X^*$ therefore forces the later output check to hit the designated source tag T^* at a distinct bridged final root transcript. The sampled-source case, item 3, of Lemma 4 bounds this probability by $\text{RootPre}_{\tau_{\text{root}}}(q_{\text{RO}}(\mathcal{B}))$. Adding the source-guess and leaf-collision events gives the theorem. $\square$

Theorem 7 (Random Oracle Everywhere Preimage Resistance of H_{TW128}). *For every query-independent target tag $T^* \in \{0, 1\}^{\tau_{\text{root}}}$ and every ePre adversary $\mathcal{B}$ in the TW128 random oracle model (Section 4.3),*

$$\text{Adv}_{\text{RO}}^{\text{ePre}}(\mathcal{B}) \leq \text{RootPre}_{\tau_{\text{root}}}(q_{\text{RO}}(\mathcal{B})) = \frac{q_{\text{RO}}(\mathcal{B}) + 1}{2^{\tau_{\text{root}}}}.$$

Proof. Run the random oracle preimage game for the query-independent target T^* , which is independent of RO_{flat} . After rejecting unless the output input $X = (K, U, A, M)$ is in the TW128 domain, the challenger performs one encryption $\text{Enc}_K^{\text{RO}}(U, A, M)$, reaching a final root transcript R . A win means $\text{H}_{\text{TW128}}(X) = T^*$, i.e. the root tag $[\text{RO}_{\text{flat}}(\text{flat}(R))]_{\tau_{\text{root}}}$ at R equals the fixed target T^* . The fixed-target case, item 2, of Lemma 4 gives $\text{Adv}_{\text{RO}}^{\text{ePre}}(\mathcal{B}) \leq \text{RootPre}_{\tau_{\text{root}}}(q_{\text{RO}}(\mathcal{B}))$. $\square$

6.4 Public Permutation Hash Bounds

The flat sponge lift transfers the preceding random oracle bounds to the public permutation model. For these one-world hash games, the lift contributes $\text{Lift}_c(N_{\text{tot}})$, and the derived random oracle adversary has $q_{\text{RO}}(\mathcal{B}) \leq N_{\text{prim}}$; challenger construction computations—including the preimage game’s source encryption—are construction-internal work counted in N .

Theorem 8 (Collision Resistance of H_{TW128}). *For every $s \geq 2$ and every adversary $\mathcal{A}$ against the s -way multi-collision resistance of H_{TW128} in the sense of [BH22], with the resources of Section 4.4,*

$$\text{Adv}_{\text{TW128}}^{\text{Coll}}(\mathcal{A}) \leq \text{Lift}_c(N_{\text{tot}}) + \text{LeafColl}_{\tau_{\text{leaf}}}(N_{\text{prim}} + N_{\text{leaf}}) + \text{RootColl}_{\tau_{\text{root}}}(N_{\text{prim}}, s).$$

Proof. By the Lift Application corollary (Corollary 9) for the collision game, the derived random oracle adversary $\mathcal{B}$ has $q_{\text{RO}}(\mathcal{B}) \leq N_{\text{prim}}$ and $n_{\text{leaf}}(\mathcal{B}) \leq N_{\text{leaf}}$. The random oracle bound of Theorem 5 is nondecreasing in both $q_{\text{RO}}(\mathcal{B})$ and $n_{\text{leaf}}(\mathcal{B})$, so the monotone form of Corollary 9 gives the stated bound. The s colliding inputs may differ in their messages, hence in their ciphertext bodies; the LeafColl term accounts for the only way two distinct inputs sharing a context (K, U, A) can reach the same final root transcript—a collision among the leaf tags it absorbs—while distinct contexts already reach distinct final root transcripts by Opening Triple Separation (Proposition 1). $\square$

Corollary 2 (Two-Way Collision Resistance of H_{TW128}). *Specializing Theorem 8 to $s = 2$ gives the classical two-way collision resistance of [RS04]: for every collision adversary $\mathcal{A}$ against H_{TW128} with the resources of Section 4.4,*

$$\text{Adv}_{\text{TW128}}^{\text{Coll}}(\mathcal{A}) \leq \text{Lift}_c(N_{\text{tot}}) + \text{LeafColl}_{\tau_{\text{leaf}}}(N_{\text{prim}} + N_{\text{leaf}}) + \text{RootColl}_{\tau_{\text{root}}}(N_{\text{prim}}, 2),$$

$$\text{with } \text{RootColl}_{\tau_{\text{root}}}(N_{\text{prim}}, 2) = \binom{N_{\text{prim}}+2}{2} / 2^{\tau_{\text{root}}}.$$

Proof. The $s = 2$ case of Theorem 8; [BH22] note that $s = 2$ recovers the classical notion of collision resistance, which under the keyed-by-the-permutation convention of Section 2.5 is the Coll notion of [RS04]. $\square$

Corollary 3 (Second-Preimage Resistance of H_{TW128}). *Fix byte lengths $0 \leq a \leq A_{\text{max}}$ and $0 \leq \mu \leq M_{\text{max}}$. For every adversary against H_{TW128} , with the resources of Section 4.4, in either the standard second-preimage (Sec) or everywhere second-preimage (eSec) sense of [RS04] on the fixed input slice $\mathcal{X}_{a,\mu}$, the corresponding advantage is bounded, independently of (a, μ) , by the two-way collision bound of Corollary 2:*

$$\text{Lift}_c(N_{\text{tot}}) + \text{LeafColl}_{\tau_{\text{leaf}}}(N_{\text{prim}} + N_{\text{leaf}}) + \text{RootColl}_{\tau_{\text{root}}}(N_{\text{prim}}, 2).$$

Proof. We use [RS04] to identify the standard and everywhere second-preimage notions with the fixed-slice notions $\text{Sec}[m_{a,\mu}]$ and $\text{eSec}[m_{a,\mu}]$ under the canonical encoding of $\mathcal{X}_{a,\mu}$; the quantitative reduction is explicit.

For a standard Sec adversary $\mathcal{A}$, define a two-way collision adversary $\mathcal{C}$ as follows. It samples $X_0 \leftarrow \$ \mathcal{X}_{a,\mu}$, runs $\mathcal{A}^{\pi, \pi^{-1}}(X_0)$, receives X_1 , and outputs (X_0, X_1) . Whenever $\mathcal{A}$

wins, X_1 is a valid input, $X_1 \neq X_0$, and $H_{TW128}(X_1) = H_{TW128}(X_0)$, so $\mathcal{C}$ wins the two-way collision game with exactly the same probability.

For eSec, fix a source $X_0 \in \mathcal{X}_{a,\mu}$ and define $\mathcal{C}_{X_0}$ to run the eSec adversary on that source and output (X_0, X_1) . The same pointwise implication holds for every X_0 ; taking the maximum over sources gives the eSec bound.

The induced collision adversary makes exactly the direct primitive queries made by $\mathcal{A}$. Its two deterministic challenger checks are the source and candidate checks, counted exactly as in the $s = 2$ collision game. Applying Corollary 2 to the induced collision adversary gives the displayed bound. $\square$

Theorem 9 (Preimage Resistance of H_{TW128}). *Fix byte lengths $0 \leq a \leq A_{\max}$ and $0 \leq \mu \leq M_{\max}$, and let $m_{a,\mu} = k + \nu + 8a + 8\mu$. For every adversary $\mathcal{A}$ against H_{TW128} , with the resources of Section 4.4, in the standard preimage (Pre) sense of [RS04] on the fixed input slice $\mathcal{X}_{a,\mu}$,*

$$\begin{aligned} \text{Adv}_{TW128}^{\text{Pre}[a,\mu]}(\mathcal{A}) \leq & \text{Lift}_c(N_{\text{tot}}) + \text{LeafColl}_{\tau_{\text{leaf}}}(N_{\text{prim}} + N_{\text{leaf}}) \\ & + \text{RootPre}_{\tau_{\text{root}}}(N_{\text{prim}}) + 2^{\tau_{\text{root}} - m_{a,\mu}}. \end{aligned}$$

Proof. By the Lift Application corollary (Corollary 9) for the preimage game, the derived random oracle adversary $\mathcal{B}$ has $q_{\text{RO}}(\mathcal{B}) \leq N_{\text{prim}}$ and $n_{\text{leaf}}(\mathcal{B}) \leq N_{\text{leaf}}$; per Section 4.4, these caps cover both challenger encryptions' final leaf transcripts, the source encryption's included; the encryptions' duplex calls are counted in N and enter through $\text{Lift}_c(N_{\text{tot}})$. The random oracle bound of Theorem 6 is nondecreasing in both resources, so the monotone form of Corollary 9 gives the stated bound. $\square$

Theorem 10 (Everywhere Preimage Resistance of H_{TW128}). *For every target tag $T^* \in \{0, 1\}^{\tau_{\text{root}}}$ fixed in advance and every ePre adversary $\mathcal{A}$ against H_{TW128} with the resources of Section 4.4,*

$$\text{Adv}_{TW128}^{\text{ePre}}(\mathcal{A}) \leq \text{Lift}_c(N_{\text{tot}}) + \text{RootPre}_{\tau_{\text{root}}}(N_{\text{prim}}) = \text{Lift}_c(N_{\text{tot}}) + \frac{N_{\text{prim}} + 1}{2^{\tau_{\text{root}}}}.$$

Proof. By the Lift Application corollary (Corollary 9) for the everywhere-preimage game—whose target T^* is fixed before the primitive sampling—the derived random oracle adversary $\mathcal{B}$ has $q_{\text{RO}}(\mathcal{B}) \leq N_{\text{prim}}$, and the random oracle bound of Theorem 7 is nondecreasing in $q_{\text{RO}}(\mathcal{B})$, so the monotone form of Corollary 9 gives the stated bound. Compared with the s -way root-collision term $\text{RootColl}_{\tau_{\text{root}}}(\cdot, s)$ of Section 4.5, the single fixed target replaces the exponent $\tau_{\text{root}}(s - 1)$ by τ_{root} . $\square$

Corollary 4 (Root Tag as a Compact Commitment). *The wrapper H_{TW128} is a secure hash of the full encryption context (K, U, A, M) in the sense of [RS04]: collision-resistant (Theorem 8), preimage-resistant on each fixed input slice (Theorem 9), second-preimage-resistant (Corollary 3), and everywhere preimage-resistant (Theorem 10). Under the implemented parameters and work cap $N_{\text{tot}} \leq 2^{63}$, these advantages are at most 2^{-128} (Corollary 7), so the single τ_{root} -bit root tag is a binding compact commitment to (K, U, A, M) with standard and everywhere preimage resistance, native to the mode.*

Proof. Collision resistance is Theorem 8; standard preimage resistance on fixed input slices is Theorem 9; second-preimage resistance follows from collision resistance by Corollary 3; and everywhere preimage resistance is Theorem 10. These guarantees quantify the adversary's success at the same τ_{root} -bit projection of the final root transcript, and Corollary 7 evaluates them at the implemented parameters. Binding is the compact-commitment property supplied by collision and second-preimage resistance. Standard preimage resistance is the property that a tag produced from a uniformly sampled fixed-length opening does not reveal an efficiently findable opening; everywhere preimage resistance is the corresponding fixed-target property for tags chosen independently of the primitive. $\square$

6.5 Committing Consequences

Theorem 11 (Random Oracle CMTDs-4). *For every $s \geq 2$ and every s -way CMTDs-4 adversary $\mathcal{B}$ in the TW128 random oracle model (Section 4.3),*

$$\mathbf{Adv}_{\text{RO}}^{\text{CMTDs-4}}(\mathcal{B}) \leq \text{RootColl}_{\tau_{\text{root}}}(q_{\text{RO}}(\mathcal{B}), s) = \frac{\binom{q_{\text{RO}}(\mathcal{B})+s}{s}}{2^{\tau_{\text{root}}(s-1)}}.$$

Proof. Run the random oracle CMTDs-4 game. After rejecting on any failed syntax, message-length, or pairwise-full-input check, the challenger performs s deterministic decryption checks $\text{Dec}_{K_i}^{\text{RO}}(U_i, A_i, C \parallel T)$ on the one common ciphertext $C \parallel T$, reaching $R_1, \dots, R_s$. On a winning execution the tuples (K_i, U_i, A_i, M_i) are pairwise distinct, and since all checks decrypt the common ciphertext, determinism of Dec makes the triples (K_i, U_i, A_i) pairwise distinct—equal triples would return equal messages, hence equal tuples. By Opening Triple Separation (Proposition 1) the candidates $\text{flat}(R_1), \dots, \text{flat}(R_s)$ are pairwise distinct (leaf tag collisions cannot merge them), and every accepting check carries the common root tag T . The Root Tag Hash and Target-Hit Bound (Lemma 4, part 1) gives $\mathbf{Adv}_{\text{RO}}^{\text{CMTDs-4}}(\mathcal{B}) \leq \text{RootColl}_{\tau_{\text{root}}}(q_{\text{RO}}(\mathcal{B}), s)$. $\square$

Corollary 5 (CMTDs-4 Committing Security). *For every $s \geq 2$ and every s -way CMTDs-4 adversary $\mathcal{A}$ against TW128 with the resources of Section 4.4,*

$$\mathbf{Adv}_{\text{TW128}}^{\text{CMTDs-4}}(\mathcal{A}) \leq \text{Lift}_c(N_{\text{tot}}) + \text{RootColl}_{\tau_{\text{root}}}(N_{\text{prim}}, s) = \text{Lift}_c(N_{\text{tot}}) + \frac{\binom{N_{\text{prim}}+s}{s}}{2^{\tau_{\text{root}}(s-1)}}.$$

Proof. A CMTDs-4 winner gives one ciphertext $C \parallel T$ and s distinct inputs (K_i, U_i, A_i, M_i) such that $\text{Dec}_{K_i}(U_i, A_i, C \parallel T) = M_i$ for every i . By TW128 tidiness (Lemma 1), each opened input re-encrypts to the same ciphertext: $\text{Enc}_{K_i}(U_i, A_i, M_i) = C \parallel T$. In particular the opened inputs share the root tag T . For the displayed no-LeafColl bound, use the dedicated root-collision derivation of Theorem 11, lifted by Corollary 9: in this all-bodies-equal case, the single shared body C forces the triples (K_i, U_i, A_i) pairwise distinct; by Opening Triple Separation (Proposition 1), a leaf tag collision cannot merge two openings and no LeafColl term arises. $\square$

Remaining committing notions. The hierarchy and s -way extension of the [BH22] committing notions are given in [BH22, §3, Figs. 4–5, and App. A]. We use those relations for the D-notions and a direct source-game argument for the CMTs- ℓ notions to preserve the resource split of Section 4.4.

Corollary 6 (Committing-Notion Root-Collision Bound). *For every $s \geq 2$, every*

$$X \in \{\text{CMTDs-1}, \text{CMTDs-3}, \text{CMTDs-4}, \text{CMTs-1}, \text{CMTs-3}, \text{CMTs-4}\},$$

and every s -way X -adversary $\mathcal{A}$ against TW128 with the resources of Section 4.4,

$$\mathbf{Adv}_{\text{TW128}}^X(\mathcal{A}) \leq \text{Lift}_c(N_{\text{tot}}) + \text{RootColl}_{\tau_{\text{root}}}(N_{\text{prim}}, s) = \text{Lift}_c(N_{\text{tot}}) + \frac{\binom{N_{\text{prim}}+s}{s}}{2^{\tau_{\text{root}}(s-1)}}.$$

Proof. CMTDs-4 is bounded directly by Corollary 5. For the remaining D-notions, the s -way hierarchy of [BH22, §3, Fig. 5 and App. A] gives, for $\ell \in \{1, 3\}$, $\mathbf{Adv}_{\text{TW128}}^{\text{CMTDs-}\ell}(\mathcal{A}) \leq \mathbf{Adv}_{\text{TW128}}^{\text{CMTDs-4}}(\mathcal{A})$ for the same output distribution. By the conservative committing-game accounting of Section 4.4, the CMTDs- ℓ and CMTDs-4 resource caps charge the same deterministic decryption checks after syntax, domain, and message-length validation, independently of their distinctness gates.

For CMTs- ℓ , the challenger performs s deterministic encryption checks, which supply the candidate final-root transcripts (Lemma 3 builds the candidate sequence for encryption

and decryption checks alike). On a winning execution all s encryptions are equal, hence of equal length and with one common root tag T . Pairwise WiC_1 -distinctness gives pairwise distinct keys, pairwise WiC_3 -distinctness gives pairwise distinct (K, U, A) triples, and pairwise WiC_4 -distinctness gives pairwise distinct full inputs; in the last case, equal triples and equal ciphertexts would decrypt deterministically to equal messages, contradicting full-input distinctness. Thus each CMTs- ℓ winner reaches pairwise distinct opening triples. Opening Triple Separation (Proposition 1) makes the s final-root candidates distinct while their τ_{root} -bit projections all equal T . The hierarchy argument does not change the output distribution or the resource counts, and the CMTs- ℓ encryption checks are the construction-internal work counted in N by Section 4.4.

By the Lift Application corollary (Corollary 9) for the CMTs- ℓ game, the derived random oracle adversary $\mathcal{B}$ has $q_{\text{RO}}(\mathcal{B}) \leq N_{\text{prim}}$. The Root Tag Hash and Target-Hit Bound (Lemma 4, part 1) then bounds its win by $\text{RootColl}_{\tau_{\text{root}}}(q_{\text{RO}}(\mathcal{B}), s)$ in the random oracle model, which the lift transfers to the public permutation setting at N_{prim} ; the encryption checks are construction-internal work counted in N , not in N_{prim} . $\square$

7 Concrete Security Summary

Sections 5 and 6 separate the security analysis into AEAD bounds and root tag hash and committing bounds. This section instantiates those bounds at the implemented TW128 parameters and records the single concrete work cap used throughout the paper.

7.1 Concrete TW128 Bounds

The implemented parameters of Section 3.2 fix

$$k = c = \nu = \tau_{\text{root}} = \tau_{\text{leaf}} = 256.$$

Substituting into Corollaries 1, 3 and 5 and Theorems 2, 4 and 8 to 10 gives the following closed-form bounds.

All-in-one AE.

$$\text{Adv}_{\text{TW128}}^{\text{ae}}(\mathcal{A}) \leq \frac{N_{\text{tot}}(N_{\text{tot}} + 1)}{2^{257}} + \frac{N_{\text{prim}}(N_{\text{prim}} + 1)}{2^{257}} + \frac{N_{\text{prim}} + Q_v + N_{\text{leaf}}}{2^{256}}.$$

Multi-user indistinguishability.

$$\text{Adv}_{\text{TW128}}^{\text{mu-ind}}(\mathcal{A}) \leq \frac{N_{\text{tot}}(N_{\text{tot}} + 1)}{2^{257}} + \frac{N_{\text{prim}}(N_{\text{prim}} + 1)}{2^{257}} + \frac{DN_{\text{prim}} + Q_v + N_{\text{leaf}}}{2^{256}} + \frac{D(D - 1)}{2^{257}}.$$

INT-RUP integrity.

$$\text{Adv}_{\text{TW128}}^{\text{INT-RUP}}(\mathcal{A}) \leq \frac{N_{\text{tot}}(N_{\text{tot}} + 1)}{2^{257}} + \frac{N_{\text{prim}} + Q_v + N_{\text{leaf}}}{2^{256}}.$$

s -way collision resistance of H_{TW128} . For every $s \geq 2$,

$$\text{Adv}_{\text{TW128}}^{\text{Coll}}(\mathcal{A}) \leq \frac{N_{\text{tot}}(N_{\text{tot}} + 1)}{2^{257}} + \frac{\binom{N_{\text{prim}} + N_{\text{leaf}}}{2}}{2^{256}} + \frac{\binom{N_{\text{prim}} + s}{s}}{2^{256(s-1)}}.$$

For $s = 2$, this is the ordinary two-input collision bound:

$$\begin{aligned} \text{Adv}_{\text{TW128}}^{\text{Coll}}(\mathcal{A}) &\leq \frac{N_{\text{tot}}(N_{\text{tot}} + 1)}{2^{257}} + \frac{(N_{\text{prim}} + N_{\text{leaf}})(N_{\text{prim}} + N_{\text{leaf}} - 1)}{2^{257}} \\ &\quad + \frac{(N_{\text{prim}} + 1)(N_{\text{prim}} + 2)}{2^{257}}. \end{aligned}$$

1378 **Second-preimage resistance of H_{TW128} .** By Corollary 3, the standard and everywhere
 1379 [RS04] second-preimage advantages of H_{TW128} are bounded by the same $s = 2$ collision
 1380 expression above.

1381 **Preimage resistance of H_{TW128} .** For fixed byte lengths a, μ , let $m_{a,\mu} = k + \nu + 8a + 8\mu$.
 1382 By Theorem 9,

$$\begin{aligned} \text{Adv}_{TW128}^{\text{Pre}[a,\mu]}(\mathcal{A}) &\leq \frac{N_{\text{tot}}(N_{\text{tot}} + 1)}{2^{257}} + \frac{(N_{\text{prim}} + N_{\text{leaf}})(N_{\text{prim}} + N_{\text{leaf}} - 1)}{2^{257}} \\ &\quad + \frac{N_{\text{prim}} + 1}{2^{256}} + 2^{256 - m_{a,\mu}}. \end{aligned}$$

1384 Under the implemented parameters $k = \nu = 256$, the slice term is $2^{-256 - 8a - 8\mu}$.

1385 **s -way CMTDs-4 committing security.** For every $s \geq 2$, the all-bodies-equal special case
 1386 drops the known-key leaf tag collision term:

$$\text{Adv}_{TW128}^{\text{CMTDs-4}}(\mathcal{A}) \leq \frac{N_{\text{tot}}(N_{\text{tot}} + 1)}{2^{257}} + \frac{\binom{N_{\text{prim}} + s}{s}}{2^{256(s-1)}}.$$

1388 *Remark 1 (Two-Opening CMTDs-4 Specialization).* Setting $s = 2$ recovers the ordinary
 1389 two-opening CMTDs-4 statement of [BH22]:

$$\text{Adv}_{TW128}^{\text{CMTDs-4}}(\mathcal{A}) \leq \frac{N_{\text{tot}}(N_{\text{tot}} + 1)}{2^{257}} + \frac{\binom{N_{\text{prim}} + 2}{2}}{2^{256}} = \frac{N_{\text{tot}}(N_{\text{tot}} + 1) + (N_{\text{prim}} + 1)(N_{\text{prim}} + 2)}{2^{257}}.$$

Everywhere preimage resistance of H_{TW128} .

$$\text{Adv}_{TW128}^{\text{ePre}}(\mathcal{A}) \leq \frac{N_{\text{tot}}(N_{\text{tot}} + 1)}{2^{257}} + \frac{N_{\text{prim}} + 1}{2^{256}}.$$

1392 7.2 Concrete Security Claim

1393 **Corollary 7** (Concrete TW128 Security). *Under the implemented parameters $k = c =$
 1394 $\nu = \tau_{\text{root}} = \tau_{\text{leaf}} = 256$, every adversary with work cap $N_{\text{tot}} \leq 2^{63}$ —and, in the multi-user
 1395 case, user cap $D \leq 2^{63}$ —has advantage at most 2^{-128} in each of TW128’s notions: multi-
 1396 user indistinguishability (Theorem 2), all-in-one AE (Corollary 1), INT-RUP integrity
 1397 (Theorem 4), the hash security of H_{TW128} — s -way collision resistance for every $s \geq 2$ in
 1398 the sense of [BH22], with the [RS04] notions of standard and everywhere second-preimage
 1399 resistance inherited from collision resistance on every fixed input slice, standard preimage
 1400 resistance on every fixed input slice, and everywhere preimage resistance (Theorems 8 to 10
 1401 and Corollary 3)—and the [BH22] CMTDs-4 committing bound it yields (Corollary 5).*

1402 *Proof.* The resource consequences of Section 4.4 give $N_{\text{prim}} \leq N_{\text{tot}}$, $Q_v \leq N_{\text{tot}}$, $N_{\text{leaf}} \leq N_{\text{tot}}$,
 1403 and $N_{\text{prim}} + N_{\text{leaf}} \leq N_{\text{tot}}$. The user cap D remains independent, since created-but-idle
 1404 users need not contribute to N .

1405 At $N_{\text{tot}} \leq 2^{63}$, each flat-sponge term—at N_{tot} or at N_{prim} —and each birthday-scale
 1406 root or leaf term at $s = 2$ is at most $2^{-131} + 2^{-192}$, while each linear root-preimage,
 1407 key-guessing, verification, or leaf tag term with no factor D is at most a 2^{-193} -scale term.
 1408 The fixed-slice Pre degradation is at most 2^{-256} , since $m_{a,\mu} \geq k + \nu = 512$. Thus the
 1409 largest non-multi-user bound is the $s = 2$ collision bound, which is below $4 \cdot 2^{-131} = 2^{-129}$;
 1410 larger s only decreases the root-collision term. The all-in-one AE bound carries its two
 1411 flat-sponge terms as its only birthday-scale terms and lies below 2^{-129} , and the INT-RUP
 1412 bound carries $\text{Lift}_c(N_{\text{tot}})$ alone. In the multi-user case, using $D \leq 2^{63}$,

$$\text{Adv}_{TW128}^{\text{mu-ind}}(\mathcal{A}) < (2^{-130} + 2^{-193}) + 2^{-130} + 2^{-192} + 2^{-131} < 2^{-128} :$$

Table 7: Concrete TW128 security at the work cap $N_{\text{tot}} \leq 2^{63}$ (with $D \leq 2^{63}$ for mu-ind), under $k = c = \nu = \tau_{\text{root}} = \tau_{\text{leaf}} = 256$ and the resource consequences $N_{\text{prim}}, Q_v, N_{\text{leaf}} \leq N_{\text{tot}}$ and $N_{\text{prim}} + N_{\text{leaf}} \leq N_{\text{tot}}$ of Section 4.4. The Coll and CMTDs-4 rows hold for every $s \geq 2$; larger s only shrinks the root-collision term.

Notion	Dominant scale	At the cap
All-in-one AE	Lift_c	$\leq 2^{-128}$
Multi-user ind.	$D \text{KeyGuess}_k(N_{\text{prim}})$	$\leq 2^{-128}$
INT-RUP	Lift_c	$\leq 2^{-128}$
H_{TW128} collision (Coll)	$\text{Lift}_c, \text{LeafColl}, \text{RootColl}$	$\leq 2^{-128}$
H_{TW128} second-preimage (Sec/eSec)	inherited from Coll	$\leq 2^{-128}$
H_{TW128} preimage (Pre)	$\text{Lift}_c, \text{LeafColl}$	$\leq 2^{-128}$
H_{TW128} everywhere preimage (ePre)	Lift_c	$\leq 2^{-128}$
CMTDs-4 [BH22]	Lift_c and RootColl	$\leq 2^{-128}$

the two flat-sponge terms sum to at most $2 \cdot (2^{-131} + 2^{-194}) = 2^{-130} + 2^{-193}$, the key-guessing term satisfies $DN_{\text{prim}}/2^{256} \leq 2^{126}/2^{256} = 2^{-130}$, the remaining linear terms satisfy $(Q_v + N_{\text{leaf}})/2^{256} \leq 2^{64}/2^{256} = 2^{-192}$, and $D(D-1)/2^{257} < 2^{-131}$; the total is below $5 \cdot 2^{-131} + 2^{-192} + 2^{-193} < 2^{-128}$. These estimates instantiate every closed form in Section 7.1; Section 7.3 examines the margin they leave. $\square$

7.3 Discussion of the Bounds

The bounds of Section 7.1 aggregate the loss terms of Section 4.5. Table 7 records the dominant terms and aggregate cap bounds; the proof of Section 7.2 supplies the individual term sizes at the work cap.

At the concrete cap. At $N_{\text{tot}} \leq 2^{63}$, the hash bounds other than ePre and the CMTDs-4 bound have flat-sponge and root or leaf birthday terms on the same 2^{-131} scale; the ePre and INT-RUP bounds retain the flat sponge loss as their only birthday-scale term, and the single-user AE bound carries its two flat-sponge terms, at N_{tot} and N_{prim} , as its only birthday-scale terms. Thus the flat sponge loss is not always the largest summand: under the loose per-term cap, the $s = 2$ root-collision term $\binom{N_{\text{prim}}+2}{2}/2^{256}$ is on the same scale as $\text{Lift}_c(N_{\text{tot}})$. The multi-user bound is the exception: its hybrid-induced term $D \text{KeyGuess}_k(N_{\text{prim}})$ can reach 2^{-130} at the extreme cap $D = N_{\text{prim}} = 2^{63}$. That corner is a property of the bound, not of any admissible attack: the matching attack of the Tightness paragraph below needs one accepted nonempty encryption per targeted user, and even the empty-AD one-block case costs two duplex calls (root initialization and closing root message). With t targeted users and $q = N_{\text{prim}}$ direct guesses, the joint cap imposes $2t + q \leq 2^{63}$, so the feasible key-guessing attack scale is about 2^{-133} , reached near $t = 2^{61}$ and $q = 2^{62}$. The remaining factor is hybrid slack from charging created-but-idle users. Expressed in bytes, $N_{\text{tot}} \leq 2^{63}$ corresponds to at most $\rho \cdot 2^{63} = 167 \cdot 2^{63} < 2^{71}$ bytes of plaintext, associated data, ciphertext, and aggregation through the duplex (Section 3.2).

Margin at the cap. Corollary 7 clears its target with a stated remainder: at the joint corner $D = N_{\text{prim}} = 2^{63}$ the multi-user estimate has exact value $5 \cdot 2^{-131} + 2^{-194}$, since the corner forces $N = 0$ and hence $Q_v = N_{\text{leaf}} = 0$; this is just over five eighths of the 2^{-128} budget. The largest of the remaining bounds, the $s = 2$ collision bound, stays below $4 \cdot 2^{-131}$ (Section 7.2). The closeness at the corner is a proof artifact rather than a consequence of the design. The corner itself is unoccupiable: it forces $N = 0$, as the

previous paragraph records, and an adversary without construction computations receives identical answers in both worlds—the primitive interface is the same real permutation (Section 4.1)—so its advantage there is exactly 0. At resource splits an adversary can occupy, the multi-user key-guessing witness reaches about the 2^{-133} scale, while the root-collision and flat-sponge tightness witnesses can still reach the 2^{-131} scale. The intervening distance to the proof bound is consumed by the two flat-sponge composition terms and the hybrid slack recorded above, the overhead of the indistinguishability route. The margin also widens away from the corner: at $D \leq 2^{32}$, the D -indexed terms drop below 2^{-160} and the multi-user bound falls to its two flat-sponge terms, below $2^{-130} + 2^{-159}$, two bits under the target. A direct keyed-duplex analysis would be expected to replace the quadratic Lift_c terms by finer construction–primitive cross terms, much as the keyed-duplex bounds synthesized by [Men23] improve on the sponge-indistinguishability bound underlying the original SpongeWrap analysis of [BDPVA11]; Section 9.5 discusses that route and the gaps it would need to close.

Reading the work cap. Corollary 7 is a success-probability statement at a fixed cap, not a statement that security is exhausted there. The closed forms of Section 7.1 remain explicit functions of the resources beyond 2^{63} , and their dominant terms are birthday-type in N_{tot} on the 256-bit capacity, so the guarantees degrade quadratically rather than abruptly: at $N_{\text{tot}} = 2^{100}$, for example, the $s = 2$ collision bound is still below 2^{-55} . Once N_{tot} reaches $2^{c/2} = 2^{128}$, the flat sponge loss exceeds $1/2$ and no longer gives meaningful security (Section 4.5); the displayed closed form itself becomes a trivial advantage bound only when it reaches 1, near $N_{\text{tot}} = 2^{(c+1)/2}$. The cap $N_{\text{tot}} \leq 2^{63}$ is therefore a data-scale choice: a single work bound at which every notion’s advantage lies below the 2^{-128} target simultaneously, not the largest workload the mode tolerates.

Tightness.

- $\text{KeyGuess}_k(Q) = Q/2^k$. Tight at constants: a direct random-guess attack on a k -bit key matches the per-query success probability $1/2^k$. The multi-user bound’s dominant term $D \text{KeyGuess}_k(N_{\text{prim}}) = DN_{\text{prim}}/2^k$ is likewise tight at the leading constant in the theorem’s free-parameter setting, using the per-user nonce discipline of Section 2.4, which permits one nonce value across users. The adversary queries every targeted user on the same nonce U , empty associated data, and a known nonempty message—construction work charged to N , which does not enter the term—so each user’s first ciphertext block reveals the keystream emitted by that user’s root initialization call (Table 4). Each direct guess is then a single forward primitive query at the root initialization transcript $p_{\text{root}} \parallel K' \parallel U$ for a fresh candidate key K' , whose output is compared against all D collected keystream blocks at once; it matches some targeted key with probability $D/2^k$, so N_{prim} guesses recover a user key—and distinguish—with probability $\approx DN_{\text{prim}}/2^k$, matching the bound. Under the joint cap the construction work for those target encryptions lowers the feasible scale to about 2^{-133} , as noted in the concrete-cap paragraph above.
- $\text{RootColl}_\tau(Q, s) = \binom{Q+s}{s}/2^{\tau(s-1)}$. Tight up to lower-order terms: it is the standard s -way multi-collision bound on a τ -bit projection of RO_{flat} , and a generic birthday-style adversary against the random oracle matches the leading-order term. It is the root term of the collision bound (H_{TW128} Coll, Theorem 8). Its N_{prim} counts every primitive query, though by Output-Point Separation (Corollary 8) only those whose induced RO_{flat} call lands at a root tag output point ($\mathcal{R}_{\text{tag}}$, or $\mathcal{R}_{\text{msg}}^+$ for single-chunk ciphertexts) can enter the candidate sequence; a class-filtered count would be no larger, and potentially smaller, but the displayed bound keeps the N_{prim} form for

- uniformity with Corollary 1 and Theorem 2, the birthday threshold at $\tau_{\text{root}} = 256$ being unaffected.
- **RootPre $_{\tau}(Q) = (Q + 1)/2^{\tau}$.** Tight at constants: a generic preimage search against a fixed τ -bit target matches the per-query $2^{-\tau}$ rate. It is the root term of the standard and everywhere preimage bounds, where a single target tag yields a linear term.
 - **Leaf tag term $N_{\text{leaf}}/2^{\tau_{\text{leaf}}}$.** Tight at constants by a nonce-respecting attack: encrypt a single two-chunk message under nonce U , obtaining ciphertext with one leaf chunk and root tag T , then submit non-replay verification queries that each replace the leaf chunk by fresh bytes of the same length, keeping T . Each query recomputes the leaf tag at the hidden slot $(K, U, 1)$ and accepts whenever that fresh tag equals the encryption’s hidden leaf tag T_1 , since the aggregation string—and hence the recomputed root tag—is then unchanged. Each verification-side final leaf transcript therefore forges with probability $2^{-\tau_{\text{leaf}}}$, matching the bound’s per-transcript rate.
 - **Known-key leaf collision LeafColl $_{\tau}(Q)$.** Genuinely birthday-tight: when keys are adversary-supplied, direct queries can target leaf tag points, and a birthday search over chosen leaf bodies yields two distinct H_{TW128} inputs sharing a leaf tag and hence colliding openings; the collision proof uses $Q = N_{\text{prim}} + N_{\text{leaf}}$.
 - **Flat sponge lift.** The term $\text{Lift}_c(N_{\text{tot}}) = N_{\text{tot}}(N_{\text{tot}} + 1)/2^{c+1}$ dominates the random-permutation differentiating bound $f_P(N_{\text{tot}})$ of [BDPVA08, Lemma 5] via Lemma 20. The differentiating bound $f_P(N_{\text{tot}})$ is itself matched at leading order by the standard inner-collision differentiator, which succeeds against every simulator because the ideal-world construction interface is the random oracle regardless of simulator strategy; any residual conservatism in $\text{Lift}_c(N_{\text{tot}})$ lies in composing indistinguishability into the specific keyed games, not in an improvable simulator bound, and is not specific to TW128. The real-vs-ideal games pay the same differentiating bound once more, at N_{prim} , to exchange the ideal-side simulator for the real permutation (Lemma 21); since $N_{\text{prim}} \leq N_{\text{tot}}$, that term never dominates the N_{tot} term.
- Cross-scheme bound-level comparisons against prior sponge AEADs and the committing AEAD line of [BH22] are deferred to Sections 9.1 to 9.3.

8 Performance Evaluation

8.1 Optimized Implementations

We analyze a Go implementation of TW128¹ with optimized assembly paths for the performance-critical permutation and absorb/squeeze loops. We compare three vectorized assembly paths: `amd64 AVX-512`, `amd64 AVX2`, and `arm64 NEON` using the Armv8.2 SHA-3 extension. The framing, domain separation, padding, tree bookkeeping, and dispatch remain in shared Go code. A pure Go path uses a scalar permutation and a scalar eight-instance loop, and serves both as a correctness oracle for the assembly cores and as a constant-time fallback on other architectures.

The tree structure of TW128 exposes two distinct permutation workloads. The root duplex—which absorbs the associated data, processes chunk 0, and absorbs the leaf tag aggregation transcript—is inherently sequential, since each call consumes the previous call’s output. It is driven by a *single-duplex* ($\times 1$) core, except for its chunk-0 message phase. The leaf duplexes operate on disjoint state and disjoint inputs (Section 3), so complete leaf chunks are batched and run through an *eight-way* ($\times 8$) core. Leftover runs of two to seven complete chunks use narrower remainder kernels, while a single leftover

¹The implementation is available at <https://github.com/codahale/treewrap>.

complete chunk, a partial trailing chunk, and the count-aggregation tail fall back to the single-duplex core.

Chunk 0 itself runs on the batched kernels rather than serially, a technique we refer to as *chunk 0 fusion*. The root duplex’s chunk-0 message phase has the same kernel-visible schedule as a leaf chunk and differs only in its initial state. The implementation therefore seeds lane 0 of the first batched call with the root state, runs chunk 0 alongside the first leaf chunks, then writes lane 0’s output state back to the root duplex before the root absorbs the leaf tags from that same call. This avoids a serial pass over chunk 0 and is the reason the cycles-per-byte curve drops as soon as complete leaves are present (Section 8.3). Section C gives the lane-major layout, remainder-kernel strategies, lane 0 writeback details, and instruction-level kernel structure.

Architecture-specific cores. On amd64 the optimized implementation selects an AVX-512 or AVX2 core once at startup from the AVX512VL CPUID feature bit. On arm64 the optimized implementation uses NEON with the Armv8.2 SHA-3 extension (FEAT_SHA3, available on Apple Silicon and Neoverse V1/V2 and N2). Both architecture-specific paths provide a single-duplex core and batched leaf kernels; the assembly stays below the mode boundary, so framing, domain separation, padding, and tree bookkeeping remain shared with the reference path.

Constant time. The scalar reference and all three vectorized implementations are constant time with respect to all inputs. The permutation and the XOR-based absorb/squeeze contain no secret-dependent branches and no secret-dependent memory or vector indices. The only indexed table is the round-constant array, addressed by the public round number; batched loads, stores, gathers, scatters, and inactive-lane handling are all controlled by public chunk counts and fixed chunk strides. Section C records the architecture-specific addressing details.

8.2 Benchmark Methodology

We measure on two hosts, one per architecture. The amd64 host is a Google Compute Engine c4-standard-4 instance with an Intel Xeon Platinum 8581C (Emerald Rapids, 2.30 GHz base, AVX-512 including AVX512VL, AVX512_VBMI2, GFNI, and VAES), two physical cores with simultaneous multithreading enabled, 48 KiB L1d and 32 KiB L1i per core, 2 MiB L2 per core, and a 260 MiB shared L3, running Debian 12 (kernel 6.1) under KVM. The guest does not expose a cpufreq interface, so the frequency governor and turbo state are not under our control; we mitigate the resulting variance statistically (below). The arm64 host is an Apple M4 Pro (10 performance cores and 4 efficiency cores; per performance core, 128 KiB L1i, 64 KiB L1d, and a 16 MiB shared L2) running macOS 26.5.1 (Darwin 25.5), with FEAT_SHA3 present. macOS exposes no userspace frequency control.

Measurements are collected by a standalone harness (cmd/cpb) that locks its goroutine to an OS thread and, for each algorithm, operation, and message length, first calibrates an iteration count that fills at least 100 ms, then records 100 samples, from which it reports the median cycles per byte and median wall-clock throughput together with the interquartile range as a dispersion measure (Table 10). Both metrics are read from the same timed loop, so the cycles-per-byte and throughput figures for a given measurement come from one interleaved run rather than two separately scheduled ones. Each sample reuses the same source and destination buffers, so the reported figures are warm-cache, steady-state rates. The single-threaded measurement leaves the SMT sibling of the amd64 host idle. All reported figures are single-threaded: they exercise the SIMD axis of TW128’s parallelism, and multicore scaling remains unmeasured. We report a sweep of seven message lengths from 1 B to 16 MiB.

The cycle counters differ by architecture, and this affects how the cycles-per-byte figures may be read. On **amd64** the harness reads `RDTSC`; on this part the time-stamp counter is invariant and ticks at the 2.30 GHz reference frequency, so the reported figures are *reference* cycles per byte and are unaffected by turbo (and convert directly to wall-clock time). On **arm64** the harness reads `CNTVCT_ELO` and scales it by the ratio of the peak core frequency to `CNTFRQ_ELO`; since Apple exposes no frequency API, the peak is auto-detected by timing a dependent integer-add chain and taking the top observed P-state, yielding *core* cycles per byte. Consequently the cycles-per-byte figures are comparable within an architecture and against the same-host AES-128-GCM baseline, but not directly across architectures; absolute throughput (Table 9) is the cross-platform metric.

The baseline is Go’s AES-128-GCM implementation (`crypto/aes` with `crypto/cipher`). It uses AES-NI and `PCLMULQDQ` on **amd64**, and the AES and `PMULL` instructions on **arm64**. The AES block cipher, GCM wrapper, and TW128 AEAD object are constructed once before timing, so the measurements exclude key schedule and GHASH key-power precomputation and compare steady-state `Seal/Open` calls under a fixed key. Each timed TW128 call still performs the root duplex initialization inside the call. Both directions are measured for each scheme; the cycles-per-byte table reports both, and the throughput table reports the encryption direction.

8.3 Throughput

Table 8 reports cycles per byte across the message-length sweep and Table 9 the corresponding wall-clock throughput in the bulk regime. The dominant feature of the TW128 rows is the descent produced by the tree-parallel leaf core: the per-byte cost falls as a message lengthens and its chunks occupy more SIMD lanes. A message of at most one chunk has no complete leaf chunk for chunk 0 to share a kernel with, so it runs entirely on the single-duplex core, and the 8 KiB column sits at that core’s rate. Past one chunk, chunk 0 fusion (Section 8.1) keeps every complete chunk on the batched kernels, and the descent tracks lane occupancy: at 32 KiB, chunk 0 and three complete leaves run four-wide in the **amd64** remainder kernels and as two full pairs on **arm64**, and 64 KiB is the first length whose eight complete chunks ($8\beta \approx 64$ KiB) fill the eight-way ($\times 8$) core. On the **amd64** AVX-512 path TW128 encryption falls $1.81 \rightarrow 0.73 \rightarrow 0.39$ across the 8, 32, and 64 KiB columns and is already at its bulk rate at 64 KiB, holding at 0.38–0.39 through 16 MiB. The AVX2 path follows the same shape at lower throughput ($2.21 \rightarrow 1.15 \rightarrow 1.10$ over the same span, thereafter roughly 1.1 cycles per byte); it runs the eight instances as two groups of four rather than in a single 8-wide register, leaving the AVX-512 core about $2.9\times$ faster in bulk. The **arm64** curve flattens earliest ($1.97 \rightarrow 0.92 \rightarrow 0.89 \rightarrow 0.89 \rightarrow 0.91$ cycles per byte from 8 KiB to 16 MiB): the eight-way core itself processes its batch as four two-instance pairs, so the fused 2-wide kernel runs at essentially the same per-chunk rate, and a 32 KiB message—chunk 0 and three leaves as two full pairs—is already close to the long-message rate. In wall-clock terms TW128 reaches 6.00 GB/s on the **amd64** AVX-512 host and 4.93 GB/s on the M4 Pro for 16 MiB messages.

Dispersion. Each cell of Tables 8 and 9 is the median of 100 samples; Table 10 summarizes their spread as the interquartile range (IQR) relative to the median. The medians are stable: across the sweep the relative IQR of the TW128 cycles-per-byte measurements is at most 0.7% on the **amd64** AVX-512 path, 0.9% on AVX2, and 1.9% on the M4 Pro, with per-path medians of 0.3%, 0.4%, and 1.2%. The full min–max range is wider at isolated points—up to 4.9%, 12.0%, and 16.7% on the three paths—because turbo, governor, and scheduler states are not under our control (Section 8.2); these are tail excursions that the IQR excludes. We therefore report medians throughout and treat the IQR as the dispersion measure.

Table 8: Median cycles per byte (100 samples). The `amd64` figures are RDTSC reference cycles at the 2.30 GHz invariant time-stamp counter; the `arm64` figures are core cycles at the auto-detected peak performance-core frequency. Cycle counts are not directly comparable across architectures (Section 8.2).

	Message length						
	1 B	64 B	8 KiB	32 KiB	64 KiB	1 MiB	16 MiB
<i>Intel Xeon 8581C (Emerald Rapids), AVX-512</i>							
TW128 encrypt	571	9.04	1.81	0.73	0.39	0.38	0.38
TW128 decrypt	605	9.55	1.82	0.73	0.39	0.38	0.38
AES-128-GCM seal	119	2.74	0.30	0.29	0.29	0.29	0.29
AES-128-GCM open	99	2.17	0.29	0.28	0.28	0.28	0.29
<i>Intel Xeon 8581C (Emerald Rapids), AVX2</i>							
TW128 encrypt	647	10.18	2.21	1.15	1.10	1.09	1.10
TW128 decrypt	682	10.70	2.21	1.06	1.10	1.08	1.06
<i>Apple M4 Pro, NEON + FEAT_SHA3</i>							
TW128 encrypt	620	9.88	1.97	0.92	0.89	0.89	0.91
TW128 decrypt	671	10.65	2.04	0.90	0.87	0.87	0.89
AES-128-GCM seal	113	2.84	0.44	0.43	0.43	0.44	0.46
AES-128-GCM open	123	2.56	0.42	0.41	0.41	0.42	0.43

Table 9: Measured wall-clock throughput (GB/s, GB = 10^9 bytes) from the same `cmd/cpb` run as Table 8, encryption direction. The `amd64` AES-128-GCM row is the AVX-512 host and is unaffected by the TW128 core selection.

	8 KiB	32 KiB	64 KiB	1 MiB	16 MiB
<i>Intel Xeon 8581C (Emerald Rapids)</i>					
TW128 (AVX-512)	1.27	3.17	5.93	6.10	6.00
TW128 (AVX2)	1.04	2.00	2.10	2.10	2.08
AES-128-GCM	7.75	8.00	8.05	8.00	7.87
<i>Apple M4 Pro</i>					
TW128 (NEON+SHA3)	2.27	4.84	5.04	5.01	4.93
AES-128-GCM	10.04	10.27	10.33	10.07	9.71

Table 10: Cycles-per-byte measurement dispersion for TW128 across the seven-length sweep (seal and open, 100 samples per point), as the interquartile range relative to the median. The final column is the widest single-point min–max range; it is excluded from the IQR as a tail effect.

Path	Relative IQR		Max
	median	max	range
Xeon 8581C, AVX-512	0.3%	0.7%	4.9%
Xeon 8581C, AVX2	0.4%	0.9%	12.0%
Apple M4 Pro	1.2%	1.9%	16.7%

8.4 Latency

For short messages the cost is set by the serial root duplex and is independent of the parallel leaf machinery. A message of at most $\rho = 167$ bytes occupies a single rate block and never enters leaf mode, so it costs exactly two permutations—one to initialize the root duplex and one to close the message block and extract the tag. The 1 B and 64 B columns of Table 8 bear this out: per-message latency is nearly flat across them at roughly 571 reference cycles on **amd64** (571 at 1 B versus $9.04 \times 64 \approx 578$ at 64 B) and 620 core cycles on **arm64** (620 versus $9.88 \times 64 \approx 633$), i.e. about 285 and 310 cycles per KECCAK- $p[1600, 12]$ respectively. Because such messages bypass the leaf core entirely, leaf-level parallelism imposes no latency penalty on them.

The latency/throughput trade-off appears instead in the intermediate regime, which chunk 0 fusion has narrowed to roughly the one-to-eight-chunk band. The worst per-byte cost in the sweep is the single-chunk regime at 8 KiB (1.8–2.2 cycles per byte across the three paths), where the message is the chunk-0 phase alone on the single-duplex core. Once the message has complete leaf chunks for chunk 0 to share a kernel with, the cost falls quickly: by 32 KiB TW128 is within a factor of two of its AVX-512 floor and close to its **arm64** floor, and the floor itself is reached once the eight-way core fills at 64 KiB. The root duplex also bounds the tail of a long message: the final tag cannot be produced until chunk 0 has been processed and every leaf tag has been absorbed into the aggregation transcript. That absorption is cheap—each rate block of the root absorbs just over five 32-byte leaf tags, so a 16 MiB message (≈ 2000 leaves) adds on the order of 400 root permutations against roughly 10^5 permutations of leaf work, under 0.5% overhead—so the serial root does not materially cap bulk throughput.

8.5 Comparison with Other AEADs

We compare against the one baseline measured under the same harness on the same hosts, the Go standard library’s hardware-accelerated AES-128-GCM (Tables 8 and 9). The harness times TW128 and AES-128-GCM back-to-back at each length and reads both metrics from the same loop, so the cycles-per-byte and throughput comparisons hold under identical conditions and agree. In the bulk regime AES-128-GCM is faster on both architectures, as expected for a primitive with dedicated AES and carryless-multiply instructions: at 16 MiB it runs at 0.29 cycles per byte on **amd64** and 0.46 on **arm64**, against 0.38 and 0.91 for TW128—about $1.3\times$ and $2.0\times$ faster. The wall-clock throughputs show the same margins: TW128 reaches 6.00 and 4.93 GB/s against AES-128-GCM’s 7.87 and 9.71 GB/s. The wider **arm64** gap reflects the strength of Apple’s AES and PMULL units. TW128 attains these rates with a permutation that has no dedicated hardware support, relying only on the SHA-3 extension on **arm64** and on general SIMD on **amd64**.

AES-128-GCM is also faster for short messages in this steady-state measurement: at 1 B it costs 119 and 113 cycles on the two hosts against TW128’s 571 and 620. This is the expected result when AES-GCM key setup is amortized outside the timed call. In the sub-bulk regime AES-128-GCM is several times faster at one chunk (0.30 versus 1.81 cycles per byte at 8 KiB on **amd64**), but the gap narrows once the fused kernels engage: about $2.5\times$ at 32 KiB and the bulk ratio from 64 KiB on. TW128’s intended value relative to AES-128-GCM lies in the tree structure and committing security properties analyzed in Sections 5 and 6 rather than in raw speed.

Table 11: Structural comparison of TW128 with representative AEADs. Sizes are in bytes; AES-GCM-SIV also admits a 32-byte key, and Kravatte-SAE’s key and nonce lengths are parameters (keys of up to 40 bytes, arbitrary-length nonces). “Misuse” is nonce-misuse resistance in the sense of [RS06]; CMT records a published committing result or attack in the relevant committing notion, with “broken” denoting a published low-cost committing attack, “hash CR” denoting the collision-resistance level of the hash used by the transform, and “open” marking a scheme with no published committing analysis. A dash means the entry depends on the base scheme. These entries are notion-specific summaries, not broad AE security claims. TW128’s concrete bounds appear in Section 7.3.

Scheme	Primitive	Par.	K/U/T	AD	Misuse	CMT
TW128	KECCAK- p [1600, 12]	yes	32/32/32	yes	no	128 b
Ascon-128a [DEMS21]	Ascon- p (320 b)	no	16/16/16	yes	no	$\approx$ 64 b
Xoodyak [DHP ⁺ 20]	Xoodoo (384 b)	no	16/16/16	yes	no	2^{17} atk.
Kravatte-SAE [BDH ⁺ 17]	KECCAK- p [1600, 6]	yes	$\leq 40/\text{var.}/16$	yes	no	open
AES-GCM [Dwo07]	AES	yes	16/12/16	yes	no	broken
AES-GCM-SIV [GLL17]	AES	yes	16/12/16	yes	yes	broken
nAE + CTX [CR22]	any nAE	—	—	yes	—	hash CR

9 Discussion

9.1 Comparison Table

Table 11 situates TW128 against representative authenticated encryption schemes along the axes that distinguish it: the underlying primitive, whether the mode is parallelizable, the key, nonce, and tag sizes, associated data support, nonce-misuse resistance, and committing security (CMT). The rows are representative rather than exhaustive: sponge/duplex relatives, a Farfalle-based parallel design, AES baselines, a misuse-resistant AES baseline, and a generic committing transform. The comparison is structural, not a total security ranking; TW128’s own concrete bounds and the regimes in which each term dominates are treated in Section 7.3, and its measured throughput against hardware-accelerated AES-128-GCM in Section 8.5. The CMT column should be read within the named committing notions only, and a negative entry there does not assert a confidentiality or integrity break of the underlying AEAD. It records the published committing result or attack most relevant to the comparison: TW128’s 128-bit CMTDs-4 bound, Ascon’s roughly 64-bit committing bound [KSW23], published low-cost attacks where known [KSW23, CR22, ADG⁺22], and the hash-collision level inherited by CTX. The table does not separately track compact (tag-only) commitment, since that property has not been systematically analyzed for the comparison schemes.

9.2 Comparison with Prior Sponge AEADs

TW128 belongs to the family of single-permutation duplex AEADs initiated by [BDPVA11]. The dedicated lightweight members of that family, Ascon [DEMS21] and Xoodyak [DHP⁺20], target 128-bit security on small permutations (320 and 384 bits) with a strictly serial duplex: each call depends on the previous one, so there is no mode-level parallelism to exploit. TW128 instead pays for a larger permutation, KECCAK- p [1600, 12], and a 256-bit capacity—the same round count as KangarooTwelve [BDP⁺16]—in exchange for a tree that processes leaf chunks in parallel and for the headroom that yields 128-bit security with a comfortable margin. The cost of this choice is visible in Section 8: on short, single-leaf inputs TW128 is slower per byte than a compact serial duplex would be, and its advantage appears only once a message spans enough leaves to fill the vector units.

TW128 reuses KangarooTwelve’s final-node/leaf split but not its Sakura coding [BDPVA14]. In KangarooTwelve, Sakura makes the final node parseable as a message hop

followed by a chaining hop: the first chunk, the sequence of leaf chaining values, the coded number of chaining values, and the final-node frame bits all have explicit roles [BDP⁺16]. In TW128, the analogous chaining values are the leaf tags, but they are not anonymous strings waiting to be interpreted by the root. Each leaf duplex is initialized with a chunk-specific state containing the leaf prefix and leaf index, while the root absorbs the leaf tags in index order and then absorbs the leaf count $\langle n \rangle_8$ (Sections 3 and 3.4). Sakura’s structural job is therefore handled by mode initialization and domain separation rather than by a tree coding. The transcript separation lemmas make this explicit (Lemma 6 and Corollary 8). The committing bounds do not depend on tree-coding soundness: distinct (K, U, A) triples already diverge before leaf aggregation (Proposition 1), and the leaf encoding enters the hash collision bound only through the explicit leaf tag collision term of Theorem 8.

Keyak [BDP⁺15], whose Motorist mode is already parallelizable, is the closest prior duplex AEAD in spirit; Section 1.2 contrasts its wire-visible, encryptor-and-decryptor-shared degree of parallelism with TW128’s run-time leaf choice. Strobe [Ham17] is a serial sponge framework aimed at whole protocols rather than at a standalone AEAD, and like Ascon and Xoodyak it does not parallelize.

Farfalle [BDH⁺17] leaves the duplex setting altogether: it is a parallel pseudorandom function whose compression layer sums masked-and-permuted input blocks into an accumulator and whose expansion layer generates output from a rolling state, with authenticated encryption supplied by the Farfalle-SAE and Farfalle-SIV modes on top (Section 1.2). Its Kravatte instance runs KECCAK- $p[1600, 6]$, and the parallelism available to an implementation is arbitrary, as in TW128. The two designs price their permutations for different adversaries. Farfalle takes its aggressive round counts from the premise that an adversary never sees both the input and the output of a permutation call, and the accumulator-collision analysis behind Kravatte’s six-round compression layer randomizes the inputs by secret masks [BDH⁺17]. TW128’s root tag claims are known-key claims: the hash and committing games hand the adversary the key, so the mode needs the hash-grade round count and margin of Section 9.4 and a proof toolkit that covers adversarial initialization—the sponge indistinguishability line of Section B. TW128 therefore keeps the strict duplex object and takes its parallelism from the tree topology, paying more permutation rounds per byte than a Farfalle instance in exchange for a tag that is a proved digest.

Committing security is not unique to TW128 in this family: a dedicated analysis of the sponge-based NIST lightweight finalists proves Ascon committing while breaking others [KSW23], and some constructions, such as Strobe and the Farfalle modes, have not been examined in that setting at all. What distinguishes TW128 is that it makes the root tag a secure compact commitment to the full input—collision-, second-preimage-, and preimage-resistant as a hash, which in particular yields the 128-bit full-input CMTDs-4 bound—as a proven, designed property of the mode without leaving the single-permutation duplex setting.

9.3 Comparison with Committing AEAD Constructions

The dominant route to committing AEAD is a generic transform over a non-committing base scheme. [BH22] give the UtC, RtC, and HtE transforms, which add key- or context-commitment to an arbitrary AEAD, and [CR22] give CTX, which makes any nonce-based AE scheme commit to its full context at the cost of a single hash over a short string and with no ciphertext expansion. These transforms are inexpensive and broadly applicable; CTX in particular is nearly free. It is nevertheless a committing-AE transform, not a compact-commitment transform: it binds the transformed ciphertext, but does not make the tag a standalone digest-like commitment to the full input. TW128 differs not by undercutting their cost but by integration: it is committing as the mode itself,

and its root tag is the compact committing object. The commitment uses the same permutation already used for confidentiality and integrity and the same tag already used for authentication, so no separate commitment primitive, key-derivation step, or auxiliary hash is invoked, and the multi-input full-commitment bound follows from a single root tag collision bound (Section 6). More strongly, Theorem 8 shows that the root tag alone is a compact commitment string even when the ciphertext bodies used as openings differ, up to the additional leaf tag birthday term. This places TW128 alongside the dedicated committing primitives—the compactly committing AEAD of [GLR17] and the encryption of [DGRW18]—which likewise build commitment in rather than bolt it on, but which were designed for message franking rather than as general nonce-based AEADs. The contrast with transformed conventional schemes is sharpened by the attacks of Section 1.2: base schemes such as GCM and OCB are not committing on their own [CR22], and a transform is the only way to make them so.

Closest to TW128 here is the recent work of [DHM⁺24], which builds nonce-based AE directly on the SHA-3 functions SHAKE and TurboSHAKE—the latter using the same round-reduced KECCAK- p permutation as TW128—and likewise achieves CMT-4 natively, from the collision resistance of the underlying function rather than from a transform. The two designs make different structural choices: [DHM⁺24] targets session-based communication, chaining a stream of messages through intermediate tags and generalizing the keyed duplex into a *duplex cipher*, whereas TW128 is a single two-level tree whose independent leaves expose run-time-selectable parallelism at constant memory (Sections 1.2 and 8.1). Because a session chain processes each message serially, TW128’s principal advantage is throughput on large messages (Section 8): its parallel leaves fill vector units a serial duplex leaves idle, at the cost of the session support [DHM⁺24] provides and TW128 does not. Both demonstrate that committing AE needs no add-on transform once it is built on a well-scrutinized KECCAK- p permutation.

9.4 Cryptanalysis of the Underlying Permutation

Every bound in this paper models KECCAK- p [1600, 12] as a uniformly random permutation (Section 4). That modeling choice is supported by cryptanalytic scrutiny rather than by proof, so this subsection summarizes the third-party record and the safety margin it indicates. The Keccak round function has been unchanged since 2008, and TW128 uses the permutation at the round count introduced by KangarooTwelve [BDP⁺16] and later exposed directly as TurboSHAKE [BDH⁺23], whose designers maintain a survey of third-party cryptanalysis. TurboSHAKE128 shares TW128’s rate and capacity ($r = 1344$, $c = 256$) as well as its permutation, so attacks on it and on round-reduced SHAKE128 transfer to TW128’s geometry directly.

In the unkeyed setting, the deepest collision attacks on any SHA-3 or SHAKE instance reach 6 of TW128’s 12 rounds: a SAT-aided attack on 6-round SHAKE128 with time $2^{123.5}$ [GLST22], and internal-differential attacks reaching 6 rounds of SHAKE256 and 5 rounds of SHA3-384 [ZHL24]. Preimage attacks reach 5 rounds [ZHL25]. Against the collision- and preimage-style attacks most relevant to the root tag hash claims of Section 6, the permutation therefore retains a margin of 6 of 12 rounds—the assessment under which the KangarooTwelve round count was chosen and later reaffirmed [BDP⁺16, BDH⁺23]. The structural distinguisher reaching the most rounds of the permutation itself, the SymSum property on 9 rounds reported in the survey of [BDH⁺23], has, like the zero-sum distinguishers before it, not been extended to an attack on any sponge or duplex mode.

The keyed setting is closer to TW128’s AE games: a hidden key, with adversary-controlled nonces and messages entering through the duplex rate. The attacks reaching the most rounds here are cube and conditional cube attacks. Dinur et al. break up to 9 rounds of Keccak-based stream ciphers and MACs faster than generic search, with keystream prediction on 8 rounds at time and data 2^{128} and on 9 rounds (of a 512-bit-key variant) at

2^{256} [DMP⁺15]. With MILP-optimized conditional cubes, nonce-respecting key recovery reaches 8 of Lake Keyak’s 12 rounds at time 2^{71} for 128-bit keys, 9 rounds at 2^{137} for 256-bit keys, and 9 of KMAC256’s 24 rounds at 2^{147} [SGSL18]. Lake Keyak is a keyed duplex over the same 1600-bit permutation whose initialization, like TW128’s root and leaf initializations, absorbs the key and nonce ahead of nonce-varied cube queries, so these results are the closest published analog to TW128’s keyed setting. They leave a margin of 3 to 4 of 12 rounds against attacks whose complexities start at 2^{71} .

In sum, the known attack frontier stands at 6 of 12 rounds for collisions and preimages and at 8 to 9 rounds for keyed-mode key recovery and structural distinguishers. TW128 uses KECCAK- p [1600, 12] at the same round count and with the same rate/capacity split as KangarooTwelve and TurboSHAKE128, so its margin assessment tracks theirs, and any future advance against the permutation applies to all three alike. What TW128 adds at the mode level and these analyses do not cover is the tree structure itself—in particular the family of leaf initializations sharing (K, U) and differing only in the leaf index—and we consider dedicated cryptanalysis of that structure a natural target for future work.

9.5 Limitations and Open Problems

Section 4 states the formal scope of the security claims; this subsection discusses the consequences of the excluded settings and the corresponding open problems.

Nonce misuse. As scoped in Section 4, TW128’s AE and mu-ind bounds require nonce-respecting encryption queries, and the mode is not misuse-resistant: a repeated nonce reuses leaf keystream and breaks confidentiality, as for any online duplex stream cipher. Schemes such as AES-GCM-SIV [GLL17], in the deterministic/misuse-resistant lineage of [RS06], tolerate nonce repetition at the price of a second pass. A misuse-resistant variant of TW128 is left open.

Unverified plaintext release. The INT-RUP theorem of Theorem 4 covers integrity when an adversary can obtain plaintext-computing oracle outputs before verification, in the release-of-unverified-plaintext setting of [ABL⁺14]. This is an integrity-only statement. It does not give plaintext awareness or privacy under release of unverified plaintext: TW128 is online and single-pass, and exposing unverified stream plaintext can reveal keystream information. Application behavior on unverified plaintext is likewise outside the provable model.

Conservative lift. Per-term tightness is itemized in Section 7.3; the flat-sponge terms are conservative because our proof follows the unkeyed sponge/hash line. The flat sponge differentiating loss $\text{Lift}_c(N_{\text{tot}}) = N_{\text{tot}}(N_{\text{tot}} + 1)/2^{c+1}$ bounds the indistinguishability loss of [BDPVA08], used here to lift the random oracle analysis to the public permutation model; the real-vs-ideal games pay it twice, at N_{tot} on the real side and at N_{prim} on the ideal side (Corollary 9). This proof route matches the role of the root tag: TW128’s hash and committing claims are consequences of hash security, and sponge hash security is supported by a mature literature.

A direct keyed-duplex treatment would have to close two gaps. First, existing keyed-duplex bounds are hidden-key AE bounds: the game samples the key, and the keyed initialization material is not part of the adversary’s output. The root tag claims here are adversarial-initialization statements. In the hash and committing games, the adversary may choose the key, nonce, associated data, message, or competing opening. Thus the proof must handle adversarially chosen initialization transcripts, not only hidden-key encryption and verification queries.

Second, the available keyed-duplex AE interface is not a direct substitute for a byte-string AEAD with arbitrary final-block length. The modern bounds synthesized by [Men23] use a duplex call phased as permutation, squeeze, then absorb or overwrite ([Men23, Sec. 3.4, Table 1]). In the authenticated-encryption application, this interface is instantiated as `MonkeySpongeWrap` [Men23, Alg. 8]. As written, however, the overwrite interface does not reconstruct the missing padding bits of a truncated final ciphertext block on decryption: encryption absorbs the padded final plaintext block after squeezing and then truncates the resulting ciphertext to the message length, while decryption receives only this truncated ciphertext, pads it independently, and uses the overwrite branch to set the rate. The decryption transcript can therefore diverge before tag recomputation.

TW128 instead stays with the BDPV11 duplex phasing: a call absorbs the actual input string, padded by the duplex rule, and then returns the requested output. At the stream layer this supports arbitrary length messages by splitting them into bounded duplex blocks, with the final short ciphertext block absorbed as it is rather than reconstructed through an overwrite branch (Section 3). A direct keyed-duplex theorem for TW128 would therefore need both this BDPV11 phasing and adversarial-initialization hash and committing games. We retain the indistinguishability lift because it covers these claims uniformly and already meets the 128-bit target (Section 7.3).

Beyond-birthday security and leakage. All bounds are birthday-type on the 256-bit capacity, delivering the 128-bit target but no more; a beyond-birthday analysis is open. The model is also single-stage and leakage-free: side-channel resistance is addressed only at the implementation level, through the constant-time properties of Section 8.1, and is not part of the provable-security claims.

10 Conclusion

TW128 starts from the premise that an AEAD tag should behave like the compact digest applications often take it to be, and builds the mode from a digest-shaped construction rather than adding commitment as an external layer. It adapts the KangarooTwelve tree-hashing pattern into a two-level tree of keyed duplexes over a single $\text{KECCAK-}p[1600, 12]$ permutation. The chunk decomposition is canonical, while the number of leaf duplexes evaluated in parallel is an implementation choice rather than a mode parameter. This gives scalar, SIMD, and multicore implementations the same ciphertexts and tags, with constant working memory and scalable parallelism.

The central security claim is that the resulting root tag satisfies the digest intuition: it is a compact commitment to the full (K, U, A, M) context. In the public permutation model, via an intermediate random oracle and the flat sponge lift, we prove concrete multi-user indistinguishability and derive all-in-one AE as its single-user corollary, and prove INT-RUP integrity under release of unverified plaintext. We also prove collision, second-preimage, and preimage security for the root tag wrapper. Those hash-security bounds yield the AEAD-facing committing results, including full-input CMTDs-4 security, at a 128-bit target for the implemented parameters.

The implementation results show why the tree-hashing origin matters. Vectorized `amd64` and `arm64` implementations exploit the scalable leaf structure, reaching bulk throughputs of 48.0 Gbit/s on the AVX-512 path, 39.4 Gbit/s on the `arm64` path, and 16.6 Gbit/s on the AVX2 path, with the AVX-512 and `arm64` paths at 51–76% of hardware-accelerated AES-128-GCM throughput.

Several questions remain open (Section 9.5): a nonce-misuse-resistant variant, privacy under release of unverified plaintext, and sharper bounds, whether through a keyed duplex treatment of the BDPV11-style phasing used here or through a beyond-birthday argument. More broadly, TW128 adds to the evidence that a single, well-scrutinized $\text{KECCAK-}p$

permutation can support authenticated encryption whose tag is digest-like without giving up scalable hashing-style performance.

Acknowledgments

A hearty thank-you to my wife, who remains reasonably mystified at my choice of hobbies.

References

- [ABL⁺14] Elena Andreeva, Andrey Bogdanov, Atul Luykx, Bart Mennink, Nicky Mouha, and Kan Yasuda. How to securely release unverified plaintext in authenticated encryption. In *Advances in Cryptology – ASIACRYPT 2014*, volume 8873 of *Lecture Notes in Computer Science*, pages 105–125. Springer, 2014.
- [ADG⁺22] Ange Albertini, Thai Duong, Shay Gueron, Stefan Kölbl, Atul Luykx, and Sophie Schmieg. How to abuse and fix authenticated encryption without key commitment. In *31st USENIX Security Symposium (USENIX Security 22)*, pages 3291–3308. USENIX Association, 2022.
- [BDH⁺17] Guido Bertoni, Joan Daemen, Seth Hoeffert, Michaël Peeters, Gilles Van Assche, and Ronny Van Keer. Farfalle: Parallel permutation-based cryptography. *IACR Transactions on Symmetric Cryptology*, 2017(4):1–38, 2017.
- [BDH⁺23] Guido Bertoni, Joan Daemen, Seth Hoeffert, Michaël Peeters, Gilles Van Assche, Ronny Van Keer, and Benoît Viguier. TurboSHAKE. Cryptology ePrint Archive, Paper 2023/342, 2023.
- [BDP⁺15] Guido Bertoni, Joan Daemen, Michaël Peeters, Gilles Van Assche, and Ronny Van Keer. CAESAR submission: Keyak v2. CAESAR Competition, Round 3 Candidate, 2015.
- [BDP⁺16] Guido Bertoni, Joan Daemen, Michaël Peeters, Gilles Van Assche, Ronny Van Keer, and Benoît Viguier. KangarooTwelve: Fast hashing based on Keccak-p. Cryptology ePrint Archive, Paper 2016/770, 2016.
- [BDPVA08] Guido Bertoni, Joan Daemen, Michaël Peeters, and Gilles Van Assche. On the indistinguishability of the sponge construction. In *Advances in Cryptology – EUROCRYPT 2008*, volume 4965 of *Lecture Notes in Computer Science*, pages 181–197. Springer, 2008.
- [BDPVA11] Guido Bertoni, Joan Daemen, Michaël Peeters, and Gilles Van Assche. Duplexing the sponge: Single-pass authenticated encryption and other applications. In *Selected Areas in Cryptography – SAC 2011*, volume 7118 of *Lecture Notes in Computer Science*, pages 320–337. Springer, 2011.
- [BDPVA14] Guido Bertoni, Joan Daemen, Michaël Peeters, and Gilles Van Assche. Sakura: A flexible coding for tree hashing. In *Applied Cryptography and Network Security – ACNS 2014*, volume 8479 of *Lecture Notes in Computer Science*, pages 217–234. Springer, 2014.
- [BH22] Mihir Bellare and Viet Tung Hoang. Efficient schemes for committing authenticated encryption. In *Advances in Cryptology – EUROCRYPT 2022*, volume 13276 of *Lecture Notes in Computer Science*, pages 845–875. Springer, 2022. Citations refer to the full version of July 23, 2024.

- [BT16] Mihir Bellare and Björn Tackmann. The multi-user security of authenticated encryption: AES-GCM in TLS 1.3. In *Advances in Cryptology – CRYPTO 2016*, volume 9814 of *Lecture Notes in Computer Science*, pages 247–276. Springer, 2016.
- [CR22] John Chan and Phillip Rogaway. On committing authenticated encryption. In *Computer Security – ESORICS 2022*, *Lecture Notes in Computer Science*. Springer, 2022.
- [DEMS21] Christoph Dobraunig, Maria Eichseder, Florian Mendel, and Martin Schläffer. Ascon v1.2: Lightweight authenticated encryption and hashing. *Journal of Cryptology*, 34(3):33, 2021.
- [DGRW18] Yevgeniy Dodis, Paul Grubbs, Thomas Ristenpart, and Joanne Woodage. Fast message franking: From invisible salamanders to encryptment. In *Advances in Cryptology – CRYPTO 2018*, volume 10991 of *Lecture Notes in Computer Science*, pages 155–186. Springer, 2018.
- [DHM⁺24] Joan Daemen, Seth Hoffert, Silvia Mella, Gilles Van Assche, and Ronny Van Keer. Shaking up authenticated encryption. *Cryptology ePrint Archive*, Paper 2024/1618, 2024.
- [DHP⁺20] Joan Daemen, Seth Hoffert, Michaël Peeters, Gilles Van Assche, and Ronny Van Keer. Xoodoo, a lightweight cryptographic scheme. *IACR Transactions on Symmetric Cryptology*, 2020(S1):60–87, 2020.
- [DMP⁺15] Itai Dinur, Paweł Morawiecki, Josef Pieprzyk, Marian Srebrny, and Michał Straus. Cube attacks and cube-attack-like cryptanalysis on the round-reduced Keccak sponge function. In *Advances in Cryptology – EUROCRYPT 2015*, volume 9056 of *Lecture Notes in Computer Science*, pages 733–761. Springer, 2015.
- [Dwo07] Morris J. Dworkin. Recommendation for block cipher modes of operation: Galois/Counter Mode (GCM) and GMAC. Technical Report Special Publication 800-38D, National Institute of Standards and Technology, 2007.
- [GLL17] Shay Gueron, Adam Langley, and Yehuda Lindell. AES-GCM-SIV: Specification and analysis. *Cryptology ePrint Archive*, Paper 2017/168, 2017.
- [GLR17] Paul Grubbs, Jiahui Lu, and Thomas Ristenpart. Message franking via committing authenticated encryption. In *Advances in Cryptology – CRYPTO 2017*, volume 10403 of *Lecture Notes in Computer Science*, pages 66–97. Springer, 2017.
- [GLST22] Jian Guo, Guozhen Liu, Ling Song, and Yi Tu. Exploring SAT for cryptanalysis: (quantum) collision attacks against 6-round SHA-3. In *Advances in Cryptology – ASIACRYPT 2022*, volume 13793 of *Lecture Notes in Computer Science*, pages 645–674. Springer, 2022.
- [Ham17] Mike Hamburg. The STROBE protocol framework. *Cryptology ePrint Archive*, Paper 2017/003, 2017.
- [KSW23] Juliane Krämer, Patrick Struck, and Maximiliane Weishäupl. Committing AE from sponges: Security analysis of the NIST LWC finalists. *Cryptology ePrint Archive*, Paper 2023/1525, 2023.

- [LGR21] Julia Len, Paul Grubbs, and Thomas Ristenpart. Partitioning oracle attacks. In *30th USENIX Security Symposium (USENIX Security 21)*, pages 195–212. USENIX Association, 2021.
- [Men23] Bart Mennink. Understanding the duplex and its security. *IACR Transactions on Symmetric Cryptology*, 2023(2):1–46, 2023.
- [MLGR23] Sanketh Menda, Julia Len, Paul Grubbs, and Thomas Ristenpart. Context discovery and commitment attacks: How to break CCM, EAX, SIV, and more. In *Advances in Cryptology – EUROCRYPT 2023*, volume 14007 of *Lecture Notes in Computer Science*, pages 379–407. Springer, 2023.
- [Nat15] National Institute of Standards and Technology. SHA-3 standard: Permutation-based hash and extendable-output functions. Federal Information Processing Standards Publication FIPS PUB 202, U.S. Department of Commerce, August 2015.
- [RS04] Phillip Rogaway and Thomas Shrimpton. Cryptographic hash-function basics: Definitions, implications, and separations for preimage resistance, second-preimage resistance, and collision resistance. In *Fast Software Encryption – FSE 2004*, volume 3017 of *Lecture Notes in Computer Science*, pages 371–388. Springer, 2004. Citations refer to the full version of July 16, 2009.
- [RS06] Phillip Rogaway and Thomas Shrimpton. A provable-security treatment of the key-wrap problem. In *Advances in Cryptology – EUROCRYPT 2006*, volume 4004 of *Lecture Notes in Computer Science*, pages 373–390. Springer, 2006.
- [SGSL18] Ling Song, Jian Guo, Danping Shi, and San Ling. New MILP modeling: Improved conditional cube attacks on Keccak-based constructions. In *Advances in Cryptology – ASIACRYPT 2018*, volume 11273 of *Lecture Notes in Computer Science*, pages 65–95. Springer, 2018.
- [ZHL24] Zhongyi Zhang, Chengan Hou, and Meicheng Liu. Probabilistic linearization: Internal differential collisions in up to 6 rounds of SHA-3. In *Advances in Cryptology – CRYPTO 2024*, volume 14923 of *Lecture Notes in Computer Science*, pages 241–272. Springer, 2024.
- [ZHL25] Zhongyi Zhang, Chengan Hou, and Meicheng Liu. Preimage attacks on up to 5 rounds of SHA-3 using internal differentials. In *Advances in Cryptology – EUROCRYPT 2025*, volume 15601 of *Lecture Notes in Computer Science*, pages 333–363. Springer, 2025.

A Supporting Lemmas

The proof core separates structural decoding from probabilistic accounting. The well-formed transcript language of Definition 2 fixes which flattened queries are valid TW128 construction prefixes. Lemmas 5, 6 and 8 then show how such a query is flattened, parsed as a typed transcript, and assigned an output point; the exact keyed decoder of Definition 4 then decodes direct queries to candidate keys, when any. Lemma 10 turns those decoded keys into an adaptive key-test bound. The remaining proof interfaces expose the events their consumers must control: encryption-side freshness, verification hidden-key exposure and root tag guessing, leaf tag collisions, and the committing candidate-sequence collision event. Thus the mu-ind and committing proofs only have to derive the relevant local event premise before applying the corresponding accounting argument.

2036 **How to read this appendix.** Figure 2 records the proof dependency map. It groups
 2037 local interfaces when the exact proof citations would obscure the spine. The Bridge &
 2038 Decode node is Lemma 8 together with the exact keyed decoder of Definition 4. The
 2039 Root/Output Separation node groups Output-Point Separation (Corollary 8), Opening
 2040 Triple Separation (Proposition 1), and Final-Root Input Separation (Proposition 2); the
 2041 use of Leaf-Transcript Recovery (Lemma 7) is folded into the final-root separation fact. The
 2042 Key-Indexed Renaming node (Lemma 9) supplies the key-field disjointness and renaming
 2043 used by hidden-key encryption and verification accounting. The Adaptive Candidate node
 2044 (Lemma 2) is the generic probabilistic engine: it bounds the Leaf Tag Collision lemma
 2045 (Lemma 11) and, through the Root Tag Candidate Sequence and Root Tag Hash and
 2046 Target-Hit Bound (Lemmas 3 and 4) of Section 6, the root tag hash and committing proofs.
 2047 The mu-ind proof reuses the same structural lemmas through the explicit hybrid game
 2048 sequence of Definition 5, adding freshness, key-test, leaf-collision, and verification-first
 2049 guessing accounting.




Figure 2: Proof dependency map for the supporting components in this appendix and their consumers. Solid boxes are lemmas or grouped local proof interfaces; dashed boxes are proof components that consume them.

2050 A.1 Duplex Flattening

2051 **Lemma 5** (Duplex Flattening). *Every TW128 root or leaf duplex output is equal to a*
 2052 *[BDPVA11] sponge output on the corresponding flattened duplex-call transcript, and the*
 2053 *ciphertext-absorbing message phase admits a flattened duplex transcript that is well defined*
 2054 *for both encryption and decryption.*

2055 *Proof.* Let D be a duplex object using permutation π , rate r , and padding $\text{pad}10^*1$. For a
 2056 sequence of duplex calls $Z_i = D.\text{duplexing}(\sigma_i, \ell_i)$ with each σ_i short enough to fit in one
 2057 duplex block after padding, [BDPVA11, Lemma 3] gives

$$2058 \quad Z_i = \text{Sponge}[\pi, \text{pad}10^*1, r](\sigma_0 \parallel \text{pad}_0 \parallel \sigma_1 \parallel \text{pad}_1 \parallel \cdots \parallel \sigma_i, \ell_i),$$

2059 where $\text{pad}_j = \text{pad}10^*1[r](|\sigma_j|)$, and Lemma 4 says that, for fixed $\text{pad}10^*1$ and r , the map
 2060 from the duplex input-string sequence to the flattened sponge input is injective.

2061 In TW128, each duplex call absorbs a byte-aligned block σ followed by a 4-bit phase
 2062 suffix $\sigma_{\text{ph}} \in \{\sigma_{\text{init}}, \sigma_{\text{ad}}^-, \dots, \sigma_{\text{agg}}^+\}$ (Section 3.4), so the duplex-call input is $\sigma \parallel \sigma_{\text{ph}}$ followed
 2063 by $\text{pad}10^*1$. With $r = 1344$ and $\rho_{\text{max}}(\text{pad}10^*1, r) = r - 2 = 1342$ bits, the maximum
 2064 byte-aligned input is $\rho = 167$ bytes (Section 3.2), giving a worst-case duplex-call input of

2065 $8 \cdot 167 + 4 = 1340 \leq 1342$ bits. The two initialization calls are shorter still: $|p_{\text{root}} \parallel K \parallel U| + 4 =$
 2066 $8 \cdot (6 + 32 + 32) + 4 = 564$ bits and $|p_{\text{leaf}} \parallel K \parallel U \parallel \langle j \rangle_8| + 4 = 8 \cdot (6 + 32 + 32 + 8) + 4 = 628$
 2067 bits. [BDPVA11, Lemmas 3 and 4] therefore apply independently to every TW128 duplex
 2068 object.

2069 The root duplex is initialized by absorbing $p_{\text{root}} \parallel K \parallel U$ under σ_{init} (Algorithm 2), then
 2070 absorbs the associated data blocks under $\sigma_{\text{ad}}^- / \sigma_{\text{ad}}^+$, the root ciphertext blocks in the message
 2071 phase under $\sigma_{\text{msg}}^- / \sigma_{\text{msg}}^+$, and the leaf tag aggregation blocks under $\sigma_{\text{agg}}^- / \sigma_{\text{agg}}^+$. By the
 2072 duplex-to-sponge equality, every root duplex output is the sponge function evaluated on
 2073 the root transcript prefix accumulated up to the relevant output point. The same holds for
 2074 each leaf duplex, initialized by $p_{\text{leaf}} \parallel K \parallel U \parallel \langle j \rangle_8$ and absorbing only its own ciphertext
 2075 blocks in the message phase. In decryption and verification, Algorithm 3 absorbs the
 2076 submitted ciphertext blocks before the root tag comparison, so the same flattened message
 2077 phase transcript is defined whether the comparison accepts or rejects. $\square$

2078 A.2 Well-formed Transcript Prefixes

2079 For a duplex-call prefix $((\sigma_0, \sigma_{\text{ph},0}), \dots, (\sigma_i, \sigma_{\text{ph},i}))$, write

$$2080 \quad \text{Flat}((\sigma_0, \sigma_{\text{ph},0}), \dots, (\sigma_i, \sigma_{\text{ph},i})) = \bigg\|_{j=0}^{i-1} (\sigma_j \parallel \sigma_{\text{ph},j} \parallel \text{pad}_j) \parallel \sigma_i \parallel \sigma_{\text{ph},i},$$

2081 where $\text{pad}_j = \text{pad10}^*1[r](|\sigma_j \parallel \sigma_{\text{ph},j}|)$, the padding of the j -th full duplex-call input. The
 2082 current call's padding pad_i is not part of Flat ; it is applied internally by the sponge
 2083 function. Equivalently, after the i -th call the duplex state is the sponge state after
 2084 absorbing $\text{Flat}(\cdot) \parallel \text{pad}_i$.

2085 **Definition 2** (Well-Formed TW128 Transcript Prefixes). A typed TW128 transcript prefix
 2086 is *well formed* if it is a prefix, ending at a construction transcript lookup point, of one of
 2087 the following two call languages.

- 2088 • A root transcript starts with $(p_{\text{root}} \parallel K \parallel U, \sigma_{\text{init}})$, with fixed-width fields $K \in \{0, 1\}^k$
 2089 and $U \in \{0, 1\}^\nu$. It then absorbs the associated data byte string (only when the
 2090 associated data is nonempty; the phase is otherwise elided, Section 3.5), the root
 2091 ciphertext byte string, and the aggregation byte string (only when $n \geq 1$; when
 2092 $n = 0$ the aggregation phase is elided, Section 3.5, and the root tag is emitted by the
 2093 final σ_{msg}^+ message call), in that order. Each present phase byte string is partitioned
 2094 by the block rule of Algorithm 1: the empty string gives one empty final block, and
 2095 a nonempty string gives zero or more ρ -byte non-final blocks followed by one final
 2096 block. The phases use respectively $\sigma_{\text{ad}}^- / \sigma_{\text{ad}}^+$, $\sigma_{\text{msg}}^- / \sigma_{\text{msg}}^+$, and $\sigma_{\text{agg}}^- / \sigma_{\text{agg}}^+$.
- 2097 • A leaf- j transcript starts with $(p_{\text{leaf}} \parallel K \parallel U \parallel \langle j \rangle_8, \sigma_{\text{init}})$, where $1 \leq j < 2^{64}$ and
 2098 the fields have fixed width. It then absorbs the leaf ciphertext byte string under
 2099 $\sigma_{\text{msg}}^- / \sigma_{\text{msg}}^+$, using the same block rule.

2100 The message byte string in this language is always the ciphertext byte string absorbed by the
 2101 construction: encryption obtains it by XORing the plaintext with the current keystream,
 2102 while decryption and verification use the submitted ciphertext. Thus a verification
 2103 transcript is well formed for every in-domain ciphertext before the tag comparison is made;
 2104 acceptance is not part of the transcript language.

2105 The root aggregation byte string has the form $T_1 \parallel \dots \parallel T_n \parallel \langle n \rangle_8$, where each T_j has
 2106 fixed byte length $\tau_{\text{leaf}}/8$. The trailing fixed-width field therefore gives a unique parse of
 2107 the leaf tag sequence. The output-point class of a well-formed counted prefix is the class
 2108 in Table 12 determined by its duplex instance and final-call suffix.

Remark 2 (Well-Formedness versus Realizability). The language of Definition 2 is deliberately larger than the set of transcript prefixes that TW128 computations realize. For example, it does not require an aggregation phase to follow a root message phase of exactly β bytes, although every real computation with $n \geq 1$ absorbs a full β -byte root chunk (Algorithm 2). The inclusion runs in the safe direction everywhere: every construction-generated prefix is well formed, so the structural lemmas of this appendix and the separation propositions of Section 6 apply to construction transcripts as restrictions of statements proved on the larger language, while the direct-query decoder of Definition 4 and the root tag candidate predicate of Lemma 3 accept the larger language and therefore only over-count adversary queries, which loosens no bound.

Name the tag-emitting transcript families:

$\mathcal{R}_{\text{tag}}$ = root transcript prefix whose sponge output gives the final root tag,

$\mathcal{L}_{\text{tag}}(j)$ = leaf- j transcript prefix whose sponge output gives the leaf tag.

The final root tag is emitted by the σ_{agg}^+ aggregation call (the $\mathcal{R}_{\text{tag}}$ family) when $n \geq 1$, and by the closing σ_{msg}^+ message call (an $\mathcal{R}_{\text{msg}}^+$ prefix, in the notation of Corollary 8) when $n = 0$ and the aggregation phase is elided. Since $n = \text{leaves}(C)$ is fixed by $\|C\|$, every computation on a given ciphertext emits its tag at the output point selected by $\|C\|$. Intermediate construction transcript lookup points (init, non-final AD, non-final message, and non-final aggregation) are also typed prefixes covered by Lemma 8; the exact keyed direct-query decoder names them in Table 12.

Definition 3 (Root Tag Output Point and Final Root Transcript Class). For any in-domain encryption or verification computation on ciphertext C , the *root tag output point* is the construction lookup that emits the root tag: $\mathcal{R}_{\text{tag}}$ when $\text{leaves}(C) \geq 1$, and the closing root message point $\mathcal{R}_{\text{msg}}^+$ when $\text{leaves}(C) = 0$ and aggregation is elided. The *final root transcript class* of such a computation is the set of construction computations that realize the same bridged final root input at this root tag output point. A class is *encryption-first* if its first construction generation is by an encryption query, and *verification-first* if its first construction generation is by a non-reply verification query.

A.3 Transcript Injectivity

Lemma 6 (Transcript Injectivity). *Well-formed TW128 flattened duplex-call transcript prefixes are injectively encoded and separated by duplex instance (root vs leaf), leaf index j , phase suffix, and duplex-call order. For final-root prefixes, equality of flattened inputs recovers the same root initialization, AD blocks, root ciphertext blocks, aggregation input, and leaf tag sequence.*

Proof. By Lemma 5, every well-formed output point can be written as a [BDPVA11] flattened sponge input for the corresponding duplex-call prefix, namely Flat as defined above. The deterministic current-call pad10^*1 is applied internally by the sponge function and is not part of the input to which [BDPVA11, Lemma 4] is applied. That lemma recovers the duplex-call sequence from the flattened sponge input; in TW128, each call input is $\sigma \parallel \sigma_{\text{ph}}$ with byte-aligned σ and fixed 4-bit suffix σ_{ph} , so recovering the call input recovers the pair $(\sigma, \sigma_{\text{ph}})$.

Definition 2 fixes the root and leaf initialization fields, the admissible phase order, and the final-call suffixes. The distinct prefixes p_{root} and p_{leaf} separate root from leaf transcripts; the fixed-width leaf field $\langle j \rangle_8$ separates leaves; and the seven suffix values of Table 3 separate phases and non-final/final calls. Thus equality of flattened inputs forces equality of the recovered instance type, key, nonce, leaf index when present, phase suffixes, and duplex-call order.

Phase-recovery sub-claim. For each present phase $P \in \{\text{ad}, \text{msg}, \text{agg}\}$ (the ad phase is present exactly when the associated data is nonempty; see Section 3.5), TW128 absorbs the phase’s input byte string s_P by partitioning into ρ -byte blocks (with the empty case yielding a single empty block, per Section 3.5) and feeding each block to the duplex with suffix σ_P^- (non-final) or σ_P^+ (final). Lemma 5 and [BDPVA11, Lemma 4] recover the corresponding $(\sigma, \sigma_{\text{ph}})$ sequence from the flattened input. Because the partition map is deterministic and uniquely invertible on its image (the empty case yielding the single empty block carrying σ_P^+), concatenating the recovered σ blocks reconstructs s_P exactly. The map from s_P to the phase’s $(\sigma, \sigma_{\text{ph}})$ sequence is therefore injective.

Applied phase by phase to a well-formed final-root transcript:

- AD. The associated data phase contributes a σ_{ad}^+ -tagged block exactly when the associated data is nonempty; an empty associated data elides the phase and contributes no block. When the phase is present the sub-claim recovers the AD byte string. Thus distinct associated data strings $A_i \neq A_j$ are separated either by the presence versus absence of the AD phase (when exactly one is empty) or, when both are nonempty, by their distinct AD σ -block sequences.
- MSG. Distinct absorbed root or leaf ciphertext byte strings produce distinct MSG σ -block sequences; the message language of Definition 2 is ciphertext-absorbing for encryption and decryption alike.
- AGG. The aggregation phase is present exactly when $n \geq 1$ (Section 3.5). When it is present, the sub-claim applied to the aggregation byte string $T_1 \parallel \cdots \parallel T_n \parallel \langle n \rangle_8$ recovers that byte string. The final eight bytes give $\langle n \rangle_8$, and each leaf tag has fixed byte length $\tau_{\text{leaf}}/8$, so the byte string has a unique parse as a leaf tag sequence. Two root aggregation encodings are therefore equal exactly when their leaf tag sequences are equal. When $n = 0$ the aggregation phase is absent and the final-root transcript ends at the recovered σ_{msg}^+ root message call; the MSG bullet recovers the root ciphertext byte string, and the absence of any $\sigma_{\text{agg}}^-/\sigma_{\text{agg}}^+$ call distinguishes such a prefix from every $n \geq 1$ final-root prefix.

Therefore the recovered final-root transcript determines the root initialization, AD blocks, root ciphertext blocks, and, when $n \geq 1$, the aggregation byte string and leaf tag sequence. $\square$

Lemma 7 (Leaf-Transcript Recovery). *Let two construction-generated final-root prefixes have equal flattened inputs, and suppose that, for each position j of their common leaf tag sequence, the two computations’ position- j final leaf transcripts are either equal or receive distinct leaf tags. Then equality of flattened inputs recovers the same leaf transcript sequence, and therefore the same full absorbed ciphertext. Both computations’ position- j final leaf transcripts are leaf- j transcripts under the common recovered (K, U) , so the hypothesis holds whenever no two distinct construction-generated final leaf transcripts with the same key, nonce, and leaf index receive the same leaf tag; when one of the two prefixes is reached by an encryption computation, it suffices that no construction-generated final leaf transcript receives the same leaf tag as a distinct encryption-reached final leaf transcript with the same key, nonce, and leaf index.*

Proof. By Transcript Injectivity (Lemma 6), the two equal flattened final-root inputs recover the same root initialization, AD blocks, root ciphertext blocks, and, when $n \geq 1$, the same aggregation byte string and leaf tag sequence; the trailing $\langle n \rangle_8$ field fixes the common leaf count n . When $n = 0$ no leaf is present and the recovered root ciphertext blocks already constitute the full absorbed ciphertext. When $n \geq 1$, the aggregation byte string parses uniquely as the leaf tag sequence $T_1 \parallel \cdots \parallel T_n$; the two leaf tag sequences are equal, so for each j the two computations’ position- j final leaf transcripts—leaf- j

transcripts carrying the common recovered (K, U) (Definition 2, Section 3.5)—receive the same leaf tag, and by hypothesis are equal, hence so are their absorbed leaf ciphertext bodies. Together with the common root ciphertext blocks this recovers the same full absorbed ciphertext. $\square$

A.4 Bridge and Typed Decoding

Lemma 8 (Bridge Injectivity). *The random oracle games used in Sections 5 and 6 and Corollary 6 are ordinary random oracle games on a single random oracle $\text{RO}_{\text{flat}} : \{0, 1\}^* \rightarrow \{0, 1\}^\infty$ over the unpadded [BDPVA08] H -input domain (Section 4.3), with TW128 root and leaf duplex outputs answered by the typed restriction*

$$\text{RF}(X) = \text{RO}_{\text{flat}}(\text{flat}(X)), \quad \text{flat}(X) = I[r, r_{\max}](\text{raw_flat}(X)),$$

for $r = r_{\max} = 1344$. Here $\text{raw_flat}(X)$ is the native flattened input $\text{Flat}(\cdot)$ of Section A.2 applied to the duplex-call sequence of X , and $I[r, r_{\max}]$ is the [BDPVA11] Theorem 3 preprocessing map. The induced map $X \mapsto \text{flat}(X)$ is injective on well-formed typed TW128 duplex transcript prefixes.

Proof. The bridge $I[r, r_{\max}]$ is described by the proof of [BDPVA11, Theorem 3] in three steps: (1) pad the input with $\text{pad10}^*1[r]$; (2) for each r -bit block, append $r_{\max} - r$ zero bits; (3) strip the trailing pad10^* from the result. The pad10^*1 padding in step (1) appends a 1, then $q = (-|\text{raw_flat}(X)| - 2) \bmod r$ zero bits, then a closing 1; the pad10^* -unpadding in step (3) strips the $r_{\max} - r$ trailing zero bits introduced by step (2) together with the closing 1 of step (1). For TW128 with $r = r_{\max}$, step (2) appends an empty string and step (3) strips zero trailing zeros and the closing 1, so the net effect of steps (1) and (3) is to map a native flattened input $W = \text{raw_flat}(X)$ to $W \parallel 1 \parallel 0^q$. Any string in this image has a unique inverse: q is the number of trailing zero bits, the preceding bit is the inserted 1, and deleting this suffix recovers W .

Injectivity of flat on well-formed typed TW128 prefixes then follows from Lemma 6: distinct well-formed typed prefixes have distinct $\text{raw_flat}(\cdot)$ images, and the bridge is injective on that image. Direct adversary queries on inputs outside the bridged TW128 well-formed transcript language still count in $q_{\text{RO}}(\mathcal{B})$ (Section 4.4) but cannot trigger bad_{key} . $\square$

The bridge injectivity established here underlies Opening Triple Separation (Proposition 1, Section 6): distinct (K, U, A) triples yield distinct bridged final root transcripts.

Corollary 8 (Output-Point Separation). *Every well-formed TW128 typed transcript prefix at a counted construction transcript lookup point belongs to exactly one of the following output-point classes, identified by the recovered duplex instance and the 4-bit suffix of its final call:*

- $\mathcal{R}_{\text{init}}(\text{root duplex}, \text{suffix } \sigma_{\text{init}})$,
- $\mathcal{L}_{\text{init}}(j)$ (leaf- j duplex, suffix σ_{init}),
- $\mathcal{R}_{\text{ad}}^- / \mathcal{R}_{\text{ad}}^+$ (root duplex, suffix σ_{ad}^- or σ_{ad}^+),
- $\mathcal{R}_{\text{msg}}^- / \mathcal{R}_{\text{msg}}^+$ (root duplex, suffix σ_{msg}^- or σ_{msg}^+),
- $\mathcal{L}_{\text{msg}}^-(j)$ (leaf- j duplex, suffix σ_{msg}^-),
- $\mathcal{L}_{\text{tag}}(j)$ (leaf- j duplex, suffix σ_{msg}^+),
- $\mathcal{R}_{\text{agg}}(\text{suffix } \sigma_{\text{agg}}^-)$,

- $\mathcal{R}_{\text{tag}}$ (root duplex, suffix σ_{agg}^+).

For any two such prefixes whose bridged flattened inputs are equal, the separation of Lemma 6 recovers the same duplex instance, leaf index j when present, (K, U) pair, output-point class, and duplex-call sequence. Equivalently, any two distinct counted transcript prefixes already differ in one of these recovered fields, and hence in their flattened inputs.

Proof. Lemma 8 recovers equality of native flattened inputs from equality of bridged inputs, after which Lemma 6 recovers the duplex-call sequence and its per-call $(\sigma, \sigma_{\text{ph}})$ pairs. The first call's σ is $p_{\text{root}} \parallel K \parallel U$ for root instances and $p_{\text{leaf}} \parallel K \parallel U \parallel \langle j \rangle_8$ for leaf instances, with fixed-width fields, so it fixes the instance type (by p_{root} vs p_{leaf}), the (K, U) pair, and the leaf index j ; within each instance type the phase suffixes are pairwise distinct, so the recovered instance type together with the final-call suffix fixes the output-point class, and each prefix therefore lies in exactly one class. Conversely, two distinct well-formed prefixes in the same class on the same instance differ in some recovered call's $(\sigma, \sigma_{\text{ph}})$ pair or in the sequence length, hence in their flattened inputs. $\square$

Definition 4 (Exact Keyed Decoder). For a direct query to RO_{flat} with input Y and finite length ℓ , define $\text{KeyedTWOOut}(Y)$ as a partial map from $\{0, 1\}^*$ to $\{0, 1\}^k \times \mathcal{C}_{\text{out}}$. The map ignores ℓ . It returns $(K, \mathcal{C})$ iff there exists a typed TW128 duplex-call prefix X such that $Y = \text{flat}(X)$, the raw flattened prefix of X parses as a well-formed TW128 transcript prefix in the class $\mathcal{C}$ listed in Table 12, and the first call of X contains the candidate key K in the fixed-width root field $p_{\text{root}} \parallel K \parallel U$ or leaf field $p_{\text{leaf}} \parallel K \parallel U \parallel \langle j \rangle_8$. If no such prefix exists, $\text{KeyedTWOOut}(Y) = \perp$. In particular, inputs with the right final suffix but a malformed phase order, noncanonical block partition, malformed aggregation string, or out-of-range leaf index are outside the well-formed language and decode to $\perp$. A non- $\perp$ decoding always fixes the candidate key K and nonce U in the fixed-width initialization field. It fixes the associated data A only for root prefixes whose recovered phase order determines the complete AD string: a final AD prefix, or any later root message or aggregation prefix, with absence of AD calls determining $A = \varepsilon$. We use the key projection $(K, \mathcal{C})$ throughout.

By Lemma 8 and Corollary 8, KeyedTWOOut is well defined: a bridged input cannot parse as two distinct counted well-formed prefixes, and an accepted input fixes a unique candidate key and output-point class. Inputs outside this well-formed transcript language decode to $\perp$.

Table 12: Counted keyed transcript classes for KeyedTWOOut .

Class $\mathcal{C}$	Instance	Final call / role
$\mathcal{R}_{\text{init}}$	root	σ_{init} , root initialization
$\mathcal{R}_{\text{ad}}^-$	root	σ_{ad}^- , non-final AD block
$\mathcal{R}_{\text{ad}}^+$	root	σ_{ad}^+ , final AD block
$\mathcal{R}_{\text{msg}}^-$	root	σ_{msg}^- , non-final root message block
$\mathcal{R}_{\text{msg}}^+$	root	σ_{msg}^+ , final root message block
$\mathcal{R}_{\text{agg}}^-$	root	σ_{agg}^- , non-final aggregation block
$\mathcal{R}_{\text{tag}}$	root	σ_{agg}^+ , final aggregation / root tag
$\mathcal{L}_{\text{init}}(j)$	leaf j	σ_{init} , leaf initialization
$\mathcal{L}_{\text{msg}}^-(j)$	leaf j	σ_{msg}^- , non-final leaf message block
$\mathcal{L}_{\text{tag}}(j)$	leaf j	σ_{msg}^+ , final leaf message block / leaf tag

The table is a transcript-lookup table, not an observation-length table. A class is counted even when the construction's requested output length at that call is zero, including root initialization, non-final AD and aggregation calls, a final AD call whose following root chunk requests no keystream bytes, and the final root message call. A direct query triggers the same decoder regardless of ℓ : short finite-output queries and zero-length queries $(Y, 0)$

still count as direct queries to the address Y . Leaves have no AD or aggregation classes; a leaf transcript enters the table only through leaf initialization, non-final leaf message blocks, or the final leaf tag point. When the associated data is empty the root transcript likewise has no $\mathcal{R}_{\text{ad}}^-$ or $\mathcal{R}_{\text{ad}}^+$ class (Section 3.5); the $\mathcal{R}_{\text{init}}$ class then supplies the first root keystream block, exactly as $\mathcal{L}_{\text{init}}(j)$ supplies the first leaf keystream block, and the decoder still ignores ℓ . Symmetrically, when $n = 0$ the root transcript has no $\mathcal{R}_{\text{agg}}$ or $\mathcal{R}_{\text{tag}}$ class (Section 3.5); the final root message call $\mathcal{R}_{\text{msg}}^+$ then supplies the root tag, exactly as $\mathcal{L}_{\text{tag}}(j)$ supplies a leaf tag, and the decoder again ignores ℓ . An $n \geq 1$ computation also reaches an $\mathcal{R}_{\text{msg}}^+$ point when closing chunk 0, but there the construction requests zero output and continues into the aggregation phase, so the $\mathcal{R}_{\text{msg}}^+$ value is exposed as a tag only on the $n = 0$ path; since $n = \text{leaves}(C)$ is fixed by $\|C\|$, the two uses never coincide on a single ciphertext. We call $\mathcal{R}_{\text{init}}$, $\mathcal{R}_{\text{ad}}^+$, $\mathcal{R}_{\text{msg}}^-$, $\mathcal{L}_{\text{init}}(j)$, and $\mathcal{L}_{\text{msg}}^-(j)$ the *stream output points*: by the schedule just described, these are the only classes at which a construction computation requests keystream bytes, while tag values are exposed only at $\mathcal{L}_{\text{tag}}(j)$ and $\mathcal{R}_{\text{tag}}$ points and, on the $n = 0$ path, at $\mathcal{R}_{\text{msg}}^+$.

The hidden-key encryption, verification, and mu-ind accounting all reindex transcript addresses by the key field. We record that reindexing once.

Lemma 9 (Key-Indexed Transcript Renaming). *For a key $K \in \{0, 1\}^k$, let $\mathcal{Y}_K$ be the set of bridged inputs $\text{flat}(X)$ of well-formed TW128 transcript prefixes X whose first-call key field is K —the counted hidden-key transcript addresses of Table 12 for key K . Then:*

1. Disjointness. *For distinct keys $K_0 \neq K_1$, the sets $\mathcal{Y}_{K_0}$ and $\mathcal{Y}_{K_1}$ are disjoint.*
2. Renaming. *The map $\Phi_{K_0 \rightarrow K_1}$ that replaces the first-call key field K_0 by K_1 and leaves every other recovered field unchanged—nonce, leaf index, phase sequence, AD blocks, ciphertext blocks, aggregated leaf tag bytes, and root/leaf instance—induces through flat a bijection $\text{flat} \circ \Phi_{K_0 \rightarrow K_1} \circ \text{flat}^{-1} : \mathcal{Y}_{K_0} \rightarrow \mathcal{Y}_{K_1}$ preserving output-point classes and equality relations among addresses.*
3. Direct-query avoidance. *A direct query input $Y \in \mathcal{Y}_K$ has $\text{KeyedTWOOut}(Y) = (K, \mathcal{C})$ for a counted class $\mathcal{C}$; hence if no prior direct query decodes to K then $\mathcal{Y}_K$ is disjoint from all prior direct-query inputs.*

Proof. By bridge injectivity (Lemma 8) and Transcript Injectivity (Lemma 6), a bridged input $\text{flat}(X)$ determines the native flattened prefix and hence the first duplex call $p_{\text{root}} \| K \| U$ or $p_{\text{leaf}} \| K \| U \| \langle j \rangle_8$, in particular its fixed-width key field K . (1) If $\text{flat}(X) = \text{flat}(X')$ with key fields K_0 and K_1 then $K_0 = K_1$, so distinct keys give disjoint address sets. (2) $\Phi_{K_0 \rightarrow K_1}$ rewrites only the key field, leaving the recovered duplex-call sequence otherwise intact, so it is a bijection on well-formed prefixes preserving the output-point class (Corollary 8) and the equality pattern among prefixes; flat transports it to the stated bijection on addresses. (3) By Lemma 8 a $Y = \text{flat}(X) \in \mathcal{Y}_K$ decodes under KeyedTWOOut to $(K, \mathcal{C})$; the contrapositive gives the avoidance claim. $\square$

A.5 Adaptive Candidate Bounds

The Adaptive Candidate Collision and Preimage lemma is stated as Lemma 2 in Section 6, where a proof sketch also appears; this subsection supplies the full proof.

Proof of Lemma 2. Both parts expose the interaction in chronological order and sample each candidate's RO_{flat} value lazily at its append step, which the premises license: when Y_j is appended its string is fixed by the prior state (predictability) and no bit of $\text{RO}_{\text{flat}}(Y_j)$ has yet been revealed (unrevealed value), so $\text{RO}_{\text{flat}}(Y_j)$ may be drawn then, uniform and independent of everything so far.

Collision. For a fixed position set $I = \{i_1 < \dots < i_s\} \subseteq \{1, \dots, L\}$, let E_I be the event that these candidate positions exist and their τ -bit projections are all equal. Condition on any positive-probability prefix immediately before position i_r is appended, for $r \geq 2$, and on the values sampled at the earlier positions in I . By the above, $\text{RO}_{\text{flat}}(Y_{i_r})$ is then uniform and independent of the conditioning, so its τ -bit projection equals the already fixed projection of Y_{i_1} with probability $2^{-\tau}$. Multiplying over $r = 2, \dots, s$ gives $\Pr[E_I] \leq 2^{-\tau(s-1)}$; if some position in I is never appended then E_I is false and the bound still holds. A union bound over the $\binom{L}{s}$ position sets gives the claim.

Preimage. For a fixed position $j \in \{1, \dots, L\}$, let E_j be the event that position j exists and $\lfloor \text{RO}_{\text{flat}}(Y_j) \rfloor_\tau = t$. Condition on any positive-probability prefix immediately before Y_j is appended; the target t is a constant. By the above, $\text{RO}_{\text{flat}}(Y_j)$ is uniform and independent of the conditioning, so its τ -bit projection equals t with probability $2^{-\tau}$. If position j is never appended, E_j is false. A union bound over the L positions gives the claim.

Second-preimage corollary. If the designated position j_0 is not appended, the event is false. Otherwise split the other candidates according to whether they occur before or after j_0 . For earlier candidates, condition on any positive-probability prefix immediately before Y_{j_0} is appended, including the already sampled projections of candidates $1, \dots, j_0 - 1$. The value $\text{RO}_{\text{flat}}(Y_{j_0})$ is then fresh and uniform, so its τ -bit projection hits one of those at most $j_0 - 1$ earlier projections with probability at most $(j_0 - 1)2^{-\tau}$. For a later candidate Y_j , $j > j_0$, condition on the full prefix immediately before it is appended, including the fixed projection $t_0 = \lfloor \text{RO}_{\text{flat}}(Y_{j_0}) \rfloor_\tau$. The value $\text{RO}_{\text{flat}}(Y_j)$ is fresh and uniform at its append step, so it hits t_0 with probability $2^{-\tau}$. A union bound over the at most $L - j_0$ later candidates gives an additional $(L - j_0)2^{-\tau}$. Adding the two sides gives $\Pr[\text{sec}_{\tau, j_0}] \leq (L - 1) \cdot 2^{-\tau}$; the designated candidate is excluded, so its trivially equal projection does not enter the count. $\square$

A.6 Adaptive Key Tests

Lemma 10 (Adaptive Key-Test Bound). *Consider a TW128 random oracle model game with a hidden key K sampled uniformly from $\{0, 1\}^k$, direct adversary access to RO_{flat} , and an immediate abort event bad_{key} that fires on a direct query (Y, ℓ) exactly when Y decodes under KeyedTWOut to the candidate key K and a counted output-point class of Table 12. Suppose that, before the abort, the following invariant holds for every direct query position i : after conditioning on the full visible execution prefix and on the set S_i of distinct decoded candidate keys from earlier direct queries, the hidden K is uniform on $\{0, 1\}^k \setminus S_i$. Then*

$$\Pr[\text{bad}_{\text{key}}] \leq \text{KeyGuess}_k(q_{\text{RO}}) = q_{\text{RO}}/2^k,$$

where q_{RO} is the execution-uniform count of direct RO_{flat} queries.

Proof. If $q_{\text{RO}} \geq 2^k$ the right-hand side $q_{\text{RO}}/2^k$ is at least one, so the claim is trivial. Assume $q_{\text{RO}} < 2^k$.

Order the direct RO_{flat} queries. By the definition of KeyedTWOut and Lemma 8, each query input either decodes to $\perp$ or decodes to a unique candidate key and counted output-point class. Queries that decode to $\perp$, and repeated queries whose decoded candidate key already lies in the current set S_i , do not change the set of possible first successful tests. Removing them can only shorten the sequence, so it suffices to analyze the subsequence of distinct fresh decoded candidate keys $K'_1, \dots, K'_m$, where $m \leq q_{\text{RO}}$.

Let F_i be the event that the first i fresh decoded candidate keys all miss the hidden K . The invariant says that, conditioned on the visible prefix before the $(i + 1)$ -st fresh test and on F_i , the hidden K is uniform over the $2^k - i$ values not yet tested. The next fresh candidate key K'_{i+1} is one of those remaining values and is determined by the adversary's

2377 next direct query, so

$$2378 \quad \Pr[K = K'_{i+1} \mid F_i] = \frac{1}{2^k - i}, \quad \Pr[F_{i+1} \mid F_i] = 1 - \frac{1}{2^k - i}.$$

2379 Multiplying the miss probabilities gives

$$2380 \quad \Pr[F_m] = \prod_{i=0}^{m-1} \left(1 - \frac{1}{2^k - i}\right) = \prod_{i=0}^{m-1} \frac{2^k - i - 1}{2^k - i} = \frac{2^k - m}{2^k}.$$

2381 Thus the probability that some fresh decoded candidate key equals the hidden K is
 2382 $m/2^k \leq q_{\text{RO}}/2^k$. This event is exactly the first bad_{key} abort, because KeyedTWOOut returns
 2383 either $\perp$ or a unique candidate key. $\square$

2384 A.7 Leaf Tag Collision Bound

2385 **Lemma 11** (Leaf Tag Collision Bound). *Consider any TW128 random oracle model game*
 2386 *with an adversary $\mathcal{B}$ that has direct access to RO_{flat} and construction computations that*
 2387 *may generate final leaf transcripts—by encryption, plaintext-computing, or verification*
 2388 *computations, whether they accept or reject. Let bad_{leaf} be the event that two distinct*
 2389 *construction-generated final leaf transcript prefixes $X \neq X'$ in leaf tag output-point classes*
 2390 *receive the same τ_{leaf} -bit projection from $\text{RF}(\cdot)$, and let $n_{\text{leaf}}(\mathcal{B})$ be the execution-uniform*
 2391 *count of distinct construction-generated final leaf transcripts (Section 4.4). The first bound*
 2392 *applies the collision case ($s = 2$, $\tau = \tau_{\text{leaf}}$) of Lemma 2 to a chronological leaf tag candidate*
 2393 *sequence; the second is a same-slot freshness argument in an all-reject game.*

- 2394 1. General count. *Unconditionally—in particular when some or all keys are adversary-*
 2395 *supplied, so direct queries may themselves land on leaf tag output points—*

$$2396 \quad \Pr[\text{bad}_{\text{leaf}}] \leq \text{LeafColl}_{\tau_{\text{leaf}}}(q_{\text{RO}}(\mathcal{B}) + n_{\text{leaf}}(\mathcal{B})) = \binom{q_{\text{RO}}(\mathcal{B}) + n_{\text{leaf}}(\mathcal{B})}{2} / 2^{\tau_{\text{leaf}}}.$$

- 2397 2. Hidden key, same slot. *Suppose the game has a hidden key, nonce-respecting en-*
 2398 *ryption queries, and all-reject verification: every non-replay verification submission*
 2399 *performs its construction computation and returns reject to the adversary. (The*
 2400 *game $\mathbf{G}_d^{\text{rj}}$ of Definition 5 and the INT-RUP continuation in the proof of Lemma 16*
 2401 *have this shape.) Say two final leaf transcripts share a slot if they recover the same*
 2402 *key, nonce, and leaf index (K, U, j) (Corollary 8). Let bad_{leaf} be the event that some*
 2403 *construction-generated final leaf transcript receives the same τ_{leaf} -bit projection as*
 2404 *a distinct final leaf transcript that shares its slot and is reached by an encryption*
 2405 *computation, and let hit_{leaf} be the event that a direct adversary query hits the bridged*
 2406 *input of a construction-generated hidden-key final leaf transcript before that transcript*
 2407 *is first created. Then, in such a game,*

$$2408 \quad \Pr[\text{bad}_{\text{leaf}} \text{ before } \text{hit}_{\text{leaf}}] \leq \frac{n_{\text{leaf}}(\mathcal{B})}{2^{\tau_{\text{leaf}}}}.$$

2409 *Proof.* Build a chronological leaf tag candidate sequence: a candidate is the bridged input
 2410 of a leaf tag output point ($\mathcal{L}_{\text{tag}}(j)$), appended the first time it is reached—by a direct query
 2411 in $\mathcal{B}$'s oracle phase or by a construction computation generating a final leaf transcript—and
 2412 not appended again. By Lemma 8, Corollary 8, and Lemma 6, two reaches of the same
 2413 bridged input come from the same final leaf transcript, so repeated evaluations add no
 2414 candidate—the lazy RO_{flat} value is reused—and distinct final leaf transcripts give distinct
 2415 candidates. No candidate's RO_{flat} value is read before its append step, so the sequence meets
 2416 the predictability and unrevealed-value premises of Lemma 2, whose collision case with

$s = 2$ and $\tau = \tau_{\text{leaf}}$ bounds the probability that two candidates share a τ_{leaf} -bit projection by $\binom{L}{2} \cdot 2^{-\tau_{\text{leaf}}}$ for a length- L sequence. Every bad_{leaf} pair is a pair of construction-generated candidates.

(1) *General count.* Direct queries may decode to a leaf tag class $\mathcal{L}_{\text{tag}}(j)$ and are appended as candidates; together with the $n_{\text{leaf}}(\mathcal{B})$ construction-generated transcripts the sequence has length at most $q_{\text{RO}}(\mathcal{B}) + n_{\text{leaf}}(\mathcal{B})$, so the case with $L = q_{\text{RO}}(\mathcal{B}) + n_{\text{leaf}}(\mathcal{B})$ gives $\Pr[\text{bad}_{\text{leaf}}] \leq \binom{q_{\text{RO}}(\mathcal{B}) + n_{\text{leaf}}(\mathcal{B})}{2} \cdot 2^{-\tau_{\text{leaf}}} = \text{LeafColl}_{\tau_{\text{leaf}}}(q_{\text{RO}}(\mathcal{B}) + n_{\text{leaf}}(\mathcal{B}))$. The count charges every direct query and every construction-generated final leaf transcript regardless of key provenance, so no assumption on the keys is used.

(2) *Hidden key, same slot.* Every bad_{leaf} pair is a pair of construction-generated hidden-key transcripts sharing a slot, with at least one member reached by an encryption computation. Nonce-respecting encryption uses each nonce in at most one encryption query, and within one encryption computation the leaf index separates final leaf transcripts, so each slot contains at most one encryption-reached final leaf transcript. That transcript may have first been created by an earlier verification or plaintext-computing computation, when the encryption's absorbed leaf body reproduces an earlier absorbed body. Charge each possible bad_{leaf} pair to its member that is not reached by an encryption computation. Each charged transcript has one slot, and that slot has at most one encryption-reached counterpart, so each charged transcript receives at most one charge. Thus at most $n_{\text{leaf}}(\mathcal{B})$ charged pairs can occur.

Each pair's collision is decided at the first creation of its second-created member. At that moment the earlier member's τ_{leaf} -bit projection is a function of the prior execution, while the later member's projection is a fresh uniform lazy sample: by Transcript Injectivity (Lemma 6) two construction reaches of the same bridged input come from the same final leaf transcript, so no construction computation has read that bridged input before the creation, and a direct query reading it before the creation would decode under **KeyedTWO**ut to the hidden key and its $\mathcal{L}_{\text{tag}}(j)$ class (Lemma 8)—precisely hit_{leaf} . Hence before hit_{leaf} , conditioned on the entire prior execution, each pair collides on its τ_{leaf} -bit projections with probability $2^{-\tau_{\text{leaf}}}$, and a union bound over the at most $n_{\text{leaf}}(\mathcal{B})$ pairs gives the claim. No assumption on what the visible interface exposes is needed: a projection revealed after its sampling cannot alter a collision already decided. By Leaf-Transcript Recovery (Lemma 7), on $\neg \text{bad}_{\text{leaf}}$ equal flattened final-root inputs of an encryption computation and any construction computation identify the same leaf transcript sequence. $\square$

A.8 Encryption Transcript Marginal Uniformity

Lemma 12 (Encryption Transcript Marginal Uniformity). *For any fixed nonce-respecting encryption query in the TW128 random oracle experiment, let $\text{hit}_{\text{fresh}}$ be the event that, before this query reaches some counted transcript point, a prior direct adversary query to RO_{flat} has already hit the bridged input of that hidden-key point. On $\neg \text{hit}_{\text{fresh}}$, every transcript output of the query that contributes to a returned ciphertext block, to an internal leaf tag, or to the final root tag is either fresh when this query first generates its bridged input, or reuses the lazy-table value sampled uniformly when that input was first generated by an earlier construction computation. The counted output points first generated by the fixed query are pairwise distinct.*

Proof. Work on an execution outside $\text{hit}_{\text{fresh}}$ for the fixed query. This rules out prior adversary reads of the hidden-key construction inputs considered below. A previous construction computation may nevertheless have generated one of these inputs, for example through a non-replay verification query using the same nonce. Such a lookup samples the lazy-table value uniformly when it first occurs; if the fixed encryption query later reaches the same bridged input, it reuses that same value. Thus it remains only to show that the counted output points newly generated by the fixed encryption query are distinct from one

another and from previous nonce-respecting encryption queries. Conditional information leaked by earlier verification comparisons of a final root tag is accounted for separately by Lemma 15, not by this lemma.

First root keystream block. The first root keystream block, when present, is produced by the root duplex after the initialization call $p_{\text{root}} \parallel K \parallel U$ and any AD absorption (elided when $A = \varepsilon$, in which case the σ_{init} call itself emits the first root keystream block). Since the nonce-respecting adversary never repeats a nonce, no previous encryption query can have used the same root initialization transcript. Within the current query, Lemma 6 separates initialization, AD, message, and aggregation calls by their suffixes and by the duplex-call encoding, so the first root keystream transcript is distinct from the other counted output points generated by this query. If the requested keystream length is zero, no ciphertext block is returned at that point and there is nothing to prove for that output.

First leaf keystream block. For each leaf chunk $j \geq 1$, the first leaf keystream block is produced after the initialization call $p_{\text{leaf}} \parallel K \parallel U \parallel \langle j \rangle_8$. The prefix separates leaves from the root, the nonce separates this encryption query from all previous encryption queries, and $\langle j \rangle_8$ separates leaves within this query. Hence every first leaf keystream transcript newly generated by the query is distinct.

Later message outputs. Argue inductively within each root or leaf message stream. Suppose the current message-output transcript is either first generated by this query or reuses a value generated by an earlier construction computation. If its output is used as a keystream block, write its RO_{flat} value as Z_i . Since the adaptive plaintext block M_i is a deterministic function of the prior view,

$$C_i = M_i \oplus Z_i$$

is now fixed. TW128 then absorbs C_i with the appropriate message phase suffix (σ_{msg}^- or σ_{msg}^+) before requesting the next keystream block. Once C_i is fixed, the next transcript is fixed; Lemma 6's injectivity and domain separation imply it has not appeared earlier unless the same duplex-call prefix has appeared earlier. Nonce uniqueness excludes this across encryption queries. By Corollary 8, the next message-keystream output point can coincide with a within-query earlier output point only if both belong to the same output-point class on the same duplex instance (root or leaf- j with matching j) and have the same duplex-call sequence; the current message-stream prefix is strictly longer than every earlier prefix in that stream, so this is excluded. Across distinct instances within the same query (root vs leaf, or two leaves with distinct j), the prefixes already differ in the first call's σ by $p_{\text{root}}/p_{\text{leaf}}$ or $\langle j \rangle_8$ respectively. Each later message-output transcript that emits keystream is therefore distinct from the other points generated by this query, unless it is a reuse of a point generated by an earlier construction computation. The terminal message-output transcript of a leaf emits the leaf tag rather than a following keystream block; the same separation argument applies to that transcript, so each internal leaf tag newly generated by the query is sampled at a distinct point. Empty inputs still induce a duplex call, but with zero keystream output they contribute no ciphertext block.

Final root tag transcript. The final root tag is sampled at the root tag output point selected by $\|C\|$: the $\mathcal{R}_{\text{tag}}$ point after the root duplex absorbs the aggregation transcript built from the leaf tags and $\langle n \rangle_8$ when $n \geq 1$, and the $\mathcal{R}_{\text{msg}}^+$ point closing the root message phase when $n = 0$ (Section 3.5). When $n \geq 1$, the leaf tags are themselves RO_{flat} outputs on final leaf transcripts; condition on their sampled values, on the sampled ciphertext, and on the adversary-visible prior view and RO_{flat} table state, so that the root aggregation transcript is fixed. When $n = 0$ there is no aggregation transcript; condition

on the sampled ciphertext and the adversary-visible prior view alone. By Corollary 8, the selected root tag output-point class is separated from every other output-point class in the same query — including root initialization (suffix σ_{init}), AD-phase outputs, the remaining message phase outputs, leaf-side output points (separated by p_{leaf}), and, in the $n \geq 1$ case, earlier $\mathcal{R}_{\text{agg}}$ -prefix output points (separated by the σ_{agg}^- vs σ_{agg}^+ suffix). Even if two encryption queries realize the same aggregation transcript (or, for $n = 0$, the same root ciphertext), the full root transcript also contains the root initialization $p_{\text{root}} \parallel K \parallel U$, and nonce uniqueness separates that initialization from all previous encryption queries. The final root tag transcript is therefore either newly generated at a distinct point of the fixed query, or is a reuse of a value first sampled by an earlier construction computation. In the newly generated case, its τ_{root} -bit RO_{flat} output is a fresh lazy-table sample; in the reuse case, it is the same uniform sample drawn when the bridged input was first generated. $\square$

A.9 Mu-Ind Hybrid Games and Verification Accounting

The direct mu-ind adjacent-hybrid proof bounds one user-replacement step of the hybrid through an explicit game sequence, defined next, using the generic key-test and leaf-collision accounting lemmas together with the verification-specific facts below. We use the root tag output point and final root transcript classes of Definition 3; in particular, computations on a common ciphertext share the same root tag output point because $\text{leaves}(C)$ is fixed by $\|C\|$.

Definition 5 (Adjacent-Hybrid Games). Fix a flat-RO mu-ind adversary $\mathcal{B}$, the pre-sampled keys $K_1, \dots, K_D$ of Section 4.2, and a user index $d \in \{1, \dots, D\}$. On the branch where two sampled keys collide, every game below outputs the same fixed bit; the remaining clauses describe the branch with pairwise distinct keys. Every game maintains one lazy table Θ implementing RO_{flat} : $\Theta[Y]$ samples a fresh uniform infinite string at the first access to Y and returns the stored string afterwards—equivalently, Θ is sampled in full at the outset and read on demand. Direct adversary queries (Y, ℓ) are answered by $[\Theta[Y]]_\ell$; the *environment* answers users $e < d$ by (Rand, Replay) and users $e > d$ by $(\text{Enc}_{K_e}^{\text{RO}}, \text{Ver}_{K_e}^{\text{RO}})$, reading construction outputs from Θ ; and the New interface and the validity and nonce-respecting checks of Section 4.2 are common to every game. Every game maintains the flags bad_{key} , bad_{leaf} , and forge , initialized false; no oracle answer or branch depends on a flag, and bad_{key} is set whenever a direct query input decodes under KeyedTWOOut to K_d and a counted output-point class of Table 12.

- Games G_d^{re} and G_d^{rj} answer user d by the shared oracle code of Algorithm 4 and differ exactly in the marked return statement: a non-replay verification whose comparison succeeds sets forge and returns accept in G_d^{re} and reject in G_d^{rj} .
- Game G_d^{id} answers user- d encryption queries with fresh uniform byte strings of length $\|M\| + \tau_{\text{root}}/8$, recorded in W , and user- d verification by the replay check alone; it performs no user- d construction computation and sets neither bad_{leaf} nor forge . G_d^{id} is the hybrid H_d^* of Lemma 17, instrumented with bad_{key} .

The flag bad_{leaf} is set, in G_d^{re} and G_d^{rj} , when a user- d construction computation first creates, or an encryption computation first reaches, a final leaf transcript whose τ_{leaf} -bit projection equals that of a distinct, already created final leaf transcript sharing its (K_d, U, j) slot, at least one of the two reached by an encryption computation—the same-slot event of Lemma 11, part 2, for user d . It guards no code difference and is used only in the accounting.

Lemma 13 (Environment Neutrality). *In each game of Definition 5, on the branch with pairwise distinct keys, the environment’s construction-internal RO_{flat} lookups—made while*

Algorithm 4 Shared user- d oracle code of the games G_d^{re} and G_d^{rj} (Definition 5). Θ is the shared lazy RO_{flat} table, read at bridged transcript prefixes, and W the user- d replay record. Inputs failing the validity or nonce-respecting checks of Section 4.2 are rejected beforehand, as in every game. The two games differ exactly in the marked lines 17 and 18.

```

1: procedure  $\text{RO}(Y, \ell)$  ▷ direct adversary query
2:   set  $\text{bad}_{\text{key}}$  if  $\text{KeyedTWOut}(Y) = (K_d, \mathcal{C})$  for a counted class  $\mathcal{C}$  (Table 12)
3:   return  $[\Theta[Y]]_\ell$ 
4: end procedure

5: procedure  $\text{ENC}(U, A, M)$ 
6:   run  $\text{Enc}_{K_d}^{\text{RO}}(U, A, M)$ , reading each requested duplex output as the corresponding
   projection of  $\Theta$  at its bridged transcript prefix (Section 4.3); let the result be  $C \parallel T$ 
7:   set  $\text{bad}_{\text{leaf}}$  if a final leaf transcript first created or first reached here completes a
   same-slot collision (Definition 5)
8:   record  $(U, A, C \parallel T)$  in  $W$ ; return  $C \parallel T$ 
9: end procedure

10: procedure  $\text{VF}(U, A, C \parallel T)$ 
11:   if  $(U, A, C \parallel T) \in W$  then return accept
12:   end if
13:   determine the verification transcript of  $(K_d, U, A, C)$ ; read from  $\Theta$  the  $\tau_{\text{leaf}}$ -bit leaf
   tags at its final leaf transcripts, setting  $\text{bad}_{\text{leaf}}$  when a transcript first created here
   completes a same-slot collision, and nothing at its stream output points
14:    $T' \leftarrow [\Theta[\text{flat}(R)]]_{\tau_{\text{root}}}$  at the root tag output point  $R$  of the computation (Defini-
   tion 3)
15:   if  $T' = T$  then
16:      $\text{forge} \leftarrow \text{true}$ 
17:     return accept ▷ in  $G_d^{\text{re}}$  only
18:     return reject ▷ in  $G_d^{\text{rj}}$  only
19:   end if
20:   return reject
21: end procedure

```

2561 *simulating $\text{Enc}_{K_e}^{\text{RO}}$ and $\text{Ver}_{K_e}^{\text{RO}}$ for users $e > d$ —neither set a user- d flag nor generate the*
 2562 *bridged input of a user- d counted transcript point.*

2563 *Proof.* Whenever such a lookup lies in a counted keyed transcript class, Lemma 8 decodes
 2564 its candidate key as K_e , which on the distinct-keys branch differs from K_d ; a lookup
 2565 outside the counted keyed language is irrelevant to the flags. By the key-field disjointness
 2566 of Lemma 9, the environment’s addresses are disjoint from user d ’s. The flags are set only
 2567 by direct adversary queries decoding to K_d (bad_{key}), by user- d final leaf transcript creations
 2568 (bad_{leaf}), and by user- d verification comparisons (forge), none of which the environment
 2569 performs. Ideal users $e < d$ perform no construction lookups at all. □

2570 **Lemma 14** (Verification Visible-Prefix Key Uniformity). *In the stopped single-key*
 2571 *TW128 random oracle encryption/verification interaction, or its INT-RUP extension with a*
 2572 *plaintext-computing oracle, with hidden key K , direct access to RO_{flat} , and immediate abort*
 2573 *on bad_{key} , order the direct RO_{flat} queries. For a direct query position i , let S_i be the set of*
 2574 *distinct candidate keys decoded by KeyedTWOut from earlier direct-query inputs. Condition*
 2575 *on any positive-probability visible execution prefix immediately before query i , on no earlier*
 2576 *bad_{key} , and on S_i . Then the hidden key K is uniform on $\{0, 1\}^k \setminus S_i$. Consequently the*
 2577 *uniformity invariant required by Lemma 10 holds in the stopped interaction.*

Proof. The visible prefix contains the adversary’s oracle queries and replies: direct RO_{flat} outputs, encryption outputs and the resulting replay table, plaintext-computing replies when that oracle is present, verification bits, and the adversary state determined by these values. Fix such a prefix and two keys $K_0, K_1 \notin S_i$.

Let $\Phi_{0 \rightarrow 1} = \Phi_{K_0 \rightarrow K_1}$ be the key-renaming map of Lemma 9 and $\Phi_{1 \rightarrow 0}$ its inverse. By that lemma, $\text{flat} \circ \Phi_{0 \rightarrow 1} \circ \text{flat}^{-1}$ is a bijection $\mathcal{Y}_{K_0} \rightarrow \mathcal{Y}_{K_1}$ between the bridged hidden-key construction addresses for K_0 and for K_1 , preserving output-point classes and equality relations among addresses.

This bijection is applied online, not to a pre-existing fixed address set. Whenever a construction computation under K_0 would first query a hidden-key address Y , the coupled computation under K_1 queries $\Phi_{0 \rightarrow 1}(Y)$ and receives the same lazy-table value; repeated queries are coupled consistently through the same map. If the queried address is a root aggregation address containing leaf tag bytes, those bytes are already coupled equal because the corresponding leaf tag lookups were either repeated or first sampled earlier under the same renaming rule. Thus adaptively generated message transcripts and root aggregation transcripts are renamed step by step with the same visible ciphertext, leaf tag, and root tag bytes.

Prior direct-query inputs are not renamed: they are adversary-chosen visible strings. Since $K_0, K_1 \notin S_i$, neither $\mathcal{Y}_{K_0}$ nor $\mathcal{Y}_{K_1}$ contains a prior direct-query input (Lemma 9, direct-query avoidance). Direct-query lazy-table values are therefore coupled identically and remain disjoint from the renamed hidden construction addresses.

Use the same adversary coins, the same direct-query lazy-table values, and the online renamed construction-internal lazy-table values. Encryption outputs have the same distribution under the coupling, because every construction lookup that determines ciphertext, leaf tags, or the final root tag is paired with a distinct renamed lookup carrying the same sampled value. Plaintext-computing computations are renamed in the same online way: a plaintext reply is the submitted ciphertext block XORed with a coupled keystream projection at a stream output point (Table 12), so it carries the same visible bytes under both keys, and the recomputed tag is not returned. Replay verification decisions depend only on the visible replay table. Non-replay verification computations are renamed in the same online way as encryption computations, and the final root tag comparison therefore gives the same accept/reject bit. Verification does not reveal plaintext on rejection or acceptance; it reveals only that bit.

Thus every visible prefix with hidden key $K_0 \notin S_i$ has the same conditional probability as the corresponding execution with hidden key $K_1 \notin S_i$. Since K is initially uniform and the condition “no earlier bad_{key} ” removes exactly the keys in S_i , all keys in $\{0, 1\}^k \setminus S_i$ remain equally likely. $\square$

Foreign-key key-test claim. In each game of Definition 5, work in the original experiment, where the key-collision branch outputs a fixed bit rather than conditioning on pairwise distinct keys. Then

$$\Pr[\text{bad}_{\text{key}}] \leq \text{KeyGuess}_k(q_{\text{RO}}(\mathcal{B})).$$

Fix the foreign keys and leave their transcript addresses unchanged. Before the first user- d key hit, the visible prefix is invariant under renaming K_d among keys that have not already been tested and are not equal to a fixed foreign key, by the same online coupling as Lemma 14 and the key-field disjointness of Lemma 9. For each fresh candidate key tested by one of $\mathcal{B}$ ’s direct RO_{flat} queries, the contributing event is contained in $\{K_d \text{ equals that candidate}\}$ and costs at most 2^{-k} in the original pre-sampled-key experiment; a candidate equal to a fixed foreign key contributes only on the fixed-output key-collision branch and therefore contributes nothing to bad_{key} . A union bound over at most $q_{\text{RO}}(\mathcal{B})$ direct queries gives the claim.

Lemma 15 (Verification-First Root Guessing). *In the game G_d^j of Definition 5, stopped at the first occurrence of $\text{bad}_{\text{key}} \cup \text{bad}_{\text{leaf}}$ —with bad_{leaf} the same-slot event of Lemma 11, part 2, for user d —let win_{vf} be the event that some non-replay user- d verification comparison succeeds within a verification-first final root transcript class. Then*

$$\Pr[\text{win}_{\text{vf}}] \leq q_v^{(d)} / 2^{\tau_{\text{root}}},$$

where $q_v^{(d)}$ is the user- d non-replay verification cap of Section 4.2.

Proof. Throughout, a *non-replay verification submission* means a valid non-replay user- d verification query; by the validity convention, an invalid verification query returns 0 and performs no construction computation, so it realizes no bridged final root input, belongs to no final root transcript class, and contributes to no G_C below.

In G_d^j , every non-replay submission performs its construction computation and returns reject (Algorithm 4), so no comparison outcome is visible to the adversary, and each class's root tag value is read from the lazy table only by the class's own comparisons and by any later encryption computation reaching the class.

Fix a verification-first final root transcript class C and follow the stopped execution until the first encryption query generating the same class, if any (the cutoff). Let G_C be the set of distinct candidate tags submitted by non-replay verification queries in class C before that cutoff; repeated tags only shrink this set. The class's correct root tag U_C is the τ_{root} -bit projection of RO_{flat} at the class's bridged root tag input. Before the stop, no direct query has read this hidden-key input: such a query would be decoded by **KeyedTWOOut** to the hidden key and would set bad_{key} ; the environment reads no user- d address (Lemma 13). Before the cutoff, no encryption query has revealed U_C as an encryption tag. (When the class's root tag point is $\mathcal{R}_{\text{msg}}^+$, i.e. $n = 0$, an encryption query on a longer message that shares chunk 0 may touch this same address while closing its chunk-0 message phase, but it requests zero output there and continues into its aggregation phase, so it does not reveal U_C ; its own tag is emitted at a later $\mathcal{R}_{\text{tag}}$ point.) The comparisons against U_C return the constant reject, so the visible execution and the set G_C are functions only of the adversary's coins, visible oracle replies, and lazy-table values at inputs other than this class's root tag address. By the lazy-sampling order exchange, conditioned on those values, U_C may equivalently be sampled after G_C is fixed, and is uniform independently of G_C . The probability that some pre-cutoff comparison in class C succeeds—that $U_C \in G_C$ —is therefore at most $|G_C|/2^{\tau_{\text{root}}}$.

After the cutoff, if it exists, class C admits no successful non-replay comparison. Let the cutoff encryption query return (U_e, A_e, C_e, T_e) . Any later verification submission in the same final root transcript class has, by the final-root determinacy of Lemma 6 and Leaf-Transcript Recovery (Lemma 7) under the absence of bad_{leaf} —which supplies the recovery hypothesis because the cutoff computation is an encryption computation—the same (U, A, C) as that encryption computation. If the submitted tag is T_e , the tuple is a recorded encryption tuple and the query is answered by the replay check before any comparison; if the submitted tag is not T_e , the comparison fails.

Each non-replay verification submission deterministically reconstructs one final root transcript at the root tag output point, so it contributes its submitted tag to at most one set G_C , and only when its class is verification-first and the submission precedes that class's cutoff. Thus $\sum_C |G_C| \leq q_v^{(d)}$, and a union bound over the verification-first classes gives the claim. $\square$

Lemma 16 (INT-RUP Root Guessing). *Let $\mathcal{B}$ be a single-key INT-RUP adversary in the TW128 random oracle model, with nonce-respecting encryption queries and the resources of Section 4.4. In its encryption/plaintext-computing/verification interaction with hidden key K , let bad_{key} be the event that some direct RO_{flat} query input Y satisfies $\text{KeyedTWOOut}(Y) = (K, C)$ for a counted class C of Table 12, and let bad_{leaf} be the same-slot event of Lemma 11,*

part 2, over the construction computations generated by encryption, plaintext-computing, and verification. Let *forge* be the event that some valid non-replay verification query accepts. Then

$$\Pr[\text{forge before } \text{bad}_{\text{key}} \cup \text{bad}_{\text{leaf}}] \leq q_v(\mathcal{B})/2^{\tau_{\text{root}}}.$$

Proof. Throughout, a *non-replay verification submission* means a valid non-replay verification query; invalid verification queries return 0, perform no construction computation, realize no bridged final root input, and contribute to no set of guessed tags.

Work in the game stopped at the first occurrence of bad_{key} or bad_{leaf} . A final root transcript class is *tag-revealed* once an encryption query has returned the root tag for that class. If a non-replay verification submission belongs to a tag-revealed class, compare it with an encryption computation that revealed the class. The two computations' final-root transcript prefixes are well-formed typed transcript prefixes with the same bridged final root input; by Lemma 8, they have the same native flattened input. Then final-root determinacy (Lemma 6) and Leaf-Transcript Recovery (Lemma 7) under $\neg \text{bad}_{\text{leaf}}$ —which supplies the recovery hypothesis because the comparand is an encryption computation—recover the same (U, A, C) . If the submitted tag is the revealed tag, the tuple is a replay-table hit; if it is any other tag, the final comparison rejects. Thus any accepting non-replay verification submission must lie in a class whose root tag has not yet been revealed by encryption.

For tag-unrevealed classes, define an all-reject continuation: encryption and plaintext-computing queries are answered normally, replay verification queries accept normally, and each non-replay verification submission performs the same construction computation but returns reject to the adversary. Plaintext-computing queries may generate the same stream and leaf tag transcripts as decryption, but they do not reveal the final root tag value. Equivalently, the final root tag sample for a tag-unrevealed class can be deferred until it is first needed for an encryption answer or a verification comparison; this does not change any plaintext reply. By Output-Point Separation (Corollary 8) and the keystream schedule recorded at Table 12, the keystream values used to compute returned plaintext blocks are projections at stream output points, not at the final root tag output point; in particular the chunk-0 $\mathcal{R}_{\text{msg}}^+$ call of an $n \geq 1$ computation requests zero output, so no plaintext reply exposes an $\mathcal{R}_{\text{msg}}^+$ tag value.

Fix a final root transcript class $\mathcal{C}$ whose first construction generation is not by an encryption query; for a class first generated by an encryption query the tag is revealed from that generation onward, so every verification submission in it is covered by the previous paragraph. Follow the all-reject path until the first encryption query that reveals the class, if any (the cutoff; any such query is later than the class's first generation). Let $G_{\mathcal{C}}$ be the set of distinct tags submitted by non-replay verification queries in class $\mathcal{C}$ before that cutoff. Before the cutoff, no encryption query has returned the class tag and no direct query has read the hidden-key root tag point without triggering bad_{key} . Conditioned on the all-reject path and on all lazy-table values except this final root tag projection, the correct tag for $\mathcal{C}$ is uniform in $\{0, 1\}^{\tau_{\text{root}}}$ and independent of $G_{\mathcal{C}}$. The probability that any pre-cutoff submitted tag in $G_{\mathcal{C}}$ equals the deferred class tag is therefore at most $|G_{\mathcal{C}}|/2^{\tau_{\text{root}}}$.

Consider the event that the first accepting non-replay verification submission in the stopped game belongs to class $\mathcal{C}$. Up to that global first accept, the adversary's visible answers match the all-reject continuation: every earlier non-replay submission rejects, replay and plaintext-computing queries are answered identically in both executions, and encryption answers are unchanged. The construction computations also agree at every address except that the continuation defers the final root tag value for class $\mathcal{C}$ until the comparison is needed. Therefore the first accepting non-replay submission can belong to class $\mathcal{C}$ only before the cutoff, and only if its submitted tag equals the deferred class tag, which has probability at most $|G_{\mathcal{C}}|/2^{\tau_{\text{root}}}$.

The execution-uniform count $q_v(\mathcal{B})$ applies to the all-reject continuation. Each valid non-replay verification query contributes its submitted tag to at most one such set $G_{\mathcal{C}}$,

and replay-table hits do not contribute. Thus $\sum_c |G_c| \leq q_v(\mathcal{B})$, and a union bound over the classes not first generated by an encryption query, each with its pre-cutoff guess set G_c , gives the claim. $\square$

A.10 Mu-Ind Adjacent-Hybrid Lemma

The multi-user indistinguishability theorem (Theorem 2) telescopes a hybrid over the D users, replacing one real user by its ideal (Rand, Replay) interface at each step. The following lemma bounds one such step through the game sequence of Definition 5: the instrumented real game G_d^e , the all-reject game G_d^j differing from it in a single return statement, and the ideal game G_d^{id} , with the divergence events set as the flags bad_{key} , bad_{leaf} , and forge by game code.

Lemma 17 (Mu-Ind Adjacent-Hybrid Bound). *Let $\mathcal{B}$ be a flat-RO mu-ind adversary in the pre-sampled-key formulation of Section 4.2, with keys $K_1, \dots, K_D$ sampled at the start of the experiment. Fix $d \in \{1, \dots, D\}$. Let H_{d-1}^* and H_d^* be the adjacent stopped hybrids that output a fixed bit if any two sampled keys collide, and otherwise run $\mathcal{B}$ with one shared lazy RO_{flat} table as follows: users $e < d$ are answered by (Rand, Replay), users $e > d$ are answered by $(\text{Enc}_{K_e}^{\text{RO}}, \text{Ver}_{K_e}^{\text{RO}})$, and user d is answered by $(\text{Enc}_{K_d}^{\text{RO}}, \text{Ver}_{K_d}^{\text{RO}})$ in H_{d-1}^* and by (Rand, Replay) in H_d^* . If $n_{\text{leaf}}^{(d)}$ and $q_v^{(d)}$ are the user- d resource caps, then*

$$|\Pr[\mathcal{B}^{H_{d-1}^*} \Rightarrow 1] - \Pr[\mathcal{B}^{H_d^*} \Rightarrow 1]| \leq \text{KeyGuess}_k(q_{\text{RO}}(\mathcal{B})) + \frac{n_{\text{leaf}}^{(d)}}{2^{\eta_{\text{leaf}}}} + \frac{q_v^{(d)}}{2^{\eta_{\text{root}}}}.$$

Proof. The fixed-output branch on a key collision is identical in H_{d-1}^* , H_d^* , and the games of Definition 5, so only the branch with pairwise distinct keys contributes; work on that branch, with the keys fixed, throughout. The proof identifies H_{d-1}^* with G_d^e (Claim 1), relates G_d^e to G_d^j pointwise (Claim 2) and G_d^j to $H_d^* = G_d^{\text{id}}$ in distribution (Claim 3), and assembles the bound by a single first-occurrence decomposition in G_d^j (Claim 4).

Claim 1: G_d^e realizes H_{d-1}^* . The user- d oracles of Algorithm 4 implement $(\text{Enc}_{K_d}^{\text{RO}}, \text{Ver}_{K_d}^{\text{RO}})$ with three distribution-preserving refactors. First, reading every construction output as a projection of the shared table Θ at its bridged transcript prefix is the lazy implementation of RO_{flat} (Section 4.3); whether Θ is sampled in full at the outset or entry by entry on first access is immaterial. Second, VF reads Θ only at leaf tag and root tag output points. This is sound because decryption absorbs the submitted ciphertext, so the verification transcript is determined by (K_d, U, A, C) independently of any stream output (Lemma 5, Section 3.6), and verification discards the recovered plaintext, so the keystream projections a real $\text{Ver}_{K_d}^{\text{RO}}$ computation reads influence neither a visible answer nor the value of any other table entry; leaving those entries unread until some later computation first accesses them leaves the joint distribution of all oracle answers unchanged. Third, the flags are instrumentation: no oracle answer or branch depends on them. Hence $\Pr[\mathcal{B}^{G_d^e} \Rightarrow 1] = \Pr[\mathcal{B}^{H_{d-1}^*} \Rightarrow 1]$.

Claim 2: G_d^e and G_d^j agree until forge. Run the two games on the same adversary coins, environment coins, and table realization Θ . The games share all code, including the comparison that sets forge , and differ in exactly one statement, the marked return line of Algorithm 4, which executes only when forge has just been set. The two executions are therefore identical—every oracle answer, table access, record in W , and flag update—up to and including the query at which forge is first set, and the first differing value is that query's returned bit. In particular, on executions in which G_d^j never sets forge , the two runs coincide, outputs included, and the flags bad_{key} and bad_{leaf} are set at the same queries in both games up to the first forge .

Claim 3: G_d^{rj} agrees with H_d^* until bad_{key} . Couple $H_d^* = G_d^{\text{id}}$ to G_d^{rj} as follows: run G_d^{rj} on the same adversary and environment coins, give the G_d^{id} coordinate the same visible oracle answers for every query answered while bad_{key} is unset, and let the G_d^{id} coordinate continue independently according to its own law from the query that sets bad_{key} onward. We claim the resulting G_d^{id} coordinate is distributed exactly as G_d^{id} . The answer rules coincide query by query except for user- d encryption: non-replay user- d verification returns reject in G_d^{rj} and under **Replay** alike, and replay queries are answered from the record W , which the coupling keeps equal in both coordinates; the environment runs identical code in both games; and direct queries read table entries that agree in both games before bad_{key} —an entry written by a user- d construction computation is read by a direct query only if that query decodes under **KeyedTWOut** to K_d and a counted class (Lemma 9), which sets bad_{key} , while entries written by the environment are produced identically in both games (Lemma 13).

It remains to show that, before bad_{key} , each user- d encryption answer of G_d^{rj} is, conditioned on the entire prior visible transcript, a fresh uniform byte string of length $\|M\| + \tau_{\text{root}}/8$ —the law of **Rand**. Fix such a query (U, A, M) . Before bad_{key} , no direct query has read the bridged input of a user- d hidden-key counted point, so the hypothesis $\neg\text{hit}_{\text{fresh}}$ of Encryption Transcript Marginal Uniformity (Lemma 12) holds, and every counted output point first generated by this query is pairwise distinct and receives a fresh uniform sample. The stream output points are always first generated by this query: user- d verification computations read no stream points (Algorithm 4), nonce-respecting encryption and Transcript Injectivity (Lemma 6) separate them from every earlier user- d encryption computation, and the environment writes only foreign-key addresses (Lemma 13). The ciphertext body C is therefore a fresh uniform string. The internal leaf tags are either fresh samples or reuse entries written by user- d verification computations; they are never returned and influence the answer only through the root tag address. The root tag T is read at the computation’s root tag output point (Definition 3), an address distinct from every stream point of the same query by Output-Point Separation (Corollary 8). If that entry is first generated here, T is a fresh uniform sample. Otherwise the entry was written by an earlier user- d verification computation—before bad_{key} , the only other writer of user- d addresses—whose reads of it produced only the constant reject answers of G_d^{rj} ; since no oracle answer or branch depends on the flags, the visible transcript is a function of the coins and of table entries other than this one, and conditioned on the visible transcript the entry remains uniform. In either case (C, T) has the **Rand** law, and at most one user- d encryption reaches any root tag address, since two would share the nonce U , contradicting nonce-respecting encryption. By induction over the query sequence, the coupled G_d^{id} coordinate has exactly the law of H_d^* , and its visible transcript agrees with G_d^{rj} ’s while bad_{key} is unset.

Claim 4: assembly. Place the three games on one probability space: shared adversary and environment coins and one table realization Θ drive G_d^{re} and G_d^{rj} pointwise, and the H_d^* coordinate is coupled to G_d^{rj} as in Claim 3. Until the first occurrence of $\text{forge} \vee \text{bad}_{\text{key}}$ in G_d^{rj} , the three visible transcripts coincide: G_d^{re} ’s by Claim 2, since forge has not yet been set, and H_d^* ’s by Claim 3, since bad_{key} has not. When neither flag is ever set, the final outputs of $\mathcal{B}$ coincide as well, so the outputs can differ only on the event $\text{forge} \vee \text{bad}_{\text{key}}$. With Claim 1 identifying $\Pr[\mathcal{B}^{H_d^*} \Rightarrow 1] = \Pr[\mathcal{B}^{G_d^{\text{re}}} \Rightarrow 1]$ and Claim 3 identifying $\Pr[\mathcal{B}^{H_d^*} \Rightarrow 1]$ with the coupled coordinate’s output probability,

$$\left| \Pr[\mathcal{B}^{H_d^*} \Rightarrow 1] - \Pr[\mathcal{B}^{H_d^*} \Rightarrow 1] \right| \leq \Pr_{G_d^{\text{rj}}}[\text{forge} \vee \text{bad}_{\text{key}}].$$

Decompose by first occurrence, inserting bad_{leaf} —which guards no code difference and enters only this accounting:

$$\begin{aligned} \Pr_{G_d^{\text{rj}}}[\text{forge} \vee \text{bad}_{\text{key}}] &\leq \Pr[\text{bad}_{\text{key}}] + \Pr[\text{bad}_{\text{leaf}} \text{ before } \text{bad}_{\text{key}} \vee \text{forge}] \\ &\quad + \Pr[\text{forge before } \text{bad}_{\text{key}} \vee \text{bad}_{\text{leaf}}]. \end{aligned}$$

The three terms are bounded in G_d^{rj} . For the first, the foreign-key key-test claim above gives $\Pr[\text{bad}_{\text{key}}] \leq \text{KeyGuess}_k(q_{\text{RO}}(\mathcal{B}))$. For the second, G_d^{rj} restricted to user d is a hidden-key game with nonce-respecting encryption and all-reject verification; before bad_{key} no direct query has read a user- d hidden final leaf point—such a query would decode to K_d and its $\mathcal{L}_{\text{tag}}(j)$ class and set bad_{key} —and the environment creates no user- d leaf transcript (Lemma 13), so the same-slot case of Lemma 11 gives $\Pr[\text{bad}_{\text{leaf}} \text{ before } \text{bad}_{\text{key}} \vee \text{forge}] \leq n_{\text{leaf}}^{(d)} / 2^{\tau_{\text{leaf}}}$. For the third, before $\text{bad}_{\text{key}} \vee \text{bad}_{\text{leaf}}$ every successful comparison lies in a verification-first class: a class first generated by a user- d encryption computation admits none, since by the final-root determinacy of Lemma 6 and Leaf-Transcript Recovery (Lemma 7)—whose hypothesis $\neg \text{bad}_{\text{leaf}}$ supplies because the comparand is the class’s encryption computation—any non-replay submission in the class carries the same (U, A, C) as that encryption and either matches the recorded tuple, in which case the replay check answers it before any comparison, or fails the comparison. Lemma 15 then gives $\Pr[\text{forge before } \text{bad}_{\text{key}} \vee \text{bad}_{\text{leaf}}] \leq q_v^{(d)} / 2^{\tau_{\text{root}}}$. Summing the three bounds gives

$$|\Pr[\mathcal{B}^{H_{d-1}^*} \Rightarrow 1] - \Pr[\mathcal{B}^{H_d^*} \Rightarrow 1]| \leq \text{KeyGuess}_k(q_{\text{RO}}(\mathcal{B})) + \frac{n_{\text{leaf}}^{(d)}}{2^{\tau_{\text{leaf}}}} + \frac{q_v^{(d)}}{2^{\tau_{\text{root}}}},$$

with the user- d key-test event bad_{key} charged a single time for the adjacent hybrid. $\square$

B The Flat Sponge Lift

This appendix supplies the flat sponge lift machinery (Lemmas 20 and 22 and Corollary 9) invoked by the AEAD security bounds of Section 5, the root tag bounds of Section 6, and the committing-notion bound of Corollary 6. The machinery applies uniformly to the games named once here.

The lift has two steps in the win-event games and three in the real-vs-ideal games. First, the [BDPVA08] indistinguishability theorem replaces the real permutation used by the TW128 construction and by direct primitive queries with the simulator over a flat random oracle, at cost $\text{Lift}_c(N_{\text{tot}})$, where $N_{\text{tot}} = N + N_{\text{prim}}$ bounds the operational primitive trace: construction-internal duplex calls induced by construction queries and deterministic challenger computations, together with direct primitive queries. The simulator-side construction trace is then identified with the flat-RO construction interface by the operational trace lemma below. Second, the simulator-world adversary is converted into a flat-RO adversary whose direct random oracle query count is at most N_{prim} . Thus the random oracle bounds are instantiated at $q_{\text{RO}} \leq N_{\text{prim}}$, while the public permutation lift pays $\text{Lift}_c(N + N_{\text{prim}})$. Third, in the real-vs-ideal games, whose ideal world exposes the real permutation rather than the simulator (Section 4.1), the ideal-side simulator introduced by the first step is exchanged back for (π, π^{-1}) at cost $\text{Lift}_c(N_{\text{prim}})$ (Lemma 21); the exchange is possible because the ideal construction oracles are independent of the primitive.

Definition 6 (Single-Stage TW128 Games). The *single-stage TW128 games* are the games of Sections 4.1 and 4.2 in which one adversary interacts with the challenger in a single stage, so that a [BDPVA08] replacement applies to an entire world’s public-permutation transcript at once, including deterministic challenger checks after the adversary’s final

output. They are all-in-one AE, mu-ind, INT-RUP, CMTDs-4, the H_{TW128} collision and preimage games (standard and everywhere), and CMTs- ℓ for $\ell \in \{1, 3, 4\}$. The win-event games among these have one world; the real-vs-ideal games—all-in-one AE and mu-ind—have two, and the lift below treats their real world by the full replacement and their ideal world by the simulator-permutation exchange of Lemma 21.

The derived adversary forwards the following game data to the corresponding flat-RO experiment.

Game family	Forwarded data
AE / mu-ind / INT-RUP	construction queries, including New where present and replay state
H_{TW128} hash games	final structured hash game output
CMTDs / CMTs	final openings and deterministic challenger checks

B.1 Simulator Convention

Throughout this appendix, $(S^{RO_{\text{flat}}}, S^{-1, RO_{\text{flat}}})$ denotes the public primitive simulator supplied by [BDPVA08, Theorem 2] — the indistinguishability theorem for the simple padded sponge over a random permutation — after applying the [BDPVA11, Theorem 3] bridge of Lemma 8 to TW128’s native `pad10*1` transcript inputs. The superscript records that both simulator directions share the same internal lazy RO_{flat} table.

The simulator is a proof device: each game of Sections 4.1 and 4.2 states its advantage over the real permutation, and the reduction to the corresponding random oracle model adversary passes through intermediate simulator worlds. In those worlds, the construction interface is the one specified by the game ((**Rand**, **Replay**), instantiated per user in mu-ind, on the ideal side) and the public primitive interface is $(S^{RO_{\text{flat}}}, S^{-1, RO_{\text{flat}}})$. The RO_{flat} table in this notation is internal to that system and is not an extra interface exposed to the application adversary.

Lemma 18 (At Most One Simulator RO Call). *For the simulator $(S^{RO_{\text{flat}}}, S^{-1, RO_{\text{flat}}})$, every primitive-interface query induces at most one simulator-side RO_{flat} call. More precisely, every inverse primitive-interface query and every forward primitive-interface query violating one of the following conditions induces no simulator RO_{flat} call. The symbols s , p , $\mathcal{R}$, $\mathcal{O}$, and $\mathcal{C}$ are the local internal notation of the BDPV simple-padded sponge simulator, not TW128 resource notation:*

(C1) New edge. *The simulator’s internal graph has no outgoing edge from the input node s .*

(C2) Unsaturated. $\mathcal{R} \cup \mathcal{O} \neq \mathcal{C}$.

(C3) Rooted input. *The input node is rooted.*

(C4) Paddable path. *The constructed path p decomposes uniquely as $p = p' \parallel 0^{r^j}$ with p' not ending in 0^r , and the prefix p' is a valid padded sponge input: equivalently, $p' = x \parallel \text{pad10}^*[r](|x|)$ for an unpadded BDPV08 random oracle input x .*

When all four conditions hold, the forward primitive-interface query induces one simulator RO_{flat} call. Consequently, each primitive-interface query induces at most one simulator-side RO_{flat} call.

Proof. This is the branch structure of the [BDPVA08] simple-padded permutation simulator. The simulator invokes the random oracle only on the forward-query branch where the input node is rooted, unsaturated, paddable to an unpadded random oracle input x , and not already extended by an outgoing edge in the simulator graph. Forward queries violating

any of these conditions and all inverse queries return outputs sampled or read from the simulator's internal graph without invoking the random oracle. Under the [BDPVA11, Theorem 3] bridge fixed in Section B.1, that simulator random oracle invocation is the simulator-side RO_{flat} call used throughout this appendix. $\square$

Lemma 19 (Operational Duplex Trace). *In any execution of a single-stage TW128 game (Definition 6), the construction computations counted by N admit an operational primitive trace of cost at most N : each root or leaf duplex call contributes one forward primitive query, and the requested output bits are the corresponding projection of the state returned by that query. When this trace is run against the [BDPVA08] simulator $(\mathcal{S}^{\text{RO}_{\text{flat}}}, \mathcal{S}^{-1, \text{RO}_{\text{flat}}})$ in an execution whose combined primitive trace has cost less than 2^c , the construction outputs are the flat-RO values $\text{RF}(X) = \text{RO}_{\text{flat}}(\text{flat}(X))$ at the same well-formed TW128 transcript prefixes X .*

Proof. A TW128 duplex call absorbs a single suffixed input block, padded by pad10^*1 , applies the permutation once, and returns at most r rate bits from the resulting state. Thus the public-permutation implementation of any root or leaf construction computation is already a primitive trace with one forward query per duplex call. The definition of N in Section 4.4 counts exactly these root initialization, associated data, message, leaf initialization, leaf message, and aggregation duplex calls, including the deterministic challenger checks of the hash and committing games.

Lemma 5 is the [BDPVA11, Lemma 3] duplex-to-sponge identity specialized to TW128: the output of each duplex call is the output of the sponge after absorbing the accumulated flattened duplex-call transcript, and no additional permutation call is made after the current absorption block when the requested output length is at most r . The bridge of Lemma 8, fixed in Section B.1, maps these pad10^*1 transcript prefixes to the [BDPVA08] padded-sponge input domain without changing the output bits at $r = r_{\text{max}}$.

When the same operational trace is run against $(\mathcal{S}^{\text{RO}_{\text{flat}}}, \mathcal{S}^{-1, \text{RO}_{\text{flat}}})$, every construction-internal forward query lies on a rooted, paddable path determined by a well-formed TW128 prefix. In an execution whose combined primitive trace has cost less than 2^c , the simulator cannot have saturated: each primitive query adds at most one capacity value to the simulator's internal $\mathcal{R} \cup \mathcal{O}$ set, so fewer than 2^c queries cannot fill the 2^c -element capacity space. The simulator's sponge-consistency invariant [BDPVA08, Lemma 2] therefore makes the rate part returned at that node agree with the corresponding lazy value $\text{RO}_{\text{flat}}(\text{flat}(X))$. Thus the simulator-side operational construction has the same construction-output law as the flat-RO construction interface used in Section 4.3. $\square$

Lemma 20 (Flat Sponge Lift Bound). *Let G be a single-stage TW128 game (Definition 6). Let $\mathcal{A}$ be a public permutation G adversary with construction interface cost at most N and at most N_{prim} direct primitive-interface queries. In the [BDPVA08] single-stage replacement, implement the TW128 construction computations by the operational duplex trace of Lemma 19 and keep the adversary's direct primitive-interface queries as primitive queries. This primitive transcript has [BDPVA08] query cost at most*

$$N_{\text{tot}} = N + N_{\text{prim}}.$$

Consequently replacing the real public permutation and TW128 construction by the simulator-world primitive interface $(\mathcal{S}^{\text{RO}_{\text{flat}}}, \mathcal{S}^{-1, \text{RO}_{\text{flat}}})$ and the corresponding flat-RO construction interface costs at most $\text{Lift}_c(N_{\text{tot}})$.

Proof. The [BDPVA08] replacement is applied before the derived-adversary construction. View the entire single-stage game, including its construction interface and deterministic challenger checks, as a primitive-only distinguisher that runs the TW128 algorithms internally. By Lemma 19, the construction side contributes at most N forward primitive queries: one per duplex call. This includes plaintext-computing computations, verification

computations whether they accept or reject, the per-user sum of construction calls in mind, the deterministic challenger decryption or encryption checks charged in the committing games, and the hash games' challenger encryptions—the collision game's s submitted-input encryptions, the standard preimage game's source encryption and post-output check, and the everywhere-preimage game's post-output check.

Each direct adversary query to (π, π^{-1}) contributes one primitive-interface query and is counted in N_{prim} . These two counts are disjoint in the public permutation experiment: N counts construction-internal use of the primitive, while N_{prim} counts only the adversary's public primitive interface queries.

Construction work is modeled by the operational duplex trace, rather than by a collection of standalone [BDPVA08] construction queries at every flattened duplex prefix. The standalone-query model would assign a nested-prefix transcript the sum of the individual sponge-evaluation costs, while the operational trace is the same stateful computation the construction actually performs, with the output of each call read from the state just returned by that call. Hence the execution-uniform caps bound the primitive-only query cost by $N + N_{\text{prim}} = N_{\text{tot}}$. If $N_{\text{tot}} \geq 2^c$ then $\text{Lift}_c(N_{\text{tot}}) \geq 1$ and the claimed bound holds trivially, so assume $N_{\text{tot}} < 2^c$, satisfying the hypothesis of [BDPVA08, Theorem 2]. Applying that theorem to this primitive-only distinguisher gives the real-to-simulator replacement, and Lemma 19 identifies the simulator-side construction outputs with the flat-RO construction interface in this nontrivial case. The theorem gives distinguishing or success-probability cost at most the random-permutation differentiating bound $f_P(N_{\text{tot}})$ of [BDPVA08, Lemma 5]. The [BDPVA08, Lemmas 4 and 5] product bounds give the random-transformation bound $f_T(N_{\text{tot}}) = 1 - \prod_{i=1}^{N_{\text{tot}}} (1 - i/2^c)$, and the factor of $f_P(N_{\text{tot}})$ corresponding to the f_T factor at index i is that factor divided by a quantity $1 - (i-1)/2^{r+c} \in (0, 1]$, hence at least as large, so $f_P(N_{\text{tot}}) \leq f_T(N_{\text{tot}})$. With $N_{\text{tot}} < 2^c$ every factor lies in $[0, 1]$ and

$$f_P(N_{\text{tot}}) \leq f_T(N_{\text{tot}}) = 1 - \prod_{i=1}^{N_{\text{tot}}} \left(1 - \frac{i}{2^c}\right) \leq \sum_{i=1}^{N_{\text{tot}}} \frac{i}{2^c} = \frac{N_{\text{tot}}(N_{\text{tot}} + 1)}{2^{c+1}} = \text{Lift}_c(N_{\text{tot}}),$$

the second inequality by the Weierstrass product inequality $\prod_i (1 - x_i) \geq 1 - \sum_i x_i$ for $x_i \in [0, 1]$. $\square$

Lemma 21 (Simulator–Permutation Proximity). *Let $\mathcal{D}$ be a distinguisher with access only to a two-directional primitive interface, making at most q queries. Then*

$$|\Pr[\mathcal{D}^{S^{\text{RO}_{\text{flat}}}, S^{-1}, \text{RO}_{\text{flat}}} \Rightarrow 1] - \Pr[\mathcal{D}^{\pi, \pi^{-1}} \Rightarrow 1]| \leq \text{Lift}_c(q),$$

where $\pi \leftarrow \$ \text{Perm}(\{0, 1\}^{1600})$ and RO_{flat} is the lazily sampled random oracle internal to the simulator.

Proof. Consider the indistinguishability distinguisher $\mathcal{D}'$ that runs $\mathcal{D}$, forwards its primitive-interface queries, and never queries the construction interface. The [BDPVA08] query cost of $\mathcal{D}'$ is at most q : each forwarded query costs one, and no construction queries arise. Because $\mathcal{D}'$ ignores the construction interface, its two views are exactly $\mathcal{D}^{\pi, \pi^{-1}}$ in the real pair and $\mathcal{D}^{S^{\text{RO}_{\text{flat}}}, S^{-1}, \text{RO}_{\text{flat}}}$ in the ideal pair of [BDPVA08, Theorem 2], applied under the [BDPVA11, Theorem 3] bridge fixed in Section B.1. If $q \geq 2^c$ then $\text{Lift}_c(q) \geq 1$ and the claim is trivial; otherwise that theorem bounds the displayed difference by the random-permutation differentiating bound $f_P(q)$ of [BDPVA08, Lemma 5], and the derivation in the proof of Lemma 20 gives $f_P(q) \leq \text{Lift}_c(q)$. $\square$

Lemma 22 (Derived-Adversary Execution Contract). *Let G be a single-stage TW128 game (Definition 6). Let $\mathcal{A}$ be a simulator-world G adversary whose public primitive interface is answered by $(S^{\text{RO}_{\text{flat}}}, S^{-1}, \text{RO}_{\text{flat}})$, with at most N_{prim} direct primitive-interface*

queries. Define $\mathcal{B}_{\text{derived}}(\mathcal{A})$ to run $\mathcal{A}$ internally, answer each primitive-interface query by running the same simulator against direct finite-output queries to its own RO_{flat} oracle, forward construction interface queries unchanged in games with construction oracles, and forward the adversary's final structured output unchanged in the output-checked H_{TW128} hash and committing games. Then the simulator-world execution of $\mathcal{A}$ and the random oracle execution of $\mathcal{B}_{\text{derived}}(\mathcal{A})$ have identical joint distributions, and

$$q_{\text{RO}}(\mathcal{B}_{\text{derived}}) \leq N_{\text{prim}}.$$

In this coupling, simulator-induced RO_{flat} calls and construction-internal RF lookups use the same lazy table. Therefore, if a simulator-induced query is the bridged input $\text{flat}(X)$ of a well-formed TW128 transcript prefix X and a forwarded construction computation later reaches the same prefix X , both interfaces read the same $\text{RO}_{\text{flat}}(\text{flat}(X))$ value, projected to the requested output length. Moreover, $\mathcal{B}_{\text{derived}}$ inherits the construction-side caps of $\mathcal{A}$:

$$\begin{aligned} q_v(\mathcal{B}_{\text{derived}}) &= q_v(\mathcal{A}) \quad (\text{AE, mu-ind, INT-RUP}), \\ n_{\text{leaf}}(\mathcal{B}_{\text{derived}}) &\leq n_{\text{leaf}}(\mathcal{A}) \quad (\text{all games above}). \end{aligned}$$

In the mu-ind game these resource inheritances hold per user, and the New-query transcript is unchanged.

Proof. Couple the two executions by using the same lazy RO_{flat} table for the simulator and for all construction-internal $\text{RF}(\cdot)$ lookups. In the simulator-world execution, $\mathcal{A}$ sees construction answers from the selected game interface and public primitive answers from $(\mathcal{S}^{\text{RO}_{\text{flat}}}, \mathcal{S}^{-1, \text{RO}_{\text{flat}}})$. In the derived random oracle execution, $\mathcal{B}_{\text{derived}}$ runs the same $\mathcal{A}$, forwards construction interface queries or the final structured output required by the selected game unchanged, and answers primitive-interface queries by invoking the same simulator code with the same lazy RO_{flat} table. The construction interface, hash challenger, or committing challenger therefore receives the same forwarded queries or final output, the simulator receives the same primitive queries in the same order, and the lazy table supplies the same values. The full views of $\mathcal{A}$ in the simulator-world execution and of $\mathcal{B}_{\text{derived}}$ in the random oracle execution are identically distributed.

The simulator's random oracle calls may be on arbitrary [BDPVA08] H -input strings, not only bridged TW128 transcript inputs. These calls are the direct RO_{flat} queries of $\mathcal{B}_{\text{derived}}$ and are counted in $q_{\text{RO}}(\mathcal{B}_{\text{derived}})$. Construction-internal $\text{RF}(\cdot)$ lookups made by Enc_K^{RO} , Dec_K^{RO} , $\text{PDec}_K^{\text{RO}}$, Ver_K^{RO} , or a committing challenger are forwarded to the same RO_{flat} table but are not counted in q_{RO} , exactly as in Section 4.4.

By Lemma 18, each primitive-interface query induces at most one direct RO_{flat} query by $\mathcal{B}_{\text{derived}}$, and inverse queries or non-RO-touching forward queries induce none. Since $\mathcal{A}$ makes at most N_{prim} direct primitive-interface queries, this gives $q_{\text{RO}}(\mathcal{B}_{\text{derived}}) \leq N_{\text{prim}}$.

Lemma 8's bridge places TW128 construction restrictions and simulator-induced RO_{flat} calls in a single lazy table. If the simulator invokes RO_{flat} on $\text{flat}(X)$ for a well-formed TW128 transcript prefix X , and a forwarded construction computation later reaches X , the construction lookup $\text{RF}(X) = \text{RO}_{\text{flat}}(\text{flat}(X))$ reads the same table entry and returns the requested projection. Since TW128 uses $r = r_{\text{max}}$, the [BDPVA11, Theorem 3] output post-processing does not alter TW128 output bits. Thus simulator-induced direct queries and construction-internal lookups at the same bridged input are transcript-consistent. In the AE-style games, a simulator call whose input is accepted by KeyedTWOut for the hidden key is still a direct RO_{flat} query and is covered by the same bad_{key} accounting used by the direct random oracle mu-ind proof. The derivation step itself does not prove a key-guessing statement; it only supplies the derived adversary and the query bound at which the random oracle bounds of Sections 5 and 6 are instantiated.

Finally, the execution-uniform caps of Section 4.4 hold along every oracle-answer trace of $\mathcal{A}$, including traces generated by the simulator on the ideal public primitive interface. The derived-adversary construction does not add construction interface encryption,

plaintext-computing, or verification queries and forwards $\mathcal{A}$'s construction interface queries unchanged along each trace. $\mathcal{B}_{\text{derived}}$ therefore inherits the relevant construction resources: $q_v(\mathcal{B}_{\text{derived}}) = q_v(\mathcal{A})$ in the AE, mu-ind, and INT-RUP games, and $n_{\text{leaf}}(\mathcal{B}_{\text{derived}}) \leq n_{\text{leaf}}(\mathcal{A})$ for the distinct construction-generated final-leaf-transcript cap. In mu-ind, construction interface forwarding preserves the New calls and these caps user by user. $\square$

B.2 Lift Application

The lift bound of Lemma 20, the simulator-permutation proximity bound of Lemma 21, and the derived-adversary construction of Lemma 22 combine into a single template, stated once and then instantiated by the bounds of Sections 5 and 6 and by the committing-notion bounds of Corollary 6. Every public-permutation bound is obtained by adding $\text{Lift}_c(N_{\text{tot}})$ —plus $\text{Lift}_c(N_{\text{prim}})$ in the real-vs-ideal games—and substituting $q_{\text{RO}} \leq N_{\text{prim}}$ in the corresponding random oracle theorem.

Corollary 9 (Lift Application). *Let G be a single-stage TW128 game (Definition 6). For every public permutation G -adversary $\mathcal{A}$ against TW128 with the resources of Section 4.4, the derived adversary $\mathcal{B} = \mathcal{B}_{\text{derived}}(\mathcal{A})$ of Lemma 22 is a random oracle G -adversary with $q_{\text{RO}}(\mathcal{B}) \leq N_{\text{prim}}$ and the construction-side caps of Lemma 22—in particular $q_v(\mathcal{B}) \leq Q_v$ and $n_{\text{leaf}}(\mathcal{B}) \leq N_{\text{leaf}}$, where these apply—such that, when G is a win-event game,*

$$\text{Adv}_{\text{TW128}}^G(\mathcal{A}) \leq \text{Lift}_c(N_{\text{tot}}) + \text{Adv}_{\text{RO}}^G(\mathcal{B}),$$

and, when G is a real-vs-ideal game (all-in-one AE or mu-ind),

$$\text{Adv}_{\text{TW128}}^G(\mathcal{A}) \leq \text{Lift}_c(N_{\text{tot}}) + \text{Adv}_{\text{RO}}^G(\mathcal{B}) + \text{Lift}_c(N_{\text{prim}}).$$

If moreover $\text{Adv}_{\text{RO}}^G(\mathcal{B}) \leq \varepsilon_G(q_{\text{RO}}(\mathcal{B}), n_{\text{leaf}}(\mathcal{B}))$ for some ε_G nondecreasing in each argument, then $\varepsilon_G(N_{\text{prim}}, N_{\text{leaf}})$ may replace $\text{Adv}_{\text{RO}}^G(\mathcal{B})$ in the corresponding bound.

Proof. By the flat sponge lift bound (Lemma 20), replacing the real public permutation and TW128 construction by $(S^{\text{RO}_{\text{flat}}}, S^{-1, \text{RO}_{\text{flat}}})$ and the flat-RO construction interface costs at most $\text{Lift}_c(N_{\text{tot}})$, where $N_{\text{tot}} = N + N_{\text{prim}}$ counts both the construction calls of N —including any construction oracle exposed during $\mathcal{A}$'s run and, in the win-event games, the challenger's construction computations, pre- and post-output (in the standard preimage game, the source encryption computing T^* as well as the post-output check)—and the N_{prim} direct primitive queries. The derived-adversary contract (Lemma 22) expresses the resulting simulator-world execution as the identically distributed random oracle execution of $\mathcal{B} = \mathcal{B}_{\text{derived}}(\mathcal{A})$, with $q_{\text{RO}}(\mathcal{B}) \leq N_{\text{prim}}$ and the stated construction-side inheritances. For a win-event game this gives the first bound.

For a real-vs-ideal game, chain four worlds: W_1 , the real construction interface with (π, π^{-1}) ; W_2 , the flat-RO real construction interface with the simulator; W_3 , the ideal construction interface with the simulator; and W_4 , the ideal construction interface with (π, π^{-1}) . The game of Sections 4.1 and 4.2 compares W_1 with W_4 . The step from W_1 to W_2 costs at most $\text{Lift}_c(N_{\text{tot}})$ as above, and under the coupling of Lemma 22 the pair (W_2, W_3) is precisely the real/ideal pair of the random oracle game for $\mathcal{B}$, so $|\Pr[\mathcal{A}^{W_2} \Rightarrow 1] - \Pr[\mathcal{A}^{W_3} \Rightarrow 1]| \leq \text{Adv}_{\text{RO}}^G(\mathcal{B})$. For the step from W_3 to W_4 , the ideal construction oracles—(Rand, Replay), instantiated per user in mu-ind, together with the New interface and the validity checks—are independent of the primitive, so a distinguisher can answer them from its own coins, run $\mathcal{A}$, and forward the at most N_{prim} primitive-interface queries; Lemma 21 bounds this step by $\text{Lift}_c(N_{\text{prim}})$. The triangle inequality combines the three estimates into the second bound.

The final claim follows by the monotonicity of ε_G . When G is parameterized by a challenge fixed before $\mathcal{A}$ runs—the everywhere-preimage target tag, or the standard

preimage source X^* , whose tag $T^* = H_{\text{TW128}}(X^*)$ each coupled world computes—the coupling holds for each fixed challenge and the challenge-independent bound holds for the maximum or average over challenges defining $\text{Adv}_{\text{TW128}}^G(\mathcal{A})$. $\square$

C Vectorized Kernel Details

This appendix details the optimized cores summarized in Section 8.1: their shared batching structure, chunk 0 fusion, instruction selection, register layout, remainder handling, and constant-time addressing.

The optimized paths share the same logical decomposition but differ in their batch widths, remainder strategy, and state writeback points:

Path	Cores and writeback
amd64 AVX-512	Wide: one eight-state ZMM group. Remainder: masked 2–7 state AVX-512 core. Single duplex: AVX-512 permutation. Chunk 0: all active states are written back through <code>state8</code> .
amd64 AVX2	Wide: two four-state YMM groups. Remainder: 2–7 states with dummy inactive lanes. Single duplex: general purpose register permutation. Chunk 0: all active states are written back through <code>state8</code> .
arm64 NEON+SHA3	Wide: four two-state NEON pairs. Remainder: two-state pair core. Single duplex: NEON/SHA-3 permutation. Chunk 0: pair (0, 1) is written back in the wide core, and both states are written back in the pair core.

Shared batching structure. The implementation separates the sequential root duplex from the independent leaf duplexes. The root duplex normally uses a single-duplex ($\times 1$) core, while complete leaf chunks are grouped into eight-instance batches. Every complete leaf chunk is exactly $\beta = 49\rho$ bytes, so all eight instances in a batch follow the same sequence of 49 full ρ -blocks: the first 48 close with `MSG_MORE`, and the last closes with `MSG_LAST`. A single public block index therefore drives the whole batch. The batched state is lane-major: for each of the 25 Keccak lanes, the kernel stores the corresponding 64-bit word of each active instance together in vector registers. Equivalently, for Keccak lane i the vector register holds $(S_0[i], \dots, S_7[i])$ across the active duplex states, so one vector operation updates the same lane in several states.

Chunk 0 fusion. Chunk 0 is part of the root transcript, but its message phase has the same kernel-visible block schedule as a complete leaf chunk. The optimized paths exploit this by placing the post-AD root state into lane 0 of the first batched call and placing the first leaf chunks in the remaining lanes. When the message has at least seven complete leaf chunks, this fused call uses the $\times 8$ core; below that, it uses the same remainder machinery that drains incomplete batches. The output state from lane 0 is written back to the root duplex, which then absorbs the leaf tags produced by the same call.

Remainder and tail handling. Messages are otherwise decomposed into full eight-chunk batches plus a remainder. On `amd64`, a register-resident remainder kernel handles two to seven complete chunks in one call, reading each active chunk in place; on `arm64`, a 2-wide kernel drains the remainder two chunks at a time. A single leftover complete chunk, a partial trailing chunk, and the count-aggregation tail use the single-duplex core. The `amd64` assembly chunk kernels handle the 48 full `MSG_MORE` blocks and leave the final `MSG_LAST` block plus leaf tag extraction to shared Go code, keeping byte-granular tail logic out of assembly. The `arm64` batched kernels process the complete chunk and store

3131 back the states needed for chunk 0 fusion: pair $(0, 1)$ in the $\times 8$ kernel, and both active
 3132 states in the 2-wide kernel.

3133 **amd64.** The **amd64** implementation provides AVX-512 and AVX2 cores and selects between
 3134 them once at startup from an AVX-512 feature probe requiring OS support for YMM/ZMM
 3135 state, **AVX512F**, and **AVX512VL**. The single-duplex permutation has two variants: a general-
 3136 purpose-register core adapted from the standard library’s **KECCAK- p** permutation (using
 3137 the lane-complement trick so that χ requires no separate negation), and an AVX-512
 3138 variant; the root duplex uses whichever matches the selected core.

3139 The eight-way AVX-512 core packs the lane-major state into 25 ZMM registers, one
 3140 per lane with all eight instances in the register’s eight 64-bit slots, and performs the
 3141 permutation with no cross-lane shuffles: rotations use **VPROLQ**, and **VPTERNLOGQ** fuses the θ
 3142 D -term formation and the χ and-not into single three-input logic instructions. The fused
 3143 chunk kernel reads the eight chunks with contiguous loads, transposes each rate block into
 3144 the lane-major layout in registers (**VPUNPCKLQDQ**/**VPUNPCKHQDQ** and **VSHUFI64X2**), absorbs
 3145 on the lane-major state, and transposes each output block back for a contiguous store. Since
 3146 $\rho = 167 = 20 \cdot 8 + 7$, the first 20 rate lanes use this transposed path and only the trailing 7
 3147 bytes in lane 20 use a masked gather. These sequential accesses are prefetcher-friendly,
 3148 whereas per-lane gathers and scatters at the chunk stride β (**VPGATHERQQ**/**VPSCATTERQQ**)
 3149 would stall on cold memory and cap the long-message rate.

3150 The register-resident kernel that drains a two-to-seven-chunk AVX-512 remainder
 3151 instead uses that per-lane gather/scatter form under a mask, reading the active chunks in
 3152 place: the transpose’s shuffle work is constant regardless of the chunk count and would be
 3153 spent partly on inactive lanes, whereas the gathers scale with the active element count.

3154 The eight-way AVX2 core, lacking 512-bit registers, runs the eight instances as two
 3155 groups of four (four 64-bit lanes per YMM register), rotating by shift-or and computing
 3156 χ with **VPANDN**. The AVX2 path drains remainders with a variant of the grouped-by-four
 3157 kernels that aims each inactive instance’s loads and stores at a dummy block inside the
 3158 kernel frame with a zero advance stride, so inactive instances touch no memory beyond
 3159 the remainder and the hot loop stays branch-free.

3160 **arm64.** The **arm64** path uses NEON plus the Armv8.2 SHA-3 extension (**FEAT_SHA3**),
 3161 available on Apple Silicon and on Neoverse V1/V2 and N2. The permutation is expressed
 3162 with the extension’s fused instructions: **EOR3** (three-input XOR) for the θ column parities,
 3163 **RAX1** (rotate-by-one and XOR) for the θ D -terms, **XAR** (XOR then rotate) to fuse the θ
 3164 accumulation with the ρ rotation, and **BCAX** (bit-clear and XOR) for χ . The single-duplex
 3165 core stores each Keccak lane in one 64-bit D lane across the 25 vector registers. The
 3166 eight-way core packs two instances per 128-bit register and processes the batch of eight
 3167 as four such pairs, using **ZIP1** and duplicating moves to interleave and extract the two
 3168 instances of each lane in the fused chunk kernel. The same pair machinery backs a 2-wide
 3169 kernel that drains a leaf-chunk remainder two chunks at a time—one instance per 64-bit
 3170 half of a 128-bit register—at roughly twice the single-duplex per-chunk rate. Both batched
 3171 kernels store the permutation state back into the batch after the closing block—the $\times 8$
 3172 kernel for its first pair, the 2-wide kernel for both instances—which is the writeback that
 3173 permits the root duplex to occupy lane 0.

3174 **Constant-time addressing.** The optimized kernels contain no secret-dependent branches
 3175 and no secret-dependent memory or vector indices. The only indexed table is the round-
 3176 constant array, addressed by the public round number. The AVX-512 steady-state transpose
 3177 uses contiguous loads and stores; the remaining gathers and scatters, for the partial rate
 3178 lane and register-resident remainder kernel, use fixed public chunk strides under masks
 3179 determined by the public chunk count. The AVX2 remainder kernel redirects inactive

instances to a dummy block in its stack frame with a zero advance stride, again selected only by the public chunk count. The `arm64` pair kernels use fixed lane interleaving and extraction patterns. None of these addresses depends on the key, nonce, plaintext, or decrypted plaintext.

D Test Vectors

The following are selected known-answer test vectors for TW128 encryption, as specified in § 3. The complete vector set, including full byte strings for the long cases and the regeneration harness, is maintained with the reference implementation at <https://github.com/codahale/treewrap>.

All byte strings are written in hexadecimal. For each vector we list the key K , nonce U , associated data A , message M , ciphertext C , and tag T ; the sealed output of encryption is $C \parallel T$. The empty string is written ε . The parameters used below are $\rho = 167$ bytes, $\beta = 8183$ bytes, and 32-byte keys, nonces, and tags. The leaf index is encoded as a little-endian 64-bit integer and the phase suffixes are appended least significant bit first.

To keep the longer vectors compact, any value exceeding 32 bytes is given by its length in bytes together with its SHA-256 digest rather than in full. For generated associated data and message inputs, byte i is defined as $x_i = (17i + 31) \bmod 256$ for $i = 0, 1, \dots$, so each such input is determined by its length alone and can be reconstructed and checked against the stated digest; a ciphertext given by its length and digest can likewise be checked by encrypting the corresponding message and hashing the result.

Vector 1. Empty associated data and empty message: the associated data and aggregation phases are both elided, and the root tag is emitted by the single closing message call with no leaves.

K	000102030405060708090a0b0c0d0e0f101112131415161718191a1b1c1d1e1f
U	202122232425262728292a2b2c2d2e2f303132333435363738393a3b3c3d3e3f
A	ε
M	ε
C	ε
T	da65bbda0b20b1e936f5326458166633d1e5a3f7aebc0b318d24b051a4efa697

Vector 2. Multi-block associated data with an empty message: the associated data spans two rate blocks and no message bytes are processed.

K	000102030405060708090a0b0c0d0e0f101112131415161718191a1b1c1d1e1f
U	202122232425262728292a2b2c2d2e2f303132333435363738393a3b3c3d3e3f
A	168 bytes, SHA-256 4c7d125a3488a7603e2f403f0010a6060b8321b60b0fd2471120ae627bbe759a
M	ε
C	ε
T	2093878418f3fca96ad3a8d1748a9be4d2a8b518e917d988fd703dcb7bf72173

Vector 3. Empty associated data and a single-block message: the root message phase only, no leaves.

K	000102030405060708090a0b0c0d0e0f101112131415161718191a1b1c1d1e1f
U	202122232425262728292a2b2c2d2e2f303132333435363738393a3b3c3d3e3f
A	ε
M	1f30415263748596a7b8c9daebfc0d1e2f405162738495a6b7c8d9eafb0c1d2e
C	676a1441e20daea0f1ee4571ee4543842c946d59e0d3df6f1dd2f92dc4f876aa
T	1873332d359ff2fe5197488224af7bf6ad821875fdbd964100de424aa77342be

Vector 4. Empty associated data and a two-block message: the root keystream is chained across two rate blocks.

3203 *K* 000102030405060708090a0b0c0d0e0f101112131415161718191a1b1c1d1e1f
U 202122232425262728292a2b2c2d2e2f303132333435363738393a3b3c3d3e3f
A ϵ
M 168 bytes, SHA-256
 4c7d125a3488a7603e2f403f0010a6060b8321b60b0fd2471120ae627bbe759a
C 168 bytes, SHA-256
 df3f0a98a91a69ed474edbc52e0dc1d8e54bfd6fc514966ce7b878d48be0bf9a
T aab2413074a202cfadd5eb784663b1831eb330025972d163d03e9ed5a964b765

Vector 5. A message of exactly one chunk: the root is full, there is still no leaf, and the aggregation phase is elided so the closing root message call emits the tag.

3204 *K* 000102030405060708090a0b0c0d0e0f101112131415161718191a1b1c1d1e1f
U 202122232425262728292a2b2c2d2e2f303132333435363738393a3b3c3d3e3f
A ϵ
M 8183 bytes, SHA-256
 6532a370a4ab6b5f4094dd8a6016512ab8e8416abe910bc1a0b532979224151d
C 8183 bytes, SHA-256
 8d58f8aa3413ea0ba1b8cc3ed4e9a334521d8dafa9694a0661fb8251049907e9
T c2201637f1116c1303b8982622756770d23acb85384f4658fcb63c37e4b629dd

Vector 6. A message of one chunk and one further byte: one root chunk and one leaf, exercising leaf tag aggregation.

3205 *K* 000102030405060708090a0b0c0d0e0f101112131415161718191a1b1c1d1e1f
U 202122232425262728292a2b2c2d2e2f303132333435363738393a3b3c3d3e3f
A ϵ
M 8184 bytes, SHA-256
 c4c383bca0808e2ce44d481ddfe56efdc7a147dcd5a14c898ef5423f2604cd7
C 8184 bytes, SHA-256
 ef1b2f7607ade503ab497fc61e62cef2b44c6328fa06e499718eb7da5fdd04d3
T c96bea48e07a8431fd19721ffbe8afe6a73047f96d953e3281f10126eeca06a1

Vector 7. A multi-leaf message of two chunks and one further byte: the full tree path with two leaves.

3206 *K* 000102030405060708090a0b0c0d0e0f101112131415161718191a1b1c1d1e1f
U 202122232425262728292a2b2c2d2e2f303132333435363738393a3b3c3d3e3f
A ϵ
M 16367 bytes, SHA-256
 17e31f3700d0602aab8a686024ecc05b8192e80a531483c67dd90d5a26557c5e
C 16367 bytes, SHA-256
 6d83d8fbdca2a7c2ef590d29a36527a45019e50d1f8c8db037959081685a46908
T cdc2e50ac5380419f06772a2f3af2146050bd0340b35f61b67f839351767804e

Vector 8. Multi-block associated data with a short message: the associated data blocks dominate the cost.

3207 *K* 000102030405060708090a0b0c0d0e0f101112131415161718191a1b1c1d1e1f
U 202122232425262728292a2b2c2d2e2f303132333435363738393a3b3c3d3e3f
A 337 bytes, SHA-256
 04ca7b24930bc6b5fe8b95ff1eb0933907657c1283cc3a7e69c5f7f02220f281
M 1f30415263748596a7b8c9daebfc0d1e
C 0dd615a0a05ed82a78c3999b9e80cbcd
T 22e4e6725f548ad0660041d66cf2f18dc45c69b447cdaae1f147d33a3b6aa1a9

Vector 9. An all-ones key and nonce with a multi-chunk message.	
<i>K</i>	ff
<i>U</i>	ff
<i>A</i>	168 bytes, SHA-256 4c7d125a3488a7603e2f403f0010a6060b8321b60b0fd2471120ae627bbe759a
<i>M</i>	8184 bytes, SHA-256 c4c383bca0808e2ce44d481ddfe56efdcd7a147dcd5a14c898ef5423f2604cd7
<i>C</i>	8184 bytes, SHA-256 49818c320ab282296639798d840cbd60751e1623f449dcbf9423371d44145dfb
<i>T</i>	310c9fff20e9651f8ed47eed9bb182a086ed1af1b3f1799e35307a4139865fc7

3208