

TW128: A Tree-Structured Duplex Authenticated Encryption Mode

Coda Hale

Hale Institute for Bicycle Appreciation

Abstract. We present TW128, a nonce-based authenticated encryption mode with associated data built as a two-level tree of keyed duplexes over the KECCAK- p [1600, 12] permutation. A serial root duplex binds the key, nonce, and associated data, encrypts the first message chunk, and aggregates a contribution from every later chunk into one authentication tag; those chunks are handled by independent leaf duplexes that, once the key and nonce are fixed, share no state and can be evaluated in parallel across SIMD lanes. We analyze TW128 in the public permutation model, proving each bound first against an intermediate random oracle and then transferring it to the real permutation by the flat sponge lift. We prove a concrete multi-user indistinguishability bound and derive all-in-one authenticated encryption as its single-user corollary, together with bounds for the commitment security of the root tag, which we obtain by proving that the wrapper returning the tag is a secure hash of the (K, N, A, M) context—collision-, second-preimage-, and preimage-resistant in the standard hash-function sense. The single root tag is thus a compact commitment; full-input committing security (CMTDs-4) follows as a corollary, and we also establish context discoverability security. At the implemented parameters these notions meet a 128-bit security target. Vectorized **amd64** and **arm64** implementations exploit the tree’s leaf parallelism, exceeding 10 Gbit/s and coming within a factor of 1.3 to 2 of hardware-accelerated AES-128-GCM.

Keywords: authenticated encryption · committing security · sponge constructions · duplex constructions · Keccak

1 Introduction

Authenticated encryption with associated data is the workhorse of applied cryptography, and the sponge and duplex constructions over a single public permutation offer an attractive route to it. The duplex of [BDPVA11] turns one permutation—here KECCAK- p [1600, 12], the round-reduced permutation underlying the [Nat15] standard—into a single-pass authenticated cipher with a small, constant memory footprint and no separate block cipher or universal hash key schedule. Its security rests on the indistinguishability of the sponge from a random oracle [BDPVA08], which reduces analysis of the full construction to a random oracle argument plus one capacity birthday loss.

This compactness comes at a structural cost. A duplex is inherently *sequential*: each call absorbs its input into a state produced by the previous call, so a single duplex cannot exploit the wide SIMD and multicore parallelism of modern processors, on which competing designs built from AES reach many bytes per cycle. A second, orthogonal gap is *commitment*. Widely deployed AEADs such as AES-GCM and ChaCha20-Poly1305 are not key-committing: an adversary can craft a single ciphertext that decrypts to valid plaintexts under two or more keys [DGRW18, ADG⁺22]. Far from a curiosity, this enables concrete abuses—the “invisible salamander” that defeats Facebook’s attachment franking scheme [DGRW18], ciphertexts that open to different well-formed files under

different keys [ADG⁺22], and partitioning oracle attacks that recover passwords from systems built on non-committing AEAD [LGR21]. The committing AE line that addresses these threats was initiated by [GLR17] and developed into formal committing notions with efficient generic transforms by [BH22]. Commitment has two facets: committing security is a collision requirement—no two contexts open one ciphertext—while its preimage counterpart, *context discoverability security* [MLGR23], forbids an adversary from *finding* even one decrypting context for a target ciphertext. The latter’s absence underlies the recent context discovery attacks that break GCM, CCM, EAX, SIV, and OCB3 [MLGR23], all of which delegate authenticity to an algebraically invertible MAC. A mode that targets these notions natively, rather than through an add-on transform, avoids the transform’s overhead and binds commitment to the same permutation it already uses for confidentiality and integrity.

TW128 (*Tree Wrap*) addresses both gaps with a two-level tree of keyed duplexes. A serial *root* duplex, initialized with the key K and nonce N , absorbs the associated data, encrypts the first message chunk, and absorbs a tag from every other chunk before emitting the final authentication tag. The remaining chunks are processed by independent *leaf* duplexes, each keyed by (K, N, j) and operating on disjoint state and disjoint input. Because the leaves share nothing once (K, N) is fixed, they can be evaluated in parallel across SIMD lanes, recovering the throughput a flat duplex forgoes; and because the root binds the associated data, the first chunk, and every leaf tag into the single output tag, that tag is the output of the sponge on the entire (K, N, A, M) input. The structure ensures only that distinct (K, N, A) triples reach distinct root transcripts; the committing and context discoverability security rests on the indistinguishability of the sponge, which makes the root tag both collision- and preimage-resistant. Because both hold for the tag projection alone, the single 256-bit root tag is itself a secure, compact cryptographic commitment to the full (K, N, A, M) context (Section 6). Every duplex still operates on a bounded 1600-bit state, and the leaf tags are absorbed into the root as the leaves complete rather than buffered together, so the memory footprint remains constant in the message length.

We analyze TW128 in the public permutation model and obtain a concrete multi-user indistinguishability bound in the sense of [BT16], with all-in-one authenticated encryption in the sense of [RS06] as its single-user corollary. We also prove commitment security of the root tag, stated as the [RS04] hash security of the tag wrapper—collision, second-preimage, and preimage resistance—from which the [BH22] full-input committing notion (CMTDs-4) follows as a corollary, alongside the [MLGR23] context discoverability notion. Each bound is proved first against an intermediate random oracle and then transferred to the real permutation by the flat sponge lift of [BDPVA08]. At the implemented parameters ($k = c = \tau_{\text{root}} = \tau_{\text{leaf}} = 256$), these guarantees meet a 128-bit security target for adversaries running up to $N_{\text{tot}} \leq 2^{63}$ KECCAK- p [1600, 12] calls—on the order of 2^{71} bytes of processed data, beyond any realistic deployment lifetime. Optimized amd64 and arm64 implementations turn the tree’s leaf parallelism into measured throughput within a small factor of hardware-accelerated AES-128-GCM (Section 7).

1.1 Our Contributions

- **The TW128 mode.** We specify TW128, a nonce-based AEAD with associated data built as a two-level tree of strict [BDPVA11] duplexes over KECCAK- p [1600, 12] (Section 3): a serial root duplex that binds (K, N, A) , the first chunk, and the leaf tags into a single authentication tag, over a set of parallel, independently keyed leaf duplexes. All phases are domain-separated within one permutation, and the state is constant-size regardless of message length.
- **Multi-user security.** We give a concrete multi-user indistinguishability bound [BT16] for D users (Theorem 2), obtained by a key-collision-stopped hybrid over the

users and a direct adjacent-user coupling in the random oracle model.

- **All-in-one AE security.** We derive a concrete bound on the all-in-one AE advantage [RS06] in the public permutation model (Corollary 1) as the $D = 1$ case of the multi-user indistinguishability theorem.
- **Commitment via a secure hash.** We show that the wrapper H returning the root tag is a secure hash of the full (K, N, A, M) input in the sense of [RS04]—collision-, second-preimage-, and preimage-resistant (Theorems 3 to 5)—so the single 256-bit root tag is a binding, preimage-resistant compact commitment to the input, matching the compactly committing AEAD goal of [GLR17] natively, with no generic transform. The [BH22] full-input committing notions follow as corollaries: s -way CMTDs-4 from collision resistance (Corollary 3), and the remaining CMT/CMTD for $\ell \in \{1, 3, 4\}$ in Corollary 6.
- **Context discoverability security.** We additionally prove that TW128 satisfies the published [MLGR23] context discoverability notion CDY^* and its per-component variants—a ciphertext-level guarantee, distinct from preimage resistance of the tag, established with a concrete bound (Theorem 6)—so no adversary, given a target ciphertext fixed in advance, can find a decrypting context. This is the facet whose absence underlies the discovery attacks of [MLGR23]; Section 1.2 contrasts TW128’s sponge-based resistance with the algebraic MAC inversion those attacks exploit.
- **Proof methodology.** We carry out the analysis in a dedicated TW128 random oracle model and transfer every bound to the public permutation setting through the [BDPVA08] flat sponge lift (Section B). A derived-adversary argument shows that each direct primitive query induces at most one random oracle query, so the random oracle bounds are instantiated at the direct-query count N_{prim} rather than the total work N_{tot} , separating the single capacity birthday loss $\text{Lift}_c(N_{\text{tot}})$ from the scheme-specific terms.
- **Optimized implementations.** We describe and evaluate vectorized amd64 (AVX-512 and AVX2) and arm64 (NEON with the SHA-3 extension) implementations whose batched leaf cores realize the tree’s parallelism in practice, exceeding 10 Gbit/s on contemporary hardware (Section 7).

1.2 Related Work

Sponge and duplex AEADs. Permutation-based authenticated encryption built on the sponge and duplex constructions of [BDPVA08, BDPVA11] is well established. Keyak [BDP⁺15] is a CAESAR candidate built on the duplex over round-reduced KECCAK- p , whose Motorist mode already admits parallel duplex instances. Ascon [DEMS21], the lightweight winner of the CAESAR competition and the basis of the NIST lightweight-cryptography standard, and Xoodyak [DHP⁺20], a finalist in that process whose Cyclist mode is a deliberately serial simplification of Motorist, are duplex AEADs over the Ascon and Xoodoo permutations. TW128 shares this single-permutation, constant-state duplex foundation but differs in structure: rather than one duplex session, it is a two-level tree whose leaves are *independently keyed* duplexes under domain-separated chunk identifiers. Like Keyak’s Motorist mode, the tree admits data-parallel evaluation across SIMD lanes; the designs differ in where the degree of parallelism lives. Keyak fixes it as a mode parameter—each named instance sets the number Π of parallel duplex instances across which a message is spread, so Π is part of the wire format and must be agreed by the encryptor and decryptor in advance. In TW128 the leaf decomposition is fixed by the mode, but the number of leaves evaluated simultaneously is not part of the transcript: an implementation chooses it at run time to match the host’s vector width, and the ciphertext

and tag are identical however it is chosen, so a wide-SIMD encryptor and a scalar decryptor interoperate.

Tree hashing. The two-level shape is inspired by tree hashing over the same permutation, in particular KangarooTwelve [BDP⁺16], which splits its input into fixed-size chunks hashed independently under KECCAK- p [1600, 12] and absorbs their chaining values into a root node. TW128 reuses the round count and the final-node/leaf split of that design and adapts it to authenticated encryption: each leaf chunk is encrypted and tagged independently, and the leaf tags are absorbed by the root in place of chaining values. It also inherits the short-message behavior of KangarooTwelve’s *kangaroo hopping*, which merges the first chunk with the hashing of the remaining chunks’ chaining values so that an input fitting in a single chunk incurs no tree overhead: in TW128 a single-chunk message engages no leaf duplexes and no aggregation phase, the root processing it directly and emitting the tag as a plain keyed duplex would. Short messages therefore pay none of the tree’s cost and retain the latency profile of an unstructured duplex (Section 7.4); the leaf layer, and its parallelism, engage only once a message spans more than one chunk.

Committing AEAD. The study of committing authenticated encryption was initiated by [GLR17], who introduced compactly committing AEAD for message franking, and by [DGRW18], who built it from a dedicated primitive (encryptment); [BH22] subsequently organized the notions into the CMT-1/3/4 and CMTD hierarchy and gave efficient generic transforms (UtC, RtC, HtE) that add commitment to a base scheme. A complementary line shows that widely used and dedicated AEADs fail these notions: the invisible salamander [DGRW18] and partitioning oracle [LGR21] attacks, multi-format ciphertexts [ADG⁺22], the context discovery and commitment attacks that break CCM, EAX, SIV, GCM, and OCB3 [MLGR23], and an $O(1)$ key-committing attack against all variants of AEGIS [WP14] that refutes the prior belief that AEGIS was committing [IR23]. [MLGR23] also introduce context discoverability, the preimage facet of commitment, and show that their discovery attacks exploit schemes whose authenticity is delegated to a non-preimage-resistant MAC such as GHASH or CMAC. TW128 instead makes the root tag itself collision- and preimage-resistant as a sponge output, rather than delegating to a transform or an invertible MAC, so the [MLGR23] discovery attacks have no analogue against it.

Closest to TW128 is the recent work of [DHM⁺24], which likewise builds nonce-based AE natively on the SHA-3 functions and obtains CMT-4 from the collision resistance of the underlying function rather than from a transform; it targets session-based communication through a sequential duplex chain, where TW128 is a parallel two-level tree. Sections 7 and 8.3 detail the structural and throughput differences.

1.3 Paper Organization

Section 2 recalls the notation, sponge and duplex constructions over KECCAK- p , the authenticated encryption, committing security [BH22], and context discoverability [MLGR23] notions used throughout, and the multi-user indistinguishability notion of [BT16]. Section 3 specifies the TW128 construction and its parameters. Section 4 fixes the public permutation games, the intermediate random oracle model, and the resource accounting that all later bounds reference. Section 5 states the headline mu-ind, AE, hash (collision, second-preimage, preimage), CMTDs-4, and context discoverability theorems and their concrete TW128-parameter forms; Section 6 proves the random oracle hash and CMTDs-4 bounds. Section 7 reports the performance of the reference and optimized implementations. Section 8 compares TW128 to prior sponge AEADs and committing constructions and discusses limitations. Section 9 concludes. Section A collects the structural decoding and accounting lemmas that form the proof core; Section B carries out the [BDPVA08] flat

190 sponge lift that converts each random oracle bound to its public permutation form; and
 191 Section D lists TW128 test vectors.

192 2 Preliminaries

193 2.1 Notation

194 **Strings.** Write $\{0, 1\}$ for the binary alphabet. For an integer $n \geq 0$, $\{0, 1\}^n$ denotes the
 195 set of n -bit strings, $\{0, 1\}^* = \bigcup_{n \geq 0} \{0, 1\}^n$ the set of all finite-length bit strings, and ε
 196 the empty string. For $x \in \{0, 1\}^*$, $|x|$ is the bit length and $x \parallel y$ is the concatenation;
 197 the bitwise XOR of two equal-length strings $x, y \in \{0, 1\}^n$ is $x \oplus y$. A *byte string* is an
 198 element of $(\{0, 1\}^8)^*$; its byte length is denoted $\|x\| = |x|/8$. For an integer $0 \leq n < 2^{64}$,
 199 $\langle n \rangle_8 \in \{0, 1\}^{64}$ is its eight-byte little-endian encoding. For a finite or infinite bit string x
 200 and $0 \leq \tau \leq |x|$ when x is finite, $\lfloor x \rfloor_\tau \in \{0, 1\}^\tau$ denotes the leading τ bits of x .

201 **Sampling and probability.** We write $x \leftarrow \$ S$ for sampling x uniformly at random from a
 202 finite set S , with implicit coins independent across samples. For an event E in a probability
 203 experiment, $\Pr[E]$ denotes its probability.

204 **Sets of maps.** For a set S , $\text{Perm}(S)$ denotes the set of permutations on S ; for a positive
 205 integer b , $\text{Perm}(\{0, 1\}^b)$ is the set of b -bit permutations.

206 **Games and advantages.** We write $\mathcal{A}^{O_1, \dots, O_k}$ for a (possibly randomized) adversary with
 207 oracle access to $O_1, \dots, O_k$, and $\Pr[\mathbf{G}^{\mathcal{A}} \Rightarrow 1]$ for the probability that a game $\mathbf{G}$ —an
 208 experiment that runs $\mathcal{A}$ against specified oracles and returns a Boolean outcome—returns
 209 1. For a distinguishing experiment between two games $\mathbf{G}_0, \mathbf{G}_1$, the advantage of $\mathcal{A}$ against
 210 a scheme Π in notion $\mathbf{X}$ is the absolute difference

$$211 \quad \text{Adv}_{\Pi}^{\mathbf{X}}(\mathcal{A}) = |\Pr[\mathbf{G}_0^{\mathcal{A}} \Rightarrow 1] - \Pr[\mathbf{G}_1^{\mathcal{A}} \Rightarrow 1]|.$$

212 For a win-event experiment without a hidden challenge bit it is instead $\Pr[\mathbf{G}^{\mathcal{A}} \text{ wins}]$, the
 213 probability the game returns 1; a bit-guessing experiment with a hidden challenge bit uses
 214 the rescaled form $|\Pr[\mathbf{G}^{\mathcal{A}} \text{ wins}] - 1|$, as in the multi-user notion of Section 2.4.

215 **Concrete security.** All bounds in this paper are concrete: each advantage is an explicit
 216 function of the resource caps introduced in Section 4.3. We make no use of asymptotic or
 217 negligibility conventions.

218 **Reference summary.** Table 1 summarizes the resource caps and loss terms used through-
 219 out. The caps are defined formally in Section 4.3, the loss terms in Section 5.1.

220 2.2 Sponges and Duplexes over Keccak- p

221 **The Keccak- p permutation.** The KECCAK- p family of permutations, introduced as part
 222 of the Keccak design and standardized within FIPS 202 [Nat15], is parameterized by
 223 a state width b and a round count n_r ; we write $\text{KECCAK-}p[b, n_r] \in \text{Perm}(\{0, 1\}^b)$ for
 224 the corresponding permutation. TW128 uses $\text{KECCAK-}p[1600, 12]$, the same primitive as
 225 KangarooTwelve [BDP⁺16], fixed throughout. In the public permutation security model of
 226 Section 4.1 and the flat sponge lift of Section B, this permutation is modeled as a uniformly
 227 random element of $\text{Perm}(\{0, 1\}^{1600})$.

Table 1: Resource caps and loss terms.

Symbol	Meaning	Defined
N	Construction interface cost (KECCAK- $p[1600, 12]$ calls)	§ 4.3
N_{prim}	Direct primitive-interface queries	§ 4.3
N_{tot}	$N + N_{\text{prim}}$ (total query-sequence cost)	§ 4.3
q_v, Q_v	Valid verification queries / cap	§ 4.3
$q_{\text{RO}}, Q_{\text{RO}}$	Direct random oracle queries / cap (RO-model)	§ 4.3
$n_{\text{leaf}}, N_{\text{leaf}}$	Distinct construction-generated final leaf transcripts / cap	§ 4.3
$\text{Lift}_c(N_{\text{tot}})$	[BDPVA08] flat sponge loss	§ 5.1
$\text{KeyGuess}_k(Q)$	Ideal system key guess	§ 5.1
$\text{RootColl}_\tau(Q, s)$	s -way root-collision	§ 5.1
$\text{RootPre}_\tau(Q)$	Root-preimage	§ 5.1
$\text{LeafColl}_\tau(Q)$	Known-key leaf-collision	§ 5.1

The sponge construction. Following [BDPVA08], fix a permutation $\pi \in \text{Perm}(\{0, 1\}^b)$, a *sponge-compliant* padding rule pad , and a rate $1 \leq r < b$ with capacity $c = b - r$. A padding rule appends $\text{pad}[r](|M|)$, a bit string determined by the input length and the rate, to an input M , extending it to a sequence of r -bit blocks; it is sponge-compliant [BDPVA11, Definition 1] if $M \parallel \text{pad}[r](|M|)$ is never empty and, for all $M \neq M'$ and all $n \geq 0$, $M \parallel \text{pad}[r](|M|) \neq M' \parallel \text{pad}[r](|M'|) \parallel 0^{nr}$. The sponge function $\text{Sponge}[\pi, \text{pad}, r] : \{0, 1\}^* \times \mathbb{N} \rightarrow \{0, 1\}^*$ maps an input M and an output length ℓ to an ℓ -bit string by absorbing $M \parallel \text{pad}[r](|M|)$ into the all-zero b -bit state in r -bit blocks (each separated by an application of π) and squeezing ℓ bits in r -bit blocks from the resulting state. [BDPVA08, Theorem 2] shows that $\text{Sponge}[\pi, \text{pad}10^*, r]$, where the simple rule $\text{pad}10^*$ appends a single 1 bit followed by the minimum number of 0 bits reaching a multiple of r , is indistinguishable from a random oracle when π is uniformly random. We bound its distinguishing advantage by $\text{Lift}_c(N_{\text{tot}}) = N_{\text{tot}}(N_{\text{tot}} + 1)/2^{c+1}$, derived in Lemma 16 from the exact differentiating probabilities of [BDPVA08, Lemmas 4 and 5].

The duplex construction. The *duplex* construction of [BDPVA11] extends the sponge to a stateful interface. A duplex object $\text{Duplex}[\pi, \text{pad}, r]$ is initialized to the all-zero state. Each call $\text{duplexing}(\sigma, \ell)$, with $\ell \leq r$, absorbs an input string σ of at most $\rho_{\max}(\text{pad}, r)$ bits via π and returns the first ℓ bits of the resulting state. Here $\rho_{\max}(\text{pad}, r) = \min\{x : x + |\text{pad}[r](x)| > r\} - 1$ is the largest input length for which the padded input fits in a single r -bit block; for the $\text{pad}10^*1$ rule, $\rho_{\max}(\text{pad}10^*1, r) = r - 2$ [BDPVA11]. [BDPVA11, Lemma 3] establishes the duplex-to-sponge equality: for any call sequence $((\sigma_0, \ell_0), \dots, (\sigma_i, \ell_i))$, the i -th output equals

$$\text{Sponge}[\pi, \text{pad}, r](\sigma_0 \parallel \text{pad}_0 \parallel \dots \parallel \sigma_i, \ell_i),$$

where $\text{pad}_j = \text{pad}[r](|\sigma_j|)$ is the padding the duplex appends in its j -th call; the flattened input carries these interior paddings explicitly, and the sponge appends only the final pad_i . The flattening map is injective by [BDPVA11, Lemma 4]: the duplex-call sequence is recoverable from the flattened sponge input.

TW128 instantiation. TW128 fixes the padding rule $\text{pad}10^*1$ and the rate $r = 1344$, so the capacity is $c = 256$ bits. The largest input length per duplex call is therefore $\rho_{\max}(\text{pad}10^*1, 1344) = 1342$ bits.

Bridge to the H -input domain. The indistinguishability theorem of [BDPVA08] is stated for the simpler $\text{pad}10^*$ padding over its unpadded H -input domain, while TW128 uses $\text{pad}10^*1$ internally. [BDPVA11, Theorem 3] supplies an injective preprocessing bridge

261 $I[r, r_{\max}]$ between the two domains. For TW128, $r = r_{\max} = 1344$, so the bridge reduces
 262 to appending $1 \parallel 0^q$ with q determined by the input length modulo $r_{\max}$; on the bridge
 263 image, q is recovered as the trailing-zero count. Section 4.2 instantiates this bridge on
 264 typed TW128 duplex transcript prefixes as the map $\text{flat}(\cdot)$ used by the random oracle
 265 proofs throughout the paper.

266 2.3 Authenticated Encryption with Associated Data

267 **Syntax.** A *nonce-based authenticated encryption scheme with associated data* (AEAD
 268 scheme) is a pair $\text{SE} = (\text{Enc}, \text{Dec})$ of deterministic algorithms with associated key space
 269 $\mathcal{K} \subseteq (\{0, 1\}^8)^*$, nonce space $\mathcal{N} \subseteq (\{0, 1\}^8)^*$, associated data space $\mathcal{A} \subseteq (\{0, 1\}^8)^*$, message
 270 space $\mathcal{M} \subseteq (\{0, 1\}^8)^*$, and ciphertext space $\mathcal{C} \subseteq (\{0, 1\}^8)^*$. The encryption algorithm
 271 $\text{Enc} : \mathcal{K} \times \mathcal{N} \times \mathcal{A} \times \mathcal{M} \rightarrow \mathcal{C}$ produces a ciphertext $C = \text{Enc}_K(N, A, M)$. The decryption
 272 algorithm $\text{Dec} : \mathcal{K} \times \mathcal{N} \times \mathcal{A} \times \mathcal{C} \rightarrow \mathcal{M} \cup \{\perp\}$ returns either a plaintext or the rejection
 273 symbol $\perp$. The *tag length in bytes* is the constant $\|\text{Enc}_K(N, A, M)\| - \|M\|$, fixed by the
 274 scheme. In this generic syntax, C denotes the full AEAD ciphertext, including any tag
 275 material. The TW128 algorithm section (Section 3) writes this same value as $C \parallel T$, with
 276 C the ciphertext body and T the root tag.

277 **Correctness.** For every $(K, N, A, M) \in \mathcal{K} \times \mathcal{N} \times \mathcal{A} \times \mathcal{M}$, $\text{Dec}_K(N, A, \text{Enc}_K(N, A, M)) =$
 278 M .

279 **All-in-one AE.** We use the all-in-one authenticated encryption notion of [RS06], instan-
 280 tiated for nonce-based AE following [BT16]. The experiment grants either the real pair
 281 $(\text{Enc}_K, \text{Ver}_K)$ or the ideal pair $(\text{Rand}, \text{Replay})$, where Rand returns a fresh uniform byte
 282 string of the ciphertext length on each encryption query and Replay accepts exactly the
 283 tuples Rand produced. The advantage $\text{Adv}_{\text{SE}}^{\text{ae}}(\mathcal{A})$ is the distinguishing advantage between
 284 these two worlds over nonce-respecting adversaries. This all-in-one notion is instantiated
 285 for TW128 in Section 4.1.

286 2.4 Multi-User Indistinguishability

287 The multi-user indistinguishability notion of [BT16, Fig. 3] extends the same real-vs-ideal
 288 oracle pair to an adversary-selected set of users. The number of users u is adversary-
 289 determined: it is the number of **New** queries the adversary makes, i.e., the final value of
 290 the user counter v in the game below. A challenge bit $b \leftarrow \{0, 1\}$ selects between the real
 291 construction interface (real ciphertexts, real verification) and an ideal interface (uniform
 292 strings of matching length, verification that accepts only previously-issued encryption
 293 tuples). The adversary also accesses a public primitive interface: the real permutation
 294 (π, π^{-1}) in the real world and the simulator $(\mathcal{S}^{\text{RO}_{\text{flat}}}, \mathcal{S}^{-1, \text{RO}_{\text{flat}}})$ in the ideal world. The
 295 adversary outputs a guess b' and wins if $b' = b$.

296 **Game and advantage.** The game $\text{G}_{\text{SE}}^{\text{mu-ind}}(\mathcal{A})$ of [BT16, Fig. 3] samples $b \leftarrow \{0, 1\}$ and
 297 maintains a user counter v , a set V of (d, N) pairs already used by encryption, and a set
 298 W of recorded encryption tuples; its **New**, **Enc**, and **Vf** oracles, together with the public
 299 primitive interface—the real permutation when $b = 1$ and the simulator when $b = 0$ —are
 300 instantiated for TW128 in Section 4.4. We write the user index as d and the associated
 301 data input as A , where [BT16] writes i and H , and take the absolute value of their signed
 302 advantage so that

$$303 \text{Adv}_{\text{SE}}^{\text{mu-ind}}(\mathcal{A}) = |2 \Pr[\text{G}_{\text{SE}}^{\text{mu-ind}}(\mathcal{A}) \text{ returns } 1] - 1| = |\Pr[\mathcal{A}^{\text{Enc}, \text{Vf}, b=1} \Rightarrow 1] - \Pr[\mathcal{A}^{\text{Enc}, \text{Vf}, b=0} \Rightarrow 1]|,$$

the second equality holding because b is uniform. Thus $\text{Adv}_{\text{SE}}^{\text{mu-ind}}$ is the two-game distinguishing advantage between the real (Enc, Ver) pair and the ideal (Rand, Replay) pair instantiated per user. The single-user case $u = 1$ recovers the all-in-one AE notion of Section 2.3, augmented with the public primitive interface (the real permutation in the real world and the simulator in the ideal world).

Per-user nonce uniqueness. Nonce uniqueness is enforced per user via the set V : the same nonce may appear under distinct users but not twice under the same user. This is strictly weaker than the global nonce-uniqueness requirement that would force every (d, N) pair across all users to be unique.

2.5 Committing and Context Discoverability Notions

[BH22] formalize a family of committing security notions for nonce-based AEAD schemes, parameterized by how much of the encryption input is committed and by whether commitment is established through decryption (the D-notions) or through encryption agreement (the E-notions). We recall their definitions and the s -way generalization, the complementary *context discoverability* notion of [MLGR23], and then frame the root tag itself as a hash function whose [RS04] security—collision, second-preimage, and preimage resistance—captures the binding and preimage facets of the tag as a compact commitment and yields the [BH22] committing notions as a consequence.

What-is-committed functions. For $\ell \in \{1, 3, 4\}$, the *what-is-committed* function WiC_ℓ projects an encryption input $(K, N, A, M) \in \mathcal{K} \times \mathcal{N} \times \mathcal{A} \times \mathcal{M}$ to the portion it should commit:

$$\begin{aligned}\text{WiC}_1(K, N, A, M) &= K, \\ \text{WiC}_3(K, N, A, M) &= (K, N, A), \\ \text{WiC}_4(K, N, A, M) &= (K, N, A, M).\end{aligned}$$

CMTDs- ℓ and CMTs- ℓ games. For $s \geq 2$, [BH22, Fig. 5] defines the s -way committing games on tuples $(K_i, N_i, A_i, M_i)_{i=1}^s$ with pairwise distinct WiC_ℓ -images, all entries in the scheme’s syntactic domain and none equal to $\perp$. The s -way CMTDs- ℓ (decryption) game has the adversary output a ciphertext C alongside the tuples and returns 1 if $\text{Dec}_{K_i}(N_i, A_i, C) = M_i$ for every i ; the s -way CMTs- ℓ (encryption) game has the adversary output the tuples alone and returns 1 if all encryptions $\text{Enc}_{K_i}(N_i, A_i, M_i)$ are equal. Advantages $\text{Adv}_{\text{SE}}^{\text{CMTDs-}\ell}(\mathcal{A})$ and $\text{Adv}_{\text{SE}}^{\text{CMTs-}\ell}(\mathcal{A})$ are the probabilities that the adversary wins the respective game; the parameter s is suppressed from the notation when clear from context. The two-tuple case $s = 2$ is the original CMTD- ℓ /CMT- ℓ pair of [BH22, Fig. 4]. These definitions are instantiated for TW128 in Section 4.1, with the public permutation lift deferred to that section.

The root tag as a hash. Many uses of commitment store, transmit, or compare a short tag rather than the full ciphertext, and ask that the short string alone be binding—the compactly committing AEAD goal introduced by [GLR17] for message franking. TW128 ciphertexts are written $C \parallel T$ with T the τ_{root} -bit root tag, and we capture this goal directly, by treating the map from the encryption input to the tag as a hash function: write $H(K, N, A, M)$ for the root tag T of $\text{Enc}_K(N, A, M) = C \parallel T$, the TW128 wrapper that discards the ciphertext body and returns only the tag. Keyed by the permutation, this is a hash family, and we ask that it satisfy the hash-function security notions of [RS04], over the TW128 input domain:

Collision resistance (Coll). No adversary finds distinct inputs $X_1, \dots, X_s$ (any $s \geq 2$) with $H(X_1) = \dots = H(X_s)$.

348 **Always second-preimage resistance (aSec).** For the fixed deployed permutation, given a
 349 challenge input X , no adversary finds $X' \neq X$ with $H(X') = H(X)$.

350 **Everywhere preimage resistance (ePre).** For a target tag T^* fixed in advance, no adver-
 351 sary finds X with $H(X) = T^*$.

352 Their advantages $\text{Adv}_{\text{TW128}}^{\text{Coll}}$, $\text{Adv}_{\text{TW128}}^{\text{aSec}}$, $\text{Adv}_{\text{TW128}}^{\text{ePre}}$ are the respective success probabilities,
 353 instantiated for TW128 in Section 4.1. Coll is the binding facet and Pre the preimage
 354 facet of the tag taken as a compact commitment to (K, N, A, M) —the two guarantees a
 355 cryptographic hash commitment provides. Of the [RS04] second-preimage notions, collision
 356 resistance already implies the everywhere and average-case forms (eSec and Sec), so we
 357 prove the always-form aSec, the one it does not imply, and with a sharper bound. Collision
 358 resistance moreover gives the [BH22] committing notions directly: a single ciphertext that
 359 two contexts both decrypt is a fortiori a pair of distinct inputs sharing a tag, so Coll
 360 bounds CMTDs- ℓ and CMTs- ℓ (Section 5.2).

361 **Relations.** [BH22, Figs. 4 and 5] record the implications among the notions. For the
 362 D-notions we use:

$$\text{CMTDs-4} \Rightarrow \text{CMTDs-}\ell \text{ for } \ell \in \{1, 3\}$$

364 (both immediate: pairwise distinct keys ($\ell = 1$) or pairwise distinct (K, N, A) triples ($\ell = 3$)
 365 are *a fortiori* pairwise distinct full tuples, so every CMTDs- ℓ winner is a CMTDs-4 winner).
 366 The E-notions are bounded directly in Corollary 6 by the same final-root candidate-collision
 367 argument, rather than through an encryption reduction.

368 **Context discoverability.** Where committing security forbids two contexts from colliding on
 369 one ciphertext, *context discoverability security* [MLGR23] forbids an adversary from *finding*
 370 a decrypting context for a target ciphertext: it is to committing security what preimage
 371 resistance is to collision resistance. [MLGR23, Fig. 5] formalize it as the game $\text{CDY}[\text{ts}, S]$,
 372 in which a query-independent context selector S —a randomized algorithm taking no input—
 373 produces a target ciphertext together with the context components named by a target
 374 specifier $\text{ts} \subseteq \{k, n, a\}$, and the adversary, given the revealed components, supplies the
 375 remaining components of a decryption context (K, N, A) and wins if the target ciphertext
 376 decrypts under it without error. The *restricted* notion $\text{CDY}^*[\text{ts}]$ takes the worst case over
 377 selectors S ; its three per-component variants CDY_k^* , CDY_n^* , CDY_a^* pin two of the three
 378 context components and have the adversary supply the third ($\text{ts} = \{n, a\}, \{k, a\}, \{k, n\}$
 379 respectively). In every variant the selector fixes the *entire* target ciphertext. [MLGR23]
 380 show that schemes which delegate authenticity to a non-preimage-resistant MAC, including
 381 CCM, EAX, SIV, GCM, and OCB3, admit efficient context discovery attacks; the preimage
 382 character of the notion is what makes TW128’s root tag, a sponge output, resist them.

383 The preimage resistance of H (ePre, above) is the tag-level preimage guarantee that
 384 the [MLGR23] context discoverability notion is the ciphertext-level counterpart of; the
 385 two are distinct, and TW128 is shown to satisfy both (Theorems 5 and 6).

366 3 The TW128 Construction

387 3.1 Design Overview

388 TW128 is a nonce-based authenticated encryption scheme with associated data, built as a
 389 two-level tree of duplex objects over the KECCAK- $p[1600, 12]$ permutation. A single *root*
 390 duplex, initialized with the key K and nonce N , absorbs the associated data A when it
 391 is nonempty, emits keystream for the first plaintext chunk M_0 of at most β bytes, and
 392 absorbs the resulting ciphertext chunk C_0 ; when A is empty the associated data phase




Figure 1: Schematic of TW128 for $n \geq 1$ leaves. The root duplex (bottom row) is initialized with (K, N) and absorbs in turn the associated data A , the root ciphertext chunk C_0 , each leaf tag $T_1, \dots, T_n$, and the eight-byte little-endian leaf count $\langle n \rangle_8$; it emits the root tag T . Each leaf duplex (top row) is initialized with (K, N, j) , emits keystream for plaintext chunk M_j , absorbs the resulting ciphertext chunk C_j , and produces T_j . Leaves can be evaluated in parallel. The ciphertext chunks $C_0, C_1, \dots, C_n$ are obtained by XORing the keystream emitted by the corresponding duplex with the plaintext chunks $M_0, M_1, \dots, M_n$. When $n = 0$ the diagram reduces to the bottom row without any T_j boxes or aggregation block: the aggregation phase is elided and the closing message phase call on C_0 emits the root tag T directly. When the associated data is empty the A block is absent: the associated data phase is elided and the initialization call emits the initial M_0 keystream directly.

393 is elided and the initialization call itself emits the initial M_0 keystream; the remaining
 394 plaintext is partitioned into *leaf* chunks $M_1, \dots, M_n$, each handled by an independent leaf
 395 duplex initialized with (K, N, j) . The leaf duplex emits the keystream for M_j , absorbs
 396 the resulting ciphertext chunk C_j , and produces a τ_{leaf} -bit leaf tag T_j . When $n \geq 1$ the
 397 leaf tags, together with the eight-byte little-endian leaf count $\langle n \rangle_8$, are absorbed into the
 398 root duplex, which emits a single τ_{root} -bit root tag T ; when $n = 0$ the aggregation phase is
 399 elided and the closing message phase call of the root chunk emits T directly, mirroring the
 400 way a leaf duplex emits its leaf tag. The construction is depicted in Figure 1.

401 Three design properties motivate the tree shape. First, the leaf duplexes operate on
 402 disjoint state and consume disjoint inputs once (K, N) is fixed, so leaves $1, \dots, n$ can be
 403 evaluated in parallel across SIMD lanes; Section 7 quantifies the resulting throughput on
 404 AMD64 and ARM64 hardware. Second, every duplex operates on a bounded 1600-bit
 405 state regardless of message length, and the leaf tags are aggregated as a stream rather
 406 than buffered in full (Section 3.5), so the working memory is constant in $\|M\|$ at a fixed
 407 degree of parallelism, scaling only with the number of leaves evaluated simultaneously.
 408 Third, the explicit aggregation of leaf tags through a single deterministic root transcript
 409 gives the construction its commitment-security guarantees: by Opening Triple Separation
 410 (Proposition 1 and § 6), distinct (K, N, A) triples force distinct final root transcripts before
 411 aggregation values matter, the structural fact behind the CMTDs-4 headline of Section 5.2.
 412 Collision resistance of the root tag wrapper (Theorem 3) treats the emitted root tag as
 413 the commitment string, distinct bodies merging only through its leaf tag collision term,
 414 and the same root tag, preimage-resistant as a sponge output, gives the preimage headline
 415 (Theorem 5).

3.2 Parameters and Domain

Table 2 lists the user-visible and derived scheme parameters. The implementation uses the same byte length for the root tag and each leaf tag, but the proofs of Sections 5 and 6 keep τ_{root} and τ_{leaf} as separate symbols.

Table 2: TW128 parameters.

Parameter	Value	Description
k	256	Key length (bits)
ν	256	Nonce length (bits)
τ_{root}	256	Root tag length (bits)
τ_{leaf}	256	Leaf tag length (bits)
β	8183	Plaintext chunk size (bytes)
M_{max}	$8183 \cdot 2^{64}$ bytes	Maximum plaintext length
A_{max}	$8183 \cdot 2^{64}$ bytes	Maximum associated data length
b	1600	Permutation state width (bits)
n_r	12	Permutation rounds
r	1344	Sponge rate (bits)
c	256	Sponge capacity (bits)
ρ	167	Maximum byte-aligned absorption per duplex call (bytes)

The bound $M_{\text{max}} = 8183 \cdot 2^{64}$ bytes is the largest plaintext whose leaf index fits in $\langle \cdot \rangle_8$: writing $\text{leaves}(M) = \max(0, \lceil \|M\|/\beta \rceil - 1)$, every supported plaintext satisfies $0 \leq \text{leaves}(M) \leq 2^{64} - 1$, so leaf indices lie in $\{1, \dots, 2^{64} - 1\}$. The associated data bound A_{max} is a fixed scheme parameter taken equal to M_{max} unless stated otherwise. Inputs exceeding these bounds are outside the scheme domain; every game in Section 4 rejects domain-violating queries.

The derived per-call absorption bound ρ follows the [BDPVA11] analysis: with $r = 1344$ and pad10^*1 padding, $\rho_{\text{max}}(\text{pad10}^*1, r) = r - 2 = 1342$ bits. The byte-aligned data field shares this allotment with the four-bit phase suffix of Section 3.4, so ρ is the largest integer satisfying $8\rho + 4 \leq 1342$, giving $\rho = 167$ data bytes (and $8\rho + 4 = 1340 \leq 1342$).

3.3 Parameter Choice Rationale

The four free parameters c , k , $\tau_{\text{root}} = \tau_{\text{leaf}}$, and β (the rate $r = b - c$ is then determined by the capacity) are chosen to give 128-bit single-key security against every term of the concrete bounds derived in Section 5.3 while keeping per-byte cost low; the nonce width ν is fixed separately, by the headroom the wide permutation affords. Specifically:

- $c = 256$. The dominant single-key term is the [BDPVA08] flat sponge loss

$$\text{Lift}_c(N_{\text{tot}}) = N_{\text{tot}}(N_{\text{tot}} + 1)/2^{c+1},$$

which reaches the 2^{-128} threshold at $N_{\text{tot}} \approx 2^{(c-127)/2} \approx 2^{64.5}$ KECCAK- $p[1600, 12]$ calls; the headline cap of $N_{\text{tot}} \leq 2^{63}$ used in Section 5.4 keeps this term at most $2^{-131} + 2^{-194}$, and hence below 2^{-130} . Setting $c = 256$ thereby delivers 128-bit single-key security on this term. Under the mu-ind bound of Section 5.2 the per-user security margin degrades by $\log D$ bits in D users; the concrete D -indexed bound appears in Section 5.4.

- $r = 1344$. With state width $b = 1600$ fixed by the KECCAK- $p[1600, 12]$ choice, $r = b - c = 1344$ is the largest rate compatible with the chosen capacity. The 12-round KECCAK- $p[1600, 12]$ permutation matches KangarooTwelve’s choice [BDP⁺16], with a substantial cryptanalytic safety margin relative to the 24-round version used in SHA-3 [Nat15].

- $\tau_{\text{root}} = \tau_{\text{leaf}} = 256$. Matching the capacity ensures that the per-tag collision terms in the authenticity and committing bounds (Section 5.3) are dominated by the capacity birthday rather than the tag length.

- $k = 256$. Matching the capacity ensures that the per-user key guess

$$\text{KeyGuess}_k(N_{\text{prim}}) = N_{\text{prim}}/2^k$$

remains below the capacity birthday at the target D of Section 5.4.

- $\beta = 8183$. The chunk size is set to $49\rho = 49 \cdot 167$ bytes, the largest multiple of the per-call absorption bound ρ not exceeding 2^{13} bytes. Taking β to be an exact multiple of ρ makes every full leaf split into exactly 49 fully utilized message duplex calls with no short final block, so a full leaf costs one initialization duplex call plus 49 message duplex calls. This amortizes the fixed per-leaf initialization cost over a large chunk while keeping the leaf state working set ($8183 < 2^{13}$ bytes) small enough for L1 cache; Section 7.3 reports the resulting cycles-per-byte figures.
- $\nu = 256$. The nonce width is not constrained by the concrete bounds, which assume a nonce-respecting adversary; it is set generously because the wide permutation makes a wide nonce nearly free. The root initialization absorbs $p_{\text{root}} \parallel K \parallel N$, 560 bits in all, in a single duplex call, comfortably within the 1336-bit ($\rho = 167$ -byte) data field (Section 3.2), so widening the nonce to 256 bits adds neither a permutation call nor any per-message cost. The width also removes any practical concern over random nonce collisions: a uniformly random 256-bit nonce repeats over q messages with probability at most $\binom{q}{2}/2^{256}$, negligible at the work cap of Section 5.4. Deployments that do not require a full-width random nonce lose nothing by the choice: a narrower nonce—a 96-bit packet counter, say—is carried in the fixed ν -bit field by zero-padding, so a single mode serves both random nonce and counter-based settings with no change of parameters.

3.4 Domain Separation

Distinct duplex instances and distinct phases within a duplex are kept separate at the transcript level by two mechanisms. First, the initialization strings carry a 6-byte ASCII prefix and, for leaves, an eight-byte little-endian chunk index:

$$p_{\text{root}} = \text{"TW128R"} \in \{0, 1\}^{48},$$

$$p_{\text{leaf}} = \text{"TW128L"} \in \{0, 1\}^{48}.$$

The root duplex's first absorbed string is $p_{\text{root}} \parallel K \parallel N$, and the j -th leaf duplex's first absorbed string is $p_{\text{leaf}} \parallel K \parallel N \parallel \langle j \rangle_8$. Together with the fixed k -bit and ν -bit widths of K and N , these prefixes separate root from leaf transcripts and distinguish leaves with different chunk indices.

Second, every byte-aligned absorption carries a four-bit phase suffix. We identify each suffix with the four-bit string given by its Table 3 value written least significant bit first, so that a duplex call absorbing a byte-aligned block σ under a phase suffix σ_{ph} feeds the duplex interface of Section 2.2 the input $\sigma \parallel \sigma_{\text{ph}}$ before pad10*1 padding. The seven phase suffixes are summarized in Table 3; they correspond to initialization, non-final and final associated data blocks, non-final and final message blocks, and non-final and final aggregation blocks. The seven values are pairwise distinct, so any two transcripts that differ in the phase of any block also differ in the flattened sponge input.

The combined effect of the initialization prefixes and the phase suffixes is that the flattened sponge input of any TW128 duplex call uniquely determines the duplex instance,

Table 3: Phase suffixes (4 bits each, appended LSB-first).

Name	Hex	Use
σ_{init}	0xC	End of initialization (the sole init call)
σ_{ad}^-	0x9	Non-final associated data block
σ_{ad}^+	0xD	Final associated data block
σ_{msg}^-	0xA	Non-final message block
σ_{msg}^+	0xE	Final message block
σ_{agg}^-	0xB	Non-final aggregation block
σ_{agg}^+	0xF	Final aggregation block (emits root tag)

the leaf index when applicable, the phase, and the duplex-call position within that phase. Definition 1 (Well-formed TW128 transcript prefixes), Lemma 4 (Transcript Injectivity), and its structural corollaries formalize this parsing fact for the proofs of Sections 5 and 6 and Corollary 6.

3.5 Encryption

Algorithm 2 specifies $\text{Enc}_K(N, A, M)$ in terms of two helper procedures, **ABSORB** and **STREAMENC**, defined in Algorithm 1. **ABSORB** drives a duplex through a sequence of blocks of a byte string with non-final σ^- and final σ^+ suffix tags, returning the output of the final call; the empty-input case yields a single empty block carrying the σ^+ tag. The streaming helpers use the same empty-input convention, so a zero-length chunk still performs one final message phase duplex call and returns a zero-length body. Algorithm 2 invokes **ABSORB** for the associated data phase only when $A \neq \varepsilon$; when $A = \varepsilon$ the phase is elided and the root σ_{init} call emits the initial M_0 keystream directly. **STREAMENC** encrypts a plaintext byte string block by block, XORing each plaintext block with the current keystream, absorbing the resulting ciphertext block back into the duplex with the appropriate phase tag, and using the duplex’s output as the next keystream. This is the duplex-level presentation of the overwrite mechanism of [BDPVA11, §6.2 and Algorithm 5]: when the current rate prefix emitted as keystream is Z , XOR-absorbing $m \oplus Z$ has the same update on the message bytes as overwriting that prefix with m , with the phase tag still carried as part of the duplex input. The decryption helper **STREAMDEC** is identical to **STREAMENC** with the roles of the plaintext and ciphertext blocks exchanged: it takes ciphertext blocks c_i as input, recovers each plaintext block as $m_i = c_i \oplus Z$, absorbs the same c_i under the same phase suffix, and returns $(m_1 \parallel \dots \parallel m_t, Z')$; we omit its pseudocode. In both helpers, each duplex call absorbs a data block followed by its phase suffix— σ^- for a non-final block and σ^+ for the final block—and requests the output bits the caller consumes next: ℓ (possibly zero) on the final call, and the next block’s keystream length on non-final calls. Each call is the duplex interface of Section 2.2 applied to the suffixed input of Section 3.4.

The encryption procedure (Algorithm 2) instantiates the root duplex and obtains the initial keystream for the root chunk M_0 : when $A \neq \varepsilon$ it absorbs the associated data A , and when $A = \varepsilon$ it elides the associated data phase and takes the keystream from the σ_{init} call itself. It then encrypts M_0 while absorbing the resulting ciphertext chunk C_0 . When $n = 0$ the closing σ_{msg}^+ call of this root chunk requests τ_{root} output bits, which are the root tag T , and the aggregation phase is elided. When $n \geq 1$ it independently encrypts each leaf chunk M_j while absorbing C_j , producing the leaf tag T_j ; the leaf tags together with $\langle n \rangle_8$ are then absorbed into the root duplex, and the final σ_{agg}^+ call emits the root tag T . The ciphertext is the concatenation $C_0 \parallel C_1 \parallel \dots \parallel C_n$. Thus this section writes C for the ciphertext body and $C \parallel T$ for the full AEAD ciphertext returned by $\text{Enc}_K(N, A, M)$.

The keystream bit length requested for the first message block of M_0 is $8 \min(\|M_0\|, \rho)$,

Algorithm 1 TW128 helper procedures.

```

1: procedure ABSORB( $D, s, \sigma^-, \sigma^+, \ell$ )
2:   Partition  $s$  into byte blocks  $s_1, \dots, s_t$  of exactly  $\rho$  bytes each, except that the last
   block may be shorter, with  $t \geq 1$  (the empty case yields the single empty block  $s_1 = \varepsilon$ ).
3:   for  $i \leftarrow 1$  to  $t - 1$  do
4:      $D.\text{duplexing}(s_i \parallel \sigma^-, 0)$ 
5:   end for
6:   return  $D.\text{duplexing}(s_t \parallel \sigma^+, \ell)$ 
7: end procedure

8: procedure STREAMENC( $D, m, Z, \sigma^-, \sigma^+, \ell$ )
9:   Partition  $m$  into byte blocks  $m_1, \dots, m_t$  of exactly  $\rho$  bytes each, except that the
   last block may be shorter, with  $t \geq 1$  (if  $m = \varepsilon$ , set  $t = 1$  and  $m_1 = \varepsilon$ ).
10:  for  $i \leftarrow 1$  to  $t - 1$  do
11:     $c_i \leftarrow m_i \oplus Z$ 
12:     $Z \leftarrow D.\text{duplexing}(c_i \parallel \sigma^-, 8\|m_{i+1}\|)$ 
13:  end for
14:   $c_t \leftarrow m_t \oplus Z$ 
15:   $Z' \leftarrow D.\text{duplexing}(c_t \parallel \sigma^+, \ell)$ 
16:  return  $(c_1 \parallel \dots \parallel c_t, Z')$ 
17: end procedure

```

532 eight times the size of the next message block. It is emitted by the closing σ_{ad}^+ call when
 533 $A \neq \varepsilon$, and by the root σ_{init} call itself when $A = \varepsilon$ and the associated data phase is elided.
 534 Each leaf's σ_{init} call likewise emits $8 \min(\|M_j\|, \rho)$ keystream bits for M_j , so the empty
 535 associated data root σ_{init} call plays exactly the role a leaf σ_{init} call already plays. When
 536 the corresponding chunk is empty, the requested keystream length is zero and the next
 537 message phase duplex call absorbs the empty ciphertext block carrying σ_{msg}^+ . The final
 538 σ_{msg}^+ call in the root chunk requests τ_{root} output bits when $n = 0$, namely the root tag T
 539 emitted by the elided-aggregation path, and zero output bits when $n \geq 1$ because the next
 540 duplex operation on the root is then the first aggregation block; the corresponding call on
 541 a leaf requests τ_{leaf} bits, namely the leaf tag. When $n \geq 1$ the final σ_{agg}^+ call requests τ_{root}
 542 bits, which is the root tag T .

543 The leaf loop is annotated **in parallel** to make explicit that the n leaf computations
 544 are mutually independent; Section 7.1 discusses the SIMD realizations.

545 **Streaming aggregation.** The string $G = T_1 \parallel \dots \parallel T_n$ in Algorithms 2 and 3 names the
 546 leaf tags only for specification. The closing ABSORB consumes $G \parallel \langle n \rangle_8$ as ρ -byte blocks in
 547 leaf order and advances the root state after each, so an implementation need not hold G
 548 in full: it can absorb each leaf tag, or each SIMD batch of leaf tags, into the root as the
 549 corresponding leaves complete, buffering only a bounded number of tags at a time. With
 550 the bounded 1600-bit state of every duplex, this keeps the working memory constant in
 551 $\|M\|$ for a fixed degree of parallelism, growing only with the number of leaves evaluated at
 552 once rather than with the message length; Section 7.1 describes the batched realization.

553 3.6 Decryption

554 Algorithm 3 specifies $\text{Dec}_K(N, A, C \parallel T)$. The procedure mirrors **Enc** block for block:
 555 each ciphertext block c_i is absorbed into the appropriate duplex (root or leaf) under the
 556 corresponding σ_{msg}^- or σ_{msg}^+ suffix, and the plaintext block is recovered as $m_i = c_i \oplus Z^{(i)}$,
 557 writing $Z^{(i)}$ for the value the running keystream variable Z of **STREAMDEC** holds when

Algorithm 2 $\text{Enc}_K(N, A, M)$.**Require:** $K \in \{0, 1\}^k$, $N \in \{0, 1\}^\nu$, A with $\|A\| \leq A_{\max}$, M with $\|M\| \leq M_{\max}$.**Ensure:** $C \parallel T$ with $\|C\| = \|M\|$ and $T \in \{0, 1\}^{\tau_{\text{root}}}$.

```

1:  $M_0 \leftarrow$  first  $\min(\|M\|, \beta)$  bytes of  $M$  (possibly  $\varepsilon$ ).
2: Partition the remaining suffix into chunks  $M_1, \dots, M_n$  of exactly  $\beta$  bytes each, except
   that the last chunk may be shorter but nonempty, so  $n = \text{leaves}(M)$ .
3:  $D_{\text{root}} \leftarrow$  new Duplex
4: if  $A = \varepsilon$  then
5:    $Z \leftarrow D_{\text{root}}.\text{duplexing}(p_{\text{root}} \parallel K \parallel N \parallel \sigma_{\text{init}}, 8 \min(\|M_0\|, \rho))$ 
6: else
7:    $D_{\text{root}}.\text{duplexing}(p_{\text{root}} \parallel K \parallel N \parallel \sigma_{\text{init}}, 0)$ 
8:    $Z \leftarrow \text{ABSORB}(D_{\text{root}}, A, \sigma_{\text{ad}}^-, \sigma_{\text{ad}}^+, 8 \min(\|M_0\|, \rho))$ 
9: end if
10: if  $n = 0$  then
11:    $(C_0, T) \leftarrow \text{STREAMENC}(D_{\text{root}}, M_0, Z, \sigma_{\text{msg}}^-, \sigma_{\text{msg}}^+, \tau_{\text{root}})$ 
12:   return  $C_0 \parallel T$ 
13: end if
14:  $(C_0, \_) \leftarrow \text{STREAMENC}(D_{\text{root}}, M_0, Z, \sigma_{\text{msg}}^-, \sigma_{\text{msg}}^+, 0)$ 
15: for  $j \leftarrow 1$  to  $n$  in parallel do
16:    $D_j \leftarrow$  new Duplex
17:    $Z_j \leftarrow D_j.\text{duplexing}(p_{\text{leaf}} \parallel K \parallel N \parallel \langle j \rangle_8 \parallel \sigma_{\text{init}}, 8 \min(\|M_j\|, \rho))$ 
18:    $(C_j, T_j) \leftarrow \text{STREAMENC}(D_j, M_j, Z_j, \sigma_{\text{msg}}^-, \sigma_{\text{msg}}^+, \tau_{\text{leaf}})$ 
19: end for
20:  $G \leftarrow T_1 \parallel \dots \parallel T_n$ .
21:  $T \leftarrow \text{ABSORB}(D_{\text{root}}, G \parallel \langle n \rangle_8, \sigma_{\text{agg}}^-, \sigma_{\text{agg}}^+, \tau_{\text{root}})$ 
22: return  $C_0 \parallel C_1 \parallel \dots \parallel C_n \parallel T$ 

```

block i is processed ($i < t$ inside the loop, $i = t$ after it), the same value the encryption procedure would have produced. The procedure absorbs ciphertext bytes before the final tag comparison, so the verification transcript is determined by (K, N, A, C) for every in-domain ciphertext, accepting or not. This property is used by the direct mu-ind proof to argue that encryption-first and verification-first transcript classes are well defined regardless of the tag-comparison outcome.

The $n = 0$ and empty-body cases mirror their encryption counterparts (Section 3.5): for $n = 0$ the aggregation phase is elided and the closing σ_{msg}^+ call of the root chunk emits the candidate tag T' directly, and an empty body still performs the single empty σ_{msg}^+ call of Algorithm 1, requesting τ_{root} output bits.

Correctness. For every $K \in \{0, 1\}^k$, $N \in \{0, 1\}^\nu$, A with $\|A\| \leq A_{\max}$, and M with $\|M\| \leq M_{\max}$, $\text{Dec}_K(N, A, \text{Enc}_K(N, A, M)) = M$. This follows by inspection: every duplex call in Dec_K receives the same input as the corresponding call in Enc_K (each ciphertext block is absorbed under the same phase tag in both directions), so the duplex states and the recomputed root tag T' exactly match. The constant-time comparison $T' = T$ succeeds, and the recovered $m_i = c_i \oplus Z^{(i)} = (m_i \oplus Z^{(i)}) \oplus Z^{(i)} = m_i$ at every position.

4 Security Model

This section fixes the security model used throughout the paper. The scheme syntax and the authenticated encryption and committing security definitions of Section 2 are instantiated here for TW128 in the lifted public permutation model (Section 4.1), and we

Algorithm 3 $\text{Dec}_K(N, A, C \parallel T)$.

Require: K, N, A as in Algorithm 2, ciphertext C with $\|C\| \leq M_{\max}$, candidate tag $T \in \{0, 1\}^{\tau_{\text{root}}}$.

Ensure: M with $\|M\| = \|C\|$, or $\perp$.

```

1:  $C_0 \leftarrow$  first  $\min(\|C\|, \beta)$  bytes of  $C$  (possibly  $\varepsilon$ ).
2: Partition the remaining suffix into chunks  $C_1, \dots, C_n$  of exactly  $\beta$  bytes each, except
   that the last chunk may be shorter but nonempty, so  $n = \text{leaves}(C)$ .
3:  $D_{\text{root}} \leftarrow$  new Duplex
4: if  $A = \varepsilon$  then
5:    $Z \leftarrow D_{\text{root}}.\text{duplexing}(p_{\text{root}} \parallel K \parallel N \parallel \sigma_{\text{init}}, 8 \min(\|C_0\|, \rho))$ 
6: else
7:    $D_{\text{root}}.\text{duplexing}(p_{\text{root}} \parallel K \parallel N \parallel \sigma_{\text{init}}, 0)$ 
8:    $Z \leftarrow \text{ABSORB}(D_{\text{root}}, A, \sigma_{\text{ad}}^-, \sigma_{\text{ad}}^+, 8 \min(\|C_0\|, \rho))$ 
9: end if
10: if  $n = 0$  then
11:    $(M_0, T') \leftarrow \text{STREAMDEC}(D_{\text{root}}, C_0, Z, \sigma_{\text{msg}}^-, \sigma_{\text{msg}}^+, \tau_{\text{root}})$ 
12:   return  $M_0$  if  $T' = T$  (compared in constant time), else  $\perp$ .
13: end if
14:  $(M_0, \_) \leftarrow \text{STREAMDEC}(D_{\text{root}}, C_0, Z, \sigma_{\text{msg}}^-, \sigma_{\text{msg}}^+, 0)$ 
15: for  $j \leftarrow 1$  to  $n$  in parallel do
16:    $D_j \leftarrow$  new Duplex
17:    $Z_j \leftarrow D_j.\text{duplexing}(p_{\text{leaf}} \parallel K \parallel N \parallel \langle j \rangle_8 \parallel \sigma_{\text{init}}, 8 \min(\|C_j\|, \rho))$ 
18:    $(M_j, T_j) \leftarrow \text{STREAMDEC}(D_j, C_j, Z_j, \sigma_{\text{msg}}^-, \sigma_{\text{msg}}^+, \tau_{\text{leaf}})$ 
19: end for
20:  $G \leftarrow T_1 \parallel \dots \parallel T_n$ .
21:  $T' \leftarrow \text{ABSORB}(D_{\text{root}}, G \parallel \langle n \rangle_8, \sigma_{\text{agg}}^-, \sigma_{\text{agg}}^+, \tau_{\text{root}})$ 
22: if  $T' = T$  (compared in constant time) then
23:   return  $M_0 \parallel M_1 \parallel \dots \parallel M_n$ 
24: else
25:   return  $\perp$ 
26: end if

```

578 define the intermediate random oracle world through which every bound is first proved
579 (Section 4.2). Section 4.3 fixes the resource accounting that all later bounds reference, and
580 Section 4.4 instantiates the multi-user indistinguishability game for TW128.

581 4.1 Public Permutation Games

582 The lifted AE-style public permutation advantages below compare a real execution, with
583 sampled $\pi \leftarrow \$ \text{Perm}(\{0, 1\}^{1600})$ and direct primitive interface (π, π^{-1}) , to a simulator-ideal
584 execution. In the simulator-ideal execution, the construction interface is the ideal interface
585 specified by the game and the public primitive interface is $(\mathcal{S}^{\text{RO}_{\text{flat}}}, \mathcal{S}^{-1, \text{RO}_{\text{flat}}})$. The simulator
586 $\mathcal{S}^{\text{RO}_{\text{flat}}}$ together with its inverse $\mathcal{S}^{-1, \text{RO}_{\text{flat}}}$ is the public primitive simulator fixed in Section B;
587 both directions share an internal lazy table for the random oracle RO_{flat} of Section 4.2. The
588 committing security games remain one-world real public permutation success probabilities;
589 the simulator appears for them only inside the reductions of Section B.

590 In the real all-in-one AE game, the challenger additionally samples a key $K \leftarrow \$$
591 $\{0, 1\}^k$ and exposes the construction oracles $\text{Enc}_K[\pi]$ (the TW128 encryption algorithm
592 of Algorithm 2 with the sampled π) and $\text{Ver}_K[\pi]$, defined by $\text{Ver}_K[\pi](N, A, C, T) = 1$
593 if $\text{Dec}_K[\pi](N, A, C \parallel T) \neq \perp$ and 0 otherwise. The committing security games have no
594 challenger-sampled key.

Validity convention. All construction interface inputs are checked against TW128's public syntax and domain before any construction computation or construction-internal primitive lookup is performed. An encryption input is valid when $N \in \{0, 1\}^\nu$, A is a byte string with $\|A\| \leq A_{\max}$, and M is a byte string with $\|M\| \leq M_{\max}$. A verification or decryption input is valid when $N \in \{0, 1\}^\nu$, A is a byte string with $\|A\| \leq A_{\max}$, and the full ciphertext parses as $C \parallel T$ with $\|C\| \leq M_{\max}$ and $T \in \{0, 1\}^{\tau_{\text{root}}}$. In the hash and committing games, each submitted input (K, N, A, M) must be a valid encryption input including $K \in \{0, 1\}^k$, and a submitted CMTDs ciphertext must be a valid verification/decryption ciphertext. Invalid encryption queries are rejected and not recorded, invalid verification queries return 0, and invalid committing-game outputs make the game return 0. Such rejected inputs contribute no construction cost and create no construction transcripts. Throughout, an *accepted* encryption query is one that passes this validity check and any game-specific nonce or user check.

All-in-one AE. The all-in-one authenticated encryption game [RS06] replaces the construction interface by the ideal pair (Rand, Replay). On each accepted encryption query (N, A, M) , Rand samples a uniformly random byte string of length $\|M\| + \tau_{\text{root}}/8$, parses it as $C \parallel T$ with $T \in \{0, 1\}^{\tau_{\text{root}}}$, records (N, A, C, T) , and returns $C \parallel T$. The oracle Replay returns 1 exactly on recorded tuples (N, A, C, T) and 0 otherwise. The adversary is nonce-respecting on encryption queries.

$$\text{Adv}_{\text{TW128}}^{\text{ae}}(\mathcal{A}) = \left| \Pr[\mathcal{A}^{\text{Enc}_K[\pi], \text{Ver}_K[\pi], \pi, \pi^{-1}} \Rightarrow 1] - \Pr[\mathcal{A}^{\text{Rand, Replay}, \mathcal{S}^{\text{ROflat}}, \mathcal{S}^{-1, \text{ROflat}}} \Rightarrow 1] \right|.$$

This lifted all-in-one AE advantage is the single-user case of the multi-user real-vs-ideal notion of Section 2.4; Corollary 1 therefore derives it from the direct mu-ind bound of Theorem 2.

Hash games. The collision, second-preimage, and preimage games for the wrapper H (Section 2.5) have no challenger-sampled key; all inputs are adversary-supplied. In the s -way collision game the adversary outputs s inputs, and the challenger returns 1 iff they are pairwise distinct, in the TW128 domain, and share a root tag:

$$\begin{aligned} & \pi \leftarrow \$ \text{Perm}(\{0, 1\}^{1600}), \quad (K_i, N_i, A_i, M_i)_{i=1}^s \leftarrow \$ \mathcal{A}^{\pi, \pi^{-1}}, \\ & \text{return 0 unless the inputs are pairwise distinct and in the TW128 domain,} \\ & \text{return } [\text{H}(K_1, N_1, A_1, M_1) = \dots = \text{H}(K_s, N_s, A_s, M_s)], \end{aligned}$$

with advantage $\text{Adv}_{\text{TW128}}^{\text{Coll}}(\mathcal{A}) = \Pr[\text{the game returns 1}]$; no nonce uniqueness is imposed. The always second-preimage (aSec) game samples a challenge input X uniformly from the TW128 domain, gives it to $\mathcal{A}$, and returns 1 iff $\mathcal{A}$ outputs $X' \neq X$ in the domain with $\text{H}(X') = \text{H}(X)$, with advantage $\text{Adv}_{\text{TW128}}^{\text{aSec}}$; following [RS04] the always-notion concerns the fixed permutation, here the deployed instance the public permutation model represents by a random π . The preimage game fixes a target tag $T^* \in \{0, 1\}^{\tau_{\text{root}}}$ before the primitive sampling and returns 1 iff $\mathcal{A}(T^*)$ outputs an input X in the domain with $\text{H}(X) = T^*$, the advantage $\text{Adv}_{\text{TW128}}^{\text{ePre}}$ taken as the worst case over query-independent targets.

Committing games. The [BH22] committing games are instantiated as in Section 2.5, with no challenger-sampled key. The s -way CMTDs- ℓ (decryption) game has the adversary output a ciphertext $C \parallel T$ and s tuples (K_i, N_i, A_i, M_i) with pairwise distinct WiC_ℓ -images, and returns 1 iff $\text{Dec}_{K_i}[\pi](N_i, A_i, C \parallel T) = M_i$ for every i , with the eager length rejection justified by Algorithm 3 ($\text{Dec}_K(N, A, C)$ returns $\perp$ or a message of length $\|C\|$); its advantage is $\text{Adv}_{\text{TW128}}^{\text{CMTDs-}\ell}(\mathcal{A})$. The CMTs- ℓ (encryption) game has the adversary output the s tuples alone and returns 1 iff the encryptions $\text{Enc}_{K_i}[\pi](N_i, A_i, M_i)$ are all equal; its

advantage is $\text{Adv}_{\text{TW128}}^{\text{CMTs-}\ell}(\mathcal{A})$. A CMTDs-4 or CMTs-4 winner's tuples share a common tag, hence are a collision of H , so Corollary 3 bounds CMTDs-4 through Theorem 7 and Corollary 6 forwards the remaining [BH22] notions to the same root-collision bound.

Context discoverability. The [MLGR23, Fig. 5] context discoverability game (Section 2.5) is a ciphertext-level preimage notion, distinct from the preimage resistance of H above. It is parameterized by a target specifier ts and a query-independent context selector S that produces a target ciphertext C^* and reveals the ts -named components rev of a decryption context; there is no challenger-sampled key.

```

648  $\pi \leftarrow \$ \text{Perm}(\{0,1\}^{1600}), \quad (C^*, \text{rev}) \leftarrow \$ S[\pi],$ 
649  $(K, N, A) \leftarrow \$ \mathcal{A}^{\pi, \pi^{-1}}(C^*, \text{rev}),$ 
650 return 0 if  $C^*$  is not a valid TW128 ciphertext, if  $(K, N, A)$  is outside the domain,
651     or if  $(K, N, A)$  disagrees with  $\text{rev}$  on its  $\text{ts}$ -components,
652 return  $[\text{Dec}_K[\pi](N, A, C^*) \neq \perp]$ .
```

The restricted advantage $\text{Adv}_{\text{TW128}}^{\text{CDY}}(\mathcal{A})$ is the worst case over query-independent selectors S , matching $\text{CDY}^*[\text{ts}]$ and, at $\text{ts} = \{n, a\}, \{k, a\}, \{k, n\}$, its per-component variants $\text{CDY}_k^*, \text{CDY}_n^*, \text{CDY}_a^*$. Because the selector holds a hidden context, the TW128 bound (Theorem 6) carries a key-test term that the keyless preimage game of H does not. Section B lifts the hash and committing bounds of Section 6 to the public permutation statements of Section 5.2.

659 4.2 The TW128 Random Oracle Model

The proofs of Section 6 and Corollary 6, together with the headline reductions of Section 5.2, are carried out through an intermediate random oracle model whose construction interfaces use the [BDPVA11] bridge of Section 2.2 to address a single random oracle. The flat sponge lift of Section B then converts the resulting bounds to the public permutation form.

The model has a single random oracle $\text{RO}_{\text{flat}} : \{0,1\}^* \rightarrow \{0,1\}^\infty$ over the unpadded H -input domain of the [BDPVA08] simple-padded sponge interface, with lazy sampling: each distinct input receives an independent infinite bit string, and repeated queries return the same string. A direct adversary query is finite-output: on input (Y, ℓ) with finite ℓ , the oracle returns the first ℓ bits of the lazy-sampled value for Y .

For a well-formed TW128 duplex transcript prefix X in the language of Definition 1, at one of the construction transcript lookup points formalized by Lemma 6 and Table 9, let $\text{flat}(X) = I[1344, 1344](\text{raw_flat}(X))$ be the bridged image of the native flattened input under the [BDPVA11, Theorem 3] preprocessing map of Section 2.2, and define

$$673 \quad \text{RF}(X) = \text{RO}_{\text{flat}}(\text{flat}(X)).$$

The construction interfaces $\text{Enc}_K^{\text{RO}}, \text{Dec}_K^{\text{RO}}$, and Ver_K^{RO} are the TW128 algorithms of Section 3 with every root or leaf duplex output replaced by the corresponding requested projection of $\text{RF}(X)$; the projection length is the duplexing call's output parameter ℓ , which may be zero. The bridge and the typed-restriction convention are formalized in Lemma 6; the direct-query key-hit predicate ignores the adversary's requested output length and is formalized as KeyedTWOOut in Lemma 6.

The random oracle model games are the construction interface variants of Sections 4.1 and 4.4: the primitive interface is replaced by direct adversary access to RO_{flat} , and each construction oracle uses the corresponding RO algorithm, the ideal oracles ($\text{Rand}, \text{Replay}$) being unchanged. Under this substitution $\text{Adv}_{\text{RO}}^{\text{ae}}$ is the all-in-one advantage of Section 4.1 with no simulator on the ideal side; the keyless hash advantages $\text{Adv}_{\text{RO}}^{\text{Coll}}(\mathcal{B})$, $\text{Adv}_{\text{RO}}^{\text{aSec}}(\mathcal{B})$,

685 $\text{Adv}_{\text{RO}}^{\text{ePre}}(\mathcal{B})$ of $\mathcal{H}$, the committing $\text{Adv}_{\text{RO}}^{\text{CMTDs-4}}(\mathcal{B})$, and, for $\ell \in \{1, 3, 4\}$, $\text{Adv}_{\text{RO}}^{\text{CMTs-}\ell}(\mathcal{B})$
 686 are the probabilities that the corresponding random oracle hash and committing games
 687 return 1, the CMTs challenger computing its deterministic encryption checks with $\text{Enc}_{\text{RO}}^{\text{RO}}$.
 688 For mu-ind, $\text{Adv}_{\text{RO}}^{\text{mu-ind}}(\mathcal{B})$ denotes the Section 4.4 game with the same New interface and
 689 per-user nonce discipline—user d using $(\text{Enc}_{K_d}^{\text{RO}}, \text{Ver}_{K_d}^{\text{RO}})$ when $b = 1$ and $(\text{Rand}, \text{Replay})$
 690 when $b = 0$ —equivalently, in the two-game notation of Section 2.4,

$$691 \quad \text{Adv}_{\text{RO}}^{\text{mu-ind}}(\mathcal{B}) = |\Pr[\mathcal{B}^{\text{Real}^{\text{mu}}, \text{RO}_{\text{flat}}} \Rightarrow 1] - \Pr[\mathcal{B}^{\text{Ideal}^{\text{mu}}, \text{RO}_{\text{flat}}} \Rightarrow 1]|,$$

692 where the superscripts name the entire multi-user construction interface in the corre-
 693 sponding world. Adversary queries to RO_{flat} are finite-output and counted in $q_{\text{RO}}(\mathcal{B})$ of
 694 Section 4.3. Construction-internal $\text{RF}(\cdot)$ lookups — those made by $\text{Enc}_K^{\text{RO}}, \text{Dec}_K^{\text{RO}}, \text{Ver}_K^{\text{RO}}$,
 695 or a committing challenger — are not.

696 4.3 Resources

697 Adversary resources are measured by fixed, execution-uniform caps on the number and
 698 size of oracle queries. An execution-uniform cap is one that holds on every execution
 699 path, including over adversary randomness, oracle randomness, and adaptive query choices.
 700 Proofs may also introduce local execution-uniform caps, such as q_e for the number of
 701 accepted encryption queries, solely to index a finite hybrid sequence. Such local caps are
 702 part of the adversary’s fixed query bounds, but they are not listed as headline resource
 703 parameters unless they affect the final bound.

704 **N :** the construction interface cost, measured as the number of $\text{KECCAK-}p[1600, 12]$ calls
 705 that the adversary’s construction interface queries induce in the real TW128 com-
 706 putation after the validity convention above has been applied. In ideal games, the
 707 same accounting applies to the corresponding real TW128 computation on the valid
 708 query, even when an ideal oracle answers by table lookup or without running the
 709 construction. N counts at per-block granularity, summing each root initialization,
 710 associated data block, root message block, leaf initialization, leaf message block,
 711 and aggregation block across the entire query sequence. Valid AE-style verification
 712 computations are counted whether they accept or reject. In the hash, committing,
 713 and discoverability games there is no adversary construction oracle: for the $\mathcal{H}$ collision
 714 game, N counts the s deterministic encryptions on the submitted inputs (and the one
 715 or two encryptions of the second-preimage and preimage games); for $\text{CMTDs-}\ell$, N
 716 counts only the deterministic decryption checks performed by the challenger on the
 717 common ciphertext; for the context discoverability game, N counts the worst-case
 718 selector’s encryption producing C^* together with the single deterministic decryp-
 719 tion check the challenger performs on C^* ; while for the CMTs- ℓ source games of
 720 Corollary 6, N counts only the deterministic encryption checks actually used by the
 721 challenger to compare the s ciphertexts.

722 **N_{prim} :** the number of direct adversary queries to the primitive interface (the pair (π, π^{-1})
 723 in public permutation games; the simulator’s RO_{flat} table is internal and is not
 724 exposed to the application adversary in those games).

725 **$N_{\text{tot}} = N + N_{\text{prim}}$:** the total query-sequence cost. The flat sponge lift of Section B is
 726 applied at N_{tot} because the indistinguishability replacement sees both the construction-
 727 internal sponge calls counted by N and the adversary’s direct primitive-interface
 728 queries counted by N_{prim} in one chronological public primitive transcript.

729 **$q_v(\mathcal{A})$, Q_v :** the number of verification queries with valid inputs, and a cap on it. Used
 730 in the verification terms of the direct mu-ind and AE bounds of Section 5.2. Every
 731 valid verification query contributes at least one verification/decryption construction
 732 call in the N accounting, so $q_v(\mathcal{A}) \leq N$.

$q_{\text{RO}}(\mathcal{B})$, Q_{RO} : the number of direct finite-output queries to RO_{flat} by a random oracle model adversary $\mathcal{B}$, and a cap on it. Repeated direct queries on the same Y are counted again. Construction-internal $\text{RF}(\cdot)$ lookups are not counted. After the derived-adversary construction, q_{RO} counts only simulator-induced RO_{flat} calls from the adversary's direct primitive-interface queries, not the construction calls already counted in N for the flat sponge lift.

$n_{\text{leaf}}(\mathcal{B})$, N_{leaf} : the number of distinct construction-generated final leaf transcripts in the execution of $\mathcal{B}$, and a cap on it. Distinct final leaf transcripts are generated by valid encryption computations, by valid AE-style verification computations (whether they accept or reject), and by hash, committing, and discoverability challenger construction checks that pass the early validity checks.

For a byte string X ,

$$\text{leaves}(X) = \max(0, \lceil \|X\|/\beta \rceil - 1)$$

counts the leaf chunks induced by a body byte string of length $\|X\|$. Thus n_{leaf} is bounded by the sum of $\text{leaves}(X)$ over the valid construction computations that create final leaf transcripts, where X is the plaintext body for encryption computations and the submitted ciphertext body for verification and committing or discoverability decryption checks. Each distinct final leaf transcript contains a final leaf duplex call counted in N , so $n_{\text{leaf}}(\mathcal{B}) \leq N$ unconditionally. The ratio n_{leaf}/N is largest when most processed bytes lie in leaf chunks, where each full leaf contributes $1 + \beta/\rho = 50$ duplex calls to N per final leaf transcript.

All theorems are stated for adversaries whose execution-uniform caps hold at the stated bounds.

4.4 Multi-User Setting

This subsection instantiates the mu-ind game of [BT16] (Section 2.4) for TW128 in the lifted public permutation model, using the same ideal-side simulator convention as the single-user AE-style games of Section 4.1. The challenger samples a challenge bit $b \leftarrow \{0, 1\}$. If $b = 1$, it samples a permutation $\pi \leftarrow \text{Perm}(\{0, 1\}^{1600})$ and exposes (π, π^{-1}) as the primitive interface; if $b = 0$, it exposes $(\mathcal{S}^{\text{RO}_{\text{flat}}}, \mathcal{S}^{-1, \text{RO}_{\text{flat}}})$ as the primitive interface. The adversary $\mathcal{A}$ creates users on demand via **New**, and the game maintains an active-user counter v . We write D for an execution-uniform cap on the number of **New** calls. Equivalently, the game may sample independent keys $K_1, \dots, K_D$ at the start of the experiment and let the d -th **New** call activate K_d ; inactive indices are never valid users.

The mu-ind construction oracles inherit the validity convention of Section 4.1. The encryption oracle on input (d, N, A, M) rejects, without recording, if (N, A, M) is outside the TW128 encryption domain, if $d \notin \{1, \dots, v\}$, or if $(d, N) \in V$; otherwise it sets $C = \text{Enc}_{K_d}[\pi](N, A, M)$ when $b = 1$, and samples a uniformly random byte string C of length $\|M\| + \tau_{\text{root}}/8$ when $b = 0$. The game records (d, N) in V and (d, N, A, C) in W , and returns C . The verification oracle on input (d, N, A, C) rejects if $d \notin \{1, \dots, v\}$ or if (N, A, C) is not a valid TW128 verification/decryption input; if $(d, N, A, C) \in W$ it returns **true**; if $b = 0$ it returns **false**; otherwise it returns $[\text{Dec}_{K_d}[\pi](N, A, C) \neq \perp]$. Nonce uniqueness is enforced per user by the set V .

After querying these oracles, $\mathcal{A}$ outputs $b' \in \{0, 1\}$. The mu-ind advantage of $\mathcal{A}$ against TW128 is

$$\text{Adv}_{\text{TW128}}^{\text{mu-ind}}(\mathcal{A}) = |2 \Pr[b' = b] - 1|.$$

Resource caps split per potential user $d \leq D$: $N^{(d)}$, $q_v^{(d)}$, $n_{\text{leaf}}^{(d)}$ are the per-user counts under Section 4.3, with inactive users contributing zero. The global totals $N = \sum_{d=1}^D N^{(d)}$, $q_v = \sum_{d=1}^D q_v^{(d)}$, $n_{\text{leaf}} = \sum_{d=1}^D n_{\text{leaf}}^{(d)}$ appear in the final bound: N through the lift term

$\text{Lift}_c(N_{\text{tot}})$, and n_{leaf} and q_v through the leaf-collision and verification-tag terms. N_{prim} remains a single shared count over the primitive interface.

The mu-ind treatment applies only to indistinguishability. The committing security games already give the adversary control of all s keys, so the s -way CMTDs- ℓ statements of Section 4.1 are themselves the multi-key committing statement and need no separate mu-cmt notion.

5 Main Results

This section states and instantiates TW128’s concrete security bounds. We first collect the recurring loss terms (Section 5.1). We then present the headline theorem sequence, including the random oracle bounds used by the lift (Section 5.2), give the concrete TW128-parameter forms (Sections 5.3 and 5.4), and close with a discussion of the bounds’ tightness (Section 5.5).

5.1 Loss Terms

Five closed-form expressions recur across the headline bounds of Section 5.2— Lift_c through-out, with KeyGuess_k , RootColl_τ , RootPre_τ , and LeafColl_τ in some. We define them here, ahead of the theorems that use them.

Flat sponge loss. The flat sponge differentiating loss is

$$\text{Lift}_c(N_{\text{tot}}) = N_{\text{tot}}(N_{\text{tot}} + 1)/2^{c+1}.$$

For $N_{\text{tot}} < 2^c$ it upper-bounds the [BDPVA08] sponge-differentiating advantage; Lemma 16 carries out the derivation from the [BDPVA08, Lemmas 4 and 5] product bounds through the Weierstrass product inequality. For $N_{\text{tot}} \geq 2^c$, $\text{Lift}_c(N_{\text{tot}}) \geq 1$ and the advantage bound is trivial without any appeal to indifferentiability, and already at $N_{\text{tot}} \geq 2^{c/2}$ the bound exceeds 1/2 and provides no meaningful security.

Key-guessing term. The ideal system key-guessing term of [BDPVA11] is

$$\text{KeyGuess}_k(Q) = Q/2^k.$$

It accounts for direct RO_{flat} queries whose inputs satisfy the decoder KeyedTWOut of Table 9. Each such query is an adaptive test of the decoded candidate key against the hidden key K , independent of the query’s requested output length. Lemma 9 bounds the probability that any of Q direct queries tests the right key by $Q/2^k$.

Root-collision term. The s -way root tag multi-collision term is

$$\text{RootColl}_\tau(Q, s) = \binom{Q+s}{s} / 2^{\tau(s-1)},$$

with the abbreviation $\text{RootColl}_{\tau_{\text{root}}}(Q, s) = \text{RootColl}_\tau(Q, s)$ at $\tau = \tau_{\text{root}}$. The $+s$ is candidate-sequence slack for the s final root transcript inputs generated by the committing challenger after the adversary outputs its tuples; these lookups are construction-internal and are not added to $q_{\text{RO}}(\mathcal{B})$. In the public permutation CMTDs-4 and direct CMTs- ℓ bounds, the committing proof builds a chronological candidate sequence and bounds only the event that some s distinct candidates have equal root tag projections. The flat sponge lift of Section B gives $Q = N_{\text{prim}}$: only the adversary’s N_{prim} direct primitive queries can induce simulator-side RO_{flat} queries that contribute to that sequence.

819 **Root-preimage term.** The root tag preimage term is

$$820 \quad \text{RootPre}_\tau(Q) = (Q + 1)/2^\tau,$$

821 with the abbreviation $\text{RootPre}_{\tau_{\text{root}}}(Q) = \text{RootPre}_\tau(Q)$ at $\tau = \tau_{\text{root}}$. It is the preimage
 822 analogue of RootColl_τ : a single query-independent target tag replaces the s -way internal
 823 collision, the $+s$ candidate-sequence slack collapses to the $+1$ of the single discoverability
 824 challenger decryption check, and the exponent drops from $\tau(s - 1)$ to τ . As in the root-
 825 collision term, the flat sponge lift gives $Q = N_{\text{prim}}$. The term bounds the query-independent
 826 target of context discoverability (Theorem 6).

827 **Leaf-collision term.** The known-key leaf tag collision term is

$$828 \quad \text{LeafColl}_\tau(Q) = \binom{Q}{2}/2^\tau.$$

829 It is the pairwise birthday bound on τ -bit leaf tag projections used by the collision-resistance
 830 and always second-preimage theorems (Theorem 3, Theorem 4). Unlike the hidden-key
 831 leaf collision term in the direct mu-ind and AE bounds, the compact committing game has
 832 adversary-supplied keys, so direct primitive queries can pre-read candidate leaf tag points.
 833 The random oracle proof therefore counts both direct leaf tag candidates and challenger-
 834 generated final leaf transcripts; after the flat sponge lift this gives $Q = N_{\text{prim}} + N_{\text{leaf}}$.

835 5.2 Headline Theorems

836 TW128 delivers its headline guarantees in the lifted public permutation model of Section 4.1:
 837 multi-user indistinguishability [BT16], all-in-one AE security [RS06] as its single-user
 838 corollary, and the commitment security of the root tag (Section 2.5). The last is stated
 839 as the [RS04] hash security of the wrapper H —collision, second-preimage, and preimage
 840 resistance—together with the [BH22] CMTDs-4 committing notion and the [MLGR23]
 841 context discoverability notion. For the AE-style notions the ideal side exposes the BDPV
 842 simulator; the committing, hash, and discoverability bounds remain real public permutation
 843 success probabilities. Theorems 1 to 6 and Corollaries 1 and 3 state these bounds in terms
 844 of the loss terms Lift_c , KeyGuess_k , RootColl_τ , LeafColl_τ , and RootPre_τ of Section 5.1.

845 The real-vs-ideal results come first: a random oracle mu-ind theorem, its public
 846 permutation lift, and the all-in-one AE corollary obtained by setting the mu-ind user cap
 847 to $D = 1$. The independent hash, committing, and discoverability results follow afterward.
 848 The CMTDs-4, collision, second-preimage, and preimage random oracle bounds are proved
 849 in Section 6. All random oracle adversaries in this subsection are flat-RO adversaries
 850 (Section 4.2) with the displayed resources, and the flat sponge lift of Section B instantiates
 851 each random oracle bound at $q_{\text{RO}} = N_{\text{prim}}$ when transferring to the public permutation
 852 setting.

853 **Theorem 1** (Random Oracle Multi-User Indistinguishability). *For every flat-RO mu-ind*
 854 *adversary $\mathcal{B}$ with execution-uniform user cap D and resource counts of Section 4.4,*

$$855 \quad \text{Adv}_{\text{RO}}^{\text{mu-ind}}(\mathcal{B}) \leq \frac{D(D-1)}{2^{k+1}} + D \cdot \text{KeyGuess}_k(q_{\text{RO}}(\mathcal{B})) + \frac{n_{\text{leaf}}(\mathcal{B})(n_{\text{leaf}}(\mathcal{B}) - 1)}{2^{\tau_{\text{leaf}}+1}} + \frac{q_v(\mathcal{B})}{2^{\tau_{\text{root}}}}.$$

856 *Proof.* Use the pre-sampled-key formulation of Section 4.4: sample keys $K_1, \dots, K_D$ at
 857 the start of the experiment, and let the d -th New call activate K_d , unused keys being
 858 ignored. Let coll denote the event that two of the pre-sampled keys $K_1, \dots, K_D$ collide;
 859 $\Pr[\text{coll}] \leq \binom{D}{2}/2^k = D(D-1)/2^{k+1}$ by union bound. In the random oracle model, use a
 860 standard hybrid over the D users: H_d answers users $1, \dots, d$ with (Rand, Replay) and users
 861 $d+1, \dots, D$ with the real construction interface, so H_0 is the real mu-ind game and H_D

is the ideal one. Let H_d^* be H_d with a fixed output bit on coll. Since coll has the same probability in all hybrids,

$$\mathbf{Adv}_{\text{RO}}^{\text{mu-ind}}(\mathcal{B}) \leq \Pr[\text{coll}] + |\Pr[\mathcal{B}^{H_0^*} \Rightarrow 1] - \Pr[\mathcal{B}^{H_D^*} \Rightarrow 1]|.$$

By the triangle inequality, it remains to bound each adjacent stopped hybrid. For each d , Lemma 14 gives

$$|\Pr[\mathcal{B}^{H_{d-1}^*} \Rightarrow 1] - \Pr[\mathcal{B}^{H_d^*} \Rightarrow 1]| \leq \text{KeyGuess}_k(q_{\text{RO}}(\mathcal{B})) + \frac{n_{\text{leaf}}^{(d)}(n_{\text{leaf}}^{(d)} - 1)}{2^{\tau_{\text{leaf}} + 1}} + \frac{q_v^{(d)}}{2^{\tau_{\text{root}}}}.$$

Summing across $d = 1, \dots, D$ gives $D \text{KeyGuess}_k(q_{\text{RO}}(\mathcal{B}))$ for the key-test contribution. The leaf-collision contribution satisfies

$$\sum_d \frac{n_{\text{leaf}}^{(d)}(n_{\text{leaf}}^{(d)} - 1)}{2^{\tau_{\text{leaf}} + 1}} = \frac{1}{2^{\tau_{\text{leaf}}}} \sum_d \binom{n_{\text{leaf}}^{(d)}}{2} \leq \frac{1}{2^{\tau_{\text{leaf}}}} \left(\sum_d n_{\text{leaf}}^{(d)} \right) = \frac{n_{\text{leaf}}(\mathcal{B})(n_{\text{leaf}}(\mathcal{B}) - 1)}{2^{\tau_{\text{leaf}} + 1}},$$

because the unordered pairs drawn within the per-user blocks are a subset of all unordered pairs among $\sum_d n_{\text{leaf}}^{(d)} = n_{\text{leaf}}(\mathcal{B})$ items. The verification-tag contribution satisfies

$$\sum_d q_v^{(d)} / 2^{\tau_{\text{root}}} = q_v(\mathcal{B}) / 2^{\tau_{\text{root}}}.$$

Combining these estimates with the coll penalty gives the stated RO bound. $\square$

Theorem 2 (Multi-User Indistinguishability). *For every mu-ind adversary $\mathcal{A}$ against TW128 with execution-uniform user cap D and the global resources of Section 4.4,*

$$\mathbf{Adv}_{\text{TW128}}^{\text{mu-ind}}(\mathcal{A}) \leq \text{Lift}_c(N_{\text{tot}}) + \frac{D(D-1)}{2^{k+1}} + D \cdot \text{KeyGuess}_k(N_{\text{prim}}) + \frac{N_{\text{leaf}}(N_{\text{leaf}} - 1)}{2^{\tau_{\text{leaf}} + 1}} + \frac{Q_v}{2^{\tau_{\text{root}}}}.$$

Proof. By the Lift Application corollary (Corollary 8) for the mu-ind game, there is a random oracle mu-ind adversary $\mathcal{B} = \mathcal{B}_{\text{derived}}(\mathcal{A})$ with $q_{\text{RO}}(\mathcal{B}) \leq N_{\text{prim}}$, $n_{\text{leaf}}(\mathcal{B}) \leq N_{\text{leaf}}$, $q_v(\mathcal{B}) \leq Q_v$, and the per-user construction-side caps of Section 4.4 such that

$$\mathbf{Adv}_{\text{TW128}}^{\text{mu-ind}}(\mathcal{A}) \leq \text{Lift}_c(N_{\text{tot}}) + \mathbf{Adv}_{\text{RO}}^{\text{mu-ind}}(\mathcal{B}).$$

Theorem 1 bounds the random oracle term. Substituting $q_{\text{RO}}(\mathcal{B}) \leq N_{\text{prim}}$, $n_{\text{leaf}}(\mathcal{B}) \leq N_{\text{leaf}}$, and $q_v(\mathcal{B}) \leq Q_v$ gives the stated bound. $\square$

Corollary 1 (All-in-One AE Security). *For every nonce-respecting AE adversary $\mathcal{A}$ against TW128 with the resources of Section 4.3,*

$$\mathbf{Adv}_{\text{TW128}}^{\text{ae}}(\mathcal{A}) \leq \text{Lift}_c(N_{\text{tot}}) + \text{KeyGuess}_k(N_{\text{prim}}) + \frac{N_{\text{leaf}}(N_{\text{leaf}} - 1)}{2^{\tau_{\text{leaf}} + 1}} + \frac{Q_v}{2^{\tau_{\text{root}}}}.$$

Proof. Build a mu-ind adversary $\mathcal{A}^\mu$ that makes one New query, runs $\mathcal{A}$, and forwards each encryption and verification query of $\mathcal{A}$ to user 1. The primitive interface is forwarded unchanged. The real and ideal views of $\mathcal{A}$ in the all-in-one AE experiment are exactly the real and ideal views of $\mathcal{A}^\mu$ in the mu-ind experiment with $D = 1$, and the resource counts are unchanged. Applying Theorem 2 with $D = 1$ eliminates the key-collision term and gives the stated bound. $\square$

Hash and committing random oracle bounds. The remaining public permutation statements use the following random oracle bounds for the root tag wrapper and committing games.

896 **CMTDs-4 bound.** The flat-RO CMTDs-4 theorem of Section 6 is the direct root-collision
 897 bound

$$898 \quad \text{Adv}_{\text{RO}}^{\text{CMTDs-4}}(\mathcal{B}) \leq \text{RootColl}_{\tau_{\text{root}}}(q_{\text{RO}}(\mathcal{B}), s) = \binom{q_{\text{RO}}(\mathcal{B}) + s}{s} / 2^{\tau_{\text{root}}(s-1)},$$

899 with no key-guessing or leaf-collision term. The keys are adversary-supplied, so a direct
 900 query on an input accepted by the keyed transcript decoder **KeyedTWO** is one of the
 901 adversary's visible RO_{flat} queries and is counted in $q_{\text{RO}}(\mathcal{B})$; leaf tag collisions cannot merge
 902 the final root transcripts whose tags are compared, since the common ciphertext body
 903 makes pairwise distinct full inputs yield pairwise distinct (K, N, A) triples; by Opening
 904 Triple Separation (Proposition 1) these reach distinct final root transcripts before any leaf
 905 tag is aggregated.

906 **Collision bound.** The flat-RO collision theorem of Section 6 is

$$907 \quad \text{Adv}_{\text{RO}}^{\text{Coll}}(\mathcal{B}) \leq \text{LeafColl}_{\tau_{\text{leaf}}}(q_{\text{RO}}(\mathcal{B}) + n_{\text{leaf}}(\mathcal{B})) + \text{RootColl}_{\tau_{\text{root}}}(q_{\text{RO}}(\mathcal{B}), s).$$

908 The root-collision part is the same final-root candidate argument as CMTDs-4. The
 909 additional known-key leaf-collision term is needed because colliding inputs may use different
 910 ciphertext bodies; absent such a leaf tag collision, pairwise distinct inputs force pairwise
 911 distinct final root transcripts before the common tag is compared.

912 **Second-preimage bound.** The flat-RO always second-preimage theorem of Section 6 is

$$913 \quad \text{Adv}_{\text{RO}}^{\text{aSec}}(\mathcal{B}) \leq \text{RootPre}_{\tau_{\text{root}}}(q_{\text{RO}}(\mathcal{B})) + \text{LeafColl}_{\tau_{\text{leaf}}}(q_{\text{RO}}(\mathcal{B}) + n_{\text{leaf}}(\mathcal{B})).$$

914 The challenge tag is a fixed second-preimage target, so the candidate-collision argument
 915 becomes a candidate-preimage argument against the designated challenge candidate; the
 916 leaf-collision term is the same adversary-supplied-key accounting as the collision bound.

917 **Preimage bound.** The flat-RO preimage theorem of Section 6 is the direct root-preimage
 918 bound

$$919 \quad \text{Adv}_{\text{RO}}^{\text{ePre}}(\mathcal{B}) \leq \text{RootPre}_{\tau_{\text{root}}}(q_{\text{RO}}(\mathcal{B})) = (q_{\text{RO}}(\mathcal{B}) + 1) / 2^{\tau_{\text{root}}},$$

920 again with no key-guessing term, by the same adversary-supplied-key accounting, and no
 921 leaf-collision term, since a single submitted input admits no pairwise leaf tag collision case.
 922 With a single context, the candidate-collision argument becomes a candidate-preimage
 923 argument against the query-independent target tag and no triple-separation is invoked.

924 Section B carries out the lift in two hops: Hop 1 (Lemma 16) replaces the real
 925 permutation by the [BDPVA08] public primitive simulator and the flat-RO construction
 926 interface at cost at most $\text{Lift}_c(N_{\text{tot}})$, with $N_{\text{tot}} = N + N_{\text{prim}}$, and Hop 2 (Lemma 17) couples
 927 the result, without loss, to a derived flat-RO adversary $\mathcal{B} = \mathcal{B}_{\text{derived}}(\mathcal{A})$ whose simulator
 928 answers induce at most one RO_{flat} call each (Lemma 15), so $q_{\text{RO}}(\mathcal{B}) \leq N_{\text{prim}}$. The random
 929 oracle bounds of Section 5.2 are therefore instantiated at N_{prim} rather than N_{tot} ; the
 930 committing, hash, and discoverability challengers' encryption and decryption checks are
 931 construction-internal and counted in N .

932 **Public permutation hash, committing, and discoverability bounds.** Applying the same
 933 lift gives the remaining headline bounds.

934 **Theorem 3** (Collision Resistance of H). *For every $s \geq 2$ and every adversary $\mathcal{A}$ against
 935 the s -way multi-collision resistance of H in the sense of [BH22], with the resources of
 936 Section 4.3,*

$$937 \quad \text{Adv}_{\text{TW128}}^{\text{Coll}}(\mathcal{A}) \leq \text{Lift}_c(N_{\text{tot}}) + \text{LeafColl}_{\tau_{\text{leaf}}}(N_{\text{prim}} + N_{\text{leaf}}) + \text{RootColl}_{\tau_{\text{root}}}(N_{\text{prim}}, s).$$

Proof. By the Lift Application corollary (Corollary 8) for the collision game, the derived random oracle adversary $\mathcal{B}$ has $q_{\text{RO}}(\mathcal{B}) \leq N_{\text{prim}}$ and $n_{\text{leaf}}(\mathcal{B}) \leq N_{\text{leaf}}$. The random oracle bound of Theorem 8 is nondecreasing in both $q_{\text{RO}}(\mathcal{B})$ and $n_{\text{leaf}}(\mathcal{B})$, so the monotone form of Corollary 8 gives the stated bound. The s colliding inputs may differ in their messages, hence in their ciphertext bodies; the LeafColl term accounts for the only way two distinct inputs sharing a context (K, N, A) can reach the same final root transcript—a collision among the leaf tags it absorbs—while distinct contexts already reach distinct final root transcripts by Opening Triple Separation (Proposition 1). $\square$

Corollary 2 (Two-Way Collision Resistance of $\mathbf{H}$). *Specializing Theorem 3 to $s = 2$ gives the classical two-way collision resistance of [RS04]: for every collision adversary $\mathcal{A}$ against $\mathbf{H}$ with the resources of Section 4.3,*

$$\text{Adv}_{\text{TW128}}^{\text{Coll}}(\mathcal{A}) \leq \text{Lift}_c(N_{\text{tot}}) + \text{LeafColl}_{\tau_{\text{leaf}}}(N_{\text{prim}} + N_{\text{leaf}}) + \text{RootColl}_{\tau_{\text{root}}}(N_{\text{prim}}, 2),$$

with $\text{RootColl}_{\tau_{\text{root}}}(N_{\text{prim}}, 2) = \binom{N_{\text{prim}}+2}{2} / 2^{\tau_{\text{root}}}$.

Proof. The $s = 2$ case of Theorem 3; [BH22] note that s -way multi-collision resistance at $s = 2$ is the classical collision resistance of [RS04]. $\square$

Theorem 4 (Always Second-Preimage Resistance of $\mathbf{H}$). *For every aSec adversary $\mathcal{A}$ against $\mathbf{H}$ with the resources of Section 4.3,*

$$\text{Adv}_{\text{TW128}}^{\text{aSec}}(\mathcal{A}) \leq \text{Lift}_c(N_{\text{tot}}) + \text{RootPre}_{\tau_{\text{root}}}(N_{\text{prim}}) + \text{LeafColl}_{\tau_{\text{leaf}}}(N_{\text{prim}} + N_{\text{leaf}}).$$

Proof. By the Lift Application corollary (Corollary 8) for the aSec game, the derived random oracle adversary $\mathcal{B}$ has $q_{\text{RO}}(\mathcal{B}) \leq N_{\text{prim}}$ and $n_{\text{leaf}}(\mathcal{B}) \leq N_{\text{leaf}}$; the random oracle bound of Theorem 10 is nondecreasing in both, so the monotone form of Corollary 8 gives the stated bound. Collision resistance (Theorem 3) already implies eSec and hence Sec in the sense of [RS04], both with the looser root-collision term; the always-notion is the one [RS04] shows collision resistance does not imply. Here the second preimage is sought against the single fixed tag $\mathbf{H}(X)$ of the challenge input X , yielding the linear root-preimage term. $\square$

Corollary 3 (CMTDs-4 Committing Security). *For every $s \geq 2$ and every s -way CMTDs-4 adversary $\mathcal{A}$ against TW128 with the resources of Section 4.3,*

$$\text{Adv}_{\text{TW128}}^{\text{CMTDs-4}}(\mathcal{A}) \leq \text{Lift}_c(N_{\text{tot}}) + \text{RootColl}_{\tau_{\text{root}}}(N_{\text{prim}}, s) = \text{Lift}_c(N_{\text{tot}}) + \frac{\binom{N_{\text{prim}}+s}{s}}{2^{\tau_{\text{root}}(s-1)}}.$$

Proof. A CMTDs-4 winner exhibits s distinct inputs (K_i, N_i, A_i, M_i) whose encryptions share one common ciphertext $C \parallel T$; in particular they share the tag T , so they are an s -way collision of $\mathbf{H}$, and $\text{Adv}_{\text{TW128}}^{\text{CMTDs-4}} \leq \text{Adv}_{\text{TW128}}^{\text{Coll}}$. For the displayed no-LeafColl bound, use the dedicated root-collision derivation of Theorem 7 for this all-bodies-equal case, lifted by Corollary 8: the single shared body C forces the triples (K_i, N_i, A_i) pairwise distinct—by Opening Triple Separation (Proposition 1)—so a leaf tag collision cannot merge two openings and no LeafColl term arises. $\square$

Theorem 5 (Everywhere Preimage Resistance of $\mathbf{H}$). *For every target tag $T^* \in \{0, 1\}^{\tau_{\text{root}}}$ fixed in advance and every preimage adversary $\mathcal{A}$ against $\mathbf{H}$ with the resources of Section 4.3,*

$$\text{Adv}_{\text{TW128}}^{\text{ePre}}(\mathcal{A}) \leq \text{Lift}_c(N_{\text{tot}}) + \text{RootPre}_{\tau_{\text{root}}}(N_{\text{prim}}) = \text{Lift}_c(N_{\text{tot}}) + \frac{N_{\text{prim}} + 1}{2^{\tau_{\text{root}}}}.$$

Proof. By the Lift Application corollary (Corollary 8) for the preimage game—whose target T^* is fixed before the primitive sampling—the derived random oracle adversary $\mathcal{B}$ has $q_{\text{RO}}(\mathcal{B}) \leq N_{\text{prim}}$, and the random oracle bound of Theorem 9 is nondecreasing in $q_{\text{RO}}(\mathcal{B})$, so the monotone form of Corollary 8 gives the stated bound. With one fixed target tag, Theorem 9 applies a candidate-preimage bound against that target and drops the exponent from $\tau_{\text{root}}(s-1)$ to τ_{root} . $\square$

Theorem 6 (Context Discoverability of [MLGR23]). *For every target specifier $\text{ts} \subsetneq \{k, n, a\}$ —so that at least one context component stays hidden—and every adversary $\mathcal{A}$ against TW128, the restricted $\text{CDY}^*[\text{ts}]$ advantage of [MLGR23, Fig. 5]—and in particular each per-component variant CDY_k^* , CDY_n^* , CDY_a^* —satisfies*

$$\text{Adv}_{\text{TW128}}^{\text{CDY}}(\mathcal{A}) \leq \text{Lift}_c(N_{\text{tot}}) + \text{RootPre}_{\tau_{\text{root}}}(N_{\text{prim}}) + \text{KeyGuess}_k(N_{\text{prim}}).$$

Proof. Full reveal $\text{ts} = \{k, n, a\}$ hides no component and is won trivially by any scheme, so the restriction to $\text{ts} \subsetneq \{k, n, a\}$ excludes no meaningful case. The restricted advantage requires $\mathcal{A}$ to win against every query-independent selector, so it is at most the success probability against the selector S^* that samples a fresh context (K_0, N_0, A_0) —each of K_0 , N_0 , and A_0 uniform on a 256-bit space—encrypts a message under it to obtain the target ciphertext $C^* = C \parallel T^*$, and reveals only the ts -components. By the Lift Application corollary (Corollary 8), the derived random oracle adversary $\mathcal{B}$ has $q_{\text{RO}}(\mathcal{B}) \leq N_{\text{prim}}$; it remains to bound the random oracle success against S^* . Here $T^* = \lfloor \text{RF}(R_0) \rfloor_{\tau_{\text{root}}}$ is the root tag of the selector’s final root transcript R_0 under (K_0, N_0, A_0) , and a winning $\mathcal{A}$ outputs a context reaching a final root transcript R with $\text{RF}(R)$ projecting to T^* .

Let Z_0 be any context component hidden by ts —one exists because $\text{ts} \subsetneq \{k, n, a\}$ —which $\mathcal{A}$ must itself supply (for the named variants, the key for $\text{ts} \in \{\emptyset, \{n, a\}\}$, the nonce for $\text{ts} = \{k, a\}$, the associated data for $\text{ts} = \{k, n\}$), and which S^* sampled uniformly on $\{0, 1\}^w$ with $w = 256 \geq k$ and which $\mathcal{A}$ never sees: the only values revealed to it, the body and tag of C^* , are RF projections independent of Z_0 . If $R = R_0$, then $\mathcal{A}$ has reconstructed the selector’s transcript, in particular its Z_0 -component, which by Lemma 6 requires a candidate decoding under KeyedTWOut_{Z_0} to the hidden Z_0 ; since Z_0 stays uniform over the values no direct query has tested, Lemma 9 (hidden component Z_0 , $w = 256$) bounds this by $\text{KeyGuess}_k(N_{\text{prim}})$. If $R \neq R_0$, the win is that R collides with R_0 on its τ_{root} -bit projection, which the Root tag Candidate Sequence lemma (Lemma 1) and the second-preimage corollary of Lemma 8, with the designated candidate $\text{flat}(R_0)$, bound by $\text{RootPre}_{\tau_{\text{root}}}(N_{\text{prim}})$. Adding the lift gives the bound, uniformly over ts . $\square$

Corollary 4 (Root Tag as a Compact Commitment). *The wrapper H is a secure hash of the full encryption context (K, N, A, M) in the sense of [RS04]: collision-resistant (Theorem 3), second-preimage-resistant (Theorem 4), and preimage-resistant (Theorem 5). At the implemented parameters all three advantages are at most 2^{-128} (Corollary 5), so the single τ_{root} -bit root tag is a binding, preimage-resistant compact commitment to (K, N, A, M) , native to the mode.*

Proof. Collision resistance is Theorem 3, second-preimage resistance Theorem 4, and preimage resistance Theorem 5; all three quantify the adversary’s success at the same τ_{root} -bit projection of the final root transcript, and Corollary 5 evaluates them at the implemented parameters. Binding (collision resistance) is the property a compact commitment must have, and preimage resistance certifies that the tag alone reveals no opening; together they are the guarantees of a hash-based commitment. $\square$

5.3 Concrete TW128 Bounds

The implemented parameters of Section 3.2 fix

$$k = c = \tau_{\text{root}} = \tau_{\text{leaf}} = 256.$$

Substituting into Corollaries 1 and 3 and Theorems 2 to 6 and using the random oracle bounds of Section 5.2 gives the following closed-form bounds.

All-in-one AE.

$$\text{Adv}_{\text{TW128}}^{\text{ae}}(\mathcal{A}) \leq \frac{N_{\text{tot}}(N_{\text{tot}} + 1)}{2^{257}} + \frac{N_{\text{prim}} + Q_v}{2^{256}} + \frac{N_{\text{leaf}}(N_{\text{leaf}} - 1)}{2^{257}}.$$

Multi-user indistinguishability.

$$\text{Adv}_{\text{TW128}}^{\text{mu-ind}}(\mathcal{A}) \leq \frac{N_{\text{tot}}(N_{\text{tot}} + 1)}{2^{257}} + \frac{DN_{\text{prim}} + Q_v}{2^{256}} + \frac{D(D - 1) + N_{\text{leaf}}(N_{\text{leaf}} - 1)}{2^{257}}.$$

s-way collision resistance of H. For every $s \geq 2$,

$$\text{Adv}_{\text{TW128}}^{\text{Coll}}(\mathcal{A}) \leq \frac{N_{\text{tot}}(N_{\text{tot}} + 1)}{2^{257}} + \frac{\binom{N_{\text{prim}} + N_{\text{leaf}}}{2}}{2^{256}} + \frac{\binom{N_{\text{prim}} + s}{s}}{2^{256(s-1)}}.$$

For $s = 2$, this is the ordinary two-input collision bound:

$$\begin{aligned} \text{Adv}_{\text{TW128}}^{\text{Coll}}(\mathcal{A}) &\leq \frac{N_{\text{tot}}(N_{\text{tot}} + 1)}{2^{257}} + \frac{(N_{\text{prim}} + N_{\text{leaf}})(N_{\text{prim}} + N_{\text{leaf}} - 1)}{2^{257}} \\ &\quad + \frac{(N_{\text{prim}} + 1)(N_{\text{prim}} + 2)}{2^{257}}. \end{aligned}$$

Always second-preimage resistance of H. The root-preimage contribution is linear in the primitive-query budget:

$$\text{Adv}_{\text{TW128}}^{\text{aSec}}(\mathcal{A}) \leq \frac{N_{\text{tot}}(N_{\text{tot}} + 1)}{2^{257}} + \frac{N_{\text{prim}} + 1}{2^{256}} + \frac{\binom{N_{\text{prim}} + N_{\text{leaf}}}{2}}{2^{256}}.$$

s-way CMTDs-4 committing security. For every $s \geq 2$, the all-bodies-equal special case drops the leaf tag term:

$$\text{Adv}_{\text{TW128}}^{\text{CMTDs-4}}(\mathcal{A}) \leq \frac{N_{\text{tot}}(N_{\text{tot}} + 1)}{2^{257}} + \frac{\binom{N_{\text{prim}} + s}{s}}{2^{256(s-1)}}.$$

Remark 1 (Two-Key CMTDs-4 Specialization). Setting $s = 2$ recovers the ordinary two-key CMTDs-4 statement of [BH22]:

$$\text{Adv}_{\text{TW128}}^{\text{CMTDs-4}}(\mathcal{A}) \leq \frac{N_{\text{tot}}(N_{\text{tot}} + 1)}{2^{257}} + \frac{\binom{N_{\text{prim}} + 2}{2}}{2^{256}} = \frac{N_{\text{tot}}(N_{\text{tot}} + 1) + (N_{\text{prim}} + 1)(N_{\text{prim}} + 2)}{2^{257}}.$$

Everywhere preimage resistance of H.

$$\text{Adv}_{\text{TW128}}^{\text{ePre}}(\mathcal{A}) \leq \frac{N_{\text{tot}}(N_{\text{tot}} + 1)}{2^{257}} + \frac{N_{\text{prim}} + 1}{2^{256}}.$$

Context discoverability ([MLGR23]).

$$\text{Adv}_{\text{TW128}}^{\text{CDY}}(\mathcal{A}) \leq \frac{N_{\text{tot}}(N_{\text{tot}} + 1)}{2^{257}} + \frac{2N_{\text{prim}} + 1}{2^{256}}.$$

Table 4: Concrete TW128 security at the work cap $N_{\text{tot}} \leq 2^{63}$ (with $D \leq 2^{63}$ for mu-ind), under $k = c = \tau_{\text{root}} = \tau_{\text{leaf}} = 256$ and the resource consequences $N_{\text{prim}}, Q_v, N_{\text{leaf}} \leq N_{\text{tot}}$ and $N_{\text{prim}} + N_{\text{leaf}} \leq N_{\text{tot}}$ of Section 4.3. The committing rows hold for every $s \geq 2$; larger s only shrinks the root-collision term.

Notion	Dominant scale	At the cap
All-in-one AE	$\text{Lift}_c(N_{\text{tot}})$	$\leq 2^{-128}$
Multi-user ind.	$D \text{KeyGuess}_k(N_{\text{prim}})$	$\leq 2^{-128}$
H collision (Coll)	$\text{Lift}_c(N_{\text{tot}})$	$\leq 2^{-128}$
H second-preimage (aSec)	$\text{Lift}_c(N_{\text{tot}})$	$\leq 2^{-128}$
H preimage (ePre)	$\text{Lift}_c(N_{\text{tot}})$	$\leq 2^{-128}$
CMTDs-4 [BH22]	$\text{Lift}_c(N_{\text{tot}})$	$\leq 2^{-128}$
CDY [MLGR23]	$\text{Lift}_c(N_{\text{tot}})$	$\leq 2^{-128}$

5.4 Concrete Security Claim

Corollary 5 (Concrete TW128 Security). *Under the implemented parameters $k = c = \tau_{\text{root}} = \tau_{\text{leaf}} = 256$, every adversary with work cap $N_{\text{tot}} \leq 2^{63}$ —and, in the multi-user case, user cap $D \leq 2^{63}$ —has advantage at most 2^{-128} in each of TW128’s notions: multi-user indistinguishability (Theorem 2), all-in-one AE (Corollary 1), the [RS04] hash security of H— s -way collision resistance for every $s \geq 2$, always second-preimage resistance, and preimage resistance (Theorems 3 to 5)—the [BH22] CMTDs-4 committing bound it yields (Corollary 3), and [MLGR23] context discoverability (Theorem 6).*

Proof. The claim instantiates the closed forms of Section 5.3 under the resource consequences of Section 4.3: $N_{\text{prim}} \leq N_{\text{tot}}$ by definition, and the pathwise bounds $q_v \leq N$ and $n_{\text{leaf}} \leq N$ let the resource caps Q_v and N_{leaf} be taken at most N_{tot} , so no separate top-level cap on Q_v or N_{leaf} is required; because $N_{\text{tot}} = N + N_{\text{prim}}$ counts construction-internal calls and direct queries disjointly, the same bound $n_{\text{leaf}} \leq N$ yields the joint constraint $N_{\text{prim}} + N_{\text{leaf}} \leq N_{\text{tot}}$. Only the mu-ind user cap D remains an independent hypothesis, since created-but-idle users need not contribute to N . At the headline work cap $N_{\text{tot}} \leq 2^{63}$ (and $D \leq 2^{63}$ for mu-ind) every bound is below 2^{-128} , with the flat sponge loss $\text{Lift}_c(N_{\text{tot}})$ the dominant scale in all but the multi-user case (Table 4).

Term by term at the cap, with $N_{\text{prim}}, Q_v, N_{\text{leaf}} \leq N_{\text{tot}}$: the flat sponge loss $\text{Lift}_c(N_{\text{tot}})$ and any leaf tag birthday term $N_{\text{leaf}}(N_{\text{leaf}} - 1)/2^{257}$ are each at most one 2^{-131} -sized term; the AE linear key-guess and verification contributions $(N_{\text{prim}} + Q_v)/2^{256}$ are at most $2 \cdot 2^{-193}$; the root-collision term $\text{RootColl}_{256}(N_{\text{prim}}, s) = \binom{N_{\text{prim}} + s}{s}/2^{256(s-1)}$ likewise matches Lift_c as a 2^{-131} -sized term in its two-key ($s = 2$) specialization and is smaller still for $s \geq 3$ under the cap; the known-key $\text{LeafColl}_{256}(N_{\text{prim}} + N_{\text{leaf}})$, which the joint cap $N_{\text{prim}} + N_{\text{leaf}} \leq N_{\text{tot}}$ bounds by $\binom{N_{\text{tot}}}{2}/2^{256}$, is likewise $\leq 2^{-131}$; and the root-preimage term $(N_{\text{prim}} + 1)/2^{256}$ is at most $2 \cdot 2^{-193}$. The single-user totals are therefore dominated by Lift_c and stay below 2^{-128} . The multi-user bound is the exception: its hybrid-induced term $D \cdot \text{KeyGuess}_k(N_{\text{prim}}) \leq DN_{\text{tot}}/2^{256}$ reaches $\approx 2\text{Lift}_c$ at the upper end and dominates—staying at most Lift_c once $D \leq N_{\text{tot}}(N_{\text{tot}} + 1)/(2N_{\text{prim}})$ (for $N_{\text{prim}} > 0$; when $N_{\text{prim}} = 0$ the hybrid-induced term vanishes), typical of any practical deployment—and the total there is at most $5 \cdot 2^{-131} + 2 \cdot 2^{-194} < 8 \cdot 2^{-131} = 2^{-128}$. Expressed in bytes, $N_{\text{tot}} \leq 2^{63}$ corresponds to at most $\rho \cdot 2^{63} = 167 \cdot 2^{63} < 2^{71}$ bytes of plaintext, associated data, ciphertext, and aggregation through the duplex (Section 3.2), well beyond any realistic deployment lifetime. $\square$

5.5 Discussion of the Bounds

The bounds of Section 5.3 aggregate the loss terms of Section 5.1, whose sizes at the work cap Table 4 records and whose dominance Section 5.4 discusses. We comment on their tightness.

Tightness.

- $\text{KeyGuess}_k(Q) = Q/2^k$. Tight at constants: a direct random-guess attack on a k -bit key matches the per-query success probability $1/2^k$. The multi-user bound's dominant term $D \text{KeyGuess}_k(N_{\text{prim}}) = DN_{\text{prim}}/2^k$ is likewise tight at the leading constant: against the D independent user keys each direct guess matches some key with probability $D/2^k$, so N_{prim} guesses recover a user key—and distinguish—with probability $\approx DN_{\text{prim}}/2^k$, matching the bound.
- $\text{RootColl}_\tau(Q, s) = \binom{Q+s}{s}/2^{\tau(s-1)}$. Tight up to lower-order terms: it is the standard s -way multi-collision bound on a τ -bit projection of RO_{flat} , and a generic birthday-style adversary against the random oracle matches the leading-order term. It is the root term of the collision bound (H Coll, Theorem 3). Its N_{prim} counts every primitive query, though by Output-Point Separation (Corollary 7) only those whose induced RO_{flat} call lands at a root tag output point ($\mathcal{R}_{\text{tag}}$, or $\mathcal{R}_{\text{msg}}^+$ for single-chunk ciphertexts) can enter the candidate sequence; a class-filtered count would be no larger, and potentially smaller, but the headline keeps the N_{prim} form for uniformity with Corollary 1 and Theorem 2, the birthday threshold at $\tau_{\text{root}} = 256$ being unaffected.
- $\text{RootPre}_\tau(Q) = (Q+1)/2^\tau$. Tight at constants: a generic preimage search against a fixed τ -bit target matches the per-query $2^{-\tau}$ rate. It is the root term of the preimage (ePre), always second-preimage (aSec), and discoverability bounds, where the single fixed target yields a linear term.
- *Leaf tag collision* $N_{\text{leaf}}(N_{\text{leaf}} - 1)/2^{\tau_{\text{leaf}}+1}$. Tight for verification security: by reusing a nonce, an adversary detects a leaf tag birthday collision through the matching root tags and splices the colliding leaf chunk into a fresh ciphertext that still verifies, forging at the leading $N_{\text{leaf}}^2/2^{\tau_{\text{leaf}}+1}$ rate. The mu-ind and all-in-one AE bounds carry the term through the verification side of their real-vs-ideal coupling; their nonce-respecting encryption rules this attack out, so there it is inherited slack.
- *Known-key leaf collision* $\text{LeafColl}_\tau(Q)$. The same birthday attack applies when keys are adversary-supplied and direct queries can target leaf tag points; the collision and always second-preimage proofs use $Q = N_{\text{prim}} + N_{\text{leaf}}$.
- *Flat sponge lift*. The term $\text{Lift}_c(N_{\text{tot}}) = N_{\text{tot}}(N_{\text{tot}} + 1)/2^{c+1}$ dominates the random-permutation differentiating bound $f_P(N_{\text{tot}})$ of [BDPVA08, Lemma 5] via Lemma 16. The remaining slack is the unimproved tightness of the [BDPVA08] simulator itself, which is an open problem on the sponge side and not specific to TW128.

Cross-scheme bound-level comparisons against prior sponge AEADs and the committing AEAD line of [BH22] are deferred to Sections 8.1 to 8.3.

6 Hash Security of the Root Tag

The tag-level guarantees of this section share one spine. The Root tag Candidate Sequence lemma (Lemma 1) collects the root tag inputs reached by a keyless game into a chronological candidate sequence, and the Root Tag Hash Bound (Lemma 2) turns a collision, preimage, or second-preimage among those candidates into the matching adaptive bound of Section A,

controlling the chance that distinct inputs land on the same τ_{root} -bit root tag projection. The four random oracle theorems—CMTDs-4 (Theorem 7), collision (Theorem 8), preimage (Theorem 9), and always second-preimage (Theorem 10)—are then short instantiations, differing only in which facet they invoke and in whether a leaf tag collision term is needed when colliding inputs may use different ciphertext bodies. Section B transfers each instance to the public permutation setting, Section 5.2 records the concrete bounds, and Corollary 6 forwards the remaining [BH22] committing notions to the same root-collision bound.

The instantiations rest on two structural separation facts, recorded first: distinct contexts (K, N, A) reach distinct final root transcripts regardless of the leaf tags they aggregate (Opening Triple Separation, Proposition 1), and—absent a collision among those leaf tags—so do distinct full inputs (Final-Root Input Separation, Proposition 2).

Proposition 1 (Opening Triple Separation). *If two TW128 root computations have distinct (K, N, A) triples, then their bridged final root transcript inputs under $\text{flat}(\cdot)$ are distinct, regardless of whether each computation is an encryption or a decryption check on a common ciphertext.*

Proof. Suppose for contradiction that $\text{flat}(R_i) = \text{flat}(R_j)$. Since the bridge $\text{flat}(\cdot)$ is injective on well-formed transcript prefixes (Lemma 6), the native flattened inputs of the two final-root prefixes are equal. Then Lemma 4 recovers the full duplex-call sequence and reads off the pair $(\sigma, \sigma_{\text{ph}})$ for each call. The seven 4-bit phase suffixes are pairwise distinct, so the σ_{init} , $\sigma_{\text{ad}}^{\pm}$, $\sigma_{\text{msg}}^{\pm}$, and (when present) $\sigma_{\text{agg}}^{\pm}$ calls in the recovered sequence are unambiguously identified. The aggregation phase is present exactly when $n \geq 1$, and $n = \text{leaves}(C)$ is determined by the common ciphertext, so the two recovered sequences either both carry an aggregation phase or both omit it; either way the recovery of (K, N, A) below uses only the σ_{init} and AD calls, which precede any message or aggregation block.

Equality of the recovered sequences forces equality of the σ_{init} call, whose σ is $p_{\text{root}} \parallel K \parallel N$ with fixed-width K and N , hence $K_i = K_j$ and $N_i = N_j$. It also forces equality of the recovered AD subsequences. By the AD bullet of Lemma 4 the associated data phase contributes a σ_{ad}^+ -tagged block exactly when the associated data is nonempty, so the recovered sequences agree on the presence of the AD phase. If exactly one of A_i, A_j were empty the recovered sequences would differ in the presence of a σ_{ad}^+ call, contradicting their equality; hence either both are empty, giving $A_i = A_j = \varepsilon$, or both are nonempty, in which case the phase-recovery sub-claim of Lemma 4 inverts the ρ -byte-block partition map on the AD byte string and equal AD σ -block sequences imply $A_i = A_j$. In all cases $(K_i, N_i, A_i) = (K_j, N_j, A_j)$, contradicting the hypothesis of distinct triples. $\square$

Triple Separation handles distinct contexts; when inputs share a context but differ in their messages, distinctness of the final root transcripts needs the absence of a leaf tag collision, recorded next.

Proposition 2 (Final-Root Input Separation). *Let $\text{bad}_{\text{leaf}}^*$ be the event that two distinct construction-generated final leaf transcripts receive the same τ_{leaf} -bit RF projection. On $\neg \text{bad}_{\text{leaf}}^*$, if two TW128 construction computations reach final root transcripts with $\text{flat}(R) = \text{flat}(R')$, then they share the same context (K, N, A) and the same full ciphertext C . Equivalently, on $\neg \text{bad}_{\text{leaf}}^*$ computations with distinct (K, N, A, C) reach distinct bridged final root transcripts.*

Proof. By bridge injectivity (Lemma 6), $\text{flat}(R) = \text{flat}(R')$ gives equal native flattened final-root inputs. Lemma 4 then recovers the same root initialization $p_{\text{root}} \parallel K \parallel N$ —hence the same K and N —the same associated data blocks—hence the same A —the same root ciphertext chunk, and the same leaf count and ordered leaf tag sequence. Under $\neg \text{bad}_{\text{leaf}}^*$, Leaf-Transcript Recovery (Lemma 5) identifies the equal leaf tag sequences with the same final leaf transcripts, hence the same leaf ciphertext bodies; with the common root ciphertext chunk this gives the same full C . $\square$

Lemma 1 (Root tag Candidate Sequence). *Consider a random oracle game of Section 4.2 in which the adversary $\mathcal{B}$ makes direct RO_{flat} queries and the challenger runs $m \geq 1$ deterministic encryption or decryption checks, the i -th reaching a final root transcript R_i whose root tag is read by a single RO_{flat} lookup at $\text{flat}(R_i)$. Every execution then admits a chronological sequence of root tag inputs such that*

1. *each distinct direct query whose input Y decodes under KeyedTWOut to a root tag output-point class— $\mathcal{R}_{\text{msg}}^+$ (the final root message block, which carries the root tag only on the no-leaf path) or $\mathcal{R}_{\text{tag}}$ (the aggregated root tag)—is appended as a candidate before its RO_{flat} lookup, and contributes at most one candidate; multi-chunk computations also reach $\mathcal{R}_{\text{msg}}^+$ with zero requested output before aggregation, so this predicate over-approximates the root tag candidate set, which only loosens the count;*
2. *if the challenger reaches its checks, each $\text{flat}(R_i)$ is in the sequence, appended—if not already present—immediately before the i -th root tag answer is sampled;*
3. *no candidate's RO_{flat} value is read before its append step, so the sequence satisfies the predictability and unrevealed-value premises of Lemma 8; and*
4. *its length is at most $q_{\text{RO}}(\mathcal{B}) + m$.*

Proof. The committing, discoverability, and hash (collision, second-preimage, and preimage) games are keyless: all keys are adversary-supplied, so a direct query to a keyed transcript is one of $\mathcal{B}$'s visible RO_{flat} queries counted in $q_{\text{RO}}(\mathcal{B})$, and there is no bad_{key} event. No construction oracle runs while $\mathcal{B}$ issues its direct queries, so each direct-query candidate is first read no earlier than its append step. The append predicate depends on Y alone through the exact decoder of Lemma 6, which ignores the requested output length, and by Output-Point Separation (Corollary 7) each input decodes to at most one output-point class; hence each distinct direct query contributes at most one candidate, giving (1).

If the challenger reaches its m checks, append $\text{flat}(R_i)$ immediately before the i -th root tag answer is sampled, unless it is already present. The intermediate RO_{flat} lookups of check i lie at non-root tag output points, so by Corollary 7 they do not read $\text{flat}(R_i)$. If $\text{flat}(R_i)$ was touched earlier—by a direct query or by an earlier check $j < i$ —then Lemma 6 and Corollary 7 make that earlier access a final-root lookup at the same address, which was therefore already appended, so the input is in the sequence and is not appended again. Otherwise its value is still unread when appended. Either way no candidate's value is read before its append step, giving (2) and (3). At most $q_{\text{RO}}(\mathcal{B})$ direct candidates and m challenger candidates are appended, giving (4). $\square$

Lemma 2 (Root Tag Hash Bound). *Consider a keyless TW128 random oracle game (Section 4.2) in which, after $\mathcal{B}$'s direct RO_{flat} queries, the challenger runs deterministic encryption or decryption checks reaching final root transcripts, with win event W . By the Root tag Candidate Sequence lemma (Lemma 1) the root tag inputs form a chronological candidate sequence meeting the premises of Lemma 8 with $\tau = \tau_{\text{root}}$. Writing $q_{\text{RO}} = q_{\text{RO}}(\mathcal{B})$:*

1. *if W forces s checks to reach pairwise distinct bridged final root transcripts carrying a common root tag, then, with $m = s$ checks, $\Pr[W] \leq \text{RootColl}_{\tau_{\text{root}}}(q_{\text{RO}}, s)$;*
2. *if W forces a single check ($m = 1$) to carry a root tag equal to a target fixed independently of RO_{flat} , then $\Pr[W] \leq \text{RootPre}_{\tau_{\text{root}}}(q_{\text{RO}})$;*
3. *if a designated check reaches R_X whose root tag is exposed before $\mathcal{B}$ outputs, and W forces a single further check ($m = 1$ besides R_X) to reach a transcript $R' \neq R_X$ carrying R_X 's root tag, then $\Pr[W] \leq \text{RootPre}_{\tau_{\text{root}}}(q_{\text{RO}})$.*

Proof. In each case the candidate sequence has the stated length and meets the premises of Lemma 8. (1) The s distinct candidates carrying a common τ_{root} -projection are an s -way collision; the collision facet with $L = q_{\text{RO}} + s$ gives $\binom{q_{\text{RO}} + s}{s} \cdot 2^{-\tau_{\text{root}}(s-1)} = \text{RootColl}_{\tau_{\text{root}}}(q_{\text{RO}}, s)$. (2) The candidate hitting the fixed target is a preimage; the preimage facet with $L = q_{\text{RO}} + 1$ gives $(q_{\text{RO}} + 1) \cdot 2^{-\tau_{\text{root}}} = \text{RootPre}_{\tau_{\text{root}}}(q_{\text{RO}})$. (3) Prepend $\text{flat}(R_X)$ as the designated candidate, its value the exposed root tag; the further candidate colliding with it on its τ_{root} -projection is a second-preimage, and the second-preimage corollary with $L = q_{\text{RO}} + 2$ gives $(q_{\text{RO}} + 1) \cdot 2^{-\tau_{\text{root}}} = \text{RootPre}_{\tau_{\text{root}}}(q_{\text{RO}})$. $\square$

Theorem 7 (Random Oracle CMTDs-4). *For every $s \geq 2$ and every s -way CMTDs-4 adversary $\mathcal{B}$ in the TW128 random oracle model (Section 4.2),*

$$\text{Adv}_{\text{RO}}^{\text{CMTDs-4}}(\mathcal{B}) \leq \text{RootColl}_{\tau_{\text{root}}}(q_{\text{RO}}(\mathcal{B}), s) = \frac{\binom{q_{\text{RO}}(\mathcal{B}) + s}{s}}{2^{\tau_{\text{root}}(s-1)}}.$$

Proof. Run the random oracle CMTDs-4 game. After rejecting on any failed syntax, message-length, or pairwise-full-input check, the challenger performs s deterministic decryption checks $\text{Dec}_{K_i}^{\text{RO}}(N_i, A_i, C \| T)$ on the one common body $C \| T$, reaching $R_1, \dots, R_s$. On a winning execution the tuples (K_i, N_i, A_i, M_i) are pairwise distinct, and since all checks decrypt the common body, determinism of Dec makes the triples (K_i, N_i, A_i) pairwise distinct—equal triples would return equal messages, hence equal tuples. By Opening Triple Separation (Proposition 1) the candidates $\text{flat}(R_1), \dots, \text{flat}(R_s)$ are pairwise distinct (leaf tag collisions cannot merge them), and every accepting check carries the common root tag T . The Root Tag Hash Bound (Lemma 2, part 1) gives $\text{Adv}_{\text{RO}}^{\text{CMTDs-4}}(\mathcal{B}) \leq \text{RootColl}_{\tau_{\text{root}}}(q_{\text{RO}}(\mathcal{B}), s)$. $\square$

Remaining committing notions. The hierarchy and s -way extension of the [BH22] committing notions are given in [BH22, §3, Figs. 4–5, and App. A]. We use those relations for the D-notions and a direct source-game argument for the CMTs- ℓ notions to preserve the resource split of Section 4.3.

Corollary 6 (Committing-Notion Root-Collision Bound). *For every $s \geq 2$, every*

$$X \in \{\text{CMTDs-1, CMTDs-3, CMTDs-4, CMTs-1, CMTs-3, CMTs-4}\},$$

and every s -way X -adversary $\mathcal{A}$ against TW128,

$$\text{Adv}_{\text{TW128}}^X(\mathcal{A}) \leq \text{Lift}_c(N_{\text{tot}}) + \text{RootColl}_{\tau_{\text{root}}}(N_{\text{prim}}, s) = \text{Lift}_c(N_{\text{tot}}) + \frac{\binom{N_{\text{prim}} + s}{s}}{2^{\tau_{\text{root}}(s-1)}}.$$

Proof. CMTDs-4 is bounded directly by Corollary 3. For $\ell \in \{1, 3\}$, the s -way hierarchy of [BH22, §3, Fig. 5 and App. A] gives $\text{Adv}_{\text{TW128}}^{\text{CMTDs-}\ell}(\mathcal{A}) \leq \text{Adv}_{\text{TW128}}^{\text{CMTDs-4}}(\mathcal{A})$ for the same output distribution, so the resource counts are unchanged.

The E-notions are bounded by the Root Tag Hash Bound, with the challenger's s deterministic *encryption* checks supplying the candidate final-root transcripts (Lemma 1 builds the candidate sequence for encryption or decryption checks alike). On a winning execution all s encryptions are equal, hence of equal length and with one common root tag T ; pairwise WiC_ℓ -distinctness forces pairwise distinct triples—for $\ell = 4$ via determinism, since equal triples with equal ciphertexts would decrypt to equal messages—so Proposition 1 makes the s final-root candidates distinct while their τ_{root} -bit projections all equal T . By the Lift Application corollary (Corollary 8) for the X game, the derived random oracle adversary $\mathcal{B}$ has $q_{\text{RO}}(\mathcal{B}) \leq N_{\text{prim}}$; the Root Tag Hash Bound (Lemma 2, part 1) then bounds its win by $\text{RootColl}_{\tau_{\text{root}}}(q_{\text{RO}}(\mathcal{B}), s)$ in the random oracle model, which the lift transfers to the public permutation setting at N_{prim} ; the encryption checks are construction-internal work counted in N , not in N_{prim} . $\square$

Theorem 8 (Random Oracle Collision Resistance of H). *For every $s \geq 2$ and every s -way collision adversary $\mathcal{B}$ in the TW128 random oracle model,*

$$\text{Adv}_{\text{RO}}^{\text{Coll}}(\mathcal{B}) \leq \text{LeafColl}_{\tau_{\text{leaf}}}(q_{\text{RO}}(\mathcal{B}) + n_{\text{leaf}}(\mathcal{B})) + \text{RootColl}_{\tau_{\text{root}}}(q_{\text{RO}}(\mathcal{B}), s).$$

Proof. Run the random oracle collision game. Let $\text{bad}_{\text{leaf}}^*$ be the event that two distinct construction-generated final leaf transcripts receive the same τ_{leaf} -bit RF projection; by the adversary-supplied-key bound of Lemma 10, $\Pr[\text{bad}_{\text{leaf}}^*] \leq \text{LeafColl}_{\tau_{\text{leaf}}}(q_{\text{RO}}(\mathcal{B}) + n_{\text{leaf}}(\mathcal{B}))$, so it remains to bound the win under $\neg \text{bad}_{\text{leaf}}^*$. After rejecting unless the s output inputs are pairwise distinct and in the TW128 domain, the challenger performs s deterministic encryptions $\text{Enc}_{K_i}^{\text{RO}}(N_i, A_i, M_i)$, reaching $R_1, \dots, R_s$, all carrying a common root tag T on a win. Pairwise distinct tuples (K_i, N_i, A_i, M_i) have pairwise distinct (K_i, N_i, A_i, C_i) : two tuples differing in their context already differ in (K, N, A) , while two sharing a context but differing in the message yield distinct ciphertexts, since for a fixed context encryption maps the message to the ciphertext bijectively. So on $\neg \text{bad}_{\text{leaf}}^*$ Final-Root Input Separation (Proposition 2) makes $\text{flat}(R_1), \dots, \text{flat}(R_s)$ pairwise distinct. The Root Tag Hash Bound (Lemma 2, part 1) bounds the win under $\neg \text{bad}_{\text{leaf}}^*$ by $\text{RootColl}_{\tau_{\text{root}}}(q_{\text{RO}}(\mathcal{B}), s)$; adding $\Pr[\text{bad}_{\text{leaf}}^*]$ gives the theorem. $\square$

Theorem 9 (Random Oracle Everywhere Preimage Resistance of H). *For every query-independent target tag $T^* \in \{0, 1\}^{\tau_{\text{root}}}$ and every preimage adversary $\mathcal{B}$ in the TW128 random oracle model (Section 4.2),*

$$\text{Adv}_{\text{RO}}^{\text{ePre}}(\mathcal{B}) \leq \text{RootPre}_{\tau_{\text{root}}}(q_{\text{RO}}(\mathcal{B})) = \frac{q_{\text{RO}}(\mathcal{B}) + 1}{2^{\tau_{\text{root}}}}.$$

Proof. Run the random oracle preimage game for the query-independent target T^* , which is independent of RO_{flat} . After rejecting unless the output input $X = (K, N, A, M)$ is in the TW128 domain, the challenger performs one encryption $\text{Enc}_K^{\text{RO}}(N, A, M)$, reaching a final root transcript R . A win means $H(X) = T^*$, i.e. the root tag $[\text{RO}_{\text{flat}}(\text{flat}(R))]_{\tau_{\text{root}}}$ at R equals the fixed target T^* . The Root Tag Hash Bound (Lemma 2, part 2, with target T^*) gives $\text{Adv}_{\text{RO}}^{\text{ePre}}(\mathcal{B}) \leq \text{RootPre}_{\tau_{\text{root}}}(q_{\text{RO}}(\mathcal{B}))$. $\square$

Theorem 10 (Random Oracle aSec). *For every aSec adversary $\mathcal{B}$ in the TW128 random oracle model (Section 4.2),*

$$\text{Adv}_{\text{RO}}^{\text{aSec}}(\mathcal{B}) \leq \text{RootPre}_{\tau_{\text{root}}}(q_{\text{RO}}(\mathcal{B})) + \text{LeafColl}_{\tau_{\text{leaf}}}(q_{\text{RO}}(\mathcal{B}) + n_{\text{leaf}}(\mathcal{B})).$$

Proof. Run the random oracle aSec game: the challenger draws the challenge input X , encrypts it to obtain $T = H(X)$ at a final root transcript R_X , and gives X to $\mathcal{B}$, which outputs $X' \neq X$, winning if $H(X') = T$. Let $\text{bad}_{\text{leaf}}^*$ be the event that two distinct construction-generated final leaf transcripts receive the same τ_{leaf} -bit RF projection; by Lemma 10, $\Pr[\text{bad}_{\text{leaf}}^*] \leq \text{LeafColl}_{\tau_{\text{leaf}}}(q_{\text{RO}}(\mathcal{B}) + n_{\text{leaf}}(\mathcal{B}))$, so it remains to bound the win under $\neg \text{bad}_{\text{leaf}}^*$. Encrypting X' reaches a final root transcript R' . If $\text{flat}(R') = \text{flat}(R_X)$ then on $\neg \text{bad}_{\text{leaf}}^*$ Final-Root Input Separation (Proposition 2) makes X and X' share (K, N, A, C) , hence the same message by determinism, so $X' = X$ —a contradiction; thus $R' \neq R_X$. The challenge root tag $T = [\text{RF}(R_X)]_{\tau_{\text{root}}}$ is exposed to $\mathcal{B}$, and a win places R' 's root tag at T . The Root Tag Hash Bound (Lemma 2, part 3, with designated R_X) bounds the win under $\neg \text{bad}_{\text{leaf}}^*$ by $\text{RootPre}_{\tau_{\text{root}}}(q_{\text{RO}}(\mathcal{B}))$; adding $\Pr[\text{bad}_{\text{leaf}}^*]$ gives the theorem. $\square$

7 Performance Evaluation

7.1 Optimized Implementations

We evaluate three implementations of TW128 sharing a common structure: a portable reference written in Go, an **amd64** implementation, and an **arm64** implementation. The two

architecture-specific implementations replace only the KECCAK- p [1600, 12] permutation and the inner absorb/squeeze loops with hand-written assembly; the framing, domain separation, padding, and tree bookkeeping remain in the shared Go code. The portable reference uses a scalar permutation and a scalar eight-instance loop, and serves both as a correctness oracle for the assembly cores and as the fallback on other architectures.

The tree structure of TW128 exposes two distinct permutation workloads, and both architectures provide a separate core for each. The root duplex—which absorbs the associated data, processes chunk 0, and absorbs the leaf tag aggregation transcript—is inherently sequential, since each of its calls consumes the previous call’s output; it is driven by a *single-duplex* ($\times 1$) core that permutes one 1600-bit state at a time, except for its chunk-0 message phase, which is fused into the first batched leaf call as described below. The leaf duplexes, by contrast, operate on disjoint state and disjoint inputs (Section 3), so the implementation batches eight complete leaf chunks and runs them in lockstep through an *eight-way* ($\times 8$) core. Because every leaf chunk is exactly $\beta = 49\rho$ bytes, the eight instances step through an identical sequence of 49 full ρ -blocks (the first 48 closed with MSG_MORE, the last with MSG_LAST), so a single shared block position drives all eight. The eight-way state is held lane-major—25 lanes each holding the corresponding 64-bit word of all eight instances—so that the batched cores can address the eight instances of a lane with a single vector load.

Messages are decomposed into eight-chunk batches run on the $\times 8$ core, and a leftover run of two to seven complete chunks is drained by a narrower kernel: on **amd64** a register-resident kernel processes the remainder in a single call, reading the active chunks in place, and on **arm64** a 2-wide kernel consumes it two chunks at a time. These remainder kernels read each chunk as a sequential stream, so they prefetch as well on the cold tail of a long message as on a cache-resident short one; on the **amd64** host they outperform both repeated single-duplex calls and padding the remainder to a full eight-chunk batch, at every remainder width. Only a single leftover complete chunk, a partial trailing chunk, and the count-aggregation tail fall back to the single-duplex ($\times 1$) core.

Chunk 0 itself runs on the batched kernels rather than serially, a technique we refer to as *lane-0 fusion*. The root duplex’s chunk-0 message phase has the same kernel-visible block schedule as a leaf chunk—the same 49 ρ -blocks with the same closing flags—and differs from a leaf only in its initial state, to which the kernels are indifferent; chunk 0 is also contiguous in the message with the first leaf chunks. The implementation therefore seeds lane 0 of the first batched call with the root state (after initialization and associated data processing) and runs the chunk-0 phase alongside up to seven leaves: in the $\times 8$ core when the message has at least seven complete leaf chunks, and in the **amd64** remainder kernel or the **arm64** 2-wide kernel below that. Lane 0’s output state is written back to the root duplex, which then absorbs the leaf tags produced by the same call; no serial pass over chunk 0 precedes the leaf work. On **amd64** the assembly kernels handle only the 48 full MSG_MORE blocks of each chunk; the final MSG_LAST block and the leaf tag extraction are performed in portable Go, shared with the reference path, which keeps the byte-granular tail logic out of assembly. The **arm64** kernels process the complete chunk, extracting the leaf tags and storing the permutation states back in assembly—the writeback that allows the root duplex to occupy lane 0.

Architecture-specific cores. On **amd64** the cores are built from AVX-512 (selected once at startup from the AVX512VL CPUID feature bit) and AVX2, holding the lane-major state in ZMM or YMM registers and permuting it with no cross-lane shuffles. The steady-state $\times 8$ kernel transposes each rate block into the lane-major layout with contiguous, prefetcher-friendly loads and stores, while the register-resident kernel that drains a two-to-seven-chunk remainder reads its active chunks in place by masked gather and scatter. On **arm64** the cores use NEON with the Armv8.2 SHA-3 extension (FEAT_SHA3, available on Apple Silicon

and Neoverse V1/V2 and N2), one instance per 64-bit lane in the single-duplex core and two instances per 128-bit register in the $\times 8$ and 2-wide kernels. Both batched paths store the permutation state back after the closing block—the writeback that lets the root duplex occupy lane 0. Section C gives the per-instruction detail, including the fused permutation instructions and the remainder-kernel strategies on each architecture.

Constant time. All three implementations are constant time with respect to the key, nonce, and plaintext. The permutation and the XOR-based absorb/squeeze contain no secret-dependent branches and no secret-dependent memory or vector indices. The only indexed table is the round-constant array, addressed by the public round number; the AVX-512 chunk transpose uses contiguous loads and stores, and the gather and scatter that remain—for the partial rate lane and the register-resident remainder kernel—are addressed by the fixed public chunk strides under a mask set by the public chunk count, and the AVX2 remainder kernel’s redirection of inactive instances to its dummy block is likewise selected by the public chunk count. None of these depend on secret data.

7.2 Benchmark Methodology

We measure on two hosts, one per architecture. The **amd64** host is a Google Compute Engine **c4-standard-4** instance with an Intel Xeon Platinum 8581C (Emerald Rapids, 2.30 GHz base, AVX-512 including AVX512VL, AVX512_VBMI2, GFNI, and VAES), two physical cores with simultaneous multithreading enabled, 48 KiB L1d and 32 KiB L1i per core, 2 MiB L2 per core, and a 260 MiB shared L3, running Debian 12 (kernel 6.1) under KVM. The guest does not expose a **cpufreq** interface, so the frequency governor and turbo state are not under our control; we mitigate the resulting variance statistically (below). The **arm64** host is an Apple M4 Pro (10 performance cores and 4 efficiency cores; per performance core, 128 KiB L1i, 64 KiB L1d, and a 16 MiB shared L2) running macOS 26.5.1 (Darwin 25.5), with **FEAT_SHA3** present. macOS exposes no userspace frequency control.

Measurements are collected by a standalone harness (**cmd/cpb**) that locks its goroutine to an OS thread and, for each algorithm, operation, and message length, first calibrates an iteration count that fills at least 100 ms, then records 100 samples, from which it reports the median cycles per byte and median wall-clock throughput together with the interquartile range as a dispersion measure (Table 7). Both metrics are read from the same timed loop, so the cycles-per-byte and throughput figures for a given measurement come from one interleaved run rather than two separately scheduled ones. Each sample reuses the same source and destination buffers, so the reported figures are warm-cache, steady-state rates. The single-threaded measurement leaves the SMT sibling of the **amd64** host idle. We report a sweep of seven message lengths from 1 B to 16 MiB.

The cycle counters differ by architecture, and this affects how the cycles-per-byte figures may be read. On **amd64** the harness reads **RDTSC**; on this part the time-stamp counter is invariant and ticks at the 2.30 GHz reference frequency, so the reported figures are *reference* cycles per byte and are unaffected by turbo (and convert directly to wall-clock time). On **arm64** the harness reads **CNTVCT_ELO** and scales it by the ratio of the peak core frequency to **CNTFRQ_ELO**; since Apple exposes no frequency API, the peak is auto-detected by timing a dependent integer-add chain and taking the top observed P-state, yielding *core* cycles per byte. Consequently the cycles-per-byte figures are comparable within an architecture and against the same-host AES-128-GCM baseline, but not directly across architectures; absolute throughput (Table 6) is the cross-platform metric.

The baseline is Go’s AES-128-GCM implementation (**crypto/aes** with **crypto/cipher**). It uses AES-NI and **PCLMULQDQ** on **amd64**, and the AES and **PMULL** instructions on **arm64**. Each timed AES-128-GCM call constructs its cipher, so its measurement includes the key schedule and the GHASH key-power precomputation; this mirrors the TW128 measurement, whose one-shot seal and open perform the root duplex initialization

Table 5: Median cycles per byte (100 samples). The **amd64** figures are RDTSC reference cycles at the 2.30 GHz invariant time-stamp counter; the **arm64** figures are core cycles at the auto-detected peak performance-core frequency. Cycle counts are not directly comparable across architectures (Section 7.2).

	Message length						
	1 B	64 B	8 KiB	32 KiB	64 KiB	1 MiB	16 MiB
<i>Intel Xeon 8581C (Emerald Rapids), AVX-512</i>							
TW128 encrypt	590	9.27	1.89	0.78	0.40	0.39	0.40
TW128 decrypt	615	9.72	1.90	0.75	0.40	0.39	0.39
AES-128-GCM seal	765	13.23	0.38	0.32	0.30	0.30	0.30
AES-128-GCM open	785	12.79	0.37	0.31	0.30	0.29	0.29
<i>Intel Xeon 8581C (Emerald Rapids), AVX2</i>							
TW128 encrypt	656	10.31	2.20	1.10	1.02	1.07	1.08
TW128 decrypt	683	10.70	2.21	1.11	1.08	1.07	1.08
<i>Apple M4 Pro, NEON + FEAT_SHA3</i>							
TW128 encrypt	620	9.77	1.95	0.91	0.88	0.88	0.89
TW128 decrypt	667	10.59	2.00	0.89	0.87	0.87	0.87
AES-128-GCM seal	1044	17.28	0.54	0.45	0.44	0.43	0.44
AES-128-GCM open	1055	17.31	0.52	0.43	0.42	0.41	0.42

(one permutation) inside the timed call. Both directions are measured for each scheme; the two directions track each other to within a few percent across the sweep, with a worst-case divergence of 8% (TW128 on **arm64** at 64 B).

7.3 Throughput

Table 5 reports cycles per byte across the message-length sweep and Table 6 the corresponding wall-clock throughput in the bulk regime. The dominant feature of the TW128 rows is the descent produced by the tree-parallel leaf core: the per-byte cost falls as a message lengthens and its chunks occupy more SIMD lanes. A message of at most one chunk has no complete leaf chunk for chunk 0 to share a kernel with, so it runs entirely on the single-duplex core, and the 8 KiB column sits at that core’s rate. Past one chunk, lane-0 fusion (Section 7.1) keeps every complete chunk on the batched kernels, and the descent tracks lane occupancy: at 32 KiB, chunk 0 and three complete leaves run four-wide in the **amd64** remainder kernels and as two full pairs on **arm64**, and 64 KiB is the first length whose eight complete chunks ($8\beta \approx 64$ KiB) fill the eight-way ($\times 8$) core. On the **amd64** AVX-512 path TW128 encryption falls $1.89 \rightarrow 0.78 \rightarrow 0.40$ across the 8, 32, and 64 KiB columns and is already at its bulk rate at 64 KiB, holding at 0.39–0.40 through 16 MiB. The AVX2 path follows the same shape at lower throughput ($2.20 \rightarrow 1.10$ over the same span, thereafter between 1.02 and 1.08); it runs the eight instances as two groups of four rather than in a single 8-wide register, leaving the AVX-512 core about $2.7\times$ faster in bulk. The **arm64** curve flattens earliest ($1.95 \rightarrow 0.91 \rightarrow 0.88 \rightarrow 0.88 \rightarrow 0.89$ cycles per byte from 8 KiB to 16 MiB): the eight-way core itself processes its batch as four two-instance pairs, so the fused 2-wide kernel runs at essentially the same per-chunk rate, and a 32 KiB message—chunk 0 and three leaves as two full pairs—is already within 3% of the 16 MiB floor. In wall-clock terms TW128 reaches 5.77 GB/s on the **amd64** AVX-512 host and 5.04 GB/s on the M4 Pro for 16 MiB messages.

Table 6: Measured wall-clock throughput (GB/s, GB = 10^9 bytes) from the same `cmd/cpb` run as Table 5, encryption direction; the decryption direction tracks it. The `amd64` AES-128-GCM row is the AVX-512 host and is unaffected by the TW128 core selection.

	8 KiB	32 KiB	64 KiB	1 MiB	16 MiB
<i>Intel Xeon 8581C (Emerald Rapids)</i>					
TW128 (AVX-512)	1.22	2.95	5.78	5.94	5.77
TW128 (AVX2)	1.05	2.09	2.25	2.15	2.14
AES-128-GCM	6.11	7.27	7.57	7.68	7.63
<i>Apple M4 Pro</i>					
TW128 (NEON+SHA3)	2.30	4.92	5.06	5.10	5.04
AES-128-GCM	8.25	9.89	10.18	10.33	10.09

Table 7: Cycles-per-byte measurement dispersion for TW128 across the seven-length sweep (seal and open, 100 samples per point), as the interquartile range relative to the median. The final column is the widest single-point min–max range; its size on the `amd64` guest reflects the turbo and governor states we cannot pin (Section 7.2) and is excluded from the IQR as a tail effect.

Path	Relative IQR		Max
	median	max	range
Xeon 8581C, AVX-512	1.2%	1.6%	16.5%
Xeon 8581C, AVX2	1.4%	2.1%	20.2%
Apple M4 Pro	0.6%	1.0%	4.7%

Dispersion. Each cell of Tables 5 and 6 is the median of 100 samples; Table 7 summarizes their spread as the interquartile range (IQR) relative to the median. The medians are stable: across the sweep the relative IQR of the TW128 cycles-per-byte measurements is at most 1.6% on the `amd64` AVX-512 path, 2.1% on AVX2, and 1.0% on the M4 Pro, with per-path medians of 1.2%, 1.4%, and 0.6%. The full min–max range is wider on the `amd64` guest—up to 16.5% (AVX-512) and 20.2% (AVX2) at isolated lengths—because its turbo and governor states are not under our control (Section 7.2); these are tail excursions that the IQR excludes, and the M4 Pro, whose frequency is steadier, keeps even its full range within 4.7%. We therefore report medians throughout and treat the IQR as the dispersion measure.

7.4 Latency

For short messages the cost is set by the serial root duplex and is independent of the parallel leaf machinery. A message of at most $\rho = 167$ bytes occupies a single rate block and never enters leaf mode, so it costs exactly two permutations—one to initialize the root duplex and one to close the message block and extract the tag. The 1 B and 64 B columns of Table 5 bear this out: per-message latency is nearly flat across them at roughly 590 reference cycles on `amd64` (590 at 1 B versus $9.27 \times 64 \approx 593$ at 64 B) and 620 core cycles on `arm64` (620 versus $9.77 \times 64 \approx 625$), i.e. about 295 and 310 cycles per KECCAK- p [1600, 12] respectively. Because such messages bypass the leaf core entirely, leaf-level parallelism imposes no latency penalty on them.

The latency/throughput trade-off appears instead in the intermediate regime, which lane-0 fusion has narrowed to roughly the one-to-eight-chunk band. The worst per-byte cost in the sweep is the single-chunk regime at 8 KiB (1.9–2.2 cycles per byte across the

three paths), where the message is the chunk-0 phase alone on the single-duplex core. Once the message has complete leaf chunks for chunk 0 to share a kernel with, the cost falls quickly: by 32 KiB TW128 is within a factor of two of its AVX-512 floor and within 3% of its `arm64` floor, and the floor itself is reached once the eight-way core fills at 64 KiB. The root duplex also bounds the tail of a long message: the final tag cannot be produced until chunk 0 has been processed and every leaf tag has been absorbed into the aggregation transcript. That absorption is cheap—each rate block of the root absorbs five 32-byte leaf tags, so a 16 MiB message (≈ 2000 leaves) adds on the order of 400 root permutations against roughly 10^5 permutations of leaf work, under 0.5% overhead—so the serial root does not materially cap bulk throughput.

7.5 Comparison with Other AEADs

We compare against the one baseline measured under the same harness on the same hosts, the Go standard library’s hardware-accelerated AES-128-GCM (Tables 5 and 6). The harness times TW128 and AES-128-GCM back-to-back at each length and reads both metrics from the same loop, so the cycles-per-byte and throughput comparisons hold under identical conditions and agree. In the bulk regime AES-128-GCM is faster on both architectures, as expected for a primitive with dedicated AES and carryless-multiply instructions: at 16 MiB it runs at 0.30 cycles per byte on `amd64` and 0.44 on `arm64`, against 0.40 and 0.89 for TW128—about $1.3\times$ and $2\times$ faster. The wall-clock throughputs show the same margins: TW128 reaches 5.77 and 5.04 GB/s against AES-128-GCM’s 7.63 and 10.09 GB/s. The wider `arm64` gap reflects the strength of Apple’s AES and PMULL units. TW128 attains these rates with a permutation that has no dedicated hardware support, relying only on the SHA-3 extension on `arm64` and on general SIMD on `amd64`.

The ordering reverses for very short messages, where AES-128-GCM’s key schedule and GHASH key-power precomputation dominate: at 1 B AES-128-GCM costs 765 and 1044 cycles on the two hosts against TW128’s 590 and 620. This comparison is setup-bound rather than representative of steady state, since each timed call rebuilds its cipher; it nonetheless reflects the per-message cost an application pays when it does not amortize key setup across messages. In the sub-bulk regime AES-128-GCM is several times faster at one chunk (0.38 versus 1.89 cycles per byte at 8 KiB on `amd64`), but the gap narrows quickly once the fused kernels engage: about $2.4\times$ at 32 KiB, and the bulk ratio from 64 KiB on. TW128’s intended value relative to AES-128-GCM lies in the tree structure and committing security properties analyzed in Sections 5 and 6 rather than in raw speed.

8 Discussion

8.1 Comparison Table

Table 8 situates TW128 against representative authenticated encryption schemes along the axes that distinguish it: the underlying primitive, whether the mode is parallelizable, the key, nonce, and tag sizes, associated data support, nonce-misuse resistance, full-input key commitment (CMT-4), and context discoverability security (CDY). The comparison is structural; TW128’s own concrete bounds and the regimes in which each term dominates are treated in Section 5.5, and its measured throughput against hardware-accelerated AES-128-GCM in Section 7.5. The CMT-4 column records full-input commitment: TW128 provides it within the mode, as does Ascon [KSW23], a committing transform supplies it for an arbitrary base scheme, and the remaining listed schemes do not provide it. The CDY column records context discoverability security in the [MLGR23] sense: TW128 provides it within the mode (Theorem 6), and the remaining entries are marked “?” because, to our knowledge, neither a context discovery attack nor a context discoverability security

Table 8: Structural comparison of TW128 with representative AEADs. Sizes are in bytes; AES-GCM-SIV also admits a 32-byte key. “Misuse-res.” is nonce-misuse resistance in the sense of [RS06]; the CMT-4 column records full-input commitment in the sense of [BH22], abstracting over its decryption (CMTD) and encryption (CMT) forms, of which TW128’s is the decryption form CMTDs-4; CDY is context discoverability in the sense of [MLGR23]. A dash means the entry is inherited from the base scheme, and “?” that no result is established in the literature. TW128’s concrete bounds appear in Section 5.5.

Scheme	Primitive	Par.	K/N/T	AD	Misuse	CMT-4	CDY
TW128	KECCAK- p [1600, 12]	yes	32/32/32	yes	no	yes	yes
Ascon-128a [DEMS21]	Ascon- p (320 b)	no	16/16/16	yes	no	yes	?
Xoodyak [DHP ⁺ 20]	Xoodoo (384 b)	no	16/16/16	yes	no	no	?
AES-GCM-SIV [GLL17]	AES	yes	16/12/16	yes	yes	no	?
nAE + CTX [CR22]	any nAE	—	—	yes	—	yes	?

proof has been published for them. Context discoverability is a recent notion [MLGR23], so its absence of entries reflects the state of the literature rather than a negative result; the schemes [MLGR23] break with context discovery attacks (GCM, CCM, EAX, SIV, OCB3) are not among the rows above. The table does not separately track compact (tag-only) commitment, since that property has not been systematically analyzed for the comparison schemes.

8.2 Comparison with Prior Sponge AEADs

TW128 belongs to the family of single-permutation duplex AEADs initiated by [BDPVA11]. The dedicated lightweight members of that family, Ascon [DEMS21] and Xoodyak [DHP⁺20], target 128-bit security on small permutations (320 and 384 bits) with a strictly serial duplex: each call depends on the previous one, so there is no mode-level parallelism to exploit. TW128 instead pays for a larger permutation, KECCAK- p [1600, 12], and a 256-bit capacity—the same round count as KangarooTwelve [BDP⁺16]—in exchange for a tree that processes leaf chunks in parallel and for the headroom that yields 128-bit security with a comfortable margin. The cost of this choice is visible in Section 7: on short, single-leaf inputs TW128 is slower per byte than a compact serial duplex would be, and its advantage appears only once a message spans enough leaves to fill the vector units.

TW128 reuses KangarooTwelve’s final-node/leaf split but not its Sakura coding, which renders that scheme’s tree hashing sound: the frame bits and coded child count of a Sakura coding make every node input unambiguously decodable, so that any collision in the hash output entails a collision in the inner function. TW128 adopts no such coding and does not need one. Its tree has a single fixed two-level shape—one root and a leaf count fixed by $\|C\|$ through $\text{leaves}(\cdot)$ —so it never distinguishes among a family of admissible topologies, the generality a Sakura coding exists to serve; the leaf count is moreover pinned by absorbing $\langle n \rangle_8$ before the root tag. The unambiguous parsing such a coding would supply is instead furnished by the per-call domain separation of Section 3.4: the prefixes p_{root} , p_{leaf} and the leaf index $\langle j \rangle_8$ fix the duplex instance, and the seven pairwise-distinct phase suffixes of Table 3 fix the phase and the final/non-final role of every call, so distinct transcripts flatten to distinct sponge inputs (Lemma 4 and Corollary 7); binding the index into each leaf is in fact a finer separation than ordering leaves within an aggregation alone. More to the point, the committing and discoverability guarantees of Section 6 do not rest on soundness of a tree coding at all: distinct (K, N, A) triples already diverge in the root initialization and associated data calls, before any leaf tag is aggregated (Proposition 1), so agreement among leaf tags cannot merge distinct contexts. The leaf encoding enters only the collision bound of Theorem 3, and there only through its explicit leaf tag collision term. A Sakura coding would re-establish parsing facts TW128 already possesses, and

would leave these bounds unchanged.

Keyak [BDP⁺15], whose Motorist mode is already parallelizable, is the closest prior duplex AEAD in spirit; Section 1.2 contrasts its wire-visible, encryptor-and-decryptor-shared degree of parallelism with TW128’s run-time leaf choice. Strobe [Ham17] is a serial sponge framework aimed at whole protocols rather than at a standalone AEAD, and like Ascon and Xoodoo it does not parallelize. Committing security is not unique to TW128 in this family: a dedicated analysis of the sponge-based NIST lightweight finalists proves Ascon committing while breaking others [KSW23], and some constructions, such as Strobe, have not been examined in that setting at all. What distinguishes TW128 is that it makes the root tag a secure compact commitment to the full input—collision-, second-preimage-, and preimage-resistant as a hash, which in particular yields the strongest full-input committing notion (CMT-4)—as a proven, designed property of the mode, and additionally resists [MLGR23] context discovery, established without leaving the single-permutation duplex setting; context discoverability security in particular is not known to hold for the other members of the family.

8.3 Comparison with Committing AEAD Constructions

The dominant route to committing AEAD is a generic transform over a non-committing base scheme. [BH22] give the UtC, RtC, and HtE transforms, which add key- or context-commitment to an arbitrary AEAD, and [CR22] give CTX, which makes any nonce-based AE scheme commit to its full context at the cost of a single hash over a short string and with no ciphertext expansion. These transforms are inexpensive and broadly applicable; CTX in particular is nearly free. TW128 differs not by undercutting their cost but by integration: it is committing as the mode itself, with the same permutation and the same tag that already provide confidentiality and integrity, so no separate commitment primitive, key-derivation step, or auxiliary hash is invoked, and the strongest s -way full-input notion follows from a single root tag collision bound (Section 6). More strongly, Theorem 3 shows that the root tag alone is a compact commitment string even when the ciphertext bodies used as openings differ, up to the additional leaf tag birthday term. The same tag, being preimage-resistant, additionally yields context discoverability security against that fixed tag, which the context-committing transforms are not known to provide. This places TW128 alongside the dedicated committing primitives—the compactly committing AEAD of [GLR17] and the encryption of [DGRW18]—which likewise build commitment in rather than bolt it on, but which were designed for message franking rather than as general nonce-based AEADs. The contrast with transformed conventional schemes is sharpened by the attacks of Section 1.2: base schemes such as GCM and OCB are not committing on their own [CR22], and a transform is the only way to make them so.

Closest to TW128 here is the recent work of [DHM⁺24], which builds nonce-based AE directly on the SHA-3 functions SHAKE and TurboSHAKE—the latter using the same round-reduced KECCAK- p permutation as TW128—and likewise achieves CMT-4 natively, from the collision resistance of the underlying function rather than from a transform. The two designs make different structural choices: [DHM⁺24] targets session-based communication, chaining a stream of messages through intermediate tags and generalizing the keyed duplex into a *duplex cipher*, whereas TW128 is a single two-level tree whose independent leaves expose run-time-selectable parallelism at constant memory (Sections 1.2 and 7.1). Because a session chain processes each message serially, TW128’s principal advantage is throughput on large messages (Section 7): its parallel leaves fill vector units a serial duplex leaves idle, at the cost of the session support [DHM⁺24] provides and TW128 does not. Both demonstrate that committing AE needs no add-on transform once it is built on a well-scrutinized KECCAK- p permutation.

8.4 Limitations and Open Problems

Nonce misuse. TW128’s AE and mu-ind bounds require nonce-respecting encryption queries, and the mode is not misuse-resistant: a repeated nonce reuses leaf keystream and breaks confidentiality, as for any online duplex stream cipher. Schemes such as AES-GCM-SIV [GLL17], in the deterministic/misuse-resistant lineage of [RS06], tolerate nonce repetition at the price of a second pass. A misuse-resistant variant of TW128 is left open.

Release of unverified plaintext. The analysis assumes that plaintext is released only after the tag verifies. TW128 is online and single-pass, so an implementation that streams decrypted leaf output before checking the root tag operates in the release-of-unverified-plaintext setting of [ABL⁺14], which our bounds do not cover; plaintext-awareness and INT-RUP security for TW128 are unstudied.

Bound tightness. Per-term tightness is itemized in Section 5.5; one term is conservative in a way that points to an open problem. The flat sponge differentiating loss $\text{Lift}_c(N_{\text{tot}}) = N_{\text{tot}}(N_{\text{tot}} + 1)/2^{c+1}$ bounds the indistinguishability loss of [BDPVA08], used here as a generic lift from the random oracle model. A direct analysis of the keyed duplex is tighter: the modern keyed-duplex bounds synthesized by [Men23], building on the full-state keyed-duplex line, bound a keyed duplex against an ideal permutation without passing through indistinguishability, and recasting TW128’s leaf and root duplexes in that framework is a promising route to a sharper capacity term. This route is not turnkey, however: the keyed duplex bundles squeezing and absorption into a single call, so its authenticated encryption mode `MonkeySpongeWrap` has the decryptor *overwrite* the rate with each ciphertext block rather than feed back the ciphertext it has just produced. On the final, padded block this overwrite no longer matches the encryptor’s XOR-absorbed padded plaintext outside the message bits, so the decryptor cannot reconstruct the final block of duplex output and recomputes the tag from a divergent state [Men23, Alg. 8]. TW128 sidesteps this by absorbing the ciphertext in both directions (Section 3), so recasting it on the keyed duplex would first require repairing the mode. We retain the indistinguishability lift for its modularity, as the resulting Lift_c already meets the 128-bit target (Section 5.5).

Beyond-birthday and leakage. All bounds are birthday-type on the 256-bit capacity, delivering the 128-bit target but no more; a beyond-birthday analysis is open. The model is also single-stage and leakage-free: side-channel resistance is addressed only at the implementation level, through the constant-time properties of Section 7.1, and is not part of the provable-security claims.

9 Conclusion

TW128 is a two-level tree of keyed duplexes over a single KECCAK- p [1600, 12] permutation that combines mode-level parallelism, constant memory, and commitment security in both its collision and preimage facets. Its concrete public permutation bounds—for multi-user indistinguishability, all-in-one AE as its single-user corollary, the secure-hash security of the root tag wrapper (collision, second-preimage, and preimage resistance) that makes the tag a compact commitment, the s -way full-input committing notion (CMTDs-4) it yields, and [MLGR23] context discoverability—meet a 128-bit target at the implemented parameters, and vectorized `amd64` and `arm64` implementations turn the tree’s leaf parallelism into multi-gigabyte per second throughput within a small factor of hardware-accelerated AES-128-GCM.

Several questions remain open (Section 8.4): a nonce-misuse-resistant variant, security under the release of unverified plaintext, and sharper bounds—whether by recasting the leaf and root duplexes on the modern keyed-duplex analysis or through a beyond-birthday argument. More broadly, TW128 adds to the growing evidence that a single, well-scrutinized KECCAK- p permutation suffices to build authenticated encryption that is at once fast, parallel, and committing.

Acknowledgments

A hearty thank-you to my wife, who remains reasonably mystified at my choice of hobbies.

References

- [ABL⁺14] Elena Andreeva, Andrey Bogdanov, Atul Luykx, Bart Mennink, Nicky Mouha, and Kan Yasuda. How to securely release unverified plaintext in authenticated encryption. In *Advances in Cryptology – ASIACRYPT 2014*, volume 8873 of *Lecture Notes in Computer Science*, pages 105–125. Springer, 2014.
- [ADG⁺22] Ange Albertini, Thai Duong, Shay Gueron, Stefan Kölbl, Atul Luykx, and Sophie Schmieg. How to abuse and fix authenticated encryption without key commitment. In *31st USENIX Security Symposium (USENIX Security 22)*, pages 3291–3308. USENIX Association, 2022.
- [BDP⁺15] Guido Bertoni, Joan Daemen, Michaël Peeters, Gilles Van Assche, and Ronny Van Keer. CAESAR submission: Keyak v2. CAESAR Competition, Round 3 Candidate, 2015.
- [BDP⁺16] Guido Bertoni, Joan Daemen, Michaël Peeters, Gilles Van Assche, Ronny Van Keer, and Benoît Viguier. KangarooTwelve: Fast hashing based on Keccak-p. *Cryptology ePrint Archive*, Paper 2016/770, 2016.
- [BDPVA08] Guido Bertoni, Joan Daemen, Michaël Peeters, and Gilles Van Assche. On the indistinguishability of the sponge construction. In *Advances in Cryptology – EUROCRYPT 2008*, volume 4965 of *Lecture Notes in Computer Science*, pages 181–197. Springer, 2008.
- [BDPVA11] Guido Bertoni, Joan Daemen, Michaël Peeters, and Gilles Van Assche. Duplexing the sponge: Single-pass authenticated encryption and other applications. In *Selected Areas in Cryptography – SAC 2011*, volume 7118 of *Lecture Notes in Computer Science*, pages 320–337. Springer, 2011.
- [BH22] Mihir Bellare and Viet Tung Hoang. Efficient schemes for committing authenticated encryption. In *Advances in Cryptology – EUROCRYPT 2022*, volume 13276 of *Lecture Notes in Computer Science*, pages 845–875. Springer, 2022.
- [BT16] Mihir Bellare and Björn Tackmann. The multi-user security of authenticated encryption: AES-GCM in TLS 1.3. In *Advances in Cryptology – CRYPTO 2016*, volume 9814 of *Lecture Notes in Computer Science*, pages 247–276. Springer, 2016.
- [CR22] John Chan and Phillip Rogaway. On committing authenticated encryption. In *Computer Security – ESORICS 2022*, *Lecture Notes in Computer Science*. Springer, 2022.

- [DEMS21] Christoph Dobraunig, Maria Eichlseder, Florian Mendel, and Martin Schl  ffer. Ascon v1.2: Lightweight authenticated encryption and hashing. *Journal of Cryptology*, 34(3):33, 2021.
- [DGRW18] Yevgeniy Dodis, Paul Grubbs, Thomas Ristenpart, and Joanne Woodage. Fast message franking: From invisible salamanders to encryption. In *Advances in Cryptology – CRYPTO 2018*, volume 10991 of *Lecture Notes in Computer Science*, pages 155–186. Springer, 2018.
- [DHM⁺24] Joan Daemen, Seth Hoffert, Silvia Mella, Gilles Van Assche, and Ronny Van Keer. Shaking up authenticated encryption. Cryptology ePrint Archive, Paper 2024/1618, 2024.
- [DHP⁺20] Joan Daemen, Seth Hoffert, Micha  l Peeters, Gilles Van Assche, and Ronny Van Keer. Xoodooak, a lightweight cryptographic scheme. *IACR Transactions on Symmetric Cryptology*, 2020(S1):60–87, 2020.
- [GLL17] Shay Gueron, Adam Langley, and Yehuda Lindell. AES-GCM-SIV: Specification and analysis. Cryptology ePrint Archive, Paper 2017/168, 2017.
- [GLR17] Paul Grubbs, Jiahui Lu, and Thomas Ristenpart. Message franking via committing authenticated encryption. In *Advances in Cryptology – CRYPTO 2017*, volume 10403 of *Lecture Notes in Computer Science*, pages 66–97. Springer, 2017.
- [Ham17] Mike Hamburg. The STROBE protocol framework. Cryptology ePrint Archive, Paper 2017/003, 2017.
- [IR23] Takanori Isobe and Mostafizar Rahman. Key committing security analysis of AEGIS. Cryptology ePrint Archive, Paper 2023/1495, 2023.
- [KSW23] Juliane Kr  mer, Patrick Struck, and Maximiliane Weish  upl. Committing AE from sponges: Security analysis of the NIST LWC finalists. Cryptology ePrint Archive, Paper 2023/1525, 2023.
- [LGR21] Julia Len, Paul Grubbs, and Thomas Ristenpart. Partitioning oracle attacks. In *30th USENIX Security Symposium (USENIX Security 21)*, pages 195–212. USENIX Association, 2021.
- [Men23] Bart Mennink. Understanding the duplex and its security. *IACR Transactions on Symmetric Cryptology*, 2023(2):1–46, 2023.
- [MLGR23] Sanketh Menda, Julia Len, Paul Grubbs, and Thomas Ristenpart. Context discovery and commitment attacks: How to break CCM, EAX, SIV, and more. In *Advances in Cryptology – EUROCRYPT 2023*, volume 14007 of *Lecture Notes in Computer Science*, pages 379–407. Springer, 2023.
- [Nat15] National Institute of Standards and Technology. SHA-3 standard: Permutation-based hash and extendable-output functions. Federal Information Processing Standards Publication FIPS PUB 202, U.S. Department of Commerce, August 2015.
- [RS04] Phillip Rogaway and Thomas Shrimpton. Cryptographic hash-function basics: Definitions, implications, and separations for preimage resistance, second-preimage resistance, and collision resistance. In *Fast Software Encryption – FSE 2004*, volume 3017 of *Lecture Notes in Computer Science*, pages 371–388. Springer, 2004.

- [RS06] Phillip Rogaway and Thomas Shrimpton. A provable-security treatment of the key-wrap problem. In *Advances in Cryptology – EUROCRYPT 2006*, volume 4004 of *Lecture Notes in Computer Science*, pages 373–390. Springer, 2006.
- [WP14] Hongjun Wu and Bart Preneel. AEGIS: A fast authenticated encryption algorithm. In *Selected Areas in Cryptography – SAC 2013*, volume 8282 of *Lecture Notes in Computer Science*, pages 185–201. Springer, 2014.

A Supporting Lemmas

The proof core separates structural decoding from probabilistic accounting. The well-formed transcript language of Definition 1 fixes which flattened queries are valid TW128 construction prefixes. Lemmas 3, 4 and 6 then show how such a query is flattened, parsed as a typed transcript, assigned an output point, and decoded to a candidate key, when any. Lemma 9 turns those decoded keys into an adaptive key-test bound. The remaining proof interfaces expose the events their consumers must control: encryption-side freshness, verification hidden-key exposure and root tag guessing, leaf tag collisions, and the committing candidate-sequence collision event. Thus the mu-ind and committing proofs only have to derive the relevant local event premise before applying the corresponding accounting argument.

A.1 Duplex Flattening

Lemma 3 (Duplex Flattening). *Every TW128 root or leaf duplex output is equal to a [BDPVA11] sponge output on the corresponding flattened duplex-call transcript, and the ciphertext-absorbing message phase admits a flattened duplex transcript that is well defined for both encryption and decryption.*

Proof. Let D be a duplex object using permutation π , rate r , and padding pad10^*1 . For a sequence of duplex calls $Z_i = D.\text{duplexing}(\sigma_i, \ell_i)$ with each σ_i short enough to fit in one duplex block after padding, [BDPVA11, Lemma 3] gives

$$Z_i = \text{Sponge}[\pi, \text{pad10}^*1, r](\sigma_0 \parallel \text{pad}_0 \parallel \sigma_1 \parallel \text{pad}_1 \parallel \cdots \parallel \sigma_i, \ell_i),$$

where $\text{pad}_j = \text{pad10}^*1[r](|\sigma_j|)$, and Lemma 4 says that, for fixed pad10^*1 and r , the map from the duplex-call sequence to the flattened sponge input is injective.

In TW128, each duplex call absorbs a byte-aligned block σ followed by a 4-bit phase suffix $\sigma_{\text{ph}} \in \{\sigma_{\text{init}}, \sigma_{\text{ad}}^-, \dots, \sigma_{\text{agg}}^+\}$ (Section 3.4), so the duplex-call input is $\sigma \parallel \sigma_{\text{ph}}$ followed by pad10^*1 . With $r = 1344$ and $\rho_{\text{max}}(\text{pad10}^*1, r) = r - 2 = 1342$ bits, the maximum byte-aligned input is $\rho = 167$ bytes (Section 3.2), giving a worst-case duplex-call input of $8 \cdot 167 + 4 = 1340 \leq 1342$ bits. The two initialization calls are shorter still: $|p_{\text{root}} \parallel K \parallel N| + 4 = 8 \cdot (6 + 32 + 32) + 4 = 564$ bits and $|p_{\text{leaf}} \parallel K \parallel N \parallel \langle j \rangle_8| + 4 = 8 \cdot (6 + 32 + 32 + 8) + 4 = 628$ bits. [BDPVA11, Lemmas 3 and 4] therefore apply independently to every TW128 duplex object.

The root duplex is initialized by absorbing $p_{\text{root}} \parallel K \parallel N$ under σ_{init} (Algorithm 2), then absorbs the associated data blocks under $\sigma_{\text{ad}}^- / \sigma_{\text{ad}}^+$, the root ciphertext blocks in the message phase under $\sigma_{\text{msg}}^- / \sigma_{\text{msg}}^+$, and the leaf tag aggregation blocks under $\sigma_{\text{agg}}^- / \sigma_{\text{agg}}^+$. By the duplex-to-sponge equality, every root duplex output is the sponge function evaluated on the root transcript prefix accumulated up to the relevant output point. The same holds for each leaf duplex, initialized by $p_{\text{leaf}} \parallel K \parallel N \parallel \langle j \rangle_8$ and absorbing only its own ciphertext blocks in the message phase. In decryption and verification, Algorithm 3 absorbs the submitted ciphertext blocks before the root tag comparison, so the same flattened message phase transcript is defined whether the comparison accepts or rejects. $\square$

A.2 Well-formed Transcript Prefixes

For a duplex-call prefix $((\sigma_0, \sigma_{\text{ph},0}), \dots, (\sigma_i, \sigma_{\text{ph},i}))$, write

$$\text{Flat}((\sigma_0, \sigma_{\text{ph},0}), \dots, (\sigma_i, \sigma_{\text{ph},i})) = \left\|_{j=0}^{i-1} (\sigma_j \parallel \sigma_{\text{ph},j} \parallel \text{pad}_j) \parallel \sigma_i \parallel \sigma_{\text{ph},i}.$$

The current call's padding pad_i is not part of Flat ; it is applied internally by the sponge function. Equivalently, after the i -th call the duplex state is the sponge state after absorbing $\text{Flat}(\cdot) \parallel \text{pad}_i$.

Definition 1 (Well-Formed TW128 Transcript Prefixes). A typed TW128 transcript prefix is *well formed* if it is a prefix, ending at a construction transcript lookup point, of one of the following two call languages.

- A root transcript starts with $(p_{\text{root}} \parallel K \parallel N, \sigma_{\text{init}})$, with fixed-width fields $K \in \{0,1\}^k$ and $N \in \{0,1\}^\nu$. It then absorbs the associated data byte string (only when the associated data is nonempty; the phase is otherwise elided, Section 3.5), the root ciphertext byte string, and the aggregation byte string (only when $n \geq 1$; when $n = 0$ the aggregation phase is elided, Section 3.5, and the root tag is emitted by the final σ_{msg}^+ message call), in that order. Each present phase byte string is partitioned by the block rule of Algorithm 1: the empty string gives one empty final block, and a nonempty string gives zero or more ρ -byte non-final blocks followed by one final block. The phases use respectively $\sigma_{\text{ad}}^-/\sigma_{\text{ad}}^+$, $\sigma_{\text{msg}}^-/\sigma_{\text{msg}}^+$, and $\sigma_{\text{agg}}^-/\sigma_{\text{agg}}^+$.
- A leaf- j transcript starts with $(p_{\text{leaf}} \parallel K \parallel N \parallel \langle j \rangle_8, \sigma_{\text{init}})$, where $1 \leq j < 2^{64}$ and the fields have fixed width. It then absorbs the leaf ciphertext byte string under $\sigma_{\text{msg}}^-/\sigma_{\text{msg}}^+$, using the same block rule.

The message byte string in this language is always the ciphertext byte string absorbed by the construction: encryption obtains it by XORing the plaintext with the current keystream, while decryption and verification use the submitted ciphertext. Thus a verification transcript is well formed for every in-domain ciphertext before the tag comparison is made; acceptance is not part of the transcript language.

The root aggregation byte string has the form $T_1 \parallel \dots \parallel T_n \parallel \langle n \rangle_8$, where each T_j has fixed byte length $\tau_{\text{leaf}}/8$. The trailing fixed-width field therefore gives a unique parse of the leaf tag sequence. The output-point class of a well-formed counted prefix is the class in Table 9 determined by its duplex instance and final-call suffix.

Name the tag-emitting transcript families:

$\mathcal{R}_{\text{tag}}$ = root transcript prefix whose sponge output gives the final root tag,

$\mathcal{L}_{\text{tag}}(j)$ = leaf- j transcript prefix whose sponge output gives the leaf tag.

The final root tag is emitted by the σ_{agg}^+ aggregation call (the $\mathcal{R}_{\text{tag}}$ family) when $n \geq 1$, and by the closing σ_{msg}^+ message call (an $\mathcal{R}_{\text{msg}}^+$ prefix, in the notation of Corollary 7) when $n = 0$ and the aggregation phase is elided. Since $n = \text{leaves}(C)$ is fixed by $\|C\|$, every computation on a given ciphertext emits its tag at the one of these two output points selected by $\|C\|$. Intermediate construction transcript lookup points (init, non-final AD, non-final message, and non-final aggregation) are also typed prefixes covered by Lemma 6; the exact keyed direct-query decoder names them in Table 9.

A.3 Transcript Injectivity

Lemma 4 (Transcript Injectivity). *Well-formed TW128 flattened duplex-call transcript prefixes are injectively encoded and separated by duplex instance (root vs leaf), leaf index*

j, phase suffix, and duplex-call order. For final-root prefixes, equality of flattened inputs recovers the same root initialization, AD blocks, root ciphertext blocks, aggregation input, and leaf tag sequence.

Proof. By Lemma 3, every well-formed output point can be written as a [BDPVA11] flattened sponge input for the corresponding duplex-call prefix, namely Flat as defined above. The deterministic current-call pad10^*1 is applied internally by the sponge function and is not part of the input to which [BDPVA11, Lemma 4] is applied. That lemma recovers the duplex-call sequence from the flattened sponge input; in TW128, each call input is $\sigma \parallel \sigma_{\text{ph}}$ with byte-aligned σ and fixed 4-bit suffix σ_{ph} , so recovering the call input recovers the pair $(\sigma, \sigma_{\text{ph}})$.

Definition 1 fixes the root and leaf initialization fields, the admissible phase order, and the final-call suffixes. The distinct prefixes p_{root} and p_{leaf} separate root from leaf transcripts; the fixed-width leaf field $\langle j \rangle_8$ separates leaves; and the seven suffix values of Table 3 separate phases and non-final/final calls. Thus equality of flattened inputs forces equality of the recovered instance type, key, nonce, leaf index when present, phase suffixes, and duplex-call order.

Phase-recovery sub-claim. For each present phase $P \in \{\text{ad}, \text{msg}, \text{agg}\}$ (the ad phase is present exactly when the associated data is nonempty; see Section 3.5), TW128 absorbs the phase’s input byte string s_P by partitioning into ρ -byte blocks (with the empty case yielding a single empty block, per Section 3.5) and feeding each block to the duplex with suffix σ_P^- (non-final) or σ_P^+ (final). Lemma 3 and [BDPVA11, Lemma 4] recover the corresponding $(\sigma, \sigma_{\text{ph}})$ sequence from the flattened input. Because the partition map is deterministic and uniquely invertible on its image (the empty case yielding the single empty block carrying σ_P^+), concatenating the recovered σ blocks reconstructs s_P exactly. The map from s_P to the phase’s $(\sigma, \sigma_{\text{ph}})$ sequence is therefore injective.

Applied phase by phase to a well-formed final-root transcript:

- AD. The associated data phase contributes a σ_{ad}^+ -tagged block exactly when the associated data is nonempty; an empty associated data elides the phase and contributes no block. When the phase is present the sub-claim recovers the AD byte string. Thus distinct associated data strings $A_i \neq A_j$ are separated either by the presence versus absence of the AD phase (when exactly one is empty) or, when both are nonempty, by their distinct AD σ -block sequences.
- MSG. Distinct absorbed root or leaf ciphertext byte strings produce distinct MSG σ -block sequences; the message language of Definition 1 is ciphertext-absorbing for encryption and decryption alike.
- AGG. The aggregation phase is present exactly when $n \geq 1$ (Section 3.5). When it is present, the sub-claim applied to the aggregation byte string $T_1 \parallel \dots \parallel T_n \parallel \langle n \rangle_8$ recovers that byte string. The final eight bytes give $\langle n \rangle_8$, and each leaf tag has fixed byte length $\tau_{\text{leaf}}/8$, so the byte string has a unique parse as a leaf tag sequence. Two root aggregation encodings are therefore equal exactly when their leaf tag sequences are equal. When $n = 0$ the aggregation phase is absent and the final-root transcript ends at the recovered σ_{msg}^+ root message call; the MSG bullet recovers the root ciphertext byte string, and the absence of any $\sigma_{\text{agg}}^-/\sigma_{\text{agg}}^+$ call distinguishes such a prefix from every $n \geq 1$ final-root prefix.

Therefore the recovered final-root transcript determines the root initialization, AD blocks, root ciphertext blocks, and, when $n \geq 1$, the aggregation byte string and leaf tag sequence. $\square$

Lemma 5 (Leaf-Transcript Recovery). *Let two construction-generated final-root prefixes have equal flattened inputs, and suppose that no two distinct construction-generated final leaf transcripts produce the same leaf tag. Then equality of flattened inputs recovers the same leaf transcript sequence, and therefore the same full absorbed ciphertext.*

Proof. By Transcript Injectivity (Lemma 4), the two equal flattened final-root inputs recover the same root initialization, AD blocks, root ciphertext blocks, and, when $n \geq 1$, the same aggregation byte string and leaf tag sequence; the trailing $\langle n \rangle_8$ field fixes the common leaf count n . When $n = 0$ no leaf is present and the recovered root ciphertext blocks already constitute the full absorbed ciphertext. When $n \geq 1$, the aggregation byte string parses uniquely as the leaf tag sequence $T_1 \parallel \dots \parallel T_n$; under the hypothesis that distinct construction-generated final leaf transcripts receive distinct leaf tags, equality of the two leaf tag sequences forces, position by position, equality of the corresponding final leaf transcripts, hence of their absorbed leaf ciphertext bodies. Together with the common root ciphertext blocks this recovers the same full absorbed ciphertext. $\square$

A.4 Bridge and Typed Decoding

Lemma 6 (Bridge and Typed Decoding). *The random oracle games used in Sections 5 and 6 and Corollary 6 are ordinary random oracle games on a single random oracle $\text{RO}_{\text{flat}} : \{0, 1\}^* \rightarrow \{0, 1\}^\infty$ over the unpadded [BDPVA08] H -input domain (Section 4.2), with TW128 root and leaf duplex outputs answered by the typed restriction*

$$\text{RF}(X) = \text{RO}_{\text{flat}}(\text{flat}(X)), \quad \text{flat}(X) = I[r, r_{\max}](\text{raw_flat}(X)),$$

for $r = r_{\max} = 1344$. Here $\text{raw_flat}(X)$ is the native flattened input $\text{Flat}(\cdot)$ of Section A.2 applied to the duplex-call sequence of X , and $I[r, r_{\max}]$ is the [BDPVA11] Theorem 3 preprocessing map. The induced map $X \mapsto \text{flat}(X)$ is injective on well-formed typed TW128 duplex transcript prefixes. The exact decoder KeyedTWOut of Definition 2 then maps any direct-query input to either $\perp$ or a unique candidate key and counted output-point class.

Proof. The bridge $I[r, r_{\max}]$ is described by the proof of [BDPVA11, Theorem 3] in three steps: (1) pad the input with $\text{pad10}^*1[r]$; (2) for each r -bit block, append $r_{\max} - r$ zero bits; (3) strip the trailing pad10^* from the result. The pad10^*1 padding in step (1) appends a 1, then $q = (-|\text{raw_flat}(X)| - 2) \bmod r$ zero bits, then a closing 1; the pad10^* -unpadding in step (3) strips the $r_{\max} - r$ trailing zero bits introduced by step (2) together with the closing 1 of step (1). For TW128 with $r = r_{\max}$, step (2) appends an empty string and step (3) strips zero trailing zeros and the closing 1, so the net effect of steps (1) and (3) is to map a native flattened input $U = \text{raw_flat}(X)$ to $U \parallel 1 \parallel 0^q$. Any string in this image has a unique inverse: q is the number of trailing zero bits, the preceding bit is the inserted 1, and deleting this suffix recovers U .

Injectivity of flat on well-formed typed TW128 prefixes then follows from Lemma 4: distinct well-formed typed prefixes have distinct $\text{raw_flat}(\cdot)$ images, and the bridge is injective on that image. Direct adversary queries on inputs outside the bridged TW128 well-formed transcript language still count in $q_{\text{RO}}(\mathcal{B})$ (Section 4.3) but cannot trigger bad_{key} .

Counted transcript lookup points are the well-formed prefixes of Definition 1 whose classes are listed in Table 9: root or leaf initialization, root AD blocks, root and leaf message blocks, root aggregation blocks, and the final root tag point. An input accepted by KeyedTWOut therefore fixes a unique candidate key—by the injectivity just shown and the fixed-width key field of Lemma 4—and a unique output-point class, by Output-Point Separation (Corollary 7); inputs outside this well-formed transcript language decode to none. $\square$

The bridge injectivity established here underlies Opening Triple Separation (Proposition 1, Section 6): distinct (K, N, A) triples yield distinct bridged final root transcripts.

Corollary 7 (Output-Point Separation). *Every well-formed TW128 typed transcript prefix at a counted construction transcript lookup point belongs to exactly one of the following output-point classes, identified by the recovered duplex instance and the 4-bit suffix of its final call:*

- $\mathcal{R}_{\text{init}}(\text{root duplex}, \text{suffix } \sigma_{\text{init}})$,
- $\mathcal{L}_{\text{init}}(j)$ (leaf- j duplex, suffix σ_{init}),
- $\mathcal{R}_{\text{ad}}^- / \mathcal{R}_{\text{ad}}^+$ (root duplex, suffix σ_{ad}^- or σ_{ad}^+),
- $\mathcal{R}_{\text{msg}}^- / \mathcal{R}_{\text{msg}}^+$ (root duplex, suffix σ_{msg}^- or σ_{msg}^+),
- $\mathcal{L}_{\text{msg}}^-(j)$ (leaf- j duplex, suffix σ_{msg}^-),
- $\mathcal{L}_{\text{tag}}(j)$ (leaf- j duplex, suffix σ_{msg}^+),
- $\mathcal{R}_{\text{agg}}(\text{suffix } \sigma_{\text{agg}}^-)$,
- $\mathcal{R}_{\text{tag}}(\text{root duplex}, \text{suffix } \sigma_{\text{agg}}^+)$.

For any two such prefixes whose bridged flattened inputs are equal, the separation of Lemma 4 recovers the same duplex instance, leaf index j when present, (K, N) pair, output-point class, and duplex-call sequence. Equivalently, any two distinct counted transcript prefixes already differ in one of these recovered fields, and hence in their flattened inputs.

Proof. Lemma 6 recovers equality of native flattened inputs from equality of bridged inputs, after which Lemma 4 recovers the duplex-call sequence and its per-call $(\sigma, \sigma_{\text{ph}})$ pairs. The first call's σ is $p_{\text{root}} \parallel K \parallel N$ for root instances and $p_{\text{leaf}} \parallel K \parallel N \parallel \langle j \rangle_8$ for leaf instances, with fixed-width fields, so it fixes the instance type (by p_{root} vs p_{leaf}), the (K, N) pair, and the leaf index j ; within each instance type the phase suffixes are pairwise distinct, so the recovered instance type together with the final-call suffix fixes the output-point class, and each prefix therefore lies in exactly one class. Conversely, two distinct well-formed prefixes in the same class on the same instance differ in some recovered call's σ or in the sequence length, hence in their flattened inputs. $\square$

Definition 2 (Exact Keyed Decoder). For a direct query to RO_{flat} with input Y and finite length ℓ , define $\text{KeyedTWOut}(Y)$ as a partial map from $\{0, 1\}^*$ to $\{0, 1\}^k \times \mathcal{C}_{\text{out}}$. The map ignores ℓ . It returns $(K, \mathcal{C})$ iff there exists a typed TW128 duplex-call prefix X such that $Y = \text{flat}(X)$, the raw flattened prefix of X parses as a well-formed TW128 transcript prefix in the class $\mathcal{C}$ listed in Table 9, and the first call of X contains the candidate key K in the fixed-width root field $p_{\text{root}} \parallel K \parallel N$ or leaf field $p_{\text{leaf}} \parallel K \parallel N \parallel \langle j \rangle_8$. If no such prefix exists, $\text{KeyedTWOut}(Y) = \perp$. In particular, inputs with the right final suffix but a malformed phase order, noncanonical block partition, malformed aggregation string, or out-of-range leaf index are outside the well-formed language and decode to $\perp$. A non- $\perp$ decoding fixes not only the candidate key K but the candidate context components present in the decoded prefix: the nonce N occupies the remaining fixed-width init-field slot, and the associated data A is read off the recovered associated data blocks. We use the key projection $(K, \mathcal{C})$ throughout; when a game hides a single context component $Z \in \{K, N, A\}$, the same decoding returns that component's candidate value, which we denote $\text{KeyedTWOut}_Z(Y)$.

The table is a transcript-lookup table, not an observation-length table. A class is counted even when the construction's requested output length at that call is zero, including root initialization, non-final AD and aggregation calls, a final AD call whose following root

Table 9: Counted keyed transcript classes for KeyedTWOOut.

Class $\mathcal{C}$	Instance	Final call / role
$\mathcal{R}_{\text{init}}$	root	σ_{init} , root initialization
$\mathcal{R}_{\text{ad}}^-$	root	σ_{ad}^- , non-final AD block
$\mathcal{R}_{\text{ad}}^+$	root	σ_{ad}^+ , final AD block
$\mathcal{R}_{\text{msg}}^-$	root	σ_{msg}^- , non-final root message block
$\mathcal{R}_{\text{msg}}^+$	root	σ_{msg}^+ , final root message block
$\mathcal{R}_{\text{agg}}^-$	root	σ_{agg}^- , non-final aggregation block
$\mathcal{R}_{\text{tag}}^+$	root	σ_{agg}^+ , final aggregation / root tag
$\mathcal{L}_{\text{init}}(j)$	leaf j	σ_{init} , leaf initialization
$\mathcal{L}_{\text{msg}}^-(j)$	leaf j	σ_{msg}^- , non-final leaf message block
$\mathcal{L}_{\text{tag}}^+(j)$	leaf j	σ_{msg}^+ , final leaf message block / leaf tag

chunk requests no keystream bytes, and the final root message call. A direct query triggers the same decoder regardless of ℓ : short finite-output queries and zero-length queries $(Y, 0)$ still count as direct queries to the address Y . Leaves have no AD or aggregation classes; a leaf transcript enters the table only through leaf initialization, non-final leaf message blocks, or the final leaf tag point. When the associated data is empty the root transcript likewise has no $\mathcal{R}_{\text{ad}}^-$ or $\mathcal{R}_{\text{ad}}^+$ class (Section 3.5); the $\mathcal{R}_{\text{init}}$ class then supplies the first root keystream block, exactly as $\mathcal{L}_{\text{init}}(j)$ supplies the first leaf keystream block, and the decoder still ignores ℓ . Symmetrically, when $n = 0$ the root transcript has no $\mathcal{R}_{\text{agg}}$ or $\mathcal{R}_{\text{tag}}$ class (Section 3.5); the final root message call $\mathcal{R}_{\text{msg}}^+$ then supplies the root tag, exactly as $\mathcal{L}_{\text{tag}}(j)$ supplies a leaf tag, and the decoder again ignores ℓ . An $n \geq 1$ computation also reaches an $\mathcal{R}_{\text{msg}}^+$ point when closing chunk 0, but there the construction requests zero output and continues into the aggregation phase, so the $\mathcal{R}_{\text{msg}}^+$ value is exposed as a tag only on the $n = 0$ path; since $n = \text{leaves}(C)$ is fixed by $\|C\|$, the two uses never coincide on a single ciphertext.

The hidden-key encryption, verification, and mu-ind accounting all reindex transcript addresses by the key field. We record that reindexing once.

Lemma 7 (Key-Indexed Transcript Renaming). *For a key $K \in \{0, 1\}^k$, let $\mathcal{Y}_K$ be the set of bridged inputs $\text{flat}(X)$ of well-formed TW128 transcript prefixes X whose first-call key field is K —the counted hidden-key transcript addresses of Table 9 for key K . Then:*

1. Disjointness. *For distinct keys $K_0 \neq K_1$, the sets $\mathcal{Y}_{K_0}$ and $\mathcal{Y}_{K_1}$ are disjoint.*
2. Renaming. *The map $\Phi_{K_0 \rightarrow K_1}$ that replaces the first-call key field K_0 by K_1 and leaves every other recovered field unchanged—nonce, leaf index, phase sequence, AD blocks, ciphertext blocks, aggregated leaf tag bytes, and root/leaf instance—induces through flat a bijection $\text{flat} \circ \Phi_{K_0 \rightarrow K_1} \circ \text{flat}^{-1} : \mathcal{Y}_{K_0} \rightarrow \mathcal{Y}_{K_1}$ preserving output-point classes and equality relations among addresses.*
3. Direct-query avoidance. *A direct query input $Y \in \mathcal{Y}_K$ has $\text{KeyedTWOOut}(Y) = (K, \mathcal{C})$ for a counted class $\mathcal{C}$; hence if no prior direct query decodes to K then $\mathcal{Y}_K$ is disjoint from all prior direct-query inputs.*

Proof. By bridge injectivity (Lemma 6) and Transcript Injectivity (Lemma 4), a bridged input $\text{flat}(X)$ determines the native flattened prefix and hence the first duplex call $p_{\text{root}} \| K \| N$ or $p_{\text{leaf}} \| K \| N \| \langle j \rangle_s$, in particular its fixed-width key field K . (1) If $\text{flat}(X) = \text{flat}(X')$ with key fields K_0 and K_1 then $K_0 = K_1$, so distinct keys give disjoint address sets. (2) $\Phi_{K_0 \rightarrow K_1}$ rewrites only the key field, leaving the recovered duplex-call sequence otherwise intact, so it is a bijection on well-formed prefixes preserving the output-point class (Corollary 7) and the equality pattern among prefixes; flat transports it to the stated bijection on

addresses. (3) By Lemma 6 a $Y = \text{flat}(X) \in \mathcal{Y}_K$ decodes under KeyedTWOOut to $(K, \mathcal{C})$; the contrapositive gives the avoidance claim. $\square$

A.5 Adaptive Candidate Bounds

Lemma 8 (Adaptive Candidate Collision and Preimage). *Consider any interaction with a lazy random oracle $\text{RO}_{\text{flat}} : \{0, 1\}^* \rightarrow \{0, 1\}^\infty$ that builds a sequence of distinct candidate inputs $Y_1, \dots, Y_m$. The process may make arbitrary non-candidate oracle queries and may choose each candidate adaptively from the full prior transcript. Whenever a new candidate Y_j is appended, two premises hold: predictability—both the decision to append and the string Y_j are determined by the transcript preceding the append step—and unrevealed value—no oracle lookup before the append step has requested a positive number of bits of $\text{RO}_{\text{flat}}(Y_j)$; zero-length lookups at Y_j are permitted. Fix a projection length τ , and suppose the sequence has at most L candidates ($m \leq L$).*

1. Collision. For every $s \geq 2$, the probability $\Pr[\text{coll}_{s,\tau}]$ that some s distinct candidates have equal τ -bit projections,

$$[\text{RO}_{\text{flat}}(Y_{i_1})]_\tau = \dots = [\text{RO}_{\text{flat}}(Y_{i_s})]_\tau \quad \text{for some } 1 \leq i_1 < \dots < i_s \leq m,$$

$$\text{satisfies } \Pr[\text{coll}_{s,\tau}] \leq \binom{L}{s} \cdot 2^{-\tau(s-1)}.$$

2. Preimage. For every target $t \in \{0, 1\}^\tau$ chosen independently of RO_{flat} , the probability $\Pr[\text{pre}_{\tau,t}]$ that some candidate's τ -bit projection equals t , i.e. $[\text{RO}_{\text{flat}}(Y_j)]_\tau = t$ for some $1 \leq j \leq m$, satisfies $\Pr[\text{pre}_{\tau,t}] \leq L \cdot 2^{-\tau}$.

The preimage facet has a second-preimage corollary: for a designated candidate position j_0 , the probability $\Pr[\text{sec}_{\tau,j_0}]$ that some candidate $j \neq j_0$ has $[\text{RO}_{\text{flat}}(Y_j)]_\tau = [\text{RO}_{\text{flat}}(Y_{j_0})]_\tau$, i.e. collides with the designated candidate on its τ -bit projection, satisfies $\Pr[\text{sec}_{\tau,j_0}] \leq (L - 1) \cdot 2^{-\tau}$.

Proof. Both parts expose the interaction in chronological order and sample each candidate's RO_{flat} value lazily at its append step, which the premises license: when Y_j is appended its string is fixed by the prior transcript (predictability) and no bit of $\text{RO}_{\text{flat}}(Y_j)$ has yet been revealed (unrevealed value), so $\text{RO}_{\text{flat}}(Y_j)$ may be drawn then, uniform and independent of everything so far.

Collision. For a fixed position set $I = \{i_1 < \dots < i_s\} \subseteq \{1, \dots, L\}$, let E_I be the event that these candidate positions exist and their τ -bit projections are all equal. Condition on any positive-probability prefix immediately before position i_r is appended, for $r \geq 2$, and on the values sampled at the earlier positions in I . By the above, $\text{RO}_{\text{flat}}(Y_{i_r})$ is then uniform and independent of the conditioning, so its τ -bit projection equals the already fixed projection of Y_{i_1} with probability $2^{-\tau}$. Multiplying over $r = 2, \dots, s$ gives $\Pr[E_I] \leq 2^{-\tau(s-1)}$; if some position in I is never appended then E_I is false and the bound still holds. A union bound over the $\binom{L}{s}$ position sets gives the claim.

Preimage. For a fixed position $j \in \{1, \dots, L\}$, let E_j be the event that position j exists and $[\text{RO}_{\text{flat}}(Y_j)]_\tau = t$. Condition on any positive-probability prefix immediately before Y_j is appended; the target t is a constant. By the above, $\text{RO}_{\text{flat}}(Y_j)$ is uniform and independent of the conditioning, so its τ -bit projection equals t with probability $2^{-\tau}$. If position j is never appended, E_j is false. A union bound over the L positions gives the claim.

Second-preimage corollary. Expose $\text{RO}_{\text{flat}}(Y_{j_0})$ at the designated candidate's append step, which predictability and unrevealed value license; this fixes $t = [\text{RO}_{\text{flat}}(Y_{j_0})]_\tau$ as a target independent of the remaining candidates. Applying the preimage part to the $L - 1$ candidates $j \neq j_0$, each still unrevealed when appended, gives $\Pr[\text{sec}_{\tau,j_0}] \leq (L - 1) \cdot 2^{-\tau}$; the designated candidate is excluded, so its trivially-equal projection does not enter the count. $\square$

A.6 Adaptive Key Tests

Lemma 9 (Adaptive Key-Test Bound). *Consider a TW128 random oracle model game with a single hidden context component Z sampled uniformly from $\{0, 1\}^w$ for some $w \geq k$ —the standard case being the hidden key, $Z = K$ and $w = k$, used in the AE-style proofs—direct adversary access to RO_{flat} , and an immediate abort event bad_{key} that fires on a direct query (Y, ℓ) exactly when Y decodes to some counted output-point class $\mathcal{C}$ of Table 9 with $\text{KeyedTWO}_{\text{out}}(Y) = Z$. Suppose that, before the abort, the following invariant holds for every direct query position i : after conditioning on the full visible execution prefix and on the set S_i of distinct decoded candidate Z -values from earlier direct queries, the hidden Z is uniform on $\{0, 1\}^w \setminus S_i$. Then*

$$\Pr[\text{bad}_{\text{key}}] \leq q_{\text{RO}}/2^w \leq \text{KeyGuess}_k(q_{\text{RO}}) = q_{\text{RO}}/2^k,$$

where q_{RO} is the execution-uniform count of direct RO_{flat} queries.

Proof. If $q_{\text{RO}} \geq 2^w$ the right-hand side $q_{\text{RO}}/2^w$ is at least one, so the claim is trivial. Assume $q_{\text{RO}} < 2^w$.

Order the direct RO_{flat} queries. By the definition of $\text{KeyedTWO}_{\text{out}}$ and Lemma 6, each query input either decodes to $\perp$ or decodes to a unique candidate Z -value and counted output-point class. Queries that decode to $\perp$, and repeated queries whose decoded Z -value already lies in the current set S_i , do not change the set of possible first successful tests. Removing them can only shorten the sequence, so it suffices to analyze the subsequence of distinct fresh decoded candidate Z -values $Z'_1, \dots, Z'_m$, where $m \leq q_{\text{RO}}$.

Let F_i be the event that the first i fresh decoded candidate values all miss the hidden Z . The invariant says that, conditioned on the visible prefix before the $(i+1)$ -st fresh test and on F_i , the hidden Z is uniform over the $2^w - i$ values not yet tested. The next fresh candidate value Z'_{i+1} is one of those remaining values and is determined by the adversary's next direct query, so

$$\Pr[Z = Z'_{i+1} \mid F_i] = \frac{1}{2^w - i}, \quad \Pr[F_{i+1} \mid F_i] = 1 - \frac{1}{2^w - i}.$$

Multiplying the miss probabilities gives

$$\Pr[F_m] = \prod_{i=0}^{m-1} \left(1 - \frac{1}{2^w - i}\right) = \prod_{i=0}^{m-1} \frac{2^w - i - 1}{2^w - i} = \frac{2^w - m}{2^w}.$$

Thus the probability that some fresh decoded candidate value equals the hidden Z is $m/2^w \leq q_{\text{RO}}/2^w \leq q_{\text{RO}}/2^k$, the last step by $w \geq k$. This event is exactly the first bad_{key} abort, because $\text{KeyedTWO}_{\text{out}}$ returns either $\perp$ or a unique candidate Z -value. $\square$

A.7 Leaf Tag Collision Bound

Lemma 10 (Leaf Tag Collision Bound). *Consider any TW128 random oracle model game with an adversary $\mathcal{B}$ that has direct access to RO_{flat} and construction computations that may generate final leaf transcripts—by encryption or by verification computations, whether they accept or reject. Let bad_{leaf} be the event that two distinct construction-generated final leaf transcript prefixes $X \neq X'$ in leaf tag output-point classes receive the same π_{leaf} -bit projection from $\text{RF}(\cdot)$, and let $n_{\text{leaf}}(\mathcal{B})$ be the execution-uniform count of distinct construction-generated final leaf transcripts (Section 4.3). Both bounds below apply the collision facet ($s = 2$, $\tau = \pi_{\text{leaf}}$) of Lemma 8 to a chronological leaf tag candidate sequence, in two regimes.*

1. Adversary-supplied keys. When the game's keys are adversary-supplied—so direct queries may themselves land on leaf tag output points—

$$\Pr[\text{bad}_{\text{leaf}}] \leq \text{LeafColl}_{\tau_{\text{leaf}}}(q_{\text{RO}}(\mathcal{B}) + n_{\text{leaf}}(\mathcal{B})) = \binom{q_{\text{RO}}(\mathcal{B}) + n_{\text{leaf}}(\mathcal{B})}{2} / 2^{\tau_{\text{leaf}}}.$$

2. Hidden key. When the game has a hidden key, let hit_{leaf} be the event that a direct adversary query hits the bridged input of a construction-generated hidden-key final leaf transcript before that transcript is first created. Then

$$\Pr[\text{bad}_{\text{leaf}} \text{ before } \text{hit}_{\text{leaf}}] \leq \binom{n_{\text{leaf}}(\mathcal{B})}{2} \cdot 2^{-\tau_{\text{leaf}}} = \frac{n_{\text{leaf}}(\mathcal{B})(n_{\text{leaf}}(\mathcal{B}) - 1)}{2^{\tau_{\text{leaf}} + 1}}.$$

Proof. Build a chronological leaf tag candidate sequence: a candidate is the bridged input of a leaf tag output point ($\mathcal{L}_{\text{tag}}(j)$), appended the first time it is reached—by a direct query in $\mathcal{B}$'s oracle phase or by a construction computation generating a final leaf transcript—and not appended again. By Lemma 6, Corollary 7, and Lemma 4, two reaches of the same bridged input come from the same final leaf transcript, so repeated evaluations add no candidate—the lazy RO_{flat} value is reused—and distinct final leaf transcripts give distinct candidates. No candidate's RO_{flat} value is read before its append step, so the sequence meets the predictability and unrevealed-value premises of Lemma 8, whose collision facet with $s = 2$ and $\tau = \tau_{\text{leaf}}$ bounds the probability that two candidates share a τ_{leaf} -bit projection by $\binom{L}{2} \cdot 2^{-\tau_{\text{leaf}}}$ for a length- L sequence. Every bad_{leaf} pair is a pair of construction-generated candidates.

(1) *Adversary-supplied keys.* Direct queries may decode to a leaf tag class $\mathcal{L}_{\text{tag}}(j)$ and are appended as candidates; together with the $n_{\text{leaf}}(\mathcal{B})$ construction-generated transcripts the sequence has length at most $q_{\text{RO}}(\mathcal{B}) + n_{\text{leaf}}(\mathcal{B})$, so the facet with $L = q_{\text{RO}}(\mathcal{B}) + n_{\text{leaf}}(\mathcal{B})$ gives $\Pr[\text{bad}_{\text{leaf}}] \leq \binom{q_{\text{RO}}(\mathcal{B}) + n_{\text{leaf}}(\mathcal{B})}{2} \cdot 2^{-\tau_{\text{leaf}}} = \text{LeafColl}_{\tau_{\text{leaf}}}(q_{\text{RO}}(\mathcal{B}) + n_{\text{leaf}}(\mathcal{B}))$.

(2) *Hidden key.* A direct query may still decode under **KeyedTWO** to a $\mathcal{L}_{\text{tag}}(j)$ class for an adversary-chosen key other than the hidden key, and is appended as a candidate. But every bad_{leaf} pair is a pair of construction-generated transcripts, and under a hidden key these are all generated with that key, so it suffices to count collisions within the subsequence of construction-generated hidden-key final leaf transcripts, of length at most $n_{\text{leaf}}(\mathcal{B})$. Before hit_{leaf} no direct query has hit the bridged input of any such transcript: by Lemma 6 such a query would decode under **KeyedTWO** to the hidden key and its $\mathcal{L}_{\text{tag}}(j)$ class, which is exactly hit_{leaf} . Each leaf tag in the subsequence is therefore fresh when its transcript is first created, so the facet applied to this length- $n_{\text{leaf}}(\mathcal{B})$ subsequence gives $\Pr[\text{bad}_{\text{leaf}} \text{ before } \text{hit}_{\text{leaf}}] \leq \binom{n_{\text{leaf}}(\mathcal{B})}{2} \cdot 2^{-\tau_{\text{leaf}}}$. By Leaf-Transcript Recovery (Lemma 5), conditioned on no leaf tag collision, equal flattened final-root inputs identify the same leaf transcript sequence. $\square$

A.8 Transcript Freshness

Lemma 11 (Transcript Freshness). *For any fixed nonce-respecting encryption query in the TW128 random oracle experiment, let $\text{hit}_{\text{fresh}}$ be the event that, before this query reaches some counted transcript point, a prior direct adversary query to RO_{flat} has already hit the bridged input of that hidden-key point. On $\neg \text{hit}_{\text{fresh}}$, every transcript output of the query that contributes to a returned ciphertext block, to an internal leaf tag, or to the final root tag is fresh at the moment its RO_{flat} value is sampled.*

Proof. Work on an execution outside $\text{hit}_{\text{fresh}}$ for the fixed query. This rules out prior adversary reads of the hidden-key construction inputs considered below. It remains to rule out repeats from previous construction interface outputs and from the current encryption query.

First root keystream block. The first root keystream block, when present, is produced by the root duplex after the initialization call $p_{\text{root}} \parallel K \parallel N$ and any AD absorption (elided when $A = \varepsilon$, in which case the σ_{init} call itself emits the first root keystream block). Since the nonce-respecting adversary never repeats a nonce, no previous encryption query can have used the same root initialization transcript. Within the current query, Lemma 4 separates initialization, AD, message, and aggregation calls by their suffixes and by the duplex-call encoding, so the first root keystream transcript is fresh. If the requested keystream length is zero, no ciphertext block is returned at that point and there is nothing to prove for that output.

First leaf keystream block. For each leaf chunk $j \geq 1$, the first leaf keystream block is produced after the initialization call $p_{\text{leaf}} \parallel K \parallel N \parallel \langle j \rangle_8$. The prefix separates leaves from the root, the nonce separates this encryption query from all previous encryption queries, and $\langle j \rangle_8$ separates leaves within this query. Hence every first leaf keystream transcript is fresh.

Later message outputs. Argue inductively within each root or leaf message stream. Suppose the current message-output transcript is fresh. If its output is used as a keystream block, its RO_{flat} value Z_i is, by freshness of the lazy sample, uniform conditioned on the full prior view; since the adaptive plaintext block M_i is a deterministic function of that prior view,

$$C_i = M_i \oplus Z_i$$

is uniform conditioned on the prior view. TW128 then absorbs C_i with the appropriate message phase suffix (σ_{msg}^- or σ_{msg}^+) before requesting the next keystream block. Once C_i is sampled, the next transcript is fixed; Lemma 4's injectivity and domain separation imply it has not appeared earlier unless the same duplex-call prefix has appeared earlier. Nonce uniqueness excludes this across encryption queries. By Corollary 7, the next message-keystream output point can coincide with a within-query earlier output point only if both belong to the same output-point class on the same duplex instance (root or leaf- j with matching j) and have the same duplex-call sequence; the current message-stream prefix is strictly longer than every earlier prefix in that stream, so this is excluded. Across distinct instances within the same query (root vs leaf, or two leaves with distinct j), the prefixes already differ in the first call's σ by $p_{\text{root}}/p_{\text{leaf}}$ or $\langle j \rangle_8$ respectively. Each later message-output transcript that emits keystream is therefore fresh when sampled. The terminal message-output transcript of a leaf emits the leaf tag rather than a following keystream block; the same separation argument applies to that transcript, so each internal leaf tag is sampled at a fresh point. Empty inputs still induce a duplex call, but with zero keystream output they contribute no ciphertext block.

Final root tag transcript. The final root tag is sampled at the root tag output point selected by $\|C\|$: the $\mathcal{R}_{\text{tag}}$ point after the root duplex absorbs the aggregation transcript built from the leaf tags and $\langle n \rangle_8$ when $n \geq 1$, and the $\mathcal{R}_{\text{msg}}^+$ point closing the root message phase when $n = 0$ (Section 3.5). When $n \geq 1$, the leaf tags are themselves RO_{flat} outputs on final leaf transcripts; condition on their sampled values, on the sampled ciphertext, and on the adversary-visible prior view and RO_{flat} table state, so that the root aggregation transcript is fixed. When $n = 0$ there is no aggregation transcript; condition on the sampled ciphertext and the adversary-visible prior view alone. By Corollary 7, the selected root tag output-point class is separated from every other output-point class in the same query — including root initialization (suffix σ_{init}), AD-phase outputs, the remaining message phase outputs, leaf-side output points (separated by p_{leaf}), and, in the $n \geq 1$ case, earlier $\mathcal{R}_{\text{agg}}$ -prefix output points (separated by the σ_{agg}^- vs σ_{agg}^+ suffix). Even if two encryption queries realize the same aggregation transcript (or, for $n = 0$, the same root ciphertext), the

full root transcript also contains the root initialization $p_{\text{root}} \parallel K \parallel N$, and nonce uniqueness separates that initialization from all previous encryption queries. The final root tag transcript is therefore fresh, and its τ_{root} -bit RO_{flat} output is uniform independently of the conditioned information. $\square$

A.9 Mu-Ind Verification Coupling

The direct mu-ind adjacent-hybrid proof uses the generic key-test and leaf-collision accounting lemmas together with the verification-specific facts below. The *root tag output point* of a computation is the output point that emits its root tag: $\mathcal{R}_{\text{tag}}$ (suffix σ_{agg}^+) when $n \geq 1$, and the closing root message point $\mathcal{R}_{\text{msg}}^+$ (suffix σ_{msg}^+) when $n = 0$ and the aggregation phase is elided (Section 3.5). Because $n = \text{leaves}(C)$ is fixed by $\|C\|$, all computations on a common ciphertext share the same root tag output point. A *final root transcript class* is the set of construction computations whose encryption or verification computation realizes the same bridged final root input at this root tag output point. Such a class is *encryption-first* if its first construction generation is by an encryption query, and *verification-first* if its first construction generation is by a non-replay verification query.

Lemma 12 (Verification Visible-Prefix Key Uniformity). *In the stopped single-key TW128 random oracle encryption/verification interaction with hidden key K , direct access to RO_{flat} , and immediate abort on bad_{key} , order the direct RO_{flat} queries. For a direct query position i , let S_i be the set of distinct candidate keys decoded by KeyedTWOut from earlier direct-query inputs. Condition on any positive-probability visible execution prefix immediately before query i , on no earlier bad_{key} , and on S_i . Then the hidden key K is uniform on $\{0, 1\}^k \setminus S_i$. Consequently the uniformity invariant required by Lemma 9 holds in the stopped interaction.*

Proof. The visible prefix contains the adversary’s oracle queries and replies: direct RO_{flat} outputs, encryption outputs and the resulting replay table, verification bits, and the adversary state determined by these values. Fix such a prefix and two keys $K_0, K_1 \notin S_i$.

Let $\Phi_{0 \rightarrow 1} = \Phi_{K_0 \rightarrow K_1}$ be the key-renaming map of Lemma 7 and $\Phi_{1 \rightarrow 0}$ its inverse. By that lemma, $\text{flat} \circ \Phi_{0 \rightarrow 1} \circ \text{flat}^{-1}$ is a bijection $\mathcal{Y}_{K_0} \rightarrow \mathcal{Y}_{K_1}$ between the bridged hidden-key construction addresses for K_0 and for K_1 , preserving output-point classes and equality relations among addresses.

This bijection is applied online, not to a pre-existing fixed address set. Whenever a construction computation under K_0 would first query a hidden-key address Y , the coupled computation under K_1 queries $\Phi_{0 \rightarrow 1}(Y)$ and receives the same lazy-table value; repeated queries are coupled consistently through the same map. If the queried address is a root aggregation address containing leaf tag bytes, those bytes are already coupled equal because the corresponding leaf tag lookups were either repeated or first sampled earlier under the same renaming rule. Thus adaptively generated message transcripts and root aggregation transcripts are renamed step by step with the same visible ciphertext, leaf tag, and root tag bytes.

Prior direct-query inputs are not renamed: they are adversary-chosen visible strings. Since $K_0, K_1 \notin S_i$, neither $\mathcal{Y}_{K_0}$ nor $\mathcal{Y}_{K_1}$ contains a prior direct-query input (Lemma 7, direct-query avoidance). Direct-query lazy-table values are therefore coupled identically and remain disjoint from the renamed hidden construction addresses.

Use the same adversary coins, the same direct-query lazy-table values, and the online renamed construction-internal lazy-table values. Encryption outputs have the same distribution under the coupling, because every construction lookup that determines ciphertext, leaf tags, or the final root tag is paired with a distinct renamed lookup carrying the same sampled value. Replay verification decisions depend only on the visible replay table. Non-replay verification computations are renamed in the same online way as encryption computations, and the final root tag comparison therefore gives the same accept/reject bit. Verification does not reveal plaintext on rejection or acceptance; it reveals only that bit.

Thus every visible prefix with hidden key $K_0 \notin S_i$ has the same conditional probability as the corresponding execution with hidden key $K_1 \notin S_i$. Since K is initially uniform and the condition “no earlier bad_{key} ” removes exactly the keys in S_i , all keys in $\{0, 1\}^k \setminus S_i$ remain equally likely. $\square$

Lemma 13 (Verification-First Root Guessing). *In the stopped single-key TW128 random oracle encryption/verification interaction before $\text{bad}_{\text{key}} \cup \text{bad}_{\text{leaf}}$, let win_{vf} be the event that some accepting non-replay verification submission belongs to a verification-first final root transcript class. Then*

$$\Pr[\text{win}_{\text{vf}}] \leq q_v(\mathcal{B})/2^{\tau_{\text{root}}}.$$

Proof. Throughout, a *non-replay verification submission* means a valid non-replay verification query; by the validity convention, an invalid verification query returns 0 and performs no construction computation, so it realizes no bridged final root input, belongs to no final root transcript class, and contributes to no G_C .

Work in the game stopped at the first occurrence of bad_{key} or bad_{leaf} . For accounting, define an all-reject continuation: encryption queries are answered normally, replay verification queries accept normally, and every non-replay verification submission performs the same construction computation but returns reject to the adversary. Lazy sampling is deferred for each verification-first class’s final root tag value: the continuation records the class address and the submitted tag, but while the class remains verification-first it does not expose that root tag value to the adversary. The execution-uniform count $q_v(\mathcal{B})$ applies to this continuation.

Fix a verification-first final root transcript class $\mathcal{C}$ and follow the all-reject path until the first later encryption query generating the same class, if any. Let G_C be the set of distinct candidate tags submitted by non-replay verification queries in class $\mathcal{C}$ before that cutoff. Repeated tags only shrink this set. The class’s correct root tag U_C is the τ_{root} -bit projection of RO_{flat} at the class’s bridged root tag input. Before the stop, no direct query has hit this hidden-key input: such a query would be decoded by KeyedTWOut to the hidden key and would trigger bad_{key} . Before the cutoff, no encryption query has revealed U_C as an encryption tag. (When the class’s root tag point is $\mathcal{R}_{\text{msg}}^+$, i.e. $n = 0$, an encryption query on a longer message that shares chunk 0 may touch this same address while closing its chunk-0 message phase, but it requests zero output there and continues into its aggregation phase, so it does not reveal U_C ; its own tag is emitted at a later $\mathcal{R}_{\text{tag}}$ point.) The all-reject path and the set G_C are therefore functions only of the adversary’s coins, visible oracle replies, and lazy-table values at inputs other than this class’s root tag address. By lazy sampling, conditioned on those values and on the fixed path, U_C may be sampled after G_C is fixed and is uniform independently of G_C .

After the cutoff, if it exists, class $\mathcal{C}$ contributes no accepting non-replay submission. Let the cutoff encryption query return (N_e, A_e, C_e, T_e) . Any later verification submission in the same final root transcript class has, by the final-root determinacy of Lemma 4 and Leaf-Transcript Recovery (Lemma 5) under the absence of bad_{leaf} , the same (N, A, C) as that encryption computation. If the submitted tag is T_e , the tuple is a recorded encryption tuple and the verification is a replay; if the submitted tag is not T_e , the final tag comparison fails.

Consider the event that the first accepting non-replay verification submission in the stopped game belongs to class $\mathcal{C}$. Up to that global first accept, the adversary’s visible answers match the all-reject continuation: every earlier non-replay submission rejects, replay queries accept in both executions, and encryption answers are unchanged. The construction computations also agree at every address except that the all-reject continuation defers the final root tag value for class $\mathcal{C}$ until the comparison is needed. Therefore the first accepting non-replay submission can belong to class $\mathcal{C}$ only before the cutoff, and only if $U_C \in G_C$, which has probability at most $|G_C|/2^{\tau_{\text{root}}}$.

Each non-replay verification submission deterministically reconstructs one final root transcript at the root tag output point. If the class is verification-first and the submission occurs before that class's encryption cutoff, the submitted tag contributes to at most one set G_C ; otherwise it contributes to none. Thus $\sum_C |G_C| \leq q_v(\mathcal{B})$, and a union bound over the verification-first classes gives the claim. $\square$

A.10 Mu-Ind Adjacent-Hybrid Lemma

The multi-user indistinguishability theorem (Theorem 2) telescopes a hybrid over the D users, replacing one real user by its ideal (Rand, Replay) interface at each step. The following lemma bounds one such step directly inside the shared random oracle environment. It couples the real and ideal user- d interfaces until one of three local events occurs: a direct hidden-key transcript hit, a hidden leaf tag collision, or an accepting non-replay verification query.

Lemma 14 (Mu-Ind Adjacent-Hybrid Bound). *Let $\mathcal{B}$ be a flat-RO mu-ind adversary in the pre-sampled-key formulation of Section 4.4, with keys $K_1, \dots, K_D$ sampled at the start of the experiment. Fix $d \in \{1, \dots, D\}$. Let H_{d-1}^* and H_d^* be the adjacent stopped hybrids that output a fixed bit if any two sampled keys collide, and otherwise run $\mathcal{B}$ with one shared lazy RO_{flat} table as follows: users $e < d$ are answered by (Rand, Replay), users $e > d$ are answered by $(\text{Enc}_{K_e}^{\text{RO}}, \text{Ver}_{K_e}^{\text{RO}})$, and user d is answered by $(\text{Enc}_{K_d}^{\text{RO}}, \text{Ver}_{K_d}^{\text{RO}})$ in H_{d-1}^* and by (Rand, Replay) in H_d^* . If $n_{\text{leaf}}^{(d)}$ and $q_v^{(d)}$ are the user- d resource caps, then*

$$|\Pr[\mathcal{B}^{H_{d-1}^*} \Rightarrow 1] - \Pr[\mathcal{B}^{H_d^*} \Rightarrow 1]| \leq \text{KeyGuess}_k(q_{\text{RO}}(\mathcal{B})) + \frac{n_{\text{leaf}}^{(d)}(n_{\text{leaf}}^{(d)} - 1)}{2^{\tau_{\text{leaf}} + 1}} + \frac{q_v^{(d)}}{2^{\tau_{\text{root}}}}.$$

Proof. The fixed-output branch on a key collision is identical in H_{d-1}^* and H_d^* , so only the branch with pairwise distinct keys contributes. On that branch, the two hybrids differ only in user d : users $e < d$ are ideal in both hybrids, users $e > d$ are real in both hybrids, and user d is answered by $(\text{Enc}_{K_d}^{\text{RO}}, \text{Ver}_{K_d}^{\text{RO}})$ in H_{d-1}^* and by (Rand, Replay) in H_d^* .

Let $\mathcal{E}_d$ denote the common surrounding environment. It runs $\mathcal{B}$, answers users $e < d$ by their (Rand, Replay) interfaces, answers users $e > d$ by $\text{Enc}_{K_e}^{\text{RO}}, \text{Ver}_{K_e}^{\text{RO}}$, forwards $\mathcal{B}$'s direct RO_{flat} queries unchanged, and shares one lazy RO_{flat} table with user d . The coupling treats user d 's encryption and verification interfaces together. Direct RO_{flat} queries that decode to K_d can affect both interfaces, so the local key-test event is charged once in the adjacent-hybrid bound.

Coupling until divergence. Run the two hybrids on shared adversary coins and one shared lazy RO_{flat} table. Introduce the user- d local events $\text{bad}_{\text{key}}^{(d)}$, that a direct RO_{flat} query decodes under KeyedTWOut to K_d and a counted output-point class; $\text{bad}_{\text{leaf}}^{(d)}$, that two distinct user- d construction-generated final leaf transcripts receive the same τ_{leaf} -bit projection; and $\text{forge}^{(d)}$, that a non-replay user- d verification query is accepted by $\text{Ver}_{K_d}^{\text{RO}}$.

Couple the two hybrids on shared adversary coins and one shared lazy RO_{flat} table, with user d 's Rand answers in H_d^* set equal to user d 's real encryption answers in H_{d-1}^* . We show the two hybrids' visible transcripts agree until $\text{bad}_{\text{key}}^{(d)} \cup \text{forge}^{(d)}$.

The coupled H_d^ has the true H_d^* law.* Each real user- d encryption output is a projection of RO_{flat} at the distinct transcript points of its computation—distinct within a computation and, by encryption nonce-uniqueness and Transcript Injectivity (Lemma 4), across distinct encryptions—so until a direct query reads such a point the outputs are jointly uniform of the prescribed lengths; this is the freshness guarantee of Lemma 11. The present environment adds a real user- d verification oracle, absent from a pure encryption-replacement game. Because only encryption queries are nonce-unique, a non-replay verification may reuse an encryption nonce and lazily read a keystream or root tag point the same-nonce encryption

later reaches; but lazy sampling draws each point uniformly when first read, so its value is unchanged and stays marginally uniform. In H_d^* the user- d verification oracle is **Replay**, which answers from the recorded ciphertexts without reading RO_{flat} . Hence the coupled H_d^* transcript has the law of the true H_d^* until $\text{bad}_{\text{key}}^{(d)}$: its encryption answers are the marginally uniform strings of **Rand** and its verification answers are the replay decisions, with no further RO_{flat} dependence. (Conditioning on a particular H_{d-1}^* verification outcome can constrain a coupled tag, but that is a property of the joint coupling; the coupled H_d^* marginal is what must match the true H_d^* , and it does.)

Agreement until $\text{bad}_{\text{key}}^{(d)} \cup \text{forge}^{(d)}$. Encryption answers are equal by construction. A verification query whose tuple is recorded returns accept in both hybrids through the shared replay-table membership check, which precedes the real-versus-ideal decryption branch (Section 4.4), so real decryption is never run on a recorded tuple. A non-replay query returns reject in H_d^* and the bit $[\text{Ver}_{K_d}^{\text{RO}} \text{ accepts}]$ in H_{d-1}^* ; the two agree unless that bit is accept, the event $\text{forge}^{(d)}$. A real verification reads hidden-key points but exposes only this bit, which is independent of the keystream values it reads—decryption absorbs the submitted ciphertext, not the recovered plaintext (Lemma 3)—and depends on an encryption output only through the root tag it recomputes. In the reused-nonce case above this correlates the H_{d-1}^* bit with that encryption’s tag, but the H_d^* verification is **Replay** and the coupled encryption answers coincide, so the two coordinates still agree on such a query unless it accepts. Direct-query answers are read from the shared table and agree until a direct query reads a user- d hidden-key point ($\text{bad}_{\text{key}}^{(d)}$). Therefore the visible transcripts agree until $\text{bad}_{\text{key}}^{(d)} \cup \text{forge}^{(d)}$, and a union bound ordered by first occurrence (inserting the leaf abort to split $\Pr[\text{forge}^{(d)}]$ before $\text{bad}_{\text{key}}^{(d)}$) gives

$$\begin{aligned} & |\Pr[\mathcal{B}^{H_{d-1}^*} \Rightarrow 1] - \Pr[\mathcal{B}^{H_d^*} \Rightarrow 1]| \\ & \leq \Pr[\text{bad}_{\text{key}}^{(d)}] + \Pr[\text{bad}_{\text{leaf}}^{(d)} \text{ before } \text{bad}_{\text{key}}^{(d)}] \\ & \quad + \Pr[\text{forge}^{(d)} \text{ before } \text{bad}_{\text{key}}^{(d)} \cup \text{bad}_{\text{leaf}}^{(d)}]. \end{aligned}$$

The environment does not enlarge the user- d events. The real users $e > d$ perform construction-internal RO_{flat} lookups while simulating $\text{Enc}_{K_e}^{\text{RO}}$ and $\text{Ver}_{K_e}^{\text{RO}}$. Whenever such a lookup lies in a counted keyed transcript class, Lemma 6 decodes its candidate key as K_e , which on the non-collision branch differs from K_d ; a lookup outside the counted keyed language is irrelevant to $\text{bad}_{\text{key}}^{(d)}$. Foreign-user lookups therefore neither trigger user d ’s key-test event nor pre-read a user- d hidden transcript point, and they create no user- d leaf tag collision or user- d non-replay verification submission, which are counted by $n_{\text{leaf}}^{(d)}$ and $q_v^{(d)}$. The three accounting events are thus local to user d ’s keyed transcript language.

Bounding the three terms. For $\Pr[\text{bad}_{\text{key}}^{(d)}]$, adapt the single-key visible-prefix renaming coupling of Lemma 12 to the environment $\mathcal{E}_d$. Its conclusion does not transfer verbatim: the foreign users $e > d$ share the lazy RO_{flat} table, so renaming K_d onto a foreign key would alter their visible transcripts, and the renaming orbit must exclude the foreign keys rather than range over all of $\{0, 1\}^k \setminus S_i$. What the coupling still delivers in $\mathcal{E}_d$ is the marginal per-candidate bound, which is all the key test needs. The key-test contribution is taken in the original pre-sampled-key experiment, not in one renormalized to the non-collision branch: fix the foreign keys and run the visible-prefix coupling with foreign-key transcript addresses left unchanged; before the first user- d key hit, the visible prefix is invariant under renaming K_d among keys not already tested and not equal to a fixed foreign key. For each fresh candidate key tested by one of $\mathcal{B}$ ’s direct RO_{flat} queries, the contributing event is contained in $\{K_d \text{ equals that candidate}\}$, costing at

most 2^{-k} , and a foreign candidate contributes nothing; with at most $q_{\text{RO}}(\mathcal{B})$ such tests, the union bound of Lemma 9 packages these into $\Pr[\text{bad}_{\text{key}}^{(d)}] \leq \text{KeyGuess}_k(q_{\text{RO}}(\mathcal{B}))$. For the leaf term, before $\text{bad}_{\text{key}}^{(d)}$ no direct query has pre-read the bridged input of a user- d hidden-key final leaf transcript—such a query would decode under KeyedTWOOut to K_d and its $\mathcal{L}_{\text{tag}}(j)$ class and would itself be $\text{bad}_{\text{key}}^{(d)}$ —so the hidden-key instance of Lemma 10 gives $\Pr[\text{bad}_{\text{leaf}}^{(d)} \text{ before } \text{bad}_{\text{key}}^{(d)}] \leq n_{\text{leaf}}^{(d)}(n_{\text{leaf}}^{(d)} - 1)/2^{\tau_{\text{leaf}}+1}$. For the forgery term, before $\text{bad}_{\text{key}}^{(d)} \cup \text{bad}_{\text{leaf}}^{(d)}$ every encryption-first final root transcript class admits no accepting non-replay submission—by the final-root determinacy of Lemma 4 and Leaf-Transcript Recovery (Lemma 5), a non-replay submission in such a class reaches the already-returned root tag and fails the tag comparison—so every accepting non-replay submission is verification-first, and Lemma 13 gives $\Pr[\text{forge}^{(d)} \text{ before } \text{bad}_{\text{key}}^{(d)} \cup \text{bad}_{\text{leaf}}^{(d)}] \leq q_v^{(d)}/2^{\tau_{\text{root}}}$.

Summing the three bounds gives

$$|\Pr[\mathcal{B}^{H_d^*} \Rightarrow 1] - \Pr[\mathcal{B}^{H_d^*} \Rightarrow 1]| \leq \text{KeyGuess}_k(q_{\text{RO}}(\mathcal{B})) + \frac{n_{\text{leaf}}^{(d)}(n_{\text{leaf}}^{(d)} - 1)}{2^{\tau_{\text{leaf}}+1}} + \frac{q_v^{(d)}}{2^{\tau_{\text{root}}}},$$

with the user- d key-test event $\text{bad}_{\text{key}}^{(d)}$ charged a single time for the adjacent hybrid. $\square$

B The Flat Sponge Lift

This appendix expands the sketch of Section 5.2 into the full argument: it supplies the flat sponge lift machinery (Lemmas 16 and 17 and Corollary 8) invoked by the headline theorems of Section 5.2 and by Corollary 6. The machinery applies uniformly to the games named once here.

Definition 3 (Single-Stage TW128 Games). The *single-stage TW128 games* are the games of Sections 4.1 and 4.4: all-in-one AE, mu-ind, CMTDs-4, the H collision, second-preimage, and preimage games, $\text{CDY}^*[\text{ts}]$, and CMTs- ℓ for $\ell \in \{1, 3, 4\}$. In $\text{CDY}^*[\text{ts}]$ the selector encryption is construction-internal, counted in N like a challenger check.

B.1 Simulator Convention

Throughout this appendix, $(\mathcal{S}^{\text{RO}_{\text{flat}}}, \mathcal{S}^{-1, \text{RO}_{\text{flat}}})$ denotes the public primitive simulator supplied by [BDPVA08, Theorem 2] — the indistinguishability theorem for the simple padded sponge over a random permutation — after applying the [BDPVA11, Theorem 3] bridge of Lemma 6 to TW128’s native pad10^*1 transcript inputs. The superscript records that both simulator directions share the same internal lazy RO_{flat} table. This is the simulator referenced by the Sections 4.1 and 4.4 definitions of $\text{Adv}_{\text{TW128}}^{\text{ae}}$ and $\text{Adv}_{\text{TW128}}^{\text{mu-ind}}$: the ideal side of each AE-style experiment exposes the construction interface specified by the game $((\text{Rand}, \text{Replay}), \text{instantiated per user in mu-ind})$ together with $(\mathcal{S}^{\text{RO}_{\text{flat}}}, \mathcal{S}^{-1, \text{RO}_{\text{flat}}})$ as the public primitive interface. The RO_{flat} table in this notation is internal to the ideal system and is not an extra interface exposed to the application adversary; the committing security advantages remain the real public permutation success probabilities of Section 4.1, with the simulator appearing only inside the reduction to the corresponding random oracle model adversary.

Lemma 15 (Simulator RO-Call Filter). *For the simulator $(\mathcal{S}^{\text{RO}_{\text{flat}}}, \mathcal{S}^{-1, \text{RO}_{\text{flat}}})$, every inverse primitive-interface query and every forward primitive-interface query violating one of the following conditions induces no simulator RO_{flat} call. The symbols s , p , $\mathcal{R}$, $\mathcal{O}$, and $\mathcal{C}$ are the local internal notation of the BDPV simple-padded-sponge simulator, not TW128 resource notation:*

2393 **(C1) New edge.** *The simulator's internal graph has no outgoing edge from the input node*
 2394 *s.*

2395 **(C2) Unsaturated.** $\mathcal{R} \cup \mathcal{O} \neq \mathcal{C}$.

2396 **(C3) Rooted input.** *The input node is rooted.*

2397 **(C4) Paddable path.** *The constructed path p decomposes uniquely as $p = p' \parallel 0^{r_j}$ with*
 2398 *p' not ending in 0^r , and the prefix p' is a valid padded sponge input: equivalently,*
 2399 *$p' = x \parallel \text{pad10}^*[r](|x|)$ for an unpadded BDPV08 random oracle input x .*

2400 *When all four conditions hold, the forward primitive-interface query induces one simulator*
 2401 *RO_{flat} call. Consequently, each primitive-interface query induces at most one simulator-side*
 2402 *RO_{flat} call.*

2403 *Proof.* This is the branch structure of the [BDPVA08] simple-padded permutation simulator.
 2404 The simulator invokes the random oracle only on the forward-query branch where the input
 2405 node is rooted, unsaturated, paddable to an unpadded random oracle input x , and not
 2406 already extended by an outgoing edge in the simulator graph. Forward queries violating
 2407 any of these conditions and all inverse queries return outputs sampled or read from the
 2408 simulator's internal graph without invoking the random oracle. Under the [BDPVA11,
 2409 Theorem 3] bridge fixed in Section B.1, that simulator random oracle invocation is the
 2410 simulator-side RO_{flat} call used throughout this appendix. $\square$

2411 **Lemma 16** (Flat Sponge Lift Bound). *Let G be a single-stage TW128 game (Definition 3).*
 2412 *Let $\mathcal{A}$ be a public permutation G adversary with construction interface cost at most N and at*
 2413 *most N_{prim} direct primitive-interface queries. In the [BDPVA08] single-stage replacement,*
 2414 *the combined construction-and-primitive transcript has length at most*

$$2415 \quad N_{\text{tot}} = N + N_{\text{prim}}.$$

2416 *Consequently replacing the real public permutation and TW128 construction by the*
 2417 *simulator-world primitive interface $(\mathcal{S}^{\text{RO}_{\text{flat}}}, \mathcal{S}^{-1, \text{RO}_{\text{flat}}})$ and the corresponding flat-RO*
 2418 *construction interface costs at most $\text{Lift}_c(N_{\text{tot}})$.*

2419 *Proof.* The [BDPVA08] replacement is applied before the derived-adversary construction,
 2420 so its transcript contains two kinds of public primitive activity. First, each construction
 2421 interface computation contributes the TW128 duplex calls it performs. By Lemmas 3
 2422 and 6, each such root or leaf output point is a typed restriction of the bridged flat sponge
 2423 construction, and by the definition of N in Section 4.3 these calls are counted at the
 2424 same per-block granularity as the underlying KECCAK- $p[1600, 12]$ calls. This includes
 2425 verification computations whether they accept or reject, the per-user sum of construction
 2426 calls in mu-ind, and the deterministic challenger decryption or encryption checks actually
 2427 performed in the committing and discoverability games.

2428 Second, each direct adversary query to (π, π^{-1}) contributes one primitive-interface
 2429 query and is counted in N_{prim} . These two counts are disjoint in the public permutation
 2430 experiment: N counts construction-internal use of the primitive, while N_{prim} counts only
 2431 the adversary's public primitive interface queries. The execution-uniform caps ensure
 2432 that every adaptive transcript seen by the single-stage replacement has length at most
 2433 $N + N_{\text{prim}} = N_{\text{tot}}$. If $N_{\text{tot}} \geq 2^c$ then $\text{Lift}_c(N_{\text{tot}}) \geq 1$ and the claimed bound holds trivially, so
 2434 assume $N_{\text{tot}} < 2^c$, satisfying the hypothesis of [BDPVA08, Theorem 2]. That theorem gives
 2435 distinguishing or success-probability cost at most the random-permutation differentiating
 2436 bound $f_P(N_{\text{tot}})$ of [BDPVA08, Lemma 5]. The [BDPVA08, Lemmas 4 and 5] product
 2437 bounds give the random-transformation bound $f_T(N_{\text{tot}}) = 1 - \prod_{i=1}^{N_{\text{tot}}} (1 - i/2^c)$, and the
 2438 factor of $f_P(N_{\text{tot}})$ corresponding to the f_T factor at index i is that factor divided by a

quantity $1 - (i - 1)/2^{r+c} \in (0, 1]$, hence at least as large, so $f_P(N_{\text{tot}}) \leq f_T(N_{\text{tot}})$. With $N_{\text{tot}} < 2^c$ every factor lies in $[0, 1]$ and

$$f_P(N_{\text{tot}}) \leq f_T(N_{\text{tot}}) = 1 - \prod_{i=1}^{N_{\text{tot}}} \left(1 - \frac{i}{2^c}\right) \leq \sum_{i=1}^{N_{\text{tot}}} \frac{i}{2^c} = \frac{N_{\text{tot}}(N_{\text{tot}} + 1)}{2^{c+1}} = \text{Lift}_c(N_{\text{tot}}),$$

the second inequality by the Weierstrass product inequality $\prod_i (1 - x_i) \geq 1 - \sum_i x_i$ for $x_i \in [0, 1]$. $\square$

Lemma 17 (Derived-Adversary Execution Contract). *Let G be a single-stage TW128 game (Definition 3). Let $\mathcal{A}$ be a simulator-world G adversary whose public primitive interface is answered by $(S^{\text{RO}_{\text{flat}}}, S^{-1, \text{RO}_{\text{flat}}})$, with at most N_{prim} direct primitive-interface queries. Define $\mathcal{B}_{\text{derived}}(\mathcal{A})$ to run $\mathcal{A}$ internally, answer each primitive-interface query by running the same simulator against direct finite-output queries to its own RO_{flat} oracle, and forward construction interface queries (or, in the committing and discoverability games, the final structured output) unchanged to the corresponding random oracle experiment. Then the simulator-world execution of $\mathcal{A}$ and the random oracle execution of $\mathcal{B}_{\text{derived}}(\mathcal{A})$ have identical joint distributions, and*

$$q_{\text{RO}}(\mathcal{B}_{\text{derived}}) \leq N_{\text{prim}}.$$

In this coupling, simulator-induced RO_{flat} calls and construction-internal RF lookups use the same lazy table. Therefore, if a simulator-induced query is the bridged input $\text{flat}(X)$ of a well-formed TW128 transcript prefix X and a forwarded construction computation later reaches the same prefix X , both interfaces read the same $\text{RO}_{\text{flat}}(\text{flat}(X))$ value, projected to the requested output length. Moreover, $\mathcal{B}_{\text{derived}}$ inherits the construction-side caps of $\mathcal{A}$:

$$\begin{aligned} q_v(\mathcal{B}_{\text{derived}}) &= q_v(\mathcal{A}) \quad (AE, \text{mu-ind}), \\ n_{\text{leaf}}(\mathcal{B}_{\text{derived}}) &\leq n_{\text{leaf}}(\mathcal{A}) \quad (\text{all games above}). \end{aligned}$$

In the mu-ind game these resource inheritances hold per user, and the New-query transcript is unchanged.

Proof. Couple the two executions by using the same lazy RO_{flat} table for the simulator and for all construction-internal $\text{RF}(\cdot)$ lookups. In the simulator-world execution, $\mathcal{A}$ sees construction answers from the selected game interface and public primitive answers from $(S^{\text{RO}_{\text{flat}}}, S^{-1, \text{RO}_{\text{flat}}})$. In the derived random oracle execution, $\mathcal{B}_{\text{derived}}$ runs the same $\mathcal{A}$, forwards construction interface queries (or the final committing- or discoverability-game output) unchanged, and answers primitive-interface queries by invoking the same simulator code with the same lazy RO_{flat} table. The construction interface or committing challenger therefore receives the same forwarded inputs, the simulator receives the same primitive queries in the same order, and the lazy table supplies the same values. The full views of $\mathcal{A}$ in the simulator-world execution and of $\mathcal{B}_{\text{derived}}$ in the random oracle execution are identically distributed.

The simulator's random oracle calls may be on arbitrary [BDPVA08] H -input strings, not only bridged TW128 transcript inputs. These calls are the direct RO_{flat} queries of $\mathcal{B}_{\text{derived}}$ and are counted in $q_{\text{RO}}(\mathcal{B}_{\text{derived}})$. Construction-internal $\text{RF}(\cdot)$ lookups made by Enc_K^{RO} , Dec_K^{RO} , Ver_K^{RO} , or a committing or discoverability challenger are forwarded to the same RO_{flat} table but are not counted in q_{RO} , exactly as in Section 4.3.

By Lemma 15, each primitive-interface query induces at most one direct RO_{flat} query by $\mathcal{B}_{\text{derived}}$, and inverse queries or non-RO-touching forward queries induce none. Since $\mathcal{A}$ makes at most N_{prim} direct primitive-interface queries, this gives $q_{\text{RO}}(\mathcal{B}_{\text{derived}}) \leq N_{\text{prim}}$.

Lemma 6's bridge places TW128 construction restrictions and simulator-induced RO_{flat} calls in a single lazy table. If the simulator invokes RO_{flat} on $\text{flat}(X)$ for a well-formed

TW128 transcript prefix X , and a forwarded construction computation later reaches X , the construction lookup $\text{RF}(X) = \text{RO}_{\text{flat}}(\text{flat}(X))$ reads the same table entry and returns the requested projection. Since TW128 uses $r = r_{\text{max}}$, the [BDPVA11, Theorem 3] output post-processing does not alter TW128 output bits. Thus simulator-induced direct queries and construction-internal lookups at the same bridged input are transcript-consistent. In the AE-style games, a simulator call whose input is accepted by KeyedTWOut for the hidden key is still a direct RO_{flat} query and is covered by the same bad_{key} accounting used by the direct random oracle mu-ind proof. The derivation step itself does not prove a key-guessing statement; it only supplies the derived adversary and the query bound at which the random oracle bounds of Section 5.2 are instantiated.

Finally, the execution-uniform caps of Section 4.3 hold along every oracle-answer trace of $\mathcal{A}$, including traces generated by the simulator on the ideal public primitive interface. The derived-adversary construction does not add construction interface encryption or verification queries and forwards $\mathcal{A}$'s construction interface queries unchanged along each trace. $\mathcal{B}_{\text{derived}}$ therefore inherits the relevant construction resources: $q_v(\mathcal{B}_{\text{derived}}) = q_v(\mathcal{A})$ in the AE and mu-ind games, and $n_{\text{leaf}}(\mathcal{B}_{\text{derived}}) \leq n_{\text{leaf}}(\mathcal{A})$ for the distinct construction-generated final-leaf-transcript cap. In mu-ind, construction interface forwarding preserves the New calls and these caps user by user. $\square$

B.2 Lift Application

The lift bound of Lemma 16 and the derived-adversary construction of Lemma 17 combine into a single template, stated once and then instantiated by the headline theorems of Section 5.2 and the committing-notion bounds of Corollary 6.

Corollary 8 (Lift Application). *Let G be a single-stage TW128 game (Definition 3). For every public permutation G -adversary $\mathcal{A}$ against TW128 with the resources of Section 4.3, the derived adversary $\mathcal{B} = \mathcal{B}_{\text{derived}}(\mathcal{A})$ of Lemma 17 is a random oracle G -adversary with $q_{\text{RO}}(\mathcal{B}) \leq N_{\text{prim}}$ and the construction-side caps of Lemma 17—in particular $q_v(\mathcal{B}) \leq Q_v$ and $n_{\text{leaf}}(\mathcal{B}) \leq N_{\text{leaf}}$, where these apply—such that*

$$\text{Adv}_{\text{TW128}}^{\mathsf{G}}(\mathcal{A}) \leq \text{Lift}_c(N_{\text{tot}}) + \text{Adv}_{\text{RO}}^{\mathsf{G}}(\mathcal{B}).$$

If moreover $\text{Adv}_{\text{RO}}^{\mathsf{G}}(\mathcal{B}) \leq \varepsilon_{\mathsf{G}}(q_{\text{RO}}(\mathcal{B}), n_{\text{leaf}}(\mathcal{B}))$ for some ε_{G} nondecreasing in each argument, then

$$\text{Adv}_{\text{TW128}}^{\mathsf{G}}(\mathcal{A}) \leq \text{Lift}_c(N_{\text{tot}}) + \varepsilon_{\mathsf{G}}(N_{\text{prim}}, N_{\text{leaf}}).$$

Proof. By the flat sponge lift bound (Lemma 16), replacing the real public permutation and TW128 construction by $(\mathcal{S}^{\text{RO}_{\text{flat}}}, \mathcal{S}^{-1, \text{RO}_{\text{flat}}})$ and the flat-RO construction interface costs at most $\text{Lift}_c(N_{\text{tot}})$, where $N_{\text{tot}} = N + N_{\text{prim}}$ counts both the construction calls of N —including any construction oracle exposed during $\mathcal{A}$'s run and, in the win-event games, the challenger's post-output decryption or encryption checks—and the N_{prim} direct primitive queries. The derived-adversary contract (Lemma 17) expresses the resulting simulator-world execution as the identically distributed random oracle execution of $\mathcal{B} = \mathcal{B}_{\text{derived}}(\mathcal{A})$, with $q_{\text{RO}}(\mathcal{B}) \leq N_{\text{prim}}$ and the stated construction-side inheritances, giving the first bound. The second follows by the monotonicity of ε_{G} . When G is parameterized by a target fixed before $\mathcal{A}$ runs—the discoverability target T^* —the coupling holds for each fixed target and the target-independent bound holds for the maximum defining $\text{Adv}_{\text{TW128}}^{\mathsf{G}}(\mathcal{A})$. $\square$

C Vectorized Kernel Details

This appendix details the architecture-specific cores summarized in Section 7.1: the instruction selection, register layout, and remainder handling of the amd64 and arm64 kernels.

amd64. The **amd64** implementation provides AVX-512 and AVX2 cores and selects between them once at startup from the **AVX512VL** CPUID feature bit. The single-duplex permutation has two variants: a general-purpose-register core adapted from the standard library’s **KECCAK- p** permutation (using the lane-complement trick so that χ requires no separate negation), and an AVX-512 variant; the root duplex uses whichever matches the selected core. The eight-way AVX-512 core packs the lane-major state into 25 ZMM registers, one per lane with all eight instances in the register’s eight 64-bit slots, and performs the permutation with no cross-lane shuffles: rotations use **VPROLQ**, and **VPTERNLOGQ** fuses the θ D -term formation and the χ and-not into single three-input logic instructions. The eight-way AVX2 core, lacking 512-bit registers, runs the eight instances as two groups of four (four 64-bit lanes per YMM register), rotating by shift-or and computing χ with **VPANDN**. The fused chunk kernels combine absorption with the permutation. The AVX-512 steady-state kernel reads the eight chunks with contiguous loads, transposes each rate block into the lane-major layout in registers (**VPUNPCKLQDQ**/**VPUNPCKHQDQ** and **VSHUFI64X2**), absorbs on the lane-major state, and transposes each output block back for a contiguous store, leaving only the 7-byte partial rate lane to a masked gather; these sequential accesses are prefetcher-friendly, whereas per-lane gathers and scatters at the chunk stride β (**VPGATHERQQ**/**VPSCATTERQQ**) would stall on cold memory and cap the long-message rate. The register-resident kernel that drains a two-to-seven-chunk remainder instead uses that per-lane gather/scatter form under a mask, reading the active chunks in place: the transpose’s shuffle work is constant regardless of the chunk count and would be spent partly on inactive lanes, whereas the gathers scale with the active element count. The AVX2 path drains the same remainders with a variant of its grouped-by-four kernels that aims each inactive instance’s loads and stores at a dummy block inside the kernel frame with a zero advance stride, so inactive instances touch no memory beyond the remainder and the hot loop stays branch-free.

arm64. The **arm64** path uses NEON plus the Armv8.2 SHA-3 extension (**FEAT_SHA3**), available on Apple Silicon and on Neoverse V1/V2 and N2. The permutation is expressed with the extension’s fused instructions: **EOR3** (three-input XOR) for the θ column parities, **RAX1** (rotate-by-one and XOR) for the θ D -terms, **XAR** (XOR then rotate) to fuse the θ accumulation with the ρ rotation, and **BCAX** (bit-clear and XOR) for χ . The single-duplex core places one instance per 64-bit D lane across the 25 vector registers; the eight-way core packs two instances per 128-bit register and processes the batch of eight as four such pairs, using **ZIP1** and duplicating moves to interleave and extract the two instances of each lane in the fused chunk kernel. The same pair machinery backs a 2-wide kernel that drains a leaf-chunk remainder two chunks at a time—one instance per 64-bit half of a 128-bit register—at roughly twice the single-duplex per-chunk rate. Both batched kernels store the permutation state back into the batch after the closing block—the $\times 8$ kernel for its first pair, the 2-wide kernel for both instances—which is the writeback that permits the root duplex to occupy lane 0.

D Test Vectors

The following test vectors cover the corner cases of TW128 encryption, as specified in § 3.

All byte strings are written in hexadecimal. For each vector we list the key K , nonce N , associated data A , message M , ciphertext C , and tag T ; the sealed output of encryption is $C \parallel T$. The empty string is written ε . The leaf index is encoded as a little-endian 64-bit integer and the phase suffixes are appended least significant bit first.

To keep the longer vectors compact, any value exceeding 32 bytes is given by its length in bytes together with its SHA-256 digest rather than in full. Every message and associated data byte is defined by $m_i = (17i + 31) \bmod 256$ for $i = 0, 1, \dots$, so each such input is

2579 determined by its length alone and can be reconstructed and checked against the stated
 2580 digest; a ciphertext given by its length and digest can likewise be checked by encrypting
 2581 the corresponding message and hashing the result.

Vector 1. Empty associated data and empty message: the associated data and aggregation phases are both elided, and the root tag is emitted by the single closing message call with no leaves.

<i>K</i>	000102030405060708090a0b0c0d0e0f101112131415161718191a1b1c1d1e1f
2582 <i>N</i>	202122232425262728292a2b2c2d2e2f303132333435363738393a3b3c3d3e3f
<i>A</i>	ϵ
<i>M</i>	ϵ
<i>C</i>	ϵ
<i>T</i>	da65bbda0b20b1e936f5326458166633d1e5a3f7aebc0b318d24b051a4efa697

Vector 2. Multi-block associated data with an empty message: the associated data spans two rate blocks and no message bytes are processed.

<i>K</i>	000102030405060708090a0b0c0d0e0f101112131415161718191a1b1c1d1e1f
2583 <i>N</i>	202122232425262728292a2b2c2d2e2f303132333435363738393a3b3c3d3e3f
<i>A</i>	168 bytes, SHA-256 4c7d125a3488a7603e2f403f0010a6060b8321b60b0fd2471120ae627bbe759a
<i>M</i>	ϵ
<i>C</i>	ϵ
<i>T</i>	2093878418f3fca96ad3a8d1748a9be4d2a8b518e917d988fd703dcb7bf72173

Vector 3. Empty associated data and a single-block message: the root message phase only, no leaves.

<i>K</i>	000102030405060708090a0b0c0d0e0f101112131415161718191a1b1c1d1e1f
2584 <i>N</i>	202122232425262728292a2b2c2d2e2f303132333435363738393a3b3c3d3e3f
<i>A</i>	ϵ
<i>M</i>	1f30415263748596a7b8c9daebfc0d1e2f405162738495a6b7c8d9eafb0c1d2e
<i>C</i>	676a1441e20daea0f1ee4571ee4543842c946d59e0d3df6f1dd2f92dc4f876aa
<i>T</i>	1873332d359ff2fe5197488224af7bf6ad821875fdbd964100de424aa77342be

Vector 4. Empty associated data and a two-block message: the root keystream is chained across two rate blocks.

<i>K</i>	000102030405060708090a0b0c0d0e0f101112131415161718191a1b1c1d1e1f
2585 <i>N</i>	202122232425262728292a2b2c2d2e2f303132333435363738393a3b3c3d3e3f
<i>A</i>	ϵ
<i>M</i>	168 bytes, SHA-256 4c7d125a3488a7603e2f403f0010a6060b8321b60b0fd2471120ae627bbe759a
<i>C</i>	168 bytes, SHA-256 df3f0a98a91a69ed474edbc52e0dc1d8e54bfd6fc514966ce7b878d48be0bf9a
<i>T</i>	aab2413074a202cfadd5eb784663b1831eb330025972d163d03e9ed5a964b765

Vector 5. A message of exactly one chunk: the root is full, there is still no leaf, and the aggregation phase is elided so the closing root message call emits the tag.

<i>K</i>	000102030405060708090a0b0c0d0e0f101112131415161718191a1b1c1d1e1f
2586 <i>N</i>	202122232425262728292a2b2c2d2e2f303132333435363738393a3b3c3d3e3f
<i>A</i>	ϵ
<i>M</i>	8183 bytes, SHA-256 6532a370a4ab6b5f4094dd8a6016512ab8e8416abe910bc1a0b532979224151d
<i>C</i>	8183 bytes, SHA-256 8d58f8aa3413ea0ba1b8cc3ed4e9a334521d8dafa9694a0661fb8251049907e9
<i>T</i>	c2201637f1116c1303b8982622756770d23acb85384f4658fcb63c37e4b629dd

Vector 6. A message of one chunk and one further byte: one root chunk and one leaf, exercising leaf tag aggregation.

<i>K</i>	000102030405060708090a0b0c0d0e0f101112131415161718191a1b1c1d1e1f
2587 <i>N</i>	202122232425262728292a2b2c2d2e2f303132333435363738393a3b3c3d3e3f
<i>A</i>	ϵ
<i>M</i>	8184 bytes, SHA-256 c4c383bca0808e2ce44d481ddfe56efdc7a147dcd5a14c898ef5423f2604cd7
<i>C</i>	8184 bytes, SHA-256 ef1b2f7607ade503ab497fc61e62cef2b44c6328fa06e499718eb7da5fdd04d3
<i>T</i>	c96bea48e07a8431fd19721ffbe8afe6a73047f96d953e3281f10126eeca06a1

Vector 7. A multi-leaf message of two chunks and one further byte: the full tree path with two leaves.

2588 *K* 000102030405060708090a0b0c0d0e0f101112131415161718191a1b1c1d1e1f
N 202122232425262728292a2b2c2d2e2f303132333435363738393a3b3c3d3e3f
A ε
M 16367 bytes, SHA-256
 17e31f3700d0602aab8a686024ecc05b8192e80a531483c67dd90d5a26557c5e
C 16367 bytes, SHA-256
 6d83d8fbdc2a7c2ef590d29a36527a45019e50d1f8c8db037959081685a46908
T cdc2e50ac5380419f06772a2f3af2146050bd0340b35f61b67f839351767804e

Vector 8. Maximum associated data with a short message: the associated data blocks dominate the cost.

2589 *K* 000102030405060708090a0b0c0d0e0f101112131415161718191a1b1c1d1e1f
N 202122232425262728292a2b2c2d2e2f303132333435363738393a3b3c3d3e3f
A 337 bytes, SHA-256
 04ca7b24930bc6b5fe8b95ff1eb0933907657c1283cc3a7e69c5f7f02220f281
M 1f30415263748596a7b8c9daebfc0d1e
C 0dd615a0a05ed82a78c3999b9e80cbed
T 22e4e6725f548ad0660041d66cf2f18dc45c69b447cdaae1f147d33a3b6aa1a9

Vector 9. An all-ones key and nonce with a multi-chunk message.

2590 *K* ff
N ff
A 168 bytes, SHA-256
 4c7d125a3488a7603e2f403f0010a6060b8321b60b0fd2471120ae627bbe759a
M 8184 bytes, SHA-256
 c4c383bca0808e2ce44d481ddfe56efdc7a147dcd5a14c898ef5423f2604cd7
C 8184 bytes, SHA-256
 49818c320ab282296639798d840cbd60751e1623f449dcbf9423371d44145dfb
T 310c9fff20e9651f8ed47eed9bb182a086ed1af1b3f1799e35307a4139865fc7




Figure 2: Dependency graph of the supporting lemmas in this appendix and their consumers. Solid boxes are lemmas; dashed boxes are the proof components that consume them. The Bridge & Decode box is Lemma 6 together with the exact keyed decoder it specifies (Definition 2); the Key-Indexed Renaming box (Lemma 7) draws on it and on Lemma 4 to supply the key-field disjointness and renaming used by the hidden-key encryption and verification accounting. The Adaptive Candidate box (Lemma 8) is the single probabilistic engine: it bounds the Leaf Tag Collision lemma (Lemma 10) and, through the Root Tag Hash Bound (Lemma 2) of Section 6, the CMTDs-4, collision, preimage, and second-preimage proofs. The edge from Lemma 4 into the committing proofs in Section 6 and the proof of Corollary 6 is mediated by Opening Triple Separation (Proposition 1) for distinct contexts and Final-Root Input Separation (Proposition 2)—which adds the full-ciphertext recovery of Lemma 5 for distinct full inputs absent a leaf tag collision—together with Corollary 7 for candidate-sequence filtering. The Leaf Tag Collision box feeds the verification side of mu-ind through its hidden-key bound and the committing collision and second-preimage proofs through its adversary-supplied-key bound. The edge from Lemma 4 into Lemma 11 uses Corollary 7; the verification-first root tag guessing accounting reached through the same separation is carried out in Section A.9 (Lemma 13) and consumed by Lemma 14. The Lemma 9 edge records the shared hidden-key bad_{key} accounting used by the mu-ind proof.